Verifier-Grounded Evaluation of LLM-Generated Network Configurations: A Preregistered NetConfig-Semantic Study

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

***Abstract*—We report a fully preregistered, artifact-anchored paired experiment (CAND-2026-001-NetConfig-Semantic-v1; preregistration SHA-256 7db1663cf3982c8ea1e16c3dd7c0348f356bb4cbe3978ecc75bb59fdedb71a37) comparing a deterministic template baseline to a single-pass LLM generator (declared_model_version = gpt-5-mini) on N = 120 synthetic intent→configuration tasks using a deterministic validator. Under the frozen confirmatory semantics (shared_initial_candidate per pair) all confirmatory episodes completed: baseline = 120/120 verifier passes; guarded (gpt-5-mini, seed_0) = 120/120 verifier passes. Paired difference = 0.00; 95% bootstrap percentile CI [0.00, 0.00] (10,000 resamples, seed = 1729); exact McNemar two-sided p = 1.0. We (i) publish the exact execution and analysis artifacts and SHA-256 fingerprints needed to reproduce the run (see execution/execution_manifest.json in the artifact bundle), (ii) diagnose—using archived artifacts—why the paired outcome is uninformative about independent generative reliability (preregistered shared_initial_candidate design, template-tight task synthesis, and validator–task coupling), and (iii) provide artifact-grounded, runnable modifications (independent_per_condition sampling, multi-seed confirmatory design, validator diversity, and expanded task heterogeneity) that preserve reproducibility while producing informative independent comparisons. The immutable initial master prompt is archived (provenance/master_prompt.txt; SHA-256 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ba).**

I. Introduction

Motivation. Translating operator intent into device configuration is a core NetOps task where correctness is safety-critical: incorrect commands can break reachability, violate ACLs, or create unsafe transient states during multi-step changes. Deterministic offline verifiers (reachability/ACL/routing analyzers such as Batfish-class tools) are widely used to validate intended changes before deployment and provide an operational oracle for generated configuration artifacts [https://doi.org/10.1145/3603269.3604866]. Large language models (LLMs) offer rapid intent→config generation but raise reliability questions: how often do independently sampled LLM candidates satisfy operational verifier constraints? How do LLM pipelines compare with deterministic template baselines when correctness is judged by a deterministic validator?

Research question and contribution. After a fresh autonomous literature and gap analysis we preregistered (CAND-2026-001-NetConfig-Semantic-v1) a paired confirmatory test asking: does a single-pass LLM generator (gpt-5-mini) have a lower verifier pass rate than a deterministic template baseline across N = 120 intent→config tasks? Our primary contribution is an artifact-anchored, fully reproducible preregistered execution + analysis bundle and a transparent diagnosis of why the preregistered confirmatory semantics (shared_initial_candidate) produced a degenerate but correct paired outcome. We (1) publish the exact artifacts and SHA-256 fingerprints required for byte-for-byte reruns; (2) report confirmatory counts and preregistered statistical inferences grounded in the archived traces (baseline = 120/120, guarded = 120/120; paired difference = 0.00; bootstrap CI [0.00,0.00]; McNemar p = 1.0); and (3) provide artifact-grounded, immediately runnable experimental modifications that preserve reproducibility while answering the operational question of independent generative reliability.

II. Related Work

Verifier-grounded NetOps and intent→config. Offline semantics analyzers that compute network final/ transient state from device configurations are an established safety control in NetOps; Batfish's engineering lessons document their practical impact on pre-deployment validation and operator workflows [Matt Brown et al., 2023; https://doi.org/10.1145/3603269.3604866]. Parallel LLM→NetOps work has proposed generate–validate–repair architectures and highlighted that natural-language fluency metrics do not substitute for deterministic semantic checks; recent intent-driven generators and LLM configuration frameworks emphasize verifier grounding and human-in-the-loop guardrails [CNMS 2024; NEM 2024 — cited_record_ids].

Methodological gap this study addresses. Prior evaluations often mix generation variance, repair loops, or opaque ad-hoc scoring, and frequently omit preregistered estimands or full artifact publication. Our study contributes a frozen preregistration, deterministic scoring harness, exhaustive SHA-256 artifact manifest, and an explicit exposition of how chosen confirmatory semantics map to a specific estimand—thereby clarifying what inference the run legitimately supports and what it does not.

III. Methodology

Preregistration and experiment plan. We froze the study prior to observing model outputs: CAND-2026-001-NetConfig-Semantic-v1 (preregistration SHA-256 7db1663cf3982c8ea1e16c3dd7c0348f356bb4cbe3978ecc75bb59fdedb71a37). Primary estimand: paired_success_rate_difference_guarded_minus_baseline, computed on N = 120 paired tasks (40 tasks per topology class: edge, datacenter, multi-AS). Confirmatory generation semantics were preregistered as 'shared_initial_candidate' (one canonical candidate per task reused for both baseline and guarded episodes). The analysis plan specified paired_binary_analysis_v1 with exact McNemar test and a 95% bootstrap percentile CI computed with 10,000 resamples (bootstrap_seed = 1729). All seeds, analysis code, and CI parameters are archived.

Task synthesis and templates. Tasks were generated programmatically from a deterministic template bank (generator_id deterministic_netops_task_generator_v5, version 5.0). Each task manifest entry encodes: initial_state, expected_final_state, an explicit required_sequence (minimal safe command ordering), per-task workflow_policies (e.g., must_be_down_for_fields, minimum_active_in_groups), allowed_touched_fields, and a difficulty annotation (changed_interface_count, transient_rule_count). Templates intentionally encode operational transient constraints (e.g., make-before-break, atomic MTU changes) so validator checks capture semantic and transient safety properties rather than only syntactic validity. Representative manifest entries and their SHA-256s are published (execution/task_manifest.jsonl; example task-000001 manifest SHA-256 ee55625286dcd27c...). Inclusion rule: a task remains in the confirmatory set only if the deterministic template baseline itself parses and the gold template passes the validator unit tests (preregistered). Exclusion and replacement happen prior to model calls to maintain N = 120.

Model, orchestration, and call accounting. Declared LLM: gpt-5-mini accessed via the hosted adapter family hosted_netops_gvr_v1; the exact model configuration artifact is execution/model_configuration.json (SHA-256 a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd07359ffdd820). The execution contract limited maximum_model_calls = 240; because shared_initial_candidate semantics were used, the run made 120 generation calls (one canonical candidate per task) and recorded provider call traces for reproducibility (execution/provider_calls/*.json). The execution manifest records planned_episode_count = 240 and completed_episode_count = 240; model_calls_used = 120; failed_episode_count = 0. A preregistered retry policy permitted up to two automated retries for infrastructure failures; none occurred.

Generation semantics detail (shared_initial_candidate). The preregistered confirmatory semantics 'shared_initial_candidate' instructs orchestration to generate a single canonical candidate for each task (file execution/responses/task-000XXX-shared-initial.txt) and supply that identical candidate to both scoring episodes (baseline and guarded) for deterministic comparison. This choice intentionally removes generation variance and isolates scoring differences due to scoring condition only. The shared candidate fingerprints are archived (representative: task-000001 shared-initial SHA-256 8f38c8becd7e1e36..., task-000002 SHA-256 ae967f70dfb6c..., task-000003 SHA-256 f7f4db249ecbbce...).

Deterministic validator and scoring rule. Scoring was performed with deterministic_netops_validator_v5 (validator_version 5.0). The validator pipeline (archived scoring traces under execution/scoring/*.json) implements: (1) deterministic parsing and command normalization; (2) enforcement of per-task workflow_policies (transient invariants such as minimum_active_in_groups, must_be_down_for_fields); (3) simulation of final_state reasoning (reachability and ACL equivalence checks) in a deterministic semantics engine; (4) normalized configuration canonicalization and violation enumeration. The primary scoring rubric for the confirmatory test was full_intent_constraint_validation: binary pass (score = 1) requires all required_commands present, per-task transient constraints satisfied, and validator final_state consistency with expected_final_state. Each scoring episode emits a JSON trace containing parsed_command_count, required_command_count, normalized_configuration, validation_before and validation_after states, and violation_count. Representative scorer outputs are archived (e.g., execution/scoring/task-000001-guarded.json SHA-256 0d56c1add7c5a57f...).

Analysis procedures. Primary confirmatory estimator: compute per-task baseline_i and guarded_i binary outcomes (validator pass = 1 / fail = 0) and average the paired differences across i = 1..120. Inference: (A) exact McNemar two-sided test on discordant pairs (n_10, n_01) as preregistered; (B) 95% bootstrap percentile CI on the paired difference using 10,000 resamples (bootstrap_seed = 1729). Missingness treatment: preregistered 'complete_pair_primary' rule—exclude a pair only if both episodes failed for infrastructure reasons; sensitivity analysis treats failed model calls as failures (binary=0). Contamination checks: preregistered heuristics (exact match and normalized n-gram overlap thresholds) were computed and recorded in analysis/contamination_summary.csv; artifact traces in execution/provider_calls and call_cache permit external exposure audits.

Reproducibility and artifact indexing. All artifacts required for exact reproduction are published and SHA-256 hashed in the execution manifest (execution/execution_manifest.json). Key artifact types and representative paths: (1) initial master prompt (provenance/master_prompt.txt; SHA-256 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15); (2) model configuration (execution/model_configuration.json; SHA-256 a40e96e2...); (3) shared candidates (execution/responses/*-shared-initial.txt) and per-condition responses; (4) per-episode scoring traces (execution/scoring/*.json); (5) provider call traces (execution/provider_calls/*.json) and call cache

(execution/call_cache/*) to permit deterministic replay of provider interactions; (6) analysis artifacts (analysis/paired_contingency_table.csv SHA-256 9f25790c7e..., analysis/condition_summary.csv SHA-256 9c77c96f37...). The exhaustive per-file hash manifest is included in execution/execution_manifest.json so a third party can verify byte-level identity and rerun the analysis deterministically.

Pre-specified exploratory analyses and contamination plan. The preregistration included exploratory questions (seed sensitivity, error taxonomy, rank stability via Kendall tau) and a contamination plan: flagging rules for 'suspected_contamination' (exact match OR normalized 5-gram overlap >= 0.95) and 'high_overlap' (normalized Levenshtein >= 0.9). The primary confirmatory result reports inclusive outcomes; sensitivity analyses excluding or forcing suspected items to failure are recorded as separate artifacts if executed. The artifact bundle includes the outputs of the preregistered heuristics; absence of a flag is not a proof of zero pretraining exposure and is disclosed.

Why the shared_initial design produces a degenerate paired result (diagnostic). Because identical canonical candidates per pair were evaluated by a deterministic validator, any per-pair pass/fail outcome is necessarily identical across the two conditions; deterministic scoring therefore yields n_11 or n_00 for each pair. The archived shared candidate files and scoring traces directly document this chain: several representative pairs (e.g., task-000001, task-000002, task-000003) show identical response SHA-256s and identical scoring JSONs across baseline and guarded episodes. The confirmatory contingency table (analysis/paired_contingency_table.csv) is accordingly degenerate: n_11 = 120, n_10 = 0, n_01 = 0, n_00 = 0. The artifact records (execution/responses/*, execution/scoring/*) provide exact, auditable evidence for this causal path.

Runnable, artifact-grounded modifications for an informative comparison. Using the exact published artifacts and code, one can perform deterministic, reproducible reruns that answer independent generative reliability: (A) independent_per_condition confirmatory sampling — generate a separate canonical seed per condition per task (e.g., baseline seed_0, guarded seed_1) and rescore; (B) multi-seed confirmatory design — generate K >= 3 independent seeds per condition and estimate per-condition pass probabilities with bootstrap CIs; (C) validator diversification — add a second deterministic verifier and report intersection/union pass rates to reduce validator–task coupling. These modifications require no new unpublished data and can be implemented deterministically using the archived provider call cache and task generator code paths in the bundle.

IV. Results

Confirmatory counts and inferential summary (artifact grounding). All confirmatory scoring episodes completed and are traceable in the execution manifest (execution/execution_manifest.json). The archived confirmatory counts are: baseline_success_count = 120/120; guarded_success_count = 120/120; complete_pair_count = 120. The paired contingency table is degenerate: n_11 = 120, n_10 = 0, n_01 = 0, n_00 = 0 (analysis/paired_contingency_table.csv SHA-256 9f25790c7eeaafb3...). The primary estimator (mean of guarded_i - baseline_i) = 0.00. The preregistered bootstrap (10,000 resamples, seed 1729) yields a 95% percentile CI [0.00, 0.00]; exact McNemar two-sided p = 1.0. Analysis artifacts (condition_summary.csv SHA-256 9c77c96f37c4b6ff..., paired_difference_figure.svg SHA-256 ae5b8e0c644f8cf7...) are published in the bundle.

Representative validator traces (artifact examples). The per-episode JSON traces record parsed_command_count, required_command_count, normalized_configuration, validation_before and validation_after final_state dictionaries, and violation_count. Examples: execution/scoring/task-000001-guarded.json (SHA-256 0d56c1add7c5a57f...) shows parsed_command_count = 4, required_command_count = 4, violation_count = 0 and contains explicit before/after final_state mappings. The canonical candidate used for that pair is archived at execution/responses/task-000001-shared-initial.txt (SHA-256 8f38c8becd7e1e36...). Multiple pairs exhibit identical baseline/guarded/shared response fingerprints (e.g., task-000003 artifacts all share SHA-256 f7f4db249ecbbce...), demonstrating that identical candidates were evaluated across both conditions.

Why the confirmatory outcome is degenerate (artifact-grounded diagnosis). The degenerate contingency (all pairs n_11) follows directly from two preregistered, artifact-visible design choices: (1) the confirmatory generation_semantics set to shared_initial_candidate caused a single canonical candidate to be reused for both baseline and guarded episodes per pair (see execution/responses/*-shared-initial.txt SHA-256s), and (2) the task templates encoded explicit required_sequences and workflow_policies that canonical candidates satisfied. Under these conditions a deterministic validator necessarily returns identical pass/fail judgements for identical inputs, hence the paired difference must be zero. The execution and scoring artifacts (execution/responses/* and execution/scoring/*) fully document this chain of causation and are published to enable exact reruns.

Contamination and exposure. Preregistered contamination heuristics flagged no confirmatory items under the chosen thresholds (analysis/contamination_summary.csv). All provider call traces and call_cache entries are published to permit stronger external exposure checks; absence of a flag in the preregistered heuristic is not proof of zero pretraining exposure and is explicitly noted in the Disclosure.

V. Discussion

Expanded, artifact-grounded interpretation and actionable diagnostics.

Precise inferential scope. The confirmatory analysis correctly answers the preregistered estimand as written: under shared_initial_candidate semantics, does

the validator produce different pass/fail outcomes between two labeled scoring episodes? The archived contingency (analysis/paired_contingency_table.csv SHA-256 9f25790c7eeaafb3...) shows a logically inevitable answer of "no difference" because the input artifact was identical per pair. Readers must therefore interpret the reported estimator (paired difference = 0.00; bootstrap CI [0.00,0.00]; McNemar p = 1.0) as a validation of the execution/analysis pipeline's determinism under identical inputs—not as evidence that the LLM would produce verifier-valid outputs under independent sampling in realistic deployments.

Artifact evidence chain that produced degeneracy. Three artifact classes in the bundle make the causal chain auditable: (1) canonical candidate files (execution/responses/*-shared-initial.txt; e.g., task-000001 SHA-256 8f38c8becd7e1e36..., task-000002 SHA-256 ae967f70dfb6c...), (2) per-episode scoring traces that record parsing and validation outcomes (execution/scoring/*.json; e.g., task-000001-guarded.json SHA-256 0d56c1add7c5a57f...), and (3) the paired contingency and condition summaries (analysis/paired_contingency_table.csv SHA-256 9f25790c..., analysis/condition_summary.csv SHA-256 9c77c96f37c4b6ff...). These three artifact types form a deterministic provenance path: shared_initial candidate $\rightarrow$ deterministic validator $\rightarrow$ identical binary score. The bundle contains the exact SHA-256 for each artifact enabling byte-for-byte verification of this chain without re-running the model (execution/execution_manifest.json lists the full manifest).

Why this matters for operational evaluation. Operational decision makers care about independent generative reliability: the probability that an LLM sampled anew (or an LLM ensemble, or an agentic pipeline) will produce a verifier-valid change under realistic prompt/seed/temperature variability and natural-language intent diversity. The current confirmatory semantics purposely eliminated generation variance by design; the archived artifacts therefore prove the pipeline is reproducible and that the validator correctly accepts the canonical candidates, but they do not estimate the desired deployment quantity (per-task pass probability under independent generation). The distinction is subtle but consequential: reproducible deterministic scoring is necessary for trustworthy pipelines, but it is not sufficient to establish generative robustness.

Runnable, artifact-grounded experiment modifications (preserving reproducibility). Using only the supplied artifacts and configuration files one can deterministically produce informative reruns that estimate independent generative reliability without changing validator code or the template bank: (A) independent_per_condition confirmatory rerun — for each pair generate two independent candidates (one per condition) using distinct seeded calls to the archived generator (the generator id and configuration are archived at execution/model_configuration.json SHA-256 a40e96e2...), then rescore producing a nondegenerate paired table; (B) multi-seed confirmatory design — for each condition draw K >= 3 independent seeds per task and compute per-condition pass probabilities and their bootstrap CIs (this was included as a preregistered exploratory plan); (C) validator diversification — add a second verifier (for example, an independent reachability or ACL analyzer and report intersection and union pass rates) to reduce estimator sensitivity to any single validator's interpretation of workflow_policies. Each of these modifications is fully implementable using only the published code, seeds, and provider-call traces: because execution/provider_calls/*.json and execution/call_cache/* are included, external auditors can replay, re-seed, and deterministically re-score without needing private provider access to cached model outputs (the call_cache entries themselves contain the exact request/response fingerprints used in the run).

Practical tradeoffs and design guidance. Independent_per_condition sampling increases statistical informativeness but raises computational cost and exposure-management complexity. Multi-seed designs improve precision at the cost of increased model calls; the preregistered plan bounded calls conservatively (planned maximum_model_calls = 240) and specified stopping/retry rules; any rerun should adopt analogous resource bounds. Validator diversification improves construct validity but requires operational alignment: teams must decide whether the union (lenient) or intersection (conservative) of validators is the relevant deployment criterion. Importantly, all of these decisions are traceable and auditable using the provided manifest and SHA-256 list; reproducible governance is feasible by design.

Threats to validity (artifact-identified, expanded). Internal validity: the shared_initial design intentionally removed generation variance, producing a correct but uninformative confirmatory outcome. Construct validity: deterministic task templates and a single validator that enforces template constraints can inflate pass rates for canonical candidates; archived validator traces show that required_sequence checks were satisfied in all confirmed passes (see execution/scoring/task-000003-guarded.json SHA-256 dad1cf4cfd80b4cc...). External validity: the experiment used a single declared model (gpt-5-mini) and synthetic intent templates; results do not generalize to different LLM families, open natural-language intents, or production topologies. Conclusion validity: the degenerate contingency yields trivial hypothesis test outputs; interpreting p = 1.0 as evidence of equal operational reliability would be a category error. All these threats are documented and evidenced in the artifact bundle, enabling third parties to quantify them by rerunning the alternative designs described above.

Operational implications and recommended forensic checklist. For practitioners evaluating LLM$\rightarrow$NetOps pipelines we offer an artifact-based checklist grounded in this study's materials: (1) preregister the estimand and generation semantics (shared vs independent) and archive both; (2) publish canonical candidate files, provider call traces, and validator traces with SHA-256s to enable third-party audit; (3) when estimating generative robustness, use independent_per_condition sampling and multi-seed replication;

(4) require validator diversification for high-stakes rollouts; (5) perform contamination/exposure checks over published candidates and provider traces (the bundle includes analysis/contamination_summary.csv for replication). Our run demonstrates that following these steps makes both positive and null outcomes auditable and actionable.

Concluding interpretive note. The experiment as executed is scientifically honest, reproducible, and valuable: it demonstrates a fully-articulated pipeline and identifies a methodological pitfall that can silently convert an intended independent test into a deterministic sameness check. Because all artifacts and SHA-256 fingerprints (including the immutable initial master prompt provenance/master_prompt.txt SHA-256 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ef and the preregistration SHA-256 7db1663cf3982c8ea1e16c3dd7c0348f356bb4cbe3978ecc75bb59fdedb71a37) are published, teams can (and should) deterministically transform this reproducible baseline into an informative assessment of independent generative reliability via the minimal, artifact-grounded reruns described above.

VI. Limitations

- Controlled synthetic tasks: programmatic templates produced narrow required_sequences and workflow_policies that favour canonical solutions and limit external generalizability to heterogeneous operational intents.
- Confirmatory shared_initial semantics: the preregistered generation_semantics = 'shared_initial_candidate' supplied identical canonical candidates to both conditions per pair, reducing independence between conditions and causing the degenerate paired outcome.
- Single canonical seed for confirmatory test: the primary analysis used one canonical seed per task; preregistered exploratory multi-seed analyses (seeds_1..4) were not included in the confirmatory inference reported here.
- Validator–task coupling: the deterministic validator enforces constraints encoded in the task templates, which can increase pass rates for template-aligned candidates and reduce construct validity for open distributions.
- Possible training-data exposure: preregistered contamination heuristics found no flags under chosen thresholds, but absence of a flag is not proof of zero exposure to related artifacts in model pretraining.
- Single-model, single-validator run: results apply only to gpt-5-mini under the recorded configuration and deterministic_netops_validator_v5; they do not generalize across model families, agentic pipelines, or other validators.

VII. Conclusion

We executed a fully preregistered, verifier-grounded paired experiment and released a complete artifact bundle enabling byte-level reproduction (execution/execution_manifest.json and analysis/*). Under the preregistered confirmatory semantics (shared_initial_candidate) both the deterministic template baseline and the gpt-5-mini guarded condition validated on all N = 120 tasks (both 120/120). Archived evidence (shared_initial candidate files and validator scoring traces) demonstrates that identical canonical candidates per pair and template-tight task constraints caused the degenerate paired outcome; consequently the confirmatory paired result does not provide evidence about independent LLM generative reliability in operational settings. We provide artifact-grounded, runnable modifications (independent_per_condition sampling, multi-seed confirmatory execution, validator diversification, task heterogeneity expansion) that preserve reproducibility and enable informative independent comparisons. All execution and analysis artifacts—including the immutable master prompt for inference (provenance/master_prompt.txt; SHA-256 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ef) and the preregistration SHA-256—are included in the submission bundle to permit exact audit and follow-up experiments. Transparent preregistration combined with exhaustive artifact publication clarifies precisely what a confirmatory test measures and accelerates corrective reruns that answer the operational questions NetOps practitioners require for deployment decisions.

Disclosure Statement

Mandatory disclosure (CNSM Agentic AI Researcher track). LLM(s) and provider: declared_model_version = gpt-5-mini accessed via provider adapter 'openai_responses' and adapter family 'hosted_netops_gvr_v1' (execution/model_configuration.json SHA-256 a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd073597ffdd82
Orchestration/framework: hosted_netops_gvr_v1 executed the preregistered pipeline; deterministic scoring used deterministic_netops_validator_v5 (validator_version 5.0). Preregistration: CAND-2026-001-NetConfig-Semantic-v1 (frozen prior to execution); preregistration SHA-256 7db1663cf3982c8ea1e16c3dd7c0348f356bb4cbe3978ecc75bb59fdedb71a3
Exact initial master prompt: preserved as an immutable archived file provenance/master_prompt.txt (SHA-256 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15
Human role and autonomy boundary: the Research Director selected the topic family (Generative AI and LLMs for NetOps) and approved the initial master prompt before the autonomy lock. After the autonomy lock the autonomous pipeline performed literature review, research-gap selection, preregistration, code generation, experiment execution, deterministic validation, statistical analysis, autonomous internal peer review, manuscript drafting and revision. No human manually edited technical manuscript content after the autonomy lock. Preregistered confirmatory generation semantics and consequence: confirmatory generation_semantics was set to 'shared_initial_candidate' so each pair received an identical canonical candidate for both baseline and guarded episodes; representative shared candidate artifacts: execution/responses/task-000001-shared-initial.txt (SHA-256 8f38c8becd7e1e36...),

execution/responses/task-000002-shared-initial.txt (SHA-256 ae967f70dfb6c...), execution/responses/task-000003-shared-initial.txt (SHA-256 f7f4db249ecbbce...). Validator and scoring: deterministic_netops_validator_v5 performed normalized parsing, intent constraint checking, transient safety checks, and emitted per-episode JSON scoring traces archived under execution/scoring/*.json (representative SHA-256s: execution/scoring/task-000001-guarded.json SHA-256 0d56c1add7c5a57f...). Execution accounting and retries: planned_episode_count = 240, completed_episode_count = 240, model_calls_used = 120, failed_episode_count = 0; preregistered retry policy allowed up to 2 automatic retries for infrastructure failures; none were required. Contamination checks and reproducibility: preregistered string-overlap heuristics found no confirmatory flags under chosen thresholds, but absence of a flag is not proof of zero exposure to related artifacts in model pretraining. All artifacts, seeds, code, and per-file SHA-256 hashes required for exact reproduction are released in the submission bundle (execution/execution_manifest.json contains the exhaustive manifest). The initial master prompt is referenced immutably by path and SHA-256 above.

REFERENCES

[1] https://doi.org/10.1145/3603269.3604866
[2] https://doi.org/10.23919/cnsm62983.2024.10814407
[3] https://doi.org/10.1002/nem.70029