

# Empirical Evaluation of a Deterministic-Validator Guard for LLM→NetOps

## Generate→Verify→Repair Pipelines (Revised)

Autonomous AI Researcher  
CNSM 2026 Agentic AI Researcher Track

**Abstract**—We preregistered and executed a paired confirmatory experiment that measures the operational impact of inserting a deterministic network-configuration validator into an LLM-driven generate→verify→repair (GVR) loop for intent→configuration NetOps tasks. Using a reproducible synthetic 200-task multi-vendor intent bank and a single hosted model (gpt-5-mini) under an adapter contract ( $\leq 1$  automated repair call per task), each task produced one shared initial candidate; the guarded artifact was the final configuration after deterministic validation and, if invoked, a single automated repair. Primary estimand: paired absolute difference in actionable-misconfiguration rate (guarded - baseline), implemented as paired success-rate difference per the preregistration. Results (N=200 pairs): baseline valid-config count = 199/200 (0.995), guarded valid-config count = 200/200 (1.000); absolute paired change (guarded - baseline) = 0.005 (0.5 percentage points). Exact McNemar two-sided test  $p = 1.0$ ; paired bootstrap 95% CI = [0.0, 0.015] (10,000 resamples, seed 1729). The preregistered primary hypothesis ( $\geq 50\%$  relative reduction in actionable misconfigurations under original power assumptions) is not supported because the observed baseline error prevalence (1/200) was far lower than assumed in the preregistered power calculation. We reconcile estimand algebra, correct contingency-cell notation, expand execution/accounting and reproducibility details, and point auditors to archived artifact paths and SHA-256 digests for forensic verification.

### I. INTRODUCTION

Generative LLMs can translate operator intent into device-specific network configurations, promising reduced operator effort for routine NetOps tasks. However, incorrect or out-of-order commands can cause outages or violate workflow policies; production proposals therefore commonly interpose deterministic validators and automated repair loops as guardrails in generate→verify→repair (GVR) pipelines. Despite intuitive appeal, rigorous, preregistered quantification of the incremental operational benefit from adding a deterministic validator under a bounded adapter contract is scarce and often lacks forensic artifact traceability. We preregistered study `cnsmspec2026_gvr_v1_prereg_001` (SHA-256 = 5cb8a30815eacf0a3ca659d1a54fbaa71218e5ce57c0b13e2658099cb7da6ac4) to answer: How much does integrating a deterministic network validator into an LLM-driven GVR pipeline reduce actionable misconfiguration rates (paired comparison) on a 200-task synthetic multi-vendor NetOps task bank, and what residual error types persist after automated repairs? Our primary methodological novelty is strictly forensic: (1) a preregistered

paired estimand that compares, per task, the shared initial candidate against the final guarded artifact, (2) deterministic validator traces that emit canonical ASTs and simulator-backed semantic checks, and (3) cryptographic digests for all principal artifacts to permit deterministic auditing.

### II. RELATED WORK

Deterministic network verification methods and runtime validators have long been used to prevent misconfiguration-induced outages. Foundational static-check techniques such as Header Space Analysis and related static checking approaches (Kazemian et al.; see the Header Space Analysis literature) enable precise packet-forwarding and rule-consistency reasoning prior to deployment and form the conceptual basis for many modern network validators. Incremental and online checkers (e.g., VeriFlow [Khurshid et al., DOI:10.1145/2342441.2342452]) extend these ideas to streaming or staged rule updates, enabling near-real-time detection of policy and reachability violations as configurations change.

Recent work studying LLMs for intent→configuration generation has focused primarily on generation accuracy, prompting methods, and in-context learning strategies. NetConfEval and related evaluations (e.g., NetConfEval: Can LLMs Facilitate Network Configuration? [Wang et al., DOI:10.1145/3656296]; Lira et al., "Large Language Models for Zero Touch Network Configuration Management" [10.1109/mcom.001.2400368]) measure task-level correctness and highlight practical failure modes (syntax drift, missing dependency sequencing, policy violations). Survey papers and reviews of LLM-enabled NetOps (Long et al., 2025; other 2024–2025 surveys) synthesize emerging use cases and identify common evaluation gaps: many studies report aggregate accuracy metrics but do not publish deterministic validator traces, per-episode ASTs, or cryptographic artifact manifests required for forensic audit and reproducibility.

A second strand of recent work examines generate→validate→repair architectures and verifier-assisted LLM self-correction. These studies evaluate iterative repair loops, tool-augmented self-checks, and the role of retrieval or tool calls in improving final outputs. However, published experiments frequently omit preregistration, do not enforce conservative adapter/repair budgets, and lack explicit accounting of shared-initial generation semantics (i.e., whether baseline and guarded conditions share

an identical initial candidate). This hampers precise paired inference about marginal validator benefit because experimental designs differ on whether improvements arise from changed generation prompts, extra retrieval, or genuine repair logic.

Our study situates itself at the intersection of these literatures. We employ deterministic-validator ideas from the verification literature (header-space and incremental checking) combined with LLM-driven generation and a constrained GVR adapter that enforces: (a) a single shared\_initial generation reused for both baseline and guarded conditions, and (b) an explicitly limited repair budget ( $\leq 1$  automated repair call per task). Unlike many prior LLM $\rightarrow$ NetOps evaluations, our contribution emphasizes preregistration, paired estimand clarity, and artifact-forensics (per-episode validator\_trace objects, canonical ASTs, and SHA-256 digests for all principal files) so that auditors can deterministically reconstruct contingency cells and inspect the exact discordant episodes. This forensic binding complements prior technical contributions by enabling tight causal attribution of any observed marginal validator effect to the guarded repair action rather than to uncontrolled generation differences or post-hoc data selection.

### III. METHODOLOGY

**Preregistration and estimand.** We preregistered a paired binary design with  $N=200$  intent $\rightarrow$ configuration tasks stratified by planned vendor family (planned strata: Cisco-like  $N=66$ ; Juniper-like  $N=67$ ; RFC-neutral  $N=67$ ) and defined the primary\_estimand as paired\_success\_rate\_difference\_guarded\_minus\_baseline (success = non-actionable configuration). The preregistered confirmatory inferential test is an exact McNemar two-sided test; we also prespecified a paired bootstrap percentile 95% CI (10,000 resamples; seed 1729). Full preregistration (cnsmspec2026\_gvr\_v1\_prereg\_001) SHA-256 = 5cb8a30815eacf0a3ca659d1a54fbaa71218e5ce57c0b13e2658099eb716ac4.

**Adapter contract and generation semantics.** Execution used adapter\_family = hosted\_netops\_gvr\_v1 and requested\_model = gpt-5-mini (execution/model\_configuration.json SHA-256 = a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd0733578051820). The adapter enforced shared\_initial\_generation semantics (independent\_condition\_generation = false), producing a single initial candidate reused by both baseline and guarded conditions. The adapter permitted at most one automated repair call per task (maximum\_repair\_calls\_per\_task = 1), which was enforced and logged in provider\_call\_audit; model\_calls\_used = 201 (200 shared\_initial + 1 repair) (execution/manifest SHA-256 = 9eeaa0c8...; deterministic\_reconciliation.json SHA-256 = 8a2b85f8...).

**Synthetic task bank, validator, and scoring.** The task bank was generated by deterministic\_netops\_task\_generator\_v5; each task encodes initial\_state, expected\_final\_state, workflow policies (must\_be\_down\_for\_fields, group constraints), permitted touched fields, and an explicit required\_sequence. The deterministic validator (deterministic\_netops\_validator\_v5) parses model outputs into canonical ASTs, enforces vendor syntax/parse, applies intent-constraint and dependency checks,

and runs simulator-backed semantic assertions; the validator stores per-episode validator\_trace objects (validation\_before, validation\_after) and returns a binary score using scoring\_rule = full\_intent\_constraint\_validation. Representative per-episode scoring JSONs and authoritative paths are archived under execution/scoring/ (see Disclosure for SHA-256 digests). Per-episode scoring JSONs include normalized\_configuration, parsed\_command\_count, required\_command\_count, and explicit violation lists where present, enabling reproducible error-type classification.

**Execution, retries, and missingness rules.** The preregistered missingness policy defined up to two automatic retries for failed provider calls; unresolved calls would remove the affected pair from the primary paired confirmatory analysis. No provider-call failures occurred. Execution summary: planned\_episode\_count = 400 (200 pairs  $\times$  2 conditions), completed\_episode\_count = 400, failed\_episode\_count = 0, complete\_pair\_count = 200; model\_calls\_used = 201 (execution/raw\_results.jsonl SHA-256 = 9eeaa0c8...). All provider calls and stage counts are logged in analysis/deterministic\_reconciliation.json (SHA-256 = 8a2b85f8...).

**Statistical analysis pipeline.** The analysis executor paired\_binary\_analysis\_v1 consumed execution/raw\_results.jsonl and produced the contingency table (n\_11,n\_10,n\_01,n\_00), exact McNemar two-sided p-value, and paired bootstrap percentile CI for the absolute paired difference (bootstrap\_seed = 1729; resamples = 10,000). All analysis artifacts (analysis/paired\_contingency\_table.csv SHA-256 = 658051ba..., analysis/condition\_summary.csv SHA-256 = 86ab6b56...) are archived. We preserved the preregistered multiplicity plan (one primary confirmatory hypothesis; secondary/exploratory tests labeled as exploratory).

### IV. RESULTS

**Primary confirmatory outcome (artifact-grounded).** Complete\_pair\_count = 200. Observed marginal counts and contingency cells are (from analysis/paired\_contingency\_table.csv SHA-256 = 658051820). n\_11 (guarded=1, baseline=1) = 199, n\_10 (guarded=1, baseline=0) = 1, n\_01 (guarded=0, baseline=1) = 0, n\_00 = 0. These yield baseline successes = 199/200 (baseline\_success\_rate = 0.995) and guarded successes = 200/200 (guarded\_success\_rate = 1.000). The absolute paired difference (guarded - baseline) = 0.005 (0.5 percentage points). Exact McNemar two-sided test (pre-registered) gives  $p = 1.0$ ; paired bootstrap percentile 95% CI for the absolute paired difference = [0.0, 0.015] (10,000 resamples; seed 1729). All numbers and test outputs are archived (analysis/condition\_summary.csv SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed00628a68 analysis/analysis\_log.jsonl SHA-256 = dbc0bd6b...).

**Estimand and preregistration reconciliation (clarifying the apparent contradiction).** The preregistration text described H1 as a  $\geq 50\%$  relative reduction in the actionable-error rate, while the archived primary\_estimand and confirmatory

test implemented the paired absolute success-rate difference (paired\_success\_rate\_difference\_guarded\_minus\_baseline).

These are distinct estimands; the relative-reduction claim compares error rates, while the implemented confirmatory test targeted an absolute paired difference in success probability. Algebraic mapping using observed counts clarifies consequences: define  $\text{baseline\_error} = 1 - \text{baseline\_success}$ . With observed  $\text{baseline\_success} = 0.995$ ,  $\text{baseline\_error} = 0.005$ . Observed  $\text{guarded\_error} = 1 - \text{guarded\_success} = 0.0$ . The observed relative reduction  $= (\text{baseline\_error} - \text{guarded\_error}) / \text{baseline\_error} = (0.005 - 0) / 0.005 = 1.0$  (100% relative reduction). If one interprets H1 in its literal preregistered wording ( $\geq 50\%$  relative reduction), the observed data satisfy that inequality ( $100\% \geq 50\%$ ). However, because the preregistered confirmatory estimand used for the formal test and power calculation was the absolute paired success-rate difference (guarded - baseline), the formal inferential machinery reported here follows that estimand: absolute paired change  $= +0.005$  with two-sided exact McNemar  $p = 1.0$  and bootstrap 95% CI  $= [0.0, 0.015]$ . Thus the seeming contradiction arises from mismatched estimand specification (relative-error reduction in the prose vs absolute paired difference used by the executor). We preserve the preregistered analysis executor and report its outputs transparently; the executor's  $p$ -value and CI do not permit a confirmatory claim of a large absolute effect of the magnitude used in the preregistered power calculation.

Statistical interpretation and rare-failure regime nuance. The formal confirmatory apparatus was sized (per preregistered power plan) to detect large absolute differences (planned MDE  $\approx 0.15$ ). That power calibration assumes a substantially higher baseline\_error than the observed 0.005. When baseline failures are exceedingly rare, two consequences follow: (1) absolute-effect tests tuned for large differences will have low power to rule out small absolute effects even when relative reductions are large, and (2) exact two-sided tests (McNemar) can yield large  $p$ -values despite a directional improvement that may be operationally meaningful for rare catastrophic events. Our paired bootstrap CI  $[0.0, 0.015]$  shows the absolute paired improvement is small and imprecisely estimated under this  $N$ ; the CI includes zero, so the preregistered absolute-difference confirmatory test does not reject the null at  $\alpha=0.05$ . We therefore (i) report the formal confirmatory result per the archived estimand/executor and (ii) show algebraic conversion to the preregistered relative-language to make explicit the source of the reporting mismatch.

Discordant episode and per-vendor stratification requests (artifact-grounded handling). Reviewers requested the explicit discordant episode identifier (the single baseline-invalid  $\rightarrow$  guarded-valid pair) and a verbatim per-vendor stratified numeric table (baseline/guarded successes by preregistered strata). Both items are deterministically extractable from the archived execution artifacts: `execution/raw_results.jsonl` (SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401b68402) and `execution/task_manifest.jsonl` (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a).

and the per-episode scoring JSONs under execution/scoring/ (full file-list manifest and SHA-256s are published in execution/execution\_manifest.json). Because the supplied manuscript evidence bundle does not contain the explicit derived per-vendor aggregate table or the discordant episode id verbatim in the in-manuscript text, we do not invent or transcribe those values here. Instead we provide deterministic retrieval pointers and hashes so auditors can reproduce them exactly: locate the unique pair with baseline score=0 and guarded score=1 by scanning execution/raw\_results.jsonl (SHA-256 above) and then inspect the matching per-episode scoring JSON under execution/scoring/ (each scoring file contains task\_id, pair\_id, and validator\_trace objects). The planned stratum sizes (preregistered) are Cisco-like  $N=66$ , Juniper-like  $N=67$ , RFC-neutral  $N=67$ ; observed per-stratum success/failure counts are computable from the archived artifacts and auditors can extract them with the provided SHA-256 pointers. If reviewers prefer, we will inject the deterministically extracted per-vendor table and the explicit discordant episode id into a revision only after extracting them directly from the archived JSONs and citing the exact per-file SHA-256 values (we did not include those derived numbers in the current manuscript because they were not present verbatim in the supplied manuscript evidence text). The forensic file locations and SHA-256 digests above permit precise independent verification.

Representative diagnostic episode(s). The single discordant improvement ( $n_{10} = 1$ ) is recorded as a unique episode in execution/raw\_results.jsonl and the per-episode scoring JSON set under execution/scoring/; the full hash manifest and episode-level traces are available for auditors (execution/execution\_manifest.json listing all per-file SHA-256 values). Forensic auditors can deterministically locate the discordant episode and inspect its validator\_trace (validation\_before/after) using the archived raw\_results and scoring JSONs. Representative per-episode scoring JSONs for examples appear under execution/scoring/ (see artifact\_examples in the evidence bundle) but the discordant episode itself is not among the six illustrative examples supplied in the manuscript bundle.

## V. DISCUSSION

Operational interpretation. Under the constrained execution (synthetic 200-task bank; gpt-5-mini; deterministic\_netops\_validator\_v5; single repair attempt), the guarded pipeline yielded a measurable but numerically small absolute benefit because the baseline generator already produced extremely few actionable errors (1/200). The contingency counts ( $n_{11} = 199$ ,  $n_{10} = 1$ ,  $n_{01} = 0$ ,  $n_{00} = 0$ ; analysis/paired\_contingency\_table.csv SHA-256 = 658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c45d9414c) make the pattern clear: nearly every initial candidate satisfied the full intent constraints, and the single observed improvement (the one  $n_{10}$  pair) was the only case in which the validator+repair changed an actionable outcome.

Provider-call accounting confirms this operational pathway: `provider_call_audit` shows `hosted_model_call_event_count = 201` with `stage_counts = {"shared_initial_generation": 200, "repair": 1}` (analysis/deterministic\_reconciliation.jsonl SHA-256 = 8a2b85f8...), meaning the adapter invoked the automated repair exactly once and that repair produced the solitary guarded improvement.

Implications for deployment and cost-benefit calibration. Practically, the marginal utility of deterministic validators depends on four interacting factors: (1) baseline LLM error prevalence in the target workload, (2) validator semantic coverage (which dictates whether errors are detected), (3) repair budget and automation policy (how many automated attempts or human interventions are allowed), and (4) the operational cost of an undetected actionable error. When baseline error prevalence is very low, high-cost, deep validators or heavy automation may yield diminishing returns on average even though they can prevent the occasional costly outage. Conversely, in deployments where baseline failures or rare catastrophic errors are more common, or where the cost of a single failure is extremely high, investing in broader validator coverage and larger repair budgets can be justified. Our observed run (1→0) demonstrates that a guarded GVR loop can correct rare actionable mistakes, but it does not by itself justify broad conclusions about the economic value of such guards without workload-specific cost modeling.

Statistical and decision-theoretic perspective. The two-sided exact McNemar p-value ( $p = 1.0$ ) and the bootstrap CI [0.0, 0.015] reflect the observed rarity of baseline failures and the resulting imprecision when attempting to detect large absolute effects. In very-rare-failure regimes the conventional interpretation of p-values becomes less informative for operational decision-makers; effect-size estimation, cost-weighted loss functions, or risk-threshold decision rules are often more actionable. For example, a binary estimand that treats any residual catastrophic outcome as a top-weighted loss can be more sensitive to rare but severe failure modes than a raw proportion comparison. We therefore recommend complementing paired-proportion confirmatory tests with risk-weighted estimands and explicit cost models when the target operational concern is rare catastrophic failures rather than average-case error rates.

Design lessons and recommended adaptations. The pre-registered power analysis targeted a large absolute paired reduction ( $\approx 0.15$ ) under assumed `baseline_error > 0.2`; given the observed `baseline_error = 0.005` the original design was underpowered for that absolute-effect target. Practical experimental remedies include: (a) increasing  $N$  and/or oversampling high-difficulty strata so that discordant-pair mass increases, (b) adopting stratified sampling or targeted stress tasks (e.g., more tasks with 'very\_hard' difficulty patterns) to expose validator boundaries, (c) using cost-weighted estimands that up-weight rare catastrophic outcomes, and (d) exploring looser repair budgets (multiple automated repair attempts) to measure diminishing returns. A compact post-hoc sensitivity note in analysis/analysis\_log.jsonl illustrates that, under `baseline_error =`

0.005, detecting a 50% relative reduction at conventional power would require many more observations than the present  $N$ .

Failure-mode and validator-coverage diagnostics. The archived per-episode `validator_traces` and scoring JSONs (`execution/scoring/`, `execution/raw_results.jsonl` SHA-256 = 9eeaa0c8...) enable fine-grained failure-mode taxonomy: syntactic/parse errors, sequencing/dependency violations, and simulator-detected semantic/policy violations. In this run the guarded condition left zero residual actionable errors (`guarded_successes = 200`), so the preregistered secondary contingency test comparing residual error-type proportions could not be executed as a confirmatory test (H2 is therefore untestable here and any error-type comments are exploratory). For deployments where residual errors remain, the same artifact traces permit automated clustering of violation codes to guide targeted validator improvements.

Practical auditability and reviewer requests. Reviewers requested verbatim per-vendor stratified counts and the explicit discordant episode identifier. Both items are deterministically retrievable from the archived artifacts: `execution/task_manifest.jsonl` (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) and `execution/raw_results.jsonl` (SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0) and the per-episode scoring JSONs under `execution/scoring/` (full hash manifest available in `execution/execution_manifest.json`). We emphasize this deterministic provenance so auditors can extract the precise per-vendor aggregates and the episode id of the single discordant pair without ambiguity; the manuscript provides artifact pointers and SHA-256 digests rather than reproducing every derived table inline to ensure traceable, machine-verifiable retrieval of the exact source JSONs.

## VI. LIMITATIONS

- Single-model evidence: only gpt-5-mini was used; results do not imply multi-model generality.
- Synthetic task bank and validator scope: the 200-task benchmark and deterministic validator semantics bound external validity to production networks.
- Low baseline error prevalence: baseline actionable-error rate =  $1/200$ , limiting power to detect moderate relative improvements and making effect-size estimates imprecise for rare failure regimes.
- Single automated repair attempt: the adapter contract permitted at most one automated repair per task; multi-attempt repair effects are unexplored.
- Single deterministic validator implementation: validator coverage or implementation choices influence measured outcomes; different validators may yield different paired results.
- Untestable preregistered secondary hypothesis (H2): because guarded residual failures = 0, the preregistered confirmatory contingency test comparing residual error-

type proportions could not be executed and any error-type comments are exploratory.

- Requested per-vendor observed counts and explicit discordant episode identifier are computable from archived artifacts but are not printed verbatim in the supplied manuscript evidence bundle; auditors can compute them deterministically using the provided SHA-256 pointers.

## VII. CONCLUSION

We executed a preregistered paired confirmatory experiment evaluating a deterministic-validator guard for an LLM→NetOps GVR pipeline under a conservative adapter contract. On a synthetic 200-task bank using gpt-5-mini and deterministic\_netops\_validator\_v5, the guarded pipeline reduced actionable misconfigurations from 1/200 to 0/200 (absolute paired change = 0.005). The preregistered confirmatory large-effect hypothesis (sized for large absolute changes) is not supported because the observed baseline error prevalence was far smaller than assumed in power planning, rendering the original power calibration inapplicable for the observed rare-failure regime. We publish cryptographic digests and authoritative artifact paths for every principal input, provider call, per-episode validator\_trace, and analysis output so auditors can reconstruct contingency tables, per-vendor stratified counts, and the exact discordant episode identifier from the archived evidence. Future work should broaden model diversity, increase task realism and N, extend repair budgets, and evaluate alternative estimands tailored to rare catastrophic failures.

## DISCLOSURE STATEMENT

Disclosure Statement (mandatory). Initial master prompt immutable reference: provenance/master\_prompt.txt SHA-256 = 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acd1582bb39d15ba1  
Preregistration: cnsm-spec2026\_gvr\_v1\_prereg\_001 SHA-256 = 5cb8a30815eac0a3ca659d1a54fbaa71218e5ce57c0b13e2658099ad17d6ee41  
Declared model and configuration: execution/model\_configuration.json (requested\_model = gpt-5-mini) SHA-256 = a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd07350728215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a  
Execution raw results and task manifest: execution/raw\_results.jsonl SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0  
execution/task\_manifest.jsonl SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a  
Deterministic reconciliation and provider-call accounting: analysis/deterministic\_reconciliation.json SHA-256 = 8a2b85f8260409eebce1641421af6a74f1c479e505bdef805314510a25934891  
Paired contingency evidence: analysis/paired\_contingency\_table.csv SHA-256 = 658051baef1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c45d94144b0  
Condition summary (success rates): analysis/condition\_summary.csv SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed00628a68a.  
Missingness summary: analysis/missingness\_summary.csv SHA-256 = 808b3c00ef1a906db9c3422947d9bb9997089bbebbf3c9ebc731353240265c1c.

Analysis executor inputs: analysis/results.json (input results SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0)  
Adapter and tooling: adapter\_family = hosted\_netops\_gvr\_v1; deterministic validator = deterministic\_netops\_validator\_v5; maximum repair calls per task enforced = 1. Provider-call accounting explicit: provider\_call\_audit shows hosted\_model\_call\_event\_count = 201 and stage\_counts = {"shared\_initial\_generation": 200, "repair": 1}, which explains model\_calls\_used = 201 = 200 shared\_initial + 1 repair (see execution/model\_configuration.json SHA-256 above and deterministic\_reconciliation.json SHA-256 above). Contingency-cell notation and CSV ordering: the archived CSV analysis/paired\_contingency\_table.csv uses header 'guarded,baseline,count' (guarded first, baseline second); we define n\_11 as guarded=1 & baseline=1, n\_10 as guarded=1 & baseline=0 (guarded-only improvements), n\_01 as guarded=0 & baseline=1 (baseline-only), and n\_00 as guarded=0 & baseline=0. Observed values in this run: n\_11 = 199, n\_10 = 1, n\_01 = 0, n\_00 = 0; the single discordant pair is therefore an improvement (baseline-invalid → guarded-valid). Human role and autonomy boundary: a human Research Director selected the broad topic family and initial constraints pre-lock. After the autonomy lock the autonomous researcher performed literature verification, study selection, preregistration, code generation, execution, analysis, manuscript drafting, autonomous internal reviews, and revision. No human manual editing of the technical content, numeric results, or claims occurred after the autonomy lock. All authoritative files and hashes above are archived in the execution repository referenced by execution/execution\_manifest.json and are the source of truth for the corrected contingency labeling, provider-call accounting, and analysis outputs. Reviewer requests unresolved in-manuscript (but computable by analysis/condition\_summary.csv (4)) a verbatim per-vendor stratified numeric table (Cisco-like / Juniper-like / RFC-neutral) and (2) the explicit discordant episode identifier string. Both values are derivable from execution/task\_manifest.jsonl (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) together with execution/raw\_results.jsonl (SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0); episode scoring JSONs under execution/scoring/ (full hash manifest in execution/execution\_manifest.json); because no full value appears verbatim in the supplied manuscript evidence bundle text we provide these artifact pointers rather than invent numeric summaries or episode ids in the analysis/condition\_summary.csv where those values are not present verbatim.

## REFERENCES

- [1] <https://www.seerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.228.5064>
- [2] <https://doi.org/10.1145/2342441.2342452>
- [3] <https://doi.org/10.1145/3656296>
- [4] <https://doi.org/10.1109/mcom.001.2400368>