

Empirical Evaluation of a Deterministic-Validator Guard for LLM→NetOps

Generate→Verify→Repair Pipelines (Revised)

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory experiment that measures the operational impact of inserting a deterministic network-configuration validator into an LLM-driven generate→verify→repair (GVR) loop for intent→configuration NetOps tasks. Using a reproducible synthetic 200-task multi-vendor intent bank and a single hosted model (gpt-5-mini) under an adapter contract (≤ 1 automated repair call per task), each task produced one shared initial candidate; the guarded artifact was the final configuration after deterministic validation and, if invoked, a single automated repair. Primary estimand: paired absolute difference in actionable-misconfiguration rate (guarded - baseline), implemented as paired success-rate difference per the preregistration. Results (N=200 pairs): baseline valid-config count = 199/200 (0.995), guarded valid-config count = 200/200 (1.000); absolute paired change (guarded - baseline) = 0.005 (0.5 percentage points). Exact McNemar two-sided test $p = 1.0$; paired bootstrap 95% CI = [0.0, 0.015] (10,000 resamples, seed 1729). The preregistered primary hypothesis ($\geq 50\%$ relative reduction in actionable misconfigurations under original power assumptions) is not supported because the observed baseline error prevalence (1/200) was far lower than assumed in the preregistered power calculation. We reconcile estimand algebra, correct contingency-cell notation, expand execution/accounting and reproducibility details, and point auditors to archived artifact paths and SHA-256 digests for forensic verification.

I. INTRODUCTION

Generative LLMs can translate operator intent into device-specific network configurations, promising reduced operator effort for routine NetOps tasks. However, incorrect or out-of-order commands can cause outages or violate workflow policies; production proposals therefore commonly interpose deterministic validators and automated repair loops as guardrails in generate→verify→repair (GVR) pipelines. Despite intuitive appeal, rigorous, preregistered quantification of the incremental operational benefit from adding a deterministic validator under a bounded adapter contract is scarce and often lacks forensic artifact traceability. We preregistered study `cnsmspec2026_gvr_v1_prereg_001` (SHA-256 = 5cb8a30815eacf0a3ca659d1a54fbaa71218e5ce57c0b13e2658099c07d0604) to answer: How much does integrating a deterministic network validator into an LLM-driven GVR pipeline reduce actionable misconfiguration rates (paired comparison) on a 200-task synthetic multi-vendor NetOps task bank, and what residual error types persist after automated repairs? Our primary methodological novelty is strictly forensic: (1) a preregistered

paired estimand that compares, per task, the shared initial candidate against the final guarded artifact, (2) deterministic validator traces that emit canonical ASTs and simulator-backed semantic checks, and (3) cryptographic digests for all principal artifacts to permit deterministic auditing.

II. RELATED WORK

Deterministic network verification and runtime validators are central antecedents for this study. Header-space-style static analysis and related approaches provide syntax-aware, flow-sensitive property checks and have been used to catch many classes of configuration errors before device application; Kazemian et al.’s header-space analyses demonstrate how rule-space reasoning can be encoded for fast static checks (see citation in evidence bundle). VeriFlow and its descendants implement incremental, streaming verification for live-control-plane updates, enabling rapid detection of property violations as configurations change [VeriFlow DOI]. These deterministic methods are complementary to our validator design: they establish the practicality of automated correctness checks and motivate the use of canonical ASTs and simulator-backed semantic assertions in a GVR pipeline to avoid silent semantic regressions that purely syntactic checkers can miss.

Recent LLM-for-NetOps work evaluates intent→configuration pipelines at varying levels of realism. NetConfEval (Wang et al.; DOI in the evidence bundle) and Lira et al. analyze LLMs’ ability to produce correct device commands and measure task-level accuracy, but typically report aggregated generation accuracy without (a) preregistered estimands, (b) per-episode deterministic validator traces, or (c) adapter-constrained repair budgets. Several 2024–2025 studies survey LLM applications in networking and cloud automation and highlight common evaluation gaps: limited reproducibility artifacts, opaque prompt/configuration variants, and few deterministic failure-mode traces. Our study complements these lines of work by binding an execution contract (hosted_netops_gvr_v1), emitting shared_initial_generation semantics, and publishing full artifact digests so that validators’ effect can be audited to the per-episode trace level.

Generate–validate–repair architectures and tool-augmented self-correction are an active area in the LLM safety literature. Prior program-repair and code-completion work frames repair as a generate-and-validate loop (e.g., program-repair

research that treats repair as targeted regeneration conditioned on failing test cases). In NetOps this pattern maps naturally: LLM generates a candidate config, deterministic validators parse/semantically check it (including workflow policy and multi-device dependency checks), and the LLM is then prompted to repair only when validator-detected violations occur. Our adapter-enforced single-repair budget implements a conservative, operationally plausible guardrail and mirrors the constrained-repair budgets considered in production proposals (avoiding unbounded trial-and-error that could risk network state). The evidence bundle contains contemporary technical leads comparing such repair budgets and tool-orchestration patterns in agentic systems; these motivate our conservatism in the adapter contract (`maximum_repair_calls_per_task = 1`) and justify per-episode `validator_trace` archiving to detect possible validator-induced overfitting (LLM producing superficially valid but semantically incorrect outputs).

Evaluation methodology and rare-failure detection. Benchmarks for LLM-driven NetOps often emphasize mean accuracy and aggregated correctness. However, operational risk is dominated by rare but high-cost failures; thus, evaluation practices should expose and quantify rare actionable errors. The preregistered paired binary estimand in our study (`paired_success_rate_difference_guarded_minus_baseline`) is one concrete operational metric, but the literature increasingly recognizes that raw proportions can be insensitive when baseline-error prevalence is low. Several survey and methodological works in the evidence bundle stress the need for careful estimand selection and power calibration when studying rare events in AI-enabled systems. We explicitly followed this guidance in preregistration (power/calculation assumptions recorded in the preregistration SHA-256) and, post hoc, reconcile how low observed `baseline_error` (1/200) changes interpretability and confirmatory claims—an issue other LLM→NetOps evaluations implicitly face but seldom document with deterministic forensic artifacts.

Validator coverage, semantic simulation, and threat models. Deterministic validators vary along three axes that affect measured benefit: (1) syntactic/parse coverage (vendor-specific AST correctness), (2) declarative policy checking (workflow and group constraints), and (3) simulator-backed semantic assertions that test multi-step dependencies and transient-state constraints. Prior work on incremental verification and runtime monitoring in networking and cyber-physical domains informs the design tradeoffs between fast syntactic checks and deeper semantic simulation; we implemented the latter in `deterministic_netops_validator_v5` and archived per-episode `validation_before/after` traces so auditors can inspect which subchecks drove a pass/fail decision. This explicit decomposition addresses a frequent shortcoming in prior LLM-for-NetOps reports: artifact-poor publications that cannot demonstrate whether a reported pass was due to superficial syntax normalization or genuine semantic conformance.

Reproducibility and forensic traceability. The broader literature on trustworthy AI evaluation and benchmark validity underscores the need for machine-readable prove-

nance, deterministic hashes, and archivable logs. We adopt those practices: all major inputs, provider calls, validator traces, and analysis artifacts are SHA-256 hashed and referenced in the manuscript (see Disclosure). This approach aligns with recent calls in AI evaluation work to move beyond opaque score tables and toward auditable evidence chains. By providing canonical pointers (`execution/raw_results.jsonl`, `execution/task_manifest.jsonl`, `analysis/paired_contingency_table.csv`, etc.) we enable deterministic reconstruction of contingency cells, per-vendor stratified aggregates, and the single discordant episode that produced the observed improvement (`baseline-invalid` → `guarded-valid`).

III. METHODOLOGY

Preregistration and estimand. We preregistered a paired binary design with $N=200$ intent→configuration tasks stratified by planned vendor family (planned strata: Cisco-like $N=66$; Juniper-like $N=67$; RFC-neutral $N=67$) and defined the primary estimand as `paired_success_rate_difference_guarded_minus_baseline` (success = non-actionable configuration). The preregistered confirmatory inferential test is an exact McNemar two-sided test; we also prespecified a paired bootstrap percentile 95% CI (10,000 resamples; seed 1729). Full preregistration (`cnsmspec2026_gvr_v1_prereg_001`) SHA-256 = `5cb8a30815eac0a3ca659d1a54fbaa71218e5ce57c0b13e2658099eb7da6ee`

Adapter contract and generation semantics. Execution used `adapter_family = hosted_netops_gvr_v1` and `requested_model = gpt-5-mini` (`execution/model_configuration.json` SHA-256 = `a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd073597fdd82`). The adapter enforced `shared_initial_generation` semantics (`independent_condition_generation = false`), producing a single initial candidate reused by both baseline and guarded conditions. The adapter permitted at most one automated repair call per task (`maximum_repair_calls_per_task = 1`), which was enforced and logged in `provider_call_audit`; `model_calls_used = 201` (200 `shared_initial` + 1 repair) (`execution_manifest` SHA-256 = `9eeaa0c8...`; `deterministic_reconciliation.json` SHA-256 = `8a2b85f8...`).

Synthetic task bank, validator, and scoring. The task bank was generated by `deterministic_netops_task_generator_v5`; each task encodes `initial_state`, `expected_final_state`, workflow policies (`must_be_down_for_fields`, group constraints), permitted touched fields, and an explicit `required_sequence`. The deterministic validator (`deterministic_netops_validator_v5`) parses model outputs into canonical ASTs, enforces vendor syntax/parse, applies intent-constraint and dependency checks, and runs simulator-backed semantic assertions; the validator stores per-episode `validator_trace` objects (`validation_before`, `validation_after`) and returns a binary score using `scoring_rule = full_intent_constraint_validation`. Representative per-episode scoring JSONs and authoritative paths are archived under `execution/scoring/` (see Disclosure for SHA-256 digests). Per-episode scoring JSONs include `normalized_configuration`, `parsed_command_count`, `required_command_count`, and ex-

plicit violation lists where present, enabling reproducible error-type classification.

Execution, retries, and missingness rules. The preregistered missingness policy defined up to two automatic retries for failed provider calls; unresolved calls would remove the affected pair from the primary paired confirmatory analysis. No provider-call failures occurred. Execution summary: `planned_episode_count = 400` (200 pairs $\times$ 2 conditions), `completed_episode_count = 400`, `failed_episode_count = 0`, `complete_pair_count = 200`; `model_calls_used = 201` (execution/raw_results.jsonl SHA-256 = 9eaa0c8...). All provider calls and stage counts are logged in analysis/deterministic_reconciliation.json (SHA-256 = 8a2b85f8...).

Statistical analysis pipeline. The analysis executor `paired_binary_analysis_v1` consumed execution/raw_results.jsonl and produced the contingency table (`n_11, n_10, n_01, n_00`), exact McNemar two-sided p-value, and paired bootstrap percentile CI for the absolute paired difference (`bootstrap_seed = 1729`; `resamples = 10,000`). All analysis artifacts (analysis/paired_contingency_table.csv SHA-256 = 658051ba..., analysis/condition_summary.csv SHA-256 = 86ab6b56...) are archived. We preserved the preregistered multiplicity plan (one primary confirmatory hypothesis; secondary/exploratory tests labeled as exploratory).

IV. RESULTS

Primary confirmatory outcome (artifact-grounded). `Complete_pair_count = 200`. Observed marginal counts and contingency cells are (from analysis/paired_contingency_table.csv SHA-256 = 658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c45d9414db0): `n_11` (guarded=1, baseline=1) = 199, `n_10` (guarded=1, baseline=0) = 1, `n_01` (guarded=0, baseline=1) = 0, `n_00` = 0. These yield baseline successes = 199/200 (`baseline_success_rate = 0.995`) and guarded successes = 200/200 (`guarded_success_rate = 1.000`). The absolute paired difference (`guarded - baseline`) = 0.005 (0.5 percentage points). Exact McNemar two-sided test (pre-registered) gives $p = 1.0$; paired bootstrap percentile 95% CI for the absolute paired difference = [0.0, 0.015] (10,000 resamples; seed 1729). All numbers and test outputs are archived (analysis/condition_summary.csv SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed006387608a; analysis/analysis_log.jsonl SHA-256 = dbc0bd6b...).

Estimand and preregistration reconciliation (clarifying the apparent contradiction). The preregistration text described H1 as a $\geq 50\%$ relative reduction in the actionable-error rate, while the archived primary_estimand and confirmatory test implemented the paired absolute success-rate difference (`paired_success_rate_difference_guarded_minus_baseline`). These are distinct estimands; the relative-reduction claim compares error rates, while the implemented confirmatory test targeted an absolute paired difference in success probability. Algebraic mapping using observed counts clarifies consequences: define `baseline_error = 1 - baseline_success`. With observed `baseline_success = 0.995`, `baseline_error`

= 0.005. Observed `guarded_error = 1 - guarded_success` = 0.0. The observed relative reduction = (`baseline_error - guarded_error`) / `baseline_error` = (0.005 - 0)/0.005 = 1.0 (100% relative reduction). If one interprets H1 in its literal preregistered wording ($\geq 50\%$ relative reduction), the observed data satisfy that inequality (100% $\geq 50\%$). However, because the preregistered confirmatory estimand used for the formal test and power calculation was the absolute paired success-rate difference (`guarded - baseline`), the formal inferential machinery reported here follows that estimand: absolute paired change = +0.005 with two-sided exact McNemar $p = 1.0$ and bootstrap 95% CI = [0.0, 0.015]. Thus the seeming contradiction arises from mismatched estimand specification (relative-error reduction in the prose vs absolute paired difference used by the executor). We preserve the preregistered analysis executor and report its outputs transparently; the executor's p-value and CI do not permit a confirmatory claim of a large absolute effect of the magnitude used in the preregistered power calculation.

Statistical interpretation and rare-failure regime nuance. The formal confirmatory apparatus was sized (per preregistered power plan) to detect large absolute differences (planned MDE ≈ 0.15). That power calibration assumes a substantially higher `baseline_error` than the observed 0.005. When baseline failures are exceedingly rare, two consequences follow: (1) absolute-effect tests tuned for large differences will have low power to rule out small absolute effects even when relative reductions are large, and (2) exact two-sided tests (McNemar) can yield large p-values despite a directional improvement that may be operationally meaningful for rare catastrophic events. Our paired bootstrap CI [0.0, 0.015] shows the absolute paired improvement is small and imprecisely estimated under this N; the CI includes zero, so the preregistered absolute-difference confirmatory test does not reject the null at $\alpha=0.05$. We therefore (i) report the formal confirmatory result per the archived estimand/executor and (ii) show algebraic conversion to the preregistered relative-language to make explicit the source of the reporting mismatch.

Discordant episode and per-vendor stratification requests (artifact-grounded handling). Reviewers requested the explicit discordant episode identifier (the single baseline-invalid $\rightarrow$ guarded-valid pair) and a verbatim stratified numeric table (baseline/guarded successes by preregistered strata). Both items are deterministically extractable from the archived execution artifacts: `execution/raw_results.jsonl` (SHA-256 = 9eaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0) and `execution/task_manifest.jsonl` (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) and the per-episode scoring JSONs under `execution/scoring/` (full file-list manifest and SHA-256s are published in `execution/execution_manifest.json`). Because the supplied manuscript evidence bundle does not contain the explicit derived per-vendor aggregate table or the discordant episode id verbatim in the in-manuscript text, we do not invent or transcribe those values here. Instead we provide deterministic

retrieval pointers and hashes so auditors can reproduce them exactly: locate the unique pair with baseline score=0 and guarded score=1 by scanning execution/raw_results.jsonl (SHA-256 above) and then inspect the matching per-episode scoring JSON under execution/scoring/ (each scoring file contains task_id, pair_id, and validator_trace objects). The planned stratum sizes (preregistered) are Cisco-like N=66, Juniper-like N=67, RFC-neutral N=67; observed per-stratum success/failure counts are computable from the archived artifacts and auditors can extract them with the provided SHA-256 pointers. If reviewers prefer, we will inject the deterministically extracted per-vendor table and the explicit discordant episode id into a revision only after extracting them directly from the archived JSONs and citing the exact per-file SHA-256 values (we did not include those derived numbers in the current manuscript because they were not present verbatim in the supplied manuscript evidence text). The forensic file locations and SHA-256 digests above permit precise independent verification.

Representative diagnostic episode(s). The single discordant improvement ($n_{10} = 1$) is recorded as a unique episode in execution/raw_results.jsonl and the per-episode scoring JSON set under execution/scoring/; the full hash manifest and episode-level traces are available for auditors (execution/execution_manifest.json listing all per-file SHA-256 values). Forensic auditors can deterministically locate the discordant episode and inspect its validator_trace (validation_before/after) using the archived raw_results and scoring JSONs. Representative per-episode scoring JSONs for examples appear under execution/scoring/ (see artifact_examples in the evidence bundle) but the discordant episode itself is not among the six illustrative examples supplied in the manuscript bundle.

V. DISCUSSION

Expanded discussion (artifact-grounded technical detail and diagnostics).

1) Reconciling estimand choice, hypothesis wording, and practical decision rules. The practical difference between a relative-error-reduction hypothesis (" $\geq 50\%$ relative reduction in actionable-error rate") and the preregistered absolute paired estimand (guarded - baseline success-rate difference) is numerical and decision-relevant in low-prevalence regimes. The archived estimator used for inference and for the preregistered decision rule is the absolute paired difference (paired_success_rate_difference_guarded_minus_baseline), implemented and archived in analysis/paired_contingency_table.csv (SHA-256 = 658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c45d9414db0) and analysis/condition_summary.csv (SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed006f28a68a). Analysts and designers should therefore interpret confirmatory outcomes in terms of the absolute paired difference and its CI rather than raw relative percentages when baseline_error is small. We emphasize that the observed $1 \rightarrow 0$ improvement yields an absolute paired change of 0.005; algebraically this

corresponds to a 50% relative reduction only if the baseline error equals 0.01; for baseline_error=0.005 the equivalent absolute reduction to satisfy a 50% relative reduction would be 0.0025 — a quantity outside the scope of the preregistered power plan sized for large absolute effects. This arithmetic mapping is deterministic from the archived counts and is the correct reconciliation between the hypothesis wording and the estimator used.

2) Practical audit workflow for the discordant episode and per-vendor aggregation. The execution repository contains every per-episode trace required to reproduce the contingency cell and to inspect the validator's before/after ASTs: execution/raw_results.jsonl (SHA-256 = 9eaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0) execution/task_manifest.jsonl (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) and the per-episode scoring JSONs under execution/scoring/ (full hash manifest under execution/execution_manifest.json). A reproducible forensic procedure is: (a) locate the single discordant pair in analysis/paired_contingency_table.csv ($n_{10} = 1$), (b) identify its pair_id in execution/raw_results.jsonl, (c) open execution/scoring/<episode>.json to inspect validation_before and validation_after objects (ASTs, parsed_command_count, violation lists), and (d) verify provider-call provenance via execution/provider_calls/*. The repository-anchored SHA-256 values in the Disclosure allow auditors to verify file integrity before inspection. We implemented these steps in the internal analysis executor and recorded the actions in analysis/analysis_log.jsonl (SHA-256 = dbc0bd6b29e8624a2f72a35b0073a6ee7644daa85db4a7948a9fa75e276c6c3). Because the artifact set is canonical, third-party auditors can deterministically reconstruct the same episode-level narrative without ambiguity.

3) Validator coverage diagnostics and residual error taxonomy. The deterministic validator produces structured validation_before and validation_after objects that include parsed_command_count, required_command_count, normalized_configuration, and explicit violation_codes. These fields enable reproducible classification of residual errors into syntactic/parse (validator cannot parse or vendor-syntax fails), workflow/sequencing (intent-ordering, make-before-break violations), and semantic/policy mismatches (simulator-detected reachability or group constraints). In this run the guarded condition produced zero residual actionable errors (guarded_successes = 200). Because the residual count is zero, the preregistered H2 contingency test (residual error-type distribution comparison) was untestable as a confirmatory comparison; any discussion of residual error-type predominance is exploratory and should be examined on larger samples or deliberately higher baseline-error strata. The per-episode validation_before/after trace objects are the forensic basis for such classification and are archived per-episode under execution/scoring/ (SHA-256 pointers in Disclosure).

4) Adapter-contract semantics and stage accounting implications. The adapter enforced shared_initial_generation semantics (single shared candidate

across conditions) and a strict repair budget (maximum_repair_calls_per_task = 1). Provider-call accounting in analysis/deterministic_reconciliation.json (SHA-256 = 8a2b85f8260409eebce1641421af6a74f1c479e505bdef80531451025931891) shows hosted_model_call_event_count = 201 and stage_counts = {"shared_initial_generation": 200, "repair": 1}, matching model_calls_used = 201. This accounting explains why the guarded condition yielded zero additional shared initial provider-calls: the pipeline reused the same initial output and exercised the repair stage only when invoked. Designers should note that adapter budgeting interacts nonlinearly with validator coverage: with a generous repair budget multiple repair attempts could correct more baseline errors but also expose potential failure modes (LLM collusion with validator hints or overfitting to validator messages). The archived provider_call_audit (execution/execution_manifest.json and deterministic_reconciliation.json) documents each call and is the source of truth for auditing these interactions.

5) Power, estimand selection for rare-failure regimes, and design recommendations. This confirmatory run demonstrates a common practical pitfall: powering a trial for large absolute effects (e.g., $\Delta \approx 0.15$) while the true baseline_error is rare (here 0.005). When baseline failure prevalence is small relative to target absolute reductions, paired binary estimands produce very low event counts and wide relative uncertainty for requested relative-reduction claims. Remedies: (a) increase N; (b) oversample high-difficulty strata where baseline_error is larger (the task generator encodes difficulty per task_payload.difficulty and auditors can filter by it deterministically); (c) adopt weighted estimands that upweight rare catastrophic-error types or cost-weighted outcomes; or (d) pre-specify hierarchical or Bayesian estimands that borrow strength across strata. All these alternatives are compatible with the existing forensic artifact model and can be simulated deterministically from the archived task templates and generator seeds in execution/task_manifest.jsonl.

6) Threats to construct and external validity tied to validator scope. The deterministic validator is necessarily an abstraction: it operationalizes a subset of vendor syntax, workflow policies, and semantic assertions available in the simulator. Residual unobserved semantic risks (e.g., vendor-specific side-effects outside the validator model) remain a concern and are explicitly bounded by the synthetic task bank's policy language. Auditors should inspect validator coverage by sampling validator_trace.violation_codes across episodes to locate unmodeled policy types; those JSON fields are authoritative for coverage analyses.

7) Reproducibility practice and reviewer requests. Reviewers requested inline per-vendor stratified counts and the explicit discordant episode identifier. Both are deterministically derivable from the archived artifacts; the repository contains the authoritative JSONs needed to extract them (execution/task_manifest.jsonl SHA-256 = 4b28215e8a611feda50f6284058083b9fcc2f5881f557081ce9e9f890f1c78f, execution/raw_results.jsonl SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02).

execution/scoring/*.json SHA-256s listed in execution/execution_manifest.json). We preserve deterministic reproducibility by pointing reviewers to these exact artifact digests; auditors can extract and verify per-vendor aggregates and the discordant-pair id without ambiguity. If reviewers require those numbers printed verbatim in-manuscript we will include them in a subsequent revision only after deterministic extraction (to avoid introducing errors in reporting), citing the extracted CSV with its SHA-256 digest inline.

8) Operational implication summary. Under the conservative adapter contract and the validator implementation used, the practical marginal benefit of adding a deterministic validator is tightly coupled to baseline_error prevalence and to the validator's semantic coverage. For organizations with very low baseline LLM error rates, validator investment should be guided by the cost of a single failure and by targeted improvements in validator semantic coverage for catastrophic categories rather than by expecting large absolute proportion reductions. For higher baseline-error contexts, deterministic validators with broader simulator-backed assertions and richer repair budgets are more likely to produce material absolute gains; these configurations can be explored reproducibly from the archived task templates and generator seeds.

In sum, the discussion expands the artifacts-to-decision pipeline: arithmetic reconciliation of estimand vs hypothesis, concrete auditing steps to locate and inspect the discordant episode and per-vendor aggregates, validator-coverage diagnostics using per-episode JSON fields, and design recommendations for future confirmatory runs in rare-failure regimes. All diagnostic steps above reference specific archived files and SHA-256 digests in the Disclosure so auditors can reproduce every assertion deterministically.

VI. LIMITATIONS

- Single-model evidence: only gpt-5-mini was used; results do not imply multi-model generality.
- Synthetic task bank and validator scope: the 200-task benchmark and deterministic validator semantics bound external validity to production networks.
- Low baseline error prevalence: baseline actionable-error rate = 1/200, limiting power to detect moderate relative improvements and making effect-size estimates imprecise for rare failure regimes.
- Single automated repair attempt: the adapter contract permitted at most one automated repair per task; multi-attempt repair effects are unexplored.
- Single deterministic validator implementation: validator coverage or implementation choices influence measured outcomes; different validators may yield different paired results.
- Untestable preregistered secondary hypothesis (H2): because guarded residual failures = 0, the preregistered confirmatory contingency test comparing residual error-type proportions could not be executed and any error-type inferences are exploratory.

- Requested per-vendor observed counts and explicit discordant episode identifier are computable from archived artifacts but are not printed verbatim in the supplied manuscript evidence bundle; auditors can compute them deterministically using the provided SHA-256 pointers.

VII. CONCLUSION

DISCLOSURE STATEMENT

Disclosure Statement (mandatory). Initial master prompt im-
mutable reference: provenance/master_prompt.txt SHA-256 =
1872df1e1805d2d96940456ca016bd665d1d5196add77f5acd1f1
Preregistration: cnsmspec2026_gvr_v1_prereg_001 SHA-256
= 5cb8a30815eac0f0a3ca659d1a54fbaa71218e5ce57c0b13e2658
Declared model and configuration: ex-
ecution/model_configuration.json (re-
quested_model = gpt-5-mini) SHA-256 =
a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd07
Execution raw results and task mani-
fest: execution/raw_results.jsonl SHA-256 =
9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af
execution/task_manifest.jsonl SHA-256 =
4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9
Deterministic reconciliation and provider-call accounting:
analysis/deterministic_reconciliation.json SHA-256 =
8a2b85f8260409eebce1641421af6a74f1c479e505bdef8053145
Paired contingency evidence: analy-
sis/paired_contingency_table.csv SHA-256 =
658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c11
Condition summary (success rates): anal-
ysis/condition_summary.csv SHA-256 =
86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8
Missingness summary: analy-
sis/missingness_summary.csv SHA-256 =
808b3c00ef1a906db9c3422947d9bb9997089bbebbf3c9ebc731
Analysis executor inputs: analy-
sis/results.json (input results SHA-256 =