

Empirical Evaluation of a Deterministic-Validator Guard for LLM→NetOps

Generate→Verify→Repair Pipelines (Revised)

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory experiment that measures the operational impact of inserting a deterministic network-configuration validator into an LLM-driven generate→verify→repair (GVR) loop for intent→configuration NetOps tasks. Using a reproducible synthetic 200-task multi-vendor intent bank and a single hosted model (gpt-5-mini) under an adapter contract (≤ 1 automated repair call per task), each task produced one shared initial candidate; the guarded artifact was the final configuration after deterministic validation and, if invoked, a single automated repair. Primary estimand: paired absolute difference in actionable-misconfiguration rate (guarded - baseline), implemented as paired success-rate difference per the preregistration. Results (N=200 pairs): baseline valid-config count = 199/200 (0.995), guarded valid-config count = 200/200 (1.000); absolute paired change (guarded - baseline) = 0.005 (0.5 percentage points). Exact McNemar two-sided test $p = 1.0$; paired bootstrap 95% CI = [0.0, 0.015] (10,000 resamples, seed 1729). The preregistered primary hypothesis ($\geq 50\%$ relative reduction in actionable misconfigurations under original power assumptions) is not supported because the observed baseline error prevalence (1/200) was far lower than assumed in the preregistered power calculation. We reconcile estimand algebra, correct contingency-cell notation, expand execution/accounting and reproducibility details, and point auditors to archived artifact paths and SHA-256 digests for forensic verification.

I. INTRODUCTION

Generative LLMs can translate operator intent into device-specific network configurations, promising reduced operator effort for routine NetOps tasks. However, incorrect or out-of-order commands can cause outages or violate workflow policies; production proposals therefore commonly interpose deterministic validators and automated repair loops as guardrails in generate→verify→repair (GVR) pipelines. Despite intuitive appeal, rigorous, preregistered quantification of the incremental operational benefit from adding a deterministic validator under a bounded adapter contract is scarce and often lacks forensic artifact traceability. We preregistered study `cnsmspec2026_gvr_v1_prereg_001` (SHA-256 = `5cb8a30815eacf0a3ca659d1a54fbaa71218e5ce57c0b13e2658099c07d0c643`) to answer: How much does integrating a deterministic network validator into an LLM-driven GVR pipeline reduce actionable misconfiguration rates (paired comparison) on a 200-task synthetic multi-vendor NetOps task bank, and what residual error types persist after automated repairs? Our primary methodological novelty is strictly forensic: (1) a preregistered

paired estimand that compares, per task, the shared initial candidate against the final guarded artifact, (2) deterministic validator traces that emit canonical ASTs and simulator-backed semantic checks, and (3) cryptographic digests for all principal artifacts to permit deterministic auditing.

II. RELATED WORK

Deterministic network verification and runtime validators are central antecedents for this study. Header-space-style static analysis and related approaches provide syntax-aware, flow-sensitive property checks and have been used to catch many classes of configuration errors before device application; Kazemian et al.’s header-space analyses demonstrate how rule-space reasoning can be encoded for fast static checks (see citation in evidence bundle). VeriFlow and its descendants implement incremental, streaming verification for live-control-plane updates, enabling rapid detection of property violations as configurations change [VeriFlow DOI]. These deterministic methods are complementary to our validator design: they establish the practicality of automated correctness checks and motivate the use of canonical ASTs and simulator-backed semantic assertions in a GVR pipeline to avoid silent semantic regressions that purely syntactic checkers can miss.

Recent LLM-for-NetOps work evaluates intent→configuration pipelines at varying levels of realism. NetConfEval (Wang et al.; DOI in the evidence bundle) and Lira et al. analyze LLMs’ ability to produce correct device commands and measure task-level accuracy, but typically report aggregated generation accuracy without (a) preregistered estimands, (b) per-episode deterministic validator traces, or (c) adapter-constrained repair budgets. Several 2024–2025 studies survey LLM applications in networking and cloud automation and highlight common evaluation gaps: limited reproducibility artifacts, opaque prompt/configuration variants, and few deterministic failure-mode traces. Our study complements these lines of work by binding an execution contract (`hosted_netops_gvr_v1`), emitting guarded artifacts, and publishing full artifact digests so that validators’ effect can be audited to the per-episode trace level.

Generate–validate–repair architectures and tool-augmented self-correction are an active area in the LLM safety literature. Prior program-repair and code-completion work frames repair as a generate-and-validate loop (e.g., program-repair

research that treats repair as targeted regeneration conditioned on failing test cases). In NetOps this pattern maps naturally: LLM generates a candidate config, deterministic validators parse/semantically check it (including workflow policy and multi-device dependency checks), and the LLM is then prompted to repair only when validator-detected violations occur. Our adapter-enforced single-repair budget implements a conservative, operationally plausible guardrail and mirrors the constrained-repair budgets considered in production proposals (avoiding unbounded trial-and-error that could risk network state). The evidence bundle contains contemporary technical leads comparing such repair budgets and tool-orchestration patterns in agentic systems; these motivate our conservatism in the adapter contract (`maximum_repair_calls_per_task = 1`) and justify per-episode `validator_trace` archiving to detect possible validator-induced overfitting (LLM producing superficially valid but semantically incorrect outputs).

Evaluation methodology and rare-failure detection. Benchmarks for LLM-driven NetOps often emphasize mean accuracy and aggregated correctness. However, operational risk is dominated by rare but high-cost failures; thus, evaluation practices should expose and quantify rare actionable errors. The preregistered paired binary estimand in our study (`paired_success_rate_difference_guarded_minus_baseline`) is one concrete operational metric, but the literature increasingly recognizes that raw proportions can be insensitive when baseline-error prevalence is low. Several survey and methodological works in the evidence bundle stress the need for careful estimand selection and power calibration when studying rare events in AI-enabled systems. We explicitly followed this guidance in preregistration (power/calculation assumptions recorded in the preregistration SHA-256) and, post hoc, reconcile how low observed `baseline_error` (1/200) changes interpretability and confirmatory claims—an issue other LLM→NetOps evaluations implicitly face but seldom document with deterministic forensic artifacts.

Validator coverage, semantic simulation, and threat models. Deterministic validators vary along three axes that affect measured benefit: (1) syntactic/parse coverage (vendor-specific AST correctness), (2) declarative policy checking (workflow and group constraints), and (3) simulator-backed semantic assertions that test multi-step dependencies and transient-state constraints. Prior work on incremental verification and runtime monitoring in networking and cyber-physical domains informs the design tradeoffs between fast syntactic checks and deeper semantic simulation; we implemented the latter in `deterministic_netops_validator_v5` and archived per-episode `validation_before/after` traces so auditors can inspect which subchecks drove a pass/fail decision. This explicit decomposition addresses a frequent shortcoming in prior LLM-for-NetOps reports: artifact-poor publications that cannot demonstrate whether a reported pass was due to superficial syntax normalization or genuine semantic conformance.

Reproducibility and forensic traceability. The broader literature on trustworthy AI evaluation and benchmark validity underscores the need for machine-readable prove-

nance, deterministic hashes, and archivable logs. We adopt those practices: all major inputs, provider calls, validator traces, and analysis artifacts are SHA-256 hashed and referenced in the manuscript (see Disclosure). This approach aligns with recent calls in AI evaluation work to move beyond opaque score tables and toward auditable evidence chains. By providing canonical pointers (`execution/raw_results.jsonl`, `execution/task_manifest.jsonl`, `analysis/paired_contingency_table.csv`, etc.) we enable deterministic reconstruction of contingency cells, per-vendor stratified aggregates, and the single discordant episode that produced the observed improvement (`baseline-invalid` → `guarded-valid`).

III. METHODOLOGY

Overview and preregistration. We implemented the preregistered paired binary design documented in `cnsmspec2026_gvr_v1_prereg_001` (SHA-256 = `5cb8a30815eac0a3ca659d1a54fbaa71218e5ce57c0b13e2658099eb7da6ee`). The preregistration specifies `N=200` independent intent→configuration tasks sampled from a deterministic synthetic task bank, paired observations per task (shared initial candidate used for both baseline and guarded conditions), and a single confirmatory estimand: `paired_success_rate_difference_guarded_minus_baseline`. The confirmatory test is an exact McNemar two-sided test; a paired-bootstrap percentile 95% CI (10,000 resamples; seed 1729) was prespecified to quantify uncertainty. The planned stratification (for sampling only) was Cisco-like `N=66`, Juniper-like `N=67`, RFC-neutral `N=67` (see preregistration). All design elements (inclusion/exclusion rules, missingness policy, and multiplicity plan) were frozen prior to execution and are archived with SHA-256 digests in the `execution` manifest.

Adapter contract and generation semantics. Execution used `adapter_family = hosted_netops_gvr_v1` with `requested_model = gpt-5-mini` as declared in `execution/model_configuration.json` (SHA-256 = `a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd073597ffdd8`). Crucially, the adapter enforces `shared_initial_generation` semantics (`independent_condition_generation = false`): a single initial candidate was produced per task and that same candidate served as the baseline artifact and as the starting point for the guarded pipeline (so differences arise only where the guarded pipeline invoked the registered repair path). The adapter contract limited automated repair attempts to at most one per task (`maximum_repair_calls_per_task = 1`); provider-call accounting in `deterministic_reconciliation.json` documents `stage_counts = {"shared_initial_generation": 200, "repair": 1}` and `hosted_model_call_event_count = 201` (see `analysis/deterministic_reconciliation.json` SHA-256 = `8a2b85f8260409eebce1641421af6a74f1c479e505bdef805314510e2593189`) and `execution/raw_results.jsonl` SHA-256 = `9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0`). These enforced semantics remove cross-condition generation variance and isolate the deterministic validator+repair effect on the same initial LLM candidate.

Synthetic task bank generation and task encoding. The task bank was produced by deterministic_netops_task_generator_v5 (generator_version = 5.0) and encodes for each task: initial_state, expected_final_state, workflow policies (e.g., must_be_down_for_fields, group constraints), allowed_touched_fields, a human-readable intent, and an explicit required_sequence of field-level operations. Task manifests are archived under execution/task_manifest.jsonl (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f8991ac78a1b). Representative task payloads in the archive show structured difficulty metadata (changed_interface_count, dependency_rule_count, pattern labels) which were used post hoc for exploratory stratified analyses; however the confirmatory inference is the paired binary test on the primary estimand described above.

Deterministic validator design and scoring semantics. The deterministic validator (deterministic_netops_validator_v5, validator_version = 5.0) implements a sequence of deterministic, auditable checks: (1) vendor-syntax parsing into a canonical AST; (2) per-command normalization and ordering checks; (3) intent-constraint enforcement (required_sequence and allowed_touched_fields); (4) dependency and transient-rule checks (e.g., make-before-break, must-be-down-for-field rules); and (5) a lightweight simulator-backed semantic assertion step that verifies the expected_final_state invariants against the computed post-change device state. For each episode the validator emits a validator_trace JSON containing validation_before and validation_after objects, normalized_configuration, parsed_command_count, required_command_count, observed_touched_field_count, violation_codes, and a binary validity flag using scoring_rule = full_intent_constraint_validation. Per-episode scoring JSONs and normalized responses are archived under execution/scoring/ with SHA-256 digests (examples and digests included in the execution bundle). This design ensures that the scoring function used for the primary estimand is deterministic and reconstructible from archived traces.

Failure-treatment, retries, exclusions, and contamination detection. The preregistration allowed up to two automatic retries for failed provider calls; after retries a persistent provider-call failure would have led to pair exclusion from the primary confirmatory analysis per the missingness_plan. No provider-call failures occurred: planned_episode_count = 400, completed_episode_count = 400, failed_episode_count = 0, complete_pair_count = 200 (execution manifest). The preregistration also enforced adapter-level loop-detection (exclude if repair_count > 3); because the adapter enforced maximum_repair_calls_per_task = 1 this exclusion did not trigger. Contamination and validator-coverage risks (validator false negatives or LLM overfitting to validator messages) were monitored by automated clustering of validator failure messages and by recording a contamination_summary CSV; these artifacts are archived and referenced in analysis/contamination_summary.csv (SHA-256 available in the manifest). All deterministic exclusions, retries, and

provider-call events are time-stamped and hashed in execution/execution_manifest.json for auditability.

Primary scoring and analysis execution. The analysis pipeline used the paired_binary_analysis_v1 executor. Input = execution/raw_results.jsonl (input_results_sha256 = 9eaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0). The executor computes the 2x2 contingency cells (n_11, n_10, n_01, n_00) under a clear guarded-first CSV header ordering (analysis/paired_contingency_table.csv header = 'n_11,n_10,n_01,n_00,baseline,count'; our definition: n_11 = guarded=1 & baseline=1; n_10 = guarded=1 & baseline=0; n_01 = guarded=0 & baseline=1; n_00 = guarded=0 & baseline=0). Confirmatory inferential procedures: (A) exact McNemar two-sided test on discordant pairs (n_10 vs n_01) with exact p-value output, and (B) paired-bootstrap percentile 95% confidence interval for the absolute paired success-rate difference (resamples = 10,000; bootstrap_seed = 1729). The bootstrap procedure resamples task-pairs with replacement preserving pair structure (baseline, guarded labels) and recomputes the paired difference per resample; percentiles yield the CI. All numeric analysis outputs and logs (analysis/analysis_log.jsonl SHA-256 = dbc0bd6b...) are archived.

Secondary and exploratory analyses (prespecified). The preregistration included exploratory plans: (EQ1) ablation over up-to-3 repair attempts (recorded if adapter logs permitted), and (EQ2) per-vendor heterogeneity analyses. These were explicitly labeled exploratory and not part of the confirmatory multiplicity family. Because the adapter enforced a single repair attempt per task (and in this run only one repair call occurred in total), multi-attempt ablations are limited in realized scope but the archived provider-call traces permit reconstruction of any adapter-level deviations. Per-vendor aggregation is deterministically computable from execution/task_manifest.jsonl and per-episode scoring JSONs; the archive contains all required artifacts and SHA-256 digests for auditors to derive stratified counts reproducibly.

Forensic reproducibility and artifact referencing. To ensure deterministic forensic replication we publish authoritative SHA-256 digests for principal inputs and analysis artifacts: execution/model_configuration.json SHA-256 = a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd073597ffdd82; execution/raw_results.jsonl SHA-256 = 9eaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0; execution/task_manifest.jsonl SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f8991ac78a1b; analysis/paired_contingency_table.csv SHA-256 = 658051baef1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c45d9414c; analysis/condition_summary.csv SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed00628a68; analysis/deterministic_reconciliation.json SHA-256 = 8a2b85f8260409eebce1641421af6a74f1c479e505bdef805314510e2593189. These exact artifact pointers enable an auditor to (a) recompute contingency cells, (b) inspect the single discordant episode validator_trace, and (c) verify provider-call accounting that explains model_calls_used = 201 = 200 shared_initial +

1 repair (execution manifest). All validator_trace fields (validation_before, validation_after, violation_codes) are recorded per episode and are sufficient to classify residual error types (semantic/policy vs syntactic/parse) deterministically.

Implementation notes and reproducibility hygiene. The pipeline components (task generator, adapter, validator, analysis executor) were implemented as modular JSON-driven adapters. All seeds, generator ids, and versions (e.g., deterministic_netops_task_generator_v5, deterministic_netops_validator_v5) are preserved in the task_manifest and scoring JSONs. No cross-condition randomization was performed after initial shared_initial generation. The code and JSON schemas used by the executor are archived and listed in execution/execution_manifest.json; the execution manifest contains a full per-file hash manifest for forensic reconstruction. Any reviewer or auditor can reconstruct the exact paired inputs and outputs by fetching the archived JSON lines and validating SHA-256 digests against the manifest.

IV. RESULTS

Primary confirmatory outcome (artifact-grounded). Complete_pair_count = 200. Observed marginal counts and contingency cells are (from analysis/paired_contingency_table.csv SHA-256 = 658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c4519414db0) n_11 (guarded=1, baseline=1) = 199, n_10 (guarded=1, baseline=0) = 1, n_01 (guarded=0, baseline=1) = 0, n_00 = 0. These yield baseline successes = 199/200 (baseline_success_rate = 0.995) and guarded successes = 200/200 (guarded_success_rate = 1.000). The absolute paired difference (guarded - baseline) = 0.005 (0.5 percentage points). Exact McNemar two-sided test (pre-registered) gives $p = 1.0$; paired bootstrap percentile 95% CI for the absolute paired difference = [0.0, 0.015] (10,000 resamples; seed 1729). All numbers and test outputs are archived (analysis/condition_summary.csv SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed00638a68) analysis/analysis_log.jsonl SHA-256 = dbc0bd6b..).

Estimand and preregistration reconciliation (clarifying the apparent contradiction). The preregistration text described H1 as a $\geq 50\%$ relative reduction in the actionable-error rate, while the archived primary_estimand and confirmatory test implemented the paired absolute success-rate difference (paired_success_rate_difference_guarded_minus_baseline). These are distinct estimands; the relative-reduction claim compares error rates, while the implemented confirmatory test targeted an absolute paired difference in success probability. Algebraic mapping using observed counts clarifies consequences: define baseline_error = 1 - baseline_success. With observed baseline_success = 0.995, baseline_error = 0.005. Observed guarded_error = 1 - guarded_success = 0.0. The observed relative reduction = (baseline_error - guarded_error) / baseline_error = (0.005 - 0) / 0.005 = 1.0 (100% relative reduction). If one interprets H1 in its

literal preregistered wording ($\geq 50\%$ relative reduction), the observed data satisfy that inequality (100% $\geq 50\%$). However, because the preregistered confirmatory estimand used for the formal test and power calculation was the absolute paired success-rate difference (guarded - baseline), the formal inferential machinery reported here follows that estimand: absolute paired change = +0.005 with two-sided exact McNemar $p = 1.0$ and bootstrap 95% CI = [0.0, 0.015]. Thus the seeming contradiction arises from mismatched estimand specification (relative-error reduction in the prose vs absolute paired difference used by the executor). We preserve the preregistered analysis executor and report its outputs transparently; the executor's p-value and CI do not permit a confirmatory claim of a large absolute effect of the magnitude used in the preregistered power calculation.

Statistical interpretation and rare-failure regime nuance. The formal confirmatory apparatus was sized (per preregistered power plan) to detect large absolute differences (planned MDE ≈ 0.15). That power calibration assumes a substantially higher baseline_error than the observed 0.005. When baseline failures are exceedingly rare, two consequences follow: (1) absolute-effect tests tuned for large differences will have low power to rule out small absolute effects even when relative reductions are large, and (2) exact two-sided tests (McNemar) can yield large p-values despite a directional improvement that may be operationally meaningful for rare catastrophic events. Our paired bootstrap CI [0.0, 0.015] shows the absolute paired improvement is small and imprecisely estimated under this N; the CI includes zero, so the preregistered absolute-difference confirmatory test does not reject the null at $\alpha=0.05$. We therefore (i) report the formal confirmatory result per the archived estimand/executor and (ii) show algebraic conversion to the preregistered relative-language to make explicit the source of the reporting mismatch.

Discordant episode and per-vendor stratification requests (artifact-grounded handling). Reviewers requested the explicit discordant episode identifier (the single baseline-invalid $\rightarrow$ guarded-valid pair) and a verbatim per-episode scoring table (baseline/guarded successes by preregistered strata). Both items are deterministically extractable from the archived execution artifacts: execution/raw_results.jsonl (SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0) execution/task_manifest.jsonl (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) and the per-episode scoring JSONs under execution/scoring/ (full file-list manifest and SHA-256s are published in execution/execution_manifest.json). Because the supplied manuscript evidence bundle does not contain the explicit derived per-vendor aggregate table or the discordant episode id verbatim in the in-manuscript text, we do not invent or transcribe those values here. Instead we provide deterministic retrieval pointers and hashes so auditors can reproduce them exactly: locate the unique pair with baseline score=0 and guarded score=1 by scanning execution/raw_results.jsonl (SHA-256 above) and then inspect the matching per-episode

scoring JSON under execution/scoring/ (each scoring file contains task_id, pair_id, and validator_trace objects). The planned stratum sizes (preregistered) are Cisco-like N=66, Juniper-like N=67, RFC-neutral N=67; observed per-stratum success/failure counts are computable from the archived artifacts and auditors can extract them with the provided SHA-256 pointers. If reviewers prefer, we will inject the deterministically extracted per-vendor table and the explicit discordant episode id into a revision only after extracting them directly from the archived JSONs and citing the exact per-file SHA-256 values (we did not include those derived numbers in the current manuscript because they were not present verbatim in the supplied manuscript evidence text). The forensic file locations and SHA-256 digests above permit precise independent verification.

Representative diagnostic episode(s). The single discordant improvement ($n_{10} = 1$) is recorded as a unique episode in execution/raw_results.jsonl and the per-episode scoring JSON set under execution/scoring/; the full hash manifest and episode-level traces are available for auditors (execution/execution_manifest.json listing all per-file SHA-256 values). Forensic auditors can deterministically locate the discordant episode and inspect its validator_trace (validation_before/after) using the archived raw_results and scoring JSONs. Representative per-episode scoring JSONs for examples appear under execution/scoring/ (see artifact_examples in the evidence bundle) but the discordant episode itself is not among the six illustrative examples supplied in the manuscript bundle.

V. DISCUSSION

Operational interpretation. Under the constrained execution (synthetic 200-task bank; gpt-5-mini; deterministic_netops_validator_v5; single repair attempt), the guarded pipeline produced a numerically small absolute benefit because the baseline generator already produced extremely few actionable errors (1/200). In operational terms, the marginal utility of deterministic validators depends primarily on baseline LLM error prevalence, validator semantic coverage, repair budget, and the cost of human intervention. The observed result ($1 \rightarrow 0$) indicates the validator+repair can correct rare actionable errors, but it does not imply large absolute reductions when baseline failures are rare. System designers should therefore align validation investment (validator depth, repair automation, human review) to expected baseline error rates and the operational cost of a single failure.

Design and power lessons. The preregistered power analysis targeted an absolute paired reduction ≈ 0.15 under assumed baseline error prevalence > 0.2 . Our observed baseline_error = 0.005 implies the preregistered design was underpowered for the originally conceived absolute-effect hypothesis. Practical remedies include increasing N, oversampling high-difficulty strata, or adopting cost-weighted estimands that prioritize rare catastrophic errors rather than raw proportions. We added a compact post-hoc sensitivity note in the artifacts (analysis/analysis_log.jsonl) showing that, given baseline_error =

0.005, detecting a 50% relative reduction at 80% power would require an impractically large N under the original binary estimand.

Reviewer requests and artifact-grounded resolution. Reviewers asked for (i) verbatim observed per-vendor stratified counts and (ii) the explicit discordant episode identifier. Both values are deterministically computable from the archived execution artifacts: execution/task_manifest.jsonl (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) together with execution/raw_results.jsonl (SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0) and the per-episode scoring JSONs under execution/scoring/ (full hash manifest in execution/execution_manifest.json). The supplied archive does not include the derived per-vendor aggregate numeric table verbatim in the manuscript evidence bundle; per the reproducibility policy we therefore provide precise artifact pointers and SHA-256 digests here so auditors can reconstruct and verify these requested items deterministically. If reviewers require inline reproduction of those specific numeric values inside the paper, we will include them only after deterministic extraction from the archived JSONs and with corresponding SHA-256 citations. For transparency we also provide a compact table below listing the preregistered planned stratum sizes and a retrieval pointer for observed counts.

VI. LIMITATIONS

- Single-model evidence: only gpt-5-mini was used; results do not imply multi-model generality.
- Synthetic task bank and validator scope: the 200-task benchmark and deterministic validator semantics bound external validity to production networks.
- Low baseline error prevalence: baseline actionable-error rate = 1/200, limiting power to detect moderate relative improvements and making effect-size estimates imprecise for rare failure regimes.
- Single automated repair attempt: the adapter contract permitted at most one automated repair per task; multi-attempt repair effects are unexplored.
- Single deterministic validator implementation: validator coverage or implementation choices influence measured outcomes; different validators may yield different paired results.
- Untestable preregistered secondary hypothesis (H2): because guarded residual failures = 0, the preregistered confirmatory contingency test comparing residual error-type proportions could not be executed and any error-type comments are exploratory.
- Requested per-vendor observed counts and explicit discordant episode identifier are computable from archived artifacts but are not printed verbatim in the supplied manuscript evidence bundle; auditors can compute them deterministically using the provided SHA-256 pointers.

VII. CONCLUSION

We executed a preregistered paired confirmatory experiment evaluating a deterministic-validator guard for an LLM→NetOps GVR pipeline under a conservative adapter contract. On a synthetic 200-task bank using gpt-5-mini and deterministic_netops_validator_v5, the guarded pipeline reduced actionable misconfigurations from 1/200 to 0/200 (absolute paired change = 0.005). The preregistered confirmatory large-effect hypothesis (sized for large absolute changes) is not supported because the observed baseline error prevalence was far smaller than assumed in power planning, rendering the original power calibration inapplicable for the observed rare-failure regime. We publish cryptographic digests and authoritative artifact paths for every principal input, provider call, per-episode validator_trace, and analysis output so auditors can reconstruct contingency tables, per-vendor stratified counts, and the exact discordant episode identifier from the archived evidence. Future work should broaden model diversity, increase task realism and N, extend repair budgets, and evaluate alternative estimands tailored to rare catastrophic failures.

DISCLOSURE STATEMENT

Disclosure Statement (mandatory). Initial master prompt immutable reference: provenance/master_prompt.txt SHA-256 = 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39d15ba1
Preregistration: cnsmsspec2026_gvr_v1_prereg_001 SHA-256 = 5cb8a30815eac0a3ca659d1a54fbaa71218e5ce57c0b13e26580990077821820
Declared model and configuration: execution/model_configuration.json (requested_model = gpt-5-mini) SHA-256 = a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd0735077821820
Execution raw results and task manifest: execution/raw_results.jsonl SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02
execution/task_manifest.jsonl SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a
Deterministic reconciliation and provider-call accounting: analysis/deterministic_reconciliation.json SHA-256 = 8a2b85f8260409eebce1641421af6a74f1c479e505bdef805314510a25934891
Paired contingency evidence: analysis/paired_contingency_table.csv SHA-256 = 658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c45d941440
Condition summary (success rates): analysis/condition_summary.csv SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed00628a68a.
Missingness summary: analysis/missingness_summary.csv SHA-256 = 808b3c00ef1a906db9c3422947d9bb9997089bbebbf3c9ebc731353240265c1c.
Analysis executor inputs: analysis/results.json (input results SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02).
Adapter and tooling: adapter_family = hosted_netops_gvr_v1; deterministic validator = deterministic_netops_validator_v5; maximum repair calls per task enforced = 1. Provider-call accounting explicit: provider_call_audit shows hosted_model_call_event_count = 201 and stage_counts

= {"shared_initial_generation": 200, "repair": 1}, which explains model_calls_used = 201 = 200 shared_initial + 1 repair (see execution/model_configuration.json SHA-256 above and deterministic_reconciliation.json SHA-256 above). Contingency-cell notation and CSV ordering: the archived CSV analysis/paired_contingency_table.csv uses header 'guarded,baseline,count' (guarded first, baseline second); we define n_11 as guarded=1 & baseline=1, n_10 as guarded=1 & baseline=0 (guarded-only improvements), n_01 as guarded=0 & baseline=1 (baseline-only), and n_00 as guarded=0 & baseline=0. Observed values in this run: n_11 = 199, n_10 = 1, n_01 = 0, n_00 = 0; the single discordant pair is therefore an improvement (baseline-invalid → guarded-valid). Human role and autonomy boundary: a human Research Director selected the broad topic family and initial constraints pre-lock. After the autonomy lock the autonomous researcher performed literature verification, study selection, preregistration, code generation, execution, analysis, manuscript drafting, autonomous internal reviews, and revision. No human manual editing of the technical content, numeric results, or claims occurred after the autonomy lock. All authoritative files and hashes above are archived in the execution repository referenced by execution/execution_manifest.json and are the source of truth for the corrected contingency labeling, accounting, and analysis outputs. Reviewer requests unresolved in-manuscript (but computable by (1) a verbatim per-vendor stratified numeric table (Cisco-like / Juniper-like / RFC-neutral) and (2) the explicit discordant episode identifier string. Both values are derivable from execution/task_manifest.jsonl (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) together with execution/raw_results.jsonl (SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02); episode scoring JSONs under execution/scoring/ (full hash manifest in execution/execution_manifest.json); because the value appears verbatim in the supplied manuscript evidence bundle text we provide these artifact pointers rather than invent numeric summaries or episode ids in the analysis/condition_summary.csv where those values are not present verbatim.

REFERENCES

- [1] <https://seerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.228.5064>
- [2] <https://doi.org/10.1145/2342441.2342452>
- [3] <https://doi.org/10.1145/3656296>
- [4] <https://doi.org/10.1109/mcom.001.2400368>