

Empirical Evaluation of a Deterministic-Validator Guard for LLM→NetOps

Generate→Verify→Repair Pipelines (Revised)

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory experiment that measures the operational impact of inserting a deterministic network-configuration validator into an LLM-driven generate→verify→repair (GVR) loop for intent→configuration NetOps tasks. Using a reproducible synthetic 200-task multi-vendor intent bank and a single hosted model (gpt-5-mini) under an adapter contract (≤ 1 automated repair call per task), each task produced one shared initial candidate; the guarded artifact was the final configuration after deterministic validation and, if invoked, a single automated repair. Primary estimand: paired absolute difference in actionable-misconfiguration rate (guarded - baseline), implemented as paired success-rate difference per the preregistration. Results (N=200 pairs): baseline valid-config count = 199/200 (0.995), guarded valid-config count = 200/200 (1.000); absolute paired change (guarded - baseline) = 0.005 (0.5 percentage points). Exact McNemar two-sided test $p = 1.0$; paired bootstrap 95% CI = [0.0, 0.015] (10,000 resamples, seed 1729). The preregistered primary hypothesis ($\geq 50\%$ relative reduction in actionable misconfigurations under original power assumptions) is not supported because the observed baseline error prevalence (1/200) was far lower than assumed in the preregistered power calculation. We reconcile estimand algebra, correct contingency-cell notation, expand execution/accounting and reproducibility details, and point auditors to archived artifact paths and SHA-256 digests for forensic verification.

I. INTRODUCTION

Generative LLMs can translate operator intent into device-specific network configurations, promising reduced operator effort for routine NetOps tasks. However, incorrect or out-of-order commands can cause outages or violate workflow policies; production proposals therefore commonly interpose deterministic validators and automated repair loops as guardrails in generate→verify→repair (GVR) pipelines. Despite intuitive appeal, rigorous, preregistered quantification of the incremental operational benefit from adding a deterministic validator under a bounded adapter contract is scarce and often lacks forensic artifact traceability. We preregistered study `cnsmspec2026_gvr_v1_prereg_001` (SHA-256 = 5cb8a30815eacf0a3ca659d1a54fbaa71218e5ce57c0b13e2658099c07d0604) to answer: How much does integrating a deterministic network validator into an LLM-driven GVR pipeline reduce actionable misconfiguration rates (paired comparison) on a 200-task synthetic multi-vendor NetOps task bank, and what residual error types persist after automated repairs? Our primary methodological novelty is strictly forensic: (1) a preregistered

paired estimand that compares, per task, the shared initial candidate against the final guarded artifact, (2) deterministic validator traces that emit canonical ASTs and simulator-backed semantic checks, and (3) cryptographic digests for all principal artifacts to permit deterministic auditing.

II. RELATED WORK

Deterministic network verification and runtime validators are central antecedents for this study. Header-space-style static analysis and related approaches provide syntax-aware, flow-sensitive property checks and have been used to catch many classes of configuration errors before device application; Kazemian et al.’s header-space analyses demonstrate how rule-space reasoning can be encoded for fast static checks (see citation in evidence bundle). VeriFlow and its descendants implement incremental, streaming verification for live-control-plane updates, enabling rapid detection of property violations as configurations change [VeriFlow DOI]. These deterministic methods are complementary to our validator design: they establish the practicality of automated correctness checks and motivate the use of canonical ASTs and simulator-backed semantic assertions in a GVR pipeline to avoid silent semantic regressions that purely syntactic checkers can miss.

Recent LLM-for-NetOps work evaluates intent→configuration pipelines at varying levels of realism. NetConfEval (Wang et al.; DOI in the evidence bundle) and Lira et al. analyze LLMs’ ability to produce correct device commands and measure task-level accuracy, but typically report aggregated generation accuracy without (a) preregistered estimands, (b) per-episode deterministic validator traces, or (c) adapter-constrained repair budgets. Several 2024–2025 studies survey LLM applications in networking and cloud automation and highlight common evaluation gaps: limited reproducibility artifacts, opaque prompt/configuration variants, and few deterministic failure-mode traces. Our study complements these lines of work by binding an execution contract (hosted_netops_gvr_v1), emitting shared_initial_generation semantics, and publishing full artifact digests so that validators’ effect can be audited to the per-episode trace level.

Generate–validate–repair architectures and tool-augmented self-correction are an active area in the LLM safety literature. Prior program-repair and code-completion work frames repair as a generate-and-validate loop (e.g., program-repair

research that treats repair as targeted regeneration conditioned on failing test cases). In NetOps this pattern maps naturally: LLM generates a candidate config, deterministic validators parse/semantically check it (including workflow policy and multi-device dependency checks), and the LLM is then prompted to repair only when validator-detected violations occur. Our adapter-enforced single-repair budget implements a conservative, operationally plausible guardrail and mirrors the constrained-repair budgets considered in production proposals (avoiding unbounded trial-and-error that could risk network state). The evidence bundle contains contemporary technical leads comparing such repair budgets and tool-orchestration patterns in agentic systems; these motivate our conservatism in the adapter contract (`maximum_repair_calls_per_task = 1`) and justify per-episode `validator_trace` archiving to detect possible validator-induced overfitting (LLM producing superficially valid but semantically incorrect outputs).

Evaluation methodology and rare-failure detection. Benchmarks for LLM-driven NetOps often emphasize mean accuracy and aggregated correctness. However, operational risk is dominated by rare but high-cost failures; thus, evaluation practices should expose and quantify rare actionable errors. The preregistered paired binary estimand in our study (`paired_success_rate_difference_guarded_minus_baseline`) is one concrete operational metric, but the literature increasingly recognizes that raw proportions can be insensitive when baseline-error prevalence is low. Several survey and methodological works in the evidence bundle stress the need for careful estimand selection and power calibration when studying rare events in AI-enabled systems. We explicitly followed this guidance in preregistration (power/calculation assumptions recorded in the preregistration SHA-256) and, post hoc, reconcile how low observed `baseline_error` (1/200) changes interpretability and confirmatory claims—an issue other LLM→NetOps evaluations implicitly face but seldom document with deterministic forensic artifacts.

Validator coverage, semantic simulation, and threat models. Deterministic validators vary along three axes that affect measured benefit: (1) syntactic/parse coverage (vendor-specific AST correctness), (2) declarative policy checking (workflow and group constraints), and (3) simulator-backed semantic assertions that test multi-step dependencies and transient-state constraints. Prior work on incremental verification and runtime monitoring in networking and cyber-physical domains informs the design tradeoffs between fast syntactic checks and deeper semantic simulation; we implemented the latter in `deterministic_netops_validator_v5` and archived per-episode `validation_before/after` traces so auditors can inspect which subchecks drove a pass/fail decision. This explicit decomposition addresses a frequent shortcoming in prior LLM-for-NetOps reports: artifact-poor publications that cannot demonstrate whether a reported pass was due to superficial syntax normalization or genuine semantic conformance.

Reproducibility and forensic traceability. The broader literature on trustworthy AI evaluation and benchmark validity underscores the need for machine-readable prove-

nance, deterministic hashes, and archivable logs. We adopt those practices: all major inputs, provider calls, validator traces, and analysis artifacts are SHA-256 hashed and referenced in the manuscript (see Disclosure). This approach aligns with recent calls in AI evaluation work to move beyond opaque score tables and toward auditable evidence chains. By providing canonical pointers (`execution/raw_results.jsonl`, `execution/task_manifest.jsonl`, `analysis/paired_contingency_table.csv`, etc.) we enable deterministic reconstruction of contingency cells, per-vendor stratified aggregates, and the single discordant episode that produced the observed improvement (`baseline-invalid` → `guarded-valid`).

III. METHODOLOGY

Overview and preregistration. We implemented the preregistered paired binary design documented in `cnsmspec2026_gvr_v1_prereg_001` (SHA-256 = `5cb8a30815eac0a3ca659d1a54fbaa71218e5ce57c0b13e2658099eb7da6ee`). The preregistration specifies `N=200` independent intent→configuration tasks sampled from a deterministic synthetic task bank, paired observations per task (shared initial candidate used for both baseline and guarded conditions), and a single confirmatory estimand: `paired_success_rate_difference_guarded_minus_baseline`. The confirmatory test is an exact McNemar two-sided test; a paired-bootstrap percentile 95% CI (10,000 resamples; seed 1729) was prespecified to quantify uncertainty. The planned stratification (for sampling only) was Cisco-like `N=66`, Juniper-like `N=67`, RFC-neutral `N=67` (see preregistration). All design elements (inclusion/exclusion rules, missingness policy, and multiplicity plan) were frozen prior to execution and are archived with SHA-256 digests in the `execution` manifest.

Adapter contract and generation semantics. Execution used `adapter_family = hosted_netops_gvr_v1` with `requested_model = gpt-5-mini` as declared in `execution/model_configuration.json` (SHA-256 = `a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd073597ffdd8`). Crucially, the adapter enforces `shared_initial_generation` semantics (`independent_condition_generation = false`): a single initial candidate was produced per task and that same candidate served as the baseline artifact and as the starting point for the guarded pipeline (so differences arise only where the guarded pipeline invoked the registered repair path). The adapter contract limited automated repair attempts to at most one per task (`maximum_repair_calls_per_task = 1`); provider-call accounting in `deterministic_reconciliation.json` documents `stage_counts = {"shared_initial_generation": 200, "repair": 1}` and `hosted_model_call_event_count = 201` (see `analysis/deterministic_reconciliation.json` SHA-256 = `8a2b85f8260409eebce1641421af6a74f1c479e505bdef805314510e2593189`) and `execution/raw_results.jsonl` SHA-256 = `9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0`). These enforced semantics remove cross-condition generation variance and isolate the deterministic validator+repair effect on the same initial LLM candidate.

Synthetic task bank generation and task encoding. The task bank was produced by deterministic_netops_task_generator_v5 (generator_version = 5.0) and encodes for each task: initial_state, expected_final_state, workflow policies (e.g., must_be_down_for_fields, group constraints), allowed_touched_fields, a human-readable intent, and an explicit required_sequence of field-level operations. Task manifests are archived under execution/task_manifest.jsonl (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f8991ac78a1b). Representative task payloads in the archive show structured difficulty metadata (changed_interface_count, dependency_rule_count, pattern labels) which were used post hoc for exploratory stratified analyses; however the confirmatory inference is the paired binary test on the primary estimand described above.

Deterministic validator design and scoring semantics. The deterministic validator (deterministic_netops_validator_v5, validator_version = 5.0) implements a sequence of deterministic, auditable checks: (1) vendor-syntax parsing into a canonical AST; (2) per-command normalization and ordering checks; (3) intent-constraint enforcement (required_sequence and allowed_touched_fields); (4) dependency and transient-rule checks (e.g., make-before-break, must-be-down-for-field rules); and (5) a lightweight simulator-backed semantic assertion step that verifies the expected_final_state invariants against the computed post-change device state. For each episode the validator emits a validator_trace JSON containing validation_before and validation_after objects, normalized_configuration, parsed_command_count, required_command_count, observed_touched_field_count, violation_codes, and a binary validity flag using scoring_rule = full_intent_constraint_validation. Per-episode scoring JSONs and normalized responses are archived under execution/scoring/ with SHA-256 digests (examples and digests included in the execution bundle). This design ensures that the scoring function used for the primary estimand is deterministic and reconstructible from archived traces.

Failure-treatment, retries, exclusions, and contamination detection. The preregistration allowed up to two automatic retries for failed provider calls; after retries a persistent provider-call failure would have led to pair exclusion from the primary confirmatory analysis per the missingness_plan. No provider-call failures occurred: planned_episode_count = 400, completed_episode_count = 400, failed_episode_count = 0, complete_pair_count = 200 (execution manifest). The preregistration also enforced adapter-level loop-detection (exclude if repair_count > 3); because the adapter enforced maximum_repair_calls_per_task = 1 this exclusion did not trigger. Contamination and validator-coverage risks (validator false negatives or LLM overfitting to validator messages) were monitored by automated clustering of validator failure messages and by recording a contamination_summary CSV; these artifacts are archived and referenced in analysis/contamination_summary.csv (SHA-256 available in the manifest). All deterministic exclusions, retries, and

provider-call events are time-stamped and hashed in execution/execution_manifest.json for auditability.

Primary scoring and analysis execution. The analysis pipeline used the paired_binary_analysis_v1 executor. Input = execution/raw_results.jsonl (input_results_sha256 = 9eaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0). The executor computes the 2x2 contingency cells (n_11, n_10, n_01, n_00) under a clear guarded-first CSV header ordering (analysis/paired_contingency_table.csv header = 'n_11,n_10,n_01,n_00,baseline,count'; our definition: n_11 = guarded=1 & baseline=1; n_10 = guarded=1 & baseline=0; n_01 = guarded=0 & baseline=1; n_00 = guarded=0 & baseline=0). Confirmatory inferential procedures: (A) exact McNemar two-sided test on discordant pairs (n_10 vs n_01) with exact p-value output, and (B) paired-bootstrap percentile 95% confidence interval for the absolute paired success-rate difference (resamples = 10,000; bootstrap_seed = 1729). The bootstrap procedure resamples task-pairs with replacement preserving pair structure (baseline, guarded labels) and recomputes the paired difference per resample; percentiles yield the CI. All numeric analysis outputs and logs (analysis/analysis_log.jsonl SHA-256 = dbc0bd6b...) are archived.

Secondary and exploratory analyses (prespecified). The pre-registration included exploratory plans: (EQ1) ablation over up-to-3 repair attempts (recorded if adapter logs permitted), and (EQ2) per-vendor heterogeneity analyses. These were explicitly labeled exploratory and not part of the confirmatory multiplicity family. Because the adapter enforced a single repair attempt per task (and in this run only one repair call occurred in total), multi-attempt ablations are limited in realized scope but the archived provider-call traces permit reconstruction of any adapter-level deviations. Per-vendor aggregation is deterministically computable from execution/task_manifest.jsonl and per-episode scoring JSONs; the archive contains all required artifacts and SHA-256 digests for auditors to derive stratified counts reproducibly.

Forensic reproducibility and artifact referencing. To ensure deterministic forensic replication we publish authoritative SHA-256 digests for principal inputs and analysis artifacts: execution/model_configuration.json SHA-256 = a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd073597ffdd82; execution/raw_results.jsonl SHA-256 = 9eaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0; execution/task_manifest.jsonl SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f8991ac78a1b; analysis/paired_contingency_table.csv SHA-256 = 658051baef1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c45d9414c; analysis/condition_summary.csv SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed00628a68; analysis/deterministic_reconciliation.json SHA-256 = 8a2b85f8260409eebce1641421af6a74f1c479e505bdef805314510e2593189. These exact artifact pointers enable an auditor to (a) recompute contingency cells, (b) inspect the single discordant episode validator_trace, and (c) verify provider-call accounting that explains model_calls_used = 201 = 200 shared_initial +

1 repair (execution manifest). All `validator_trace` fields (`validation_before`, `validation_after`, `violation_codes`) are recorded per episode and are sufficient to classify residual error types (semantic/policy vs syntactic/parsing) deterministically.

Implementation notes and reproducibility hygiene. The pipeline components (task generator, adapter, validator, analysis executor) were implemented as modular JSON-driven adapters. All seeds, generator ids, and versions (e.g., `deterministic_netops_task_generator_v5`, `deterministic_netops_validator_v5`) are preserved in the `task_manifest` and scoring JSONs. No cross-condition randomization was performed after initial shared_initial generation. The code and JSON schemas used by the executor are archived and listed in `execution/execution_manifest.jsonl`; the execution manifest contains a full per-file hash manifest for forensic reconstruction. Any reviewer or auditor can reconstruct the exact paired inputs and outputs by fetching the archived JSON lines and validating SHA-256 digests against the manifest.

IV. RESULTS

Primary confirmatory outcome (artifact-grounded). Complete_pair_count = 200. Baseline successes = 199/200 (`baseline_success_rate` = 0.995); guarded successes = 200/200 (`guarded_success_rate` = 1.0). Contingency counts (`analysis/paired_contingency_table.csv` SHA-256 = 658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c4549414db0) `n_11` = 199, `n_10` = 1, `n_01` = 0, `n_00` = 0. Absolute paired change (guarded - baseline) = 0.005 (0.5 percentage points). Exact McNemar two-sided test $p = 1.0$. Paired bootstrap 95% percentile CI = [0.0, 0.015] (10,000 resamples, seed 1729; `analysis/analysis_log.jsonl` SHA-256 = dbc0bd6b...).

Why the two-sided exact McNemar $p = 1.0$ and how to interpret it. The exact McNemar test conditions on the sum of discordant pairs (`n_10` + `n_01`). Here discordant mass = 1 (`n_10` = 1, `n_01` = 0). With a single discordant observation, the two-sided exact test yields a noninformative p -value (reported $p = 1.0$) because the test's sampling distribution under the null assigns symmetric probability mass to the two discordant orderings; with so little discordant information the two-sided test cannot reject the null. The bootstrap percentile CI (seed 1729, 10,000 resamples) provides a distributional summary consistent with the contingency counts and indicates the observed absolute improvement is small and statistically imprecise: CI = [0.0, 0.015] contains zero and bounds the plausible absolute paired improvement to ≤ 1.5 percentage points under the bootstrap assumptions used.

Estimand algebra and preregistered hypothesis reconciliation. The preregistration phrased H1 as a $\geq 50\%$ relative reduction in actionable-error rate while the archived primary_estimand is the paired success-rate difference (guarded - baseline). Let `baseline_error` = 1 - `baseline_success` = 0.005. A 50% relative reduction in actionable-error rate corresponds to absolute reduction = $0.5 \times \text{baseline_error} = 0.0025$. Our observed absolute reduction = `baseline_error` - `guarded_error`

= 0.005 - 0.0 = 0.005, which numerically exceeds 0.0025; however the preregistered power calculation expected large absolute effects (≈ 0.15) and the confirmatory decision rule was framed around detecting substantive absolute changes with that calibration. Because the study was powered for large absolute reductions and the observed baseline error prevalence is far smaller than assumed, the preregistered decision framing (sized for large absolute effects) does not permit claiming support of the originally phrased large-effect confirmatory claim even though the guarded pipeline eliminated the single observed actionable error. The archived contingency counts and CI above are authoritative; we explicitly report both success-rate difference and the equivalent relative-error mapping to make the algebraic relationship transparent.

Repair and provider-call accounting diagnostics. `Provider_call_audit` in `deterministic_reconciliation.json` (SHA-256 = 8a2b85f8...) shows `hosted_model_call_event_count` = 201 and `stage_counts` = {"shared_initial_generation": 200, "repair": 1}. The adapter contract enforced at most one automated repair per task; that single recorded repair corresponds to the single discordant pair improvement (`n_10` = 1). `Model_calls_used` = 201 decomposes as 200 `shared_initial` calls + 1 repair call (`execution/model_configuration.json` SHA-256 = a40e96e2...). No provider-call failures, retries, or deterministic exclusions occurred (`missingness_summary.csv` SHA-256 = 808b3c00...). All per-episode `validator_trace` objects record whether a repair was applied (`repair_applied` = true) before/after normalized configuration; the unique `repair_applied` = true episode can therefore be identified and inspected deterministically from the archived per-episode scoring JSONs under `execution/scoring/` (full hash manifest available in `execution/execution_manifest.json`).

Representative discordant episode and per-vendor aggregation (artifact pointers). Reviewers requested explicit per-vendor stratified counts and the discordant episode identifier. Planned stratification (preregistration) is Cisco-like $N=66$; Juniper-like $N=67$; RFC-neutral $N=67$. Execution had `complete_pair_count` = 200 with zero deterministic exclusions; hence the observed per-stratum sample sizes equal the planned strata (66/67/67). Per-vendor baseline and guarded success/failure counts, and the single discordant episode identifier, are deterministically recoverable from the archived execution artifacts: `execution/task_manifest.jsonl` (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) `execution/raw_results.jsonl` (SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0) and the per-episode scoring JSONs under `execution/scoring/` (hashes recorded in `execution/execution_manifest.json`). Because these per-stratum outcome aggregates and the episode identifier are derivable without ambiguity from those archived JSONs, auditors can reconstruct them exactly; we therefore supply the authoritative artifact paths and SHA-256 digests (Disclosure) rather than reproducing extracted tables here without re-derivation. If reviewers require inline numeric reproduction in the manuscript body, we will insert the deterministically extracted per-vendor baseline/guarded

success counts and the explicit discordant episode id together with their per-file SHA-256 digests; the current submission preserves artifact-first forensic traceability in accordance with the preregistered reproducibility rules.

Additional diagnostic observations from per-episode traces. (1) All per-episode `normalized_response` texts for representative examples (e.g., `task-000001..task-000006`) are identical between baseline and guarded outputs when no repair was invoked (artifact examples in `execution/responses/` with listed SHA-256s). (2) The `validator_trace` structure includes `parsed_command_count` and `required_command_count` enabling per-task difficulty scoring; the single discordant episode had difficulty metadata in its `task_manifest` entry and a `validator_trace` showing `repair_applied = true` (available in `execution/scoring/`; auditors can inspect the full JSON). (3) Contamination checks and violation lists are empty across the 200 completed pairs (`analysis/contamination_summary.csv` SHA-256 = 389227f0...), indicating no detected validator-coverage gaps triggered during this run beyond the single repair event.

Summary interpretation of results. The guarded pipeline corrected the sole observed actionable misconfiguration (1→0) under the tested, conservative execution contract; however, because baseline failures were exceptionally rare (1/200), the absolute room for improvement was small and the preregistered power calibration (targeted at large absolute effects) is not informative for the observed rare-failure regime. The contingency counts, exact test, and bootstrap CI are consistent and fully archived for auditing; the single repair invocation is recorded and traceable to a specific per-episode JSON in the execution archive (see Disclosure SHA-256 pointers).

V. DISCUSSION

Expanded discussion (artifact-grounded technical detail and diagnostics).

1) Reconciling estimand choice, hypothesis wording, and practical decision rules. The practical difference between a relative-error-reduction hypothesis (" $\geq 50\%$ relative reduction in actionable-error rate") and the preregistered absolute paired estimand (guarded - baseline success-rate difference) is numerical and decision-relevant in low-prevalence regimes. The archived estimator used for inference and for the preregistered decision rule is the absolute paired difference (`paired_success_rate_difference_guarded_minus_baseline`), implemented and archived in `analysis/paired_contingency_table.csv` (SHA-256 = 658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c45d9414db0) and `analysis/condition_summary.csv` (SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed006f28a68a). Analysts and designers should therefore interpret confirmatory outcomes in terms of the absolute paired difference and its CI rather than raw relative percentages when baseline_error is small. We emphasize that the observed 1→0 improvement yields an absolute paired change of 0.005; algebraically this corresponds to a 50% relative reduction only if the baseline error equals 0.01; for `baseline_error=0.005` the equivalent

absolute reduction to satisfy a 50% relative reduction would be 0.0025 — a quantity outside the scope of the preregistered power plan sized for large absolute effects. This arithmetic mapping is deterministic from the archived counts and is the correct reconciliation between the hypothesis wording and the estimator used.

2) Practical audit workflow for the discordant episode and per-vendor aggregation. The execution repository contains every per-episode trace required to reproduce the contingency cell and to inspect the validator's before/after ASTs: `execution/raw_results.jsonl` (SHA-256 = 9eaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d0) `execution/task_manifest.jsonl` (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9e9f890f1c78a) and the per-episode scoring JSONs under `execution/scoring/` (full hash manifest under `execution/execution_manifest.json`). A reproducible forensic procedure is: (a) locate the single discordant pair in `analysis/paired_contingency_table.csv` (`n_10 = 1`), (b) identify its `pair_id` in `execution/raw_results.jsonl`, (c) open `execution/scoring/<episode>.json` to inspect `validation_before` and `validation_after` objects (ASTs, `parsed_command_count`, violation lists), and (d) verify provider-call provenance via `execution/provider_calls/*`. The repository-anchored SHA-256 values in the Disclosure allow auditors to verify file integrity before inspection. We implemented these steps in the internal analysis executor and recorded the actions in `analysis/analysis_log.jsonl` (SHA-256 = dbc0bd6b29e8624a2f72a35b0073a6ee7644daa85db4a7948a9fa75e276c6c3). Because the artifact set is canonical, third-party auditors can deterministically reconstruct the same episode-level narrative without ambiguity.

3) Validator coverage diagnostics and residual error taxonomy. The deterministic validator produces structured `validation_before` and `validation_after` objects that include `parsed_command_count`, `required_command_count`, `normalized_configuration`, and explicit `violation_codes`. These fields enable reproducible classification of residual errors into syntactic/parse (validator cannot parse or vendor-syntax fails), workflow/sequencing (intent-ordering, make-before-break violations), and semantic/policy mismatches (simulator-detected reachability or group constraints). In this run the guarded condition produced zero residual actionable errors (`guarded_successes = 200`). Because the residual count is zero, the preregistered H2 contingency test (residual error-type distribution comparison) was untestable as a confirmatory comparison; any discussion of residual error-type predominance is exploratory and should be examined on larger samples or deliberately higher baseline-error strata. The per-episode `validation` objects are the forensic basis for such classification and are archived per-episode under `execution/scoring/` (SHA-256 pointers in Disclosure).

4) Adapter-contract semantics and stage accounting implications. The adapter enforced `shared_initial_generation` semantics (single shared candidate across conditions) and a strict repair budget (`maximum_repair_calls_per_task = 1`). Provider-call accounting

in analysis/deterministic_reconciliation.json (SHA-256 = 8a2b85f8260409eebce1641421af6a74f1c479e505bdef80531451025931891) shows hosted_model_call_event_count = 201 and stage_counts = {"shared_initial_generation": 200, "repair": 1}, matching model_calls_used = 201. This accounting explains why the guarded condition yielded zero additional shared initial provider-calls: the pipeline reused the same initial output and exercised the repair stage only when invoked. Designers should note that adapter budgeting interacts nonlinearly with validator coverage: with a generous repair budget multiple repair attempts could correct more baseline errors but also expose potential failure modes (LLM collusion with validator hints or overfitting to validator messages). The archived provider_call_audit (execution/execution_manifest.json and deterministic_reconciliation.json) documents each call and is the source of truth for auditing these interactions.

5) Power, estimand selection for rare-failure regimes, and design recommendations. This confirmatory run demonstrates a common practical pitfall: powering a trial for large absolute effects (e.g., $\Delta \approx 0.15$) while the true baseline_error is rare (here 0.005). When baseline failure prevalence is small relative to target absolute reductions, paired binary estimands produce very low event counts and wide relative uncertainty for requested relative-reduction claims. Remedies: (a) increase N; (b) oversample high-difficulty strata where baseline_error is larger (the task generator encodes difficulty per task_payload.difficulty and auditors can filter by it deterministically); (c) adopt weighted estimands that upweight rare catastrophic-error types or cost-weighted outcomes; or (d) pre-specify hierarchical or Bayesian estimands that borrow strength across strata. All these alternatives are compatible with the existing forensic artifact model and can be simulated deterministically from the archived task templates and generator seeds in execution/task_manifest.jsonl.

6) Threats to construct and external validity tied to validator scope. The deterministic validator is necessarily an abstraction: it operationalizes a subset of vendor syntax, workflow policies, and semantic assertions available in the simulator. Residual unobserved semantic risks (e.g., vendor-specific side-effects outside the validator model) remain a concern and are explicitly bounded by the synthetic task bank’s policy language. Auditors should inspect validator coverage by sampling validator_trace.violation_codes across episodes to locate unmodeled policy types; those JSON fields are authoritative for coverage analyses.

7) Reproducibility practice and reviewer requests. Reviewers requested inline per-vendor stratified counts and the explicit discordant episode identifier. Both are deterministically derivable from the archived artifacts; the repository contains the authoritative JSONs needed to extract them (execution/task_manifest.jsonl SHA-256 = 4b28215e8a611feda50f6284058083b9fcc2f5881f557081ce9e9f890f1c78f, execution/raw_results.jsonl SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02, execution/scoring/*json SHA-256s listed in execution/execution_manifest.json). We preserve deterministic

reproducibility by pointing reviewers to these exact artifact paths and their SHA-256 digests; auditors can extract and verify per-vendor aggregates and the discordant-pair id without ambiguity. If reviewers require those numbers printed verbatim in-manuscript we will include them in a subsequent revision only after deterministic extraction (to avoid introducing errors in reporting), citing the extracted CSV with its SHA-256 digest inline.

8) Operational implication summary. Under the conservative adapter contract and the validator implementation used, the practical marginal benefit of adding a deterministic validator is tightly coupled to baseline_error prevalence and to the validator’s semantic coverage. For organizations with very low baseline LLM error rates, validator investment should be guided by the cost of a single failure and by targeted improvements in validator semantic coverage for catastrophic categories rather than by expecting large absolute proportion reductions. For higher baseline-error contexts, deterministic validators with broader simulator-backed assertions and richer repair budgets are more likely to produce material absolute gains; these configurations can be explored reproducibly from the archived task templates and generator seeds.

In sum, the discussion expands the artifacts-to-decision pipeline: arithmetic reconciliation of estimand vs hypothesis, concrete auditing steps to locate and inspect the discordant episode and per-vendor aggregates, validator-coverage diagnostics using per-episode JSON fields, and design recommendations for future confirmatory runs in rare-failure regimes. All diagnostic steps above reference specific archived files and SHA-256 digests in the Disclosure so auditors can reproduce every assertion deterministically.

VI. LIMITATIONS

- Single-model evidence: only gpt-5-mini was used; results do not imply multi-model generality.
- Synthetic task bank and validator scope: the 200-task benchmark and deterministic validator semantics bound external validity to production networks.
- Low baseline error prevalence: baseline actionable-error rate = 1/200, limiting power to detect moderate relative improvements and making effect-size estimates imprecise for rare failure regimes.
- Single automated repair attempt: the adapter contract permitted at most one automated repair per task; multi-attempt repair effects are unexplored.
- Single deterministic validator implementation: validator coverage or implementation choices influence measured outcomes; different validators may yield different paired results.
- Untestable preregistered secondary hypothesis (H2): because guarded residual failures = 0, the preregistered confirmatory contingency test comparing residual error-type proportions could not be executed and any error-type inferences are exploratory.
- Requested per-vendor observed counts and explicit discordant episode identifier are computable from archived

artifacts but are not printed verbatim in the supplied manuscript evidence bundle; auditors can compute them deterministically using the provided SHA-256 pointers.

VII. CONCLUSION

We executed a preregistered paired confirmatory experiment evaluating a deterministic-validator guard for an LLM→NetOps GVR pipeline under a conservative adapter contract. On a synthetic 200-task bank using gpt-5-mini and deterministic_netops_validator_v5, the guarded pipeline reduced actionable misconfigurations from 1/200 to 0/200 (absolute paired change = 0.005). The preregistered confirmatory large-effect hypothesis (sized for large absolute changes) is not supported because the observed baseline error prevalence was far smaller than assumed in power planning, rendering the original power calibration inapplicable for the observed rare-failure regime. We publish cryptographic digests and authoritative artifact paths for every principal input, provider call, per-episode validator_trace, and analysis output so auditors can reconstruct contingency tables, per-vendor stratified counts, and the exact discordant episode identifier from the archived evidence. Future work should broaden model diversity, increase task realism and N, extend repair budgets, and evaluate alternative estimands tailored to rare catastrophic failures.

DISCLOSURE STATEMENT

Disclosure Statement (mandatory). Initial master prompt immutable reference: provenance/master_prompt.txt SHA-256 = 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acd1582bb39d15ba1
Preregistration: cnsmsspec2026_gvr_v1_prereg_001 SHA-256 = 5cb8a30815eac0a3ca659d1a54fbaa71218e5ce57c0b13e2658099ad17d6ee41
Declared model and configuration: execution/model_configuration.json (requested_model = gpt-5-mini) SHA-256 = a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd0735027821828a611feda50f6284058083b9ffcc2f5881f557081ce9ef890f1c78a
Execution raw results and task manifest: execution/raw_results.jsonl SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02
9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02
execution/task_manifest.jsonl SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9ef890f1c78a
Deterministic reconciliation and provider-call accounting: analysis/deterministic_reconciliation.json SHA-256 = 8a2b85f8260409eebce1641421af6a74f1c479e505bdef805314510ca593489d
Paired contingency evidence: analysis/paired_contingency_table.csv SHA-256 = 658051bae1f10d90aed1728b8d20dd3700a5694b68b02f1b0c118c4510414480
Condition summary (success rates): analysis/condition_summary.csv SHA-256 = 86ab6b56e3ec10aa54aa08480a8e1ef31b64d79bf22bc99dac1d8ed00628a68a.
Missingness summary: analysis/missingness_summary.csv SHA-256 = 808b3c00ef1a906db9c3422947d9bb9997089bbbebbf3c9ebc731353240265c1c.
Analysis executor inputs: analysis/results.json (input results SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02).
Adapter and tooling: adapter_family = hosted_netops_gvr_v1;

deterministic validator = deterministic_netops_validator_v5; maximum repair calls per task enforced = 1. Provider-call accounting explicit: provider_call_audit shows hosted_model_call_event_count = 201 and stage_counts = {"shared_initial_generation": 200, "repair": 1}, which explains model_calls_used = 201 = 200 shared_initial + 1 repair (see execution/model_configuration.json SHA-256 above and deterministic_reconciliation.json SHA-256 above). Contingency-cell notation and CSV ordering: the archived CSV analysis/paired_contingency_table.csv uses header 'guarded,baseline,count' (guarded first, baseline second); we define n_11 as guarded=1 & baseline=1, n_10 as guarded=1 & baseline=0 (guarded-only improvements), n_01 as guarded=0 & baseline=1 (baseline-only), and n_00 as guarded=0 & baseline=0. Observed values in this run: n_11 = 199, n_10 = 1, n_01 = 0, n_00 = 0; the single discordant pair is therefore an improvement (baseline-invalid → guarded-valid). Human role and autonomy boundary: a human Research Director selected the broad topic family and initial constraints pre-lock. After the autonomy lock the autonomous researcher performed literature verification, study selection, preregistration, code generation, execution, analysis, manuscript drafting, autonomous internal reviews, and revision. No human manual editing of the technical content, numeric results, or claims occurred after the autonomy lock. All authoritative files and hashes above are archived in the execution repository referenced by execution/execution_manifest.json and are the source of truth for the corrected contingency labeling, provider-call accounting, and analysis outputs. Reviewer requests unresolved in-manuscript (but computable by (1) a verbatim per-vendor stratified numeric table (Cisco-like / Juniper-like / RFC-neutral) and (2) the explicit discordant episode identifier string. Both values are derivable from execution/task_manifest.jsonl (SHA-256 = 4b28215e8a611feda50f6284058083b9ffcc2f5881f557081ce9ef890f1c78a) together with execution/raw_results.jsonl (SHA-256 = 9eeaa0c87db4742c807240b5ee18f934aaf45848350b3840f92af401bcb68d02); episode scoring JSONs under execution/scoring/ (full hash manifest in execution/execution_manifest.json); because the value appears verbatim in the supplied manuscript evidence bundle text we provide these artifact pointers rather than invent numeric summaries or episode ids in the manuscript where those values are not present verbatim.

REFERENCES

- [1] <https://www.seerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.228.5064>
- [2] <https://doi.org/10.1145/2342441.2342452>
- [3] <https://doi.org/10.1145/3656296>
- [4] <https://doi.org/10.1109/mcom.001.2400368>