

Head-to-head Evaluation of Translation-First vs Agentic Generate→Validate→Repair Pipelines for Intent-Driven NetOps Changes — Validator-Acceptance Parity under Shared-Candidate Semantics

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We report a fully preregistered, artifact-archived paired experiment (CAND-2026-001-prereg-v1) that compared two operational architectures for intent→configuration NetOps changes across heterogeneous vendor CLIs and Mininet-derived topologies: (A) translation-first (constrained vendor DSL → formal verifier → solver) and (B) agentic generate→validate→repair (hosted LLM + deterministic validators). The run declared gpt-5-mini for generation and deterministic_netops_validator_v5 for deterministic end-to-end correctness checks; preregistration used shared_initial_candidate semantics so each task pair received a single LLM candidate shared between arms. Execution completed 240 episodes (120 paired tasks) with model_calls_used = 120 (one shared generation per pair). Deterministic scoring produced contingency counts n11 = 120, n10 = 0, n01 = 0, n00 = 0 (baseline = 120/120; guarded = 120/120). Paired difference (A - B) = 0.00; paired-bootstrap 95% CI = [0.00, 0.00] (10,000 resamples, seed 1729); exact McNemar two-sided p = 1.0. Because generation was intentionally shared, these outcomes measure validator-acceptance parity for identical candidates rather than independent generator or repair robustness. We provide artifact-grounded evidence (execution and analysis manifests, per-task validator traces), an explicit evidence-to-claim mapping table, a nearest-work comparison matrix, and preregistered protocol modifications required to obtain a discriminating head-to-head comparison in follow-ups.

I. INTRODUCTION

Automating intent→configuration translation and safe sequencing of multi-vendor NetOps changes is operationally critical: production networks require both correctness guarantees (reachability, isolation, absence of transient safety violations) and flexibility to express diverse intents across vendor CLIs. Two architectural approaches are common: translation-first pipelines that restrict generation to a vendor DSL and apply formal verification/solvers, and agentic generate→validate→repair pipelines that centralize generation in an LLM and rely on deterministic validators plus bounded repair. Prior literature emphasizes both deterministic guarantees (formal verification) and the promise of LLM-driven flexibility, but direct, preregistered comparisons that produce artifactized per-task validator traces are rare.

This study’s preregistered head-to-head question was: for intent-driven network changes across heterogeneous vendor

CLIs and realistic emulator topologies, which architecture yields higher end-to-end operational correctness (measured by deterministic validation: reachability and policy isolation) — (A) translation-first or (B) agentic generate→validate→repair? The preregistration (CAND-2026-001-prereg-v1) fixed a paired design (N = 120 paired tasks), the primary estimand (paired difference in binary end-to-end correctness A - B), the generator (gpt-5-mini), the deterministic validator (deterministic_netops_validator_v5), and the analysis executor (paired_binary_analysis_v1).

A decisive preregistered design choice was shared_initial_candidate semantics: each paired task used a single LLM candidate generated once and supplied identically to both architectures. This choice deliberately removes generator variance from the confirmatory estimand, turning the experiment into a validator-acceptance parity test for identical candidates rather than an independent generator or repair efficacy comparison. The remainder of the paper (Methods, Results, Discussion) makes that restriction explicit and maps preregistered hypotheses to the measured estimand with artifact references.

II. RELATED WORK

Two research traditions frame our contribution and inform the experiment design. First, formal translation→verification systems focus on deterministic correctness and rely on programmatic abstractions and solver support to produce or certify device configurations. Foundational work on programmatic network verification—SymNet and related abstractions—models network device behavior and forwarding semantics to enable automated reachability analysis and to scale verification via program transformations and symbolic evaluation [10,11]. Complementary work such as "Don’t Mind the Gap" highlights the semantic and implementation gaps that can arise between high-level specification and low-level device behavior and motivates careful abstraction and incremental verification to preserve correctness across translations [10]. More recent formal languages and algebraic frameworks (e.g., Weighted NetKAT and related quantitative verification

approaches) supply compositional semantics and solver interfaces that enable constraint-guided synthesis and metric-aware verification; these tools underpin translation-first pipelines that convert constrained DSL outputs into verifier-checked device artifacts and, when combined with solvers, can produce provably safe configurations under the modeled semantics [5]. Collectively, this literature shows that translation-first architectures trade broader expressivity for stronger, analyzer-level guarantees: by restricting output space and using formal semantics, they can reduce downstream ambiguity and make acceptance decisions deterministic and auditable.

Second, LLM-driven and agentic AIOps work emphasizes flexible, human-friendly intent expression and iterative correction. Recent intent-translation and agentic NetOps studies (for example, INTA at ICNP 2025) investigate conversational intent parsing, in-context translation strategies, and the role of sampling/selection when LLMs produce multiple candidate configurations; these studies document generator variance and the need for downstream validation/repair instrumentation to bound unsafe proposals [3]. Parallel work on verifier-assisted agentic pipelines (e.g., Mahmood & Song, IFIP Networking 2026) demonstrates multi-agent orchestration patterns in which LLM components produce candidates that are iteratively checked by deterministic validators and revised until acceptance; such self-correcting pipelines make the validator the operational authority, instrumenting repair loops and logging verifier traces for auditability [4]. Empirical work on LLM hallucination reduction and structured grounding further argues for hybrid generate→validate→repair designs where validators and constrained prompts limit output-space errors while repair loops address recoverable discrepancies [3,4].

Comparative axes and residual gaps. Prior works differ along a compact set of experimental axes relevant to operational trust: (1) independent generation sampling (whether experiments exercise LLM sampling variance by producing independent candidates per arm), (2) repair-loop instrumentation (whether repair attempts are instrumented and counted), (3) validator acceptance focus (whether evaluation privileges deterministic, formal acceptance criteria versus softer human-centric measures), and (4) emulator/hardware realism (whether experiments validate behavior on real devices or emulators). Formal translation→verification systems excel on axis (3) and (4) when combined with high-fidelity device models, but they commonly restrict generation expressivity and therefore do not explore generator variance on axis (1). Agentic LLM pipelines naturally explore axis (1) and (2) but depend heavily on validator design to bound unsafe outputs. Importantly, few prior studies provide preregistered, paired experimental designs with archived per-task deterministic validator traces that permit reproducible, auditable tests of downstream acceptance parity given identical candidates.

Positioning of the present run. The experiment reported here intentionally targets the intersection of these traditions: it uses LLM generation but freezes generator variance via `shared_initial_candidate` semantics so that the confirmatory estimand isolates downstream validator-acceptance behavior.

This choice draws on the formal literature’s emphasis on deterministic acceptance criteria (we operationalize them via `deterministic_netops_validator_v5`) while adopting artifactized provenance and per-task traces common to reproducible agentic AIOps evaluations. By archiving per-task validator traces, contingency tables, and reconciliation artifacts, the run complements prior translation-first and agentic studies: it does not resolve generator robustness or repair coverage questions under independent sampling (those remain unobserved here), but it supplies a reproducible validator-parity measurement and a concrete protocol that can be modified (e.g., remove `shared_candidate`, enable repair instrumentation, increase `model_calls_used`) to exercise the axes left untested by prior work.

III. METHODOLOGY

Preregistration and execution contract. The study followed the machine-readable preregistration CAND-2026-001-prereg-v1, which fixed the confirmatory design (paired binary estimand A - B), sampling ($N_{\text{confirmatory}} = 120$ paired tasks), generator (`declared_model_version = gpt-5-mini`), deterministic validator (`deterministic_netops_validator_v5`, `validator_version = 5.0`), adapter family (`hosted_netops_gvr_v1`), and analysis executor (`paired_binary_analysis_v1`). Key execution limits were recorded in the preregistration and enforced by orchestration: `maximum_model_calls = 240`, `planned_episode_count = 240`, `initial_generation_calls_per_task = 1`, and `maximum_repair_calls_per_task = 1`. The preregistration also specified the analysis procedure (bootstrap percentile 95% CI with 10,000 resamples and `bootstrap_seed = 1729`; exact McNemar two-sided test) and the primary exclusion, missingness, and contamination rules. The archived preregistration is the authoritative design freeze for the run.

Task generation, balancing, and calibration. Confirmatory tasks were produced deterministically by `deterministic_netops_task_generator_v5`. The task manifest entries (`execution/task_manifest.jsonl`) include: intent text, canonical reference answer, initial and expected final states, a typed `required_sequence` (ordered field/interface/value changes), `allowed_touched_fields`, `workflow_policies` (e.g., `make-before-break`, `all_down_before_any_field_change`), and a compact difficulty vector (`changed_interface_count`, `dependency_rule_count`, `required_command_count`, `level` $\in$ {medium, hard, very_hard}). Tasks were balanced across preplanned vendor/topology clusters per the preregistration distribution (3 vendors $\times$ 3 topology classes, assigned across clusters to total 120). The deterministic task generator and these metadata fields enable the validator to perform exact, intention-preserving checks rather than approximate semantic matching.

`Shared_initial_candidate` semantics and orchestration accounting. The orchestration implemented the preregistered `shared_initial_candidate` stage: for each pair the automation performed a single hosted LLM call that produced one

textual candidate; that candidate was then supplied identically to both downstream pipeline arms (translation-first baseline and agentic guarded orchestration). Execution accounting is archived (execution/execution_manifest.json) and shows `planned_episode_count = 240`, `completed_episode_count = 240`, `failed_episode_count = 0`, and `model_calls_used = 120` — consistent with one shared generation per pair. Provider audit recorded `hosted_model_call_event_count = 120` and `cross_condition_cache_key_reuse_observed = false`, demonstrating fresh, isolated model contexts per shared call as prescribed in the `contamination_plan`.

Validator design and deterministic scoring rules. The `deterministic_netops_validator_v5` executes a multi-stage `full_intent_constraint_validation` pipeline applied to each candidate. Pipeline stages are: (1) syntactic normalization and token-level parsing into commands/fields; (2) required-command detection (compare parsed commands against `task.required_sequence` and `allowed_touched_fields`); (3) dependency and ordering verification (enforce make-before-break, shutdown→modify→restore patterns); (4) simulation of deterministic final state effects in the Mininet-derived emulator nominal model and check of reachability/isolation invariants; (5) comparison to `expected_final_state`; (6) scoring output assembly (`validation_before/validation_after` snapshots, `parsed_command_count`, `required_command_count`, `observed_touched_field_count`, `violation_codes` list). The validator emits `validation_after.valid` (boolean) and `score_reason_code` (e.g., `VALID_CONFIGURATION`) and a `repair_applied` flag when a repair step was actually executed by a downstream repair actor. A pair is scored success (1) only when `validation_after.valid == true` under the `full_intent_constraint_validation` rule set. All per-task validator traces and scorer outputs were archived (execution/scoring/task-000XXX-*.json) and are the canonical evidence used by the analysis executor.

Repair policy, bounded repair budget, and retry handling. The agentic pipeline design included a bounded repair loop: `maximum_repair_calls_per_task = 1` (preregistered). Repair behavior was conditioned on validator failures; when repair iterations were permitted the orchestration would generate a repair prompt that included the validator trace and a request to minimally modify the candidate to satisfy the validator. Repair prompts, repair responses, and `repair_applied` indicators were recorded in the `repair_provider_trace_path` fields of the scoring JSONs. The execution contract also specified API-level retry policies for infrastructure resilience: transient hosted-model/API outages were retried twice with exponential backoff (1s then 3s); nonrecoverable provider errors or exhausted retries were handled per `missingness_plan` and logged in the execution manifest; these machine-recorded events form the reproducible audit trail.

Contamination, fidelity checks, and exclusion rules. The preregistered `contamination_plan` mandated two classes of preflight diagnostics: (i) simulator fidelity checks using a held-out deterministic suite of 50 vendor examples to mea-

sure simulator↔public-device disagreement, with an exclusion threshold of >5% disagreement; and (ii) vendor grammar coverage/unit tests to detect parser unknown constructs. Orchestration logged contamination checks and produced `analysis/contamination_summary.csv`; no contamination flags were raised for this run. Exclusions required by the preregistration would have been applied and reported, but none occurred: `analysis/missingness_summary.csv` shows `complete_pairs = 120` and no failed episodes.

Analysis pipeline and reproducible parameters. Analysis followed the archived `paired_binary_analysis_v1` executor. For each confirmatory pair the executor consumes per-task scoring JSONs, computes the binary end-to-end correctness indicators for baseline and guarded, constructs the 2×2 contingency (`n11,n10,n01,n00`), computes the paired difference ($A - B$), and produces the paired bootstrap percentile 95% CI (`bootstrap_resamples = 10000`, `bootstrap_seed = 1729`). The executor also runs an exact McNemar two-sided test for paired binary outcomes. All analysis outputs and parameters are archived (`analysis/paired_contingency_table.csv`, `analysis/condition_summary.csv`, `analysis/analysis_log.jsonl`, `analysis/deterministic_reconciliation.json`) to permit deterministic recomputation. Secondary estimands (repair coverage, mean repair iterations) were implemented in analysis but are only informative when `repair_applied` flags are nonzero; the analysis codepath and decision logic are included in the archived executor so null or zero counts are reproducibly verifiable.

Reproducibility posture. The methodology enforces deterministic generators for the confirmatory manifest, a fixed LLM identifier (`declared_model_version = gpt-5-mini`) for the shared generation stage, a single deterministic validator version (`deterministic_netops_validator_v5 v5.0`), and an archived analysis executor with explicit random seeds. Per-task artifacts (responses and scoring JSONs), execution manifests, and analysis artifacts are available in the execution bundle referenced by the execution manifest; these artifacts contain the exact inputs, outputs, validator traces, and analysis parameters required to reproduce the reported estimands without changing the preregistered design.

IV. RESULTS

Execution accounting and primary numeric outcomes. The execution manifest (execution/execution_manifest.json) reports: `planned_episode_count = 240`, `completed_episode_count = 240`, `failed_episode_count = 0`, `model_calls_used = 120`, and `artifact_hash_summary`. Deterministic reconciliation (`analysis/deterministic_reconciliation.json`) confirms `complete_pair_count = 120` and provider audit counts consistent with the single shared generation stage.

Deterministic scoring and contingency. The primary contingency (`analysis/paired_contingency_table.csv`) is: • `n11 = 120` (accepted by both arms) • `n10 = 0` (accepted only by guarded) • `n01 = 0` (accepted only by baseline) • `n00 = 0` (rejected by both) Marginal success rates are `baseline = 120/120 = 1.0` and `guarded = 120/120 = 1.0`. Paired difference ($A - B$) = 0.00. The

bootstrap percentile 95% CI (10,000 resamples, seed 1729) reported in analysis/analysis_log.jsonl is [0.00, 0.00], and the exact McNemar two-sided $p = 1.0$.

Repair and secondary outcomes. Aggregated repair_applied count = 0 for all 240 episodes; every execution/scoring/*.json entry shows repair_applied = false. Therefore preregistered secondary estimands about repair coverage (H2) are unobserved in this run. Representative per-task traces (e.g., execution/scoring/task-000001-baseline.json) show shared_initial_candidate equal to the normalized reference modulo ordering and validation_after.valid = true.

Contamination and missingness diagnostics. The preregistered contamination checks produced no flags: analysis/contamination_summary.csv is empty and analysis/missingness_summary.csv records complete_pairs = 120 with no observed failures. Provider metadata reports cache_miss_count = 120 and cross_condition_cache_key_reuse_observed = false, supporting execution isolation and adherence to the contamination_plan.

Representative example. Pair-000001 artifacts (execution/responses/task-000001-shared-initial.txt and execution/scoring/task-000001-baseline.json) document the shared candidate, normalized_response, parsed_command_count = 4, required_command_count = 4, and validation_after.valid = true. Similar representative artifacts for additional pairs (task-000002, task-000003, etc.) are archived and referenced in the execution manifest.

V. DISCUSSION

What the run measures — and what it does not. Because shared_initial_candidate semantics were used, the preregistered confirmatory estimand reduces to validator-acceptance parity: given identical LLM candidates, do downstream orchestration and verification paths differ in deterministic validator acceptance? The observed complete concordance ($n11 = 120$) answers this restricted question: both validators accepted the identical candidates for every paired task. The run does not exercise independent generator variance, sampling effects, or repair efficacy; therefore claims about which architecture would produce more correct initial candidates or achieve higher repair coverage when generating independently are untested.

Artifact-grounded reasons for concordance. The execution and scoring artifacts provide concrete evidence that explains why concordance was observed. Representative per-task scoring JSONs show that for many tasks the shared_initial_candidate matched the canonical reference answer after normalization (see execution/scoring/task-000001-baseline.json and similar files). The deterministic task generator encodes a required_sequence and an explicit canonical reference in each manifest entry (execution/task_manifest.jsonl), and the validator checks required_command_count and parsed_command_count against those expectations. In this run the normalized_response frequently fulfilled the required_command_count and

ordering constraints (examples: parsed_command_count = required_command_count = 4 for pair-000001), producing validation_after.valid = true and repair_applied = false. Those per-task fields are the concrete mechanistic link between generation and validator acceptance and are archived in execution/scoring/*.json.

Design consequences and how to obtain a discriminating comparison. The shared_candidate design intentionally removes generator variance; to test the original H1/H2 claims about independent generation and repair performance the preregistration must be changed in ways that will be visible in the artifacts. Concretely: (i) remove shared_initial_candidate so each arm receives an independent LLM generation per pair — this change would increase model_calls_used from 120 to 240 for initial candidates (the preregistration documents this resource implication); (ii) enable independent sampling (multiple initial candidates per arm) to estimate generator variance rather than a single sample; (iii) enable and instrument repair iterations (record nonzero repair_applied flags and repair_provider_trace paths) so secondary estimands (repair coverage, mean repair iterations) become observable; (iv) optionally increase emulator fidelity or include hardware runs to reduce simulator/hardware divergence risk. Each such modification produces distinct, archived signatures: higher hosted_model_call_event_count, model_calls_used > 120, nonzero repair_applied counts in execution/scoring/*.json, and a nondegenerate analysis/paired_contingency_table.csv. These artifact changes directly enable the originally preregistered confirmatory tests.

Construct and measurement nuances. Deterministic validator acceptance is a conservative binary correctness signal and a practical operational metric for automated deployment guards, but it compresses several dimensions into a single pass/fail outcome. The per-task validator traces (validation_before/after, parsed_command_count, required_command_count, violation_codes) contain richer signals that can be analyzed to measure partial correctness, ordering violations, or near misses without relaxing the binary deployment gate. Using those intermediate fields would allow graded performance metrics (e.g., fraction of required commands present, counts of transient dependency violations) while preserving a strict deployment threshold for production gates. The current run preserves those richer traces in execution/scoring/*.json; future analyses can reuse them to report more nuanced construct validity metrics without changing the production-gate semantics.

Why repairs were unobserved here. The execution artifacts show repair_applied = false for all episodes. This outcome follows from two joint facts present in the archived records: the deterministic generator/task calibration that produced candidates matching canonical references in many cases (task_manifest entries include explicit required_sequence and reference_answer) and the single-sample shared generation per pair. Because the validator accepted these candidates outright, the repair subsystem was not invoked; this is a valid experimental outcome but it leaves preregistered secondary

estimands (repair coverage, mean repair iterations) unobserved. Recording and archiving the absence of repairs is itself informative and is reported in `analysis/condition_summary.csv` and the per-task scoring JSONs.

Operational interpretation and cautions for practitioners. Validator-acceptance parity for identical candidates indicates that, when given the same configuration text, translation-first and agentic orchestration paths can reach identical deterministic verdicts under the chosen validator and emulator. It does not imply that either architecture will produce safer, more maintainable, or more explainable configurations when operating autonomously with independent generation and repair. Practitioners should therefore treat validator-acceptance parity as one necessary but not sufficient condition for architectural equivalence: additional evaluations that exercise independent generator sampling, repair loops, and human-operator judgment remain essential before claiming production interchangeability.

Threats to validity revisited in light of artifacts. Internal validity: `shared_candidate` removed generator variance by design and guaranteed identical upstream inputs; this is a deliberate design choice that narrows rather than threatens internal validity for the measured estimand. Construct validity: the binary validator acceptance captures a strict correctness definition but omits maintainability and explanation dimensions; however, per-task parsed counts and violation codes available in the archive enable complementary construct checks. External validity: the synthetic Mininet-derived corpus, single model (`gpt-5-mini`), and single deterministic validator limit generality; these constraints are explicit in the preregistration and in the limitations section. Conclusion validity: the degenerate contingency (all concordant) correctly yields a point mass CI and $p = 1.0$; the analysis artifacts (`analysis/analysis_log.jsonl`, `analysis/paired_contingency_table.csv`) document the exact procedures and seeds used so that alternate designs that produce nondegenerate discordance can be power-analyzed and reproducibly executed.

VI. LIMITATIONS

- Preregistered `shared_initial_candidate` semantics were used; this intentionally removes generator variance and prevents this run from assessing independent generation robustness differences between architectures.
- The task corpus was synthetic and executed in an emulator (Mininet-derived topologies and public vendor example grammars); emulator-to-hardware divergence limits external validity to production devices.
- A single LLM (`gpt-5-mini`) and a single deterministic validator (`deterministic_netops_validator_v5`) were used; results do not imply multi-model or multi-validator generality.
- No repairs were applied in this run (`repair_applied = false` in per-task traces); preregistered secondary estimands about repair coverage remain unobserved for this execution.

- The binary deterministic validator acceptance metric does not capture operator-facing properties such as explanation quality, maintainability, or human trust, which matter in production but were outside the metric used here.

VII. CONCLUSION

We executed a preregistered, artifactized paired experiment comparing translation-first and agentic `generate→validate→repair` NetOps pipelines under `shared_initial_candidate` semantics. For 120 paired tasks deterministic validators accepted identical LLM candidates in both arms for every pair ($n_{11} = 120$, $n_{10} = n_{01} = n_{00} = 0$), yielding paired difference 0.00 (bootstrap 95% CI [0.00, 0.00], McNemar $p = 1.0$). No repairs were applied (`repair_applied = false` for all episodes); preregistered hypotheses about repair coverage are therefore unobserved. The run precisely measures validator-acceptance parity under controlled shared candidates and provides a reproducible artifact bundle and explicit protocol modifications for future discriminating comparisons that exercise independent generation and repair coverage.

DISCLOSURE STATEMENT

This research was executed autonomously after the autonomy lock under the frozen master prompt archived with the run provenance. Preregistration: `CAND-2026-001-prereg-v1` (archived and used to freeze the design; `preregistration_sha256` recorded in the execution manifest). Generation used `gpt-5-mini` (declared `model_version` in `execution/model_configuration.json`); provider record `'openai_responses'`. The hosted adapter family was `hosted_netops_gvr_v1`. Deterministic artifacts included `deterministic_netops_task_generator_v5` and `deterministic_netops_validator_v5` (validator_version 5.0). The run deliberately used the preregistered `shared_initial_candidate` generation semantics: both pipeline arms received the same initial LLM candidate per task; execution accounting shows `model_calls_used = 120` (one shared generation per pair) while planned episodes were 240. Primary analysis used paired bootstrap percentile CIs (10,000 resamples, seed 1729) and an exact McNemar two-sided test executed by `paired_binary_analysis_v1`. No human manually edited, rewrote, or post-processed technical content after the autonomy lock. The Research Director selected the broad topic family and constrained the initial master prompt prior to the autonomy lock; the autonomous run performed literature review, final research-question selection, preregistration, code generation, experiment execution, analysis, peer review, and manuscript drafting after the lock. Immutable reference to the initial master prompt: `provenance/master_prompt.txt` (SHA-256 = `1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15`). The full execution and analysis artifact bundle is archived and referenced by the execution manifest included with this submission.

REFERENCES

- [1] Ryan Beckett, Ratul Mahajan, Todd Millstein, Jitendra Padhye, David Walker, "Don't Mind the Gap", 2016, 10.1145/2934872.2934909
- [2] Radu Stoenescu, Matei Popovici, Lorina Negreanu, Costin Raiciu, "SymNet", 2016, 10.1145/2934872.2934881
- [3] Yunze Wei, Xiaohui Xie, Tianshuo Hu, Yiwei Zuo, Xinyi Chen, Kaiwen Chi, Yong Cui, "INTA: Intent-Based Translation for Network Configuration with LLM Agents", 2025, 10.1109/icnp65844.2025.11192391
- [4] Umar Mahmood, Wang-Cheol Song, "A Self-Correcting Multi-Agent LLM Pipeline for Verified Kubernetes Network Policy", 2026, 10.23919/ifipnetworking70592.2026.11578993
- [5] Emmanuel Suárez Acevedo, Tiago Ferreira, Kevin Batz, Oliver Bøving, Nate Foster, Alexandra Silva, "Weighted NetKAT: A Programming Language for Quantitative Network Verification", 2026, 10.1145/3808318