A Preregistered Audit of Dataset Contamination Controls in LLM-for-NetOps Benchmarks

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

***Abstract*—We preregistered and executed a paired, confirmatory audit to test whether conservative deterministic contamination detectors (exact-match + fuzzy 5-gram + provenance heuristics) flag items in a reproducible 150-item NetOps sample and whether applying preregistered contamination controls changes validator-based success rates. Crucially, the executed sample used provenance-stable synthetic tasks (source_identifier prefix "synthetic-netops:") and shared generation semantics (independent_condition_generation = false), recorded in the archived model and task manifests; these two execution facts materially limited the audit's sensitivity to generation-level contamination. No preregistered high-confidence contamination flags occurred and therefore no guarded exclusions or replacements were performed. All 150 baseline and matched guarded episodes passed the deterministic validator (baseline = 150/150, guarded = 150/150). The paired success-rate difference = 0.0 percentage points (95% CI [0.0, 0.0]); exact McNemar two-sided p = 1.0. We reconcile the preregistration→execution gap with artifact-grounded diagnosis, explain why the run is degenerate for detecting public-benchmark leakage, and provide concrete, artifact-anchored recommendations for audit designs that would be sensitive to contamination in web-sourced benchmarks.**

I. Introduction

Motivation. Large language models (LLMs) are increasingly applied to NetOps tasks such as intent-to-configuration translation and automated configuration repair. Operational decisions based on benchmark results require that measured performance reflect model capability rather than dataset contamination or templated item leakage into model training. Meta-analyses and recent surveys call for contamination-aware evaluations and validator-backed oracles to avoid inflated reliability claims.

Research question and contribution. We preregistered and executed an artifact-anchored audit to answer: under conservative deterministic contamination detectors (exact normalized token equality; fuzzy 5-gram shingle overlap with preregistered thresholds; provenance timestamp heuristics) and a guarded exclusion policy, how many high-confidence flags arise in a reproducible 150-item NetOps sample, and how much does applying those controls change a single hosted LLM's validator pass rate? Our contributions are: (1) a fully preregistered confirmatory experiment (study_id = contamination-audit-netops-2026-v1) executed end-to-end and archived; (2) exact, artifact-verified confirmatory statistics on Npairs = 150 with deterministic validator traces; (3) an explicit preregistration→execution reconciliation that identifies why the run produced a degenerate null and prescribes artifact-grounded design changes to increase audit sensitivity.

Scope and bounded claims. The audit intentionally targeted a methodological question about preregistered contamination controls, not an empirical prevalence estimate for public web-sourced benchmarks. The executed confirmatory sample used provenance-stable synthetic tasks (source_identifier prefix "synthetic-netops:") produced by the deterministic task generator and employed shared_initial_candidate generation semantics (independent_condition_generation = false). Because both facts are recorded in the archived model and task manifests, they limit inference: the observed zero-flag, zero-difference outcome applies to this synthetic sampling and generation regime and does not imply absence of contamination in independent public benchmarks.

II. Related Work

LLM-driven NetOps systems and verifier integration. Recent system and evaluation work has investigated LLMs for intent-driven configuration synthesis and for proposing configuration repairs; these studies highlight practical value but also recurring failure modes that motivated tool-assisted or agentic architectures. Representative intent→configuration systems show that LLMs can produce syntactically valid device artifacts when guided by structured prompts or in-context examples, but that semantic correctness (reachability, policy preservation) typically requires an external validator or iterative refine-and-repair loop [https://openalex.org/W4404602270, 10.1109/vtc2025-fall65116.2025.11310185]. Benchmarks and controlled experiments on LLM-assisted repair report improved average repair quality when the pipeline integrates a deterministic checker (or formal verifier) with generation and repair steps; however, per-case regressions remain common and depend strongly on the verifier's property coverage and on dataset realism [https://openalex.org/W7157506260, https://openalex.org/W7163636922].

Formal verification and deterministic oracles. The NetOps literature provides mature deterministic verification techniques (Header Space Analysis and derived tools, localized subspecification approaches, and NetKAT/related formalizations) that are practical oracles for reachability and many policy properties; these tools form the empirical basis for validator-backed evaluation in NetOps generate-validate pipelines [http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.300.8659, https://openalex.org/W2742515467, https://openalex.org/W1882012874]. Studies integrating such verifiers show that the choice of checks (which fields and temporal/workflow constraints are modeled) materially

changes which generation errors are visible to the auditor—an observation we incorporate by using a deterministic validator (deterministic_netops_validator_v5) that emits parse→state traces and explicit violation codes for every episode.

Agentic validate–repair architectures and their limits. Multiple contemporaneous projects evaluate agentic or multi-stage designs where an LLM proposes a candidate, a verifier checks it, and a repair generator attempts fixes when violations are found; these generate–validate–repair patterns generally improve mean repair rates versus single-pass generation but do not eliminate regressions and depend on the repair-policy and localization strategy [https://openalex.org/W7162218766, https://openalex.org/W7163636922]. Empirical studies also show scale sensitivity: effectiveness reported on curated or synthetic misconfiguration tasks often degrades when confronted with more heterogeneous, multi-vendor, or production-sourced corpora. These findings motivate two audit recommendations we make below: (1) use provenance-rich sampling to exercise contamination detectors, and (2) persist richer validator/tracer outputs to analyze subtle regression modes.

Benchmark validity, contamination, and detector strategies. Recent surveys and meta-evaluations emphasize two recurring evaluation validity risks: dataset contamination, where benchmark items (or templated prompts) are present in LLM training data and inflate measured accuracy, and construct invalidity, where tasks are overly templated or low complexity and fail to exercise operational difficulty [https://openalex.org/W4404534210, https://openalex.org/W7166900660]. The literature recommends deterministic fingerprinting (canonicalized exact matching), fuzzy n-gram overlap, and provenance/timestamp tracing as complementary heuristic detectors; these techniques form the basis for many proposed audit workflows because they are computationally tractable and interpretable. However, prior work also highlights tradeoffs: exact fingerprints miss paraphrased memorization; n-gram overlap thresholds require sensitivity analysis; and provenance heuristics depend on available source metadata and archival coverage. Our preregistration implemented a conservative combination (exact normalized equality OR fuzzy 5-gram overlap >= 0.80 for high confidence, with lower thresholds for sensitivity analyses) precisely because the literature documents these tradeoffs and recommends multimodal flagging and post-hoc adjudication [https://openalex.org/W4385681611, https://openalex.org/W7166900660].

How our study differs from and complements prior work. Unlike broader capability benchmarks that measure absolute LLM performance across many models, or agentic repair evaluations that emphasize mean repair rates, our study is a preregistered contamination-control audit: it asks whether a conservative, preregistered detector ensemble would produce high-confidence flags in a reproducible 150-item sample and whether deterministic guarded exclusions would alter validator pass rates. Prior benchmark and system papers typically do not publish preregistered exclusion rules or the exact generation semantics used for paired comparisons; that omission can obscure whether observed generation discordance arises from detector sensitivity or from sampling/generation choices. By freezing detector thresholds, exclusion policies, sampling seeds, and generation semantics in a preregistration and archiving the validator traces, our experiment complements the literature by making the audit contract explicit and by demonstrating (artifact-grounded) how deterministic sampling and shared generation semantics can collapse a paired confirmatory test into a degenerate outcome. This reconciliation—linking verifier traces, provider call accounting, and the detector workflow to the resulting paired contingency—is the primary methodological contribution that complements prior empirical and system-level findings [https://openalex.org/W7157506260, https://openalex.org/W7162218766, https://openalex.org/W4404602270].

III. Methodology

Preregistration and primary estimand. The experiment was preregistered (study_id = contamination-audit-netops-2026-v1) and froze analysis and exclusion rules before execution. The primary estimand is the paired success-rate difference (guarded - baseline) across item×model pairs, tested by exact McNemar two-sided test and summarized with a paired bootstrap percentile 95% CI (10,000 resamples, seed = 1729) as specified in the preregistered paired_binary_analysis_v1 executor.

Detectors and guarded policy (preregistered). The preregistered detectors were: (1) exact normalized token equality after canonicalization; (2) fuzzy 5-gram shingle overlap (Jaccard/overlap) with preregistered high-confidence threshold >= 0.80 and sensitivity thresholds at 0.65 and 0.50; and (3) provenance timestamp heuristics comparing item timestamps to model cutoff. The guarded policy deterministically excludes only high-confidence flagged items (exact match OR fuzzy overlap >= 0.80) and replaces them via deterministic retemplating/topology augmentation from the same generator with fixed seeds; replacements must meet an embedding similarity constraint per the preregistration.

Sampling, generation semantics, and the critical execution choice. The preregistration allowed deterministic task generators for reproducibility and permitted shared generation semantics when appropriate. The executed run used deterministic_netops_task_generator_v5 to produce synthetic tasks (source_identifier prefix "synthetic-netops:") for the confirmatory sample. Execution settings used shared_initial_candidate semantics (independent_condition_generation = false): when no exclusions occur, one shared model generation per item is reused for both baseline and guarded episodes. These execution settings are recorded in the archived model_configuration and task_manifest artifacts. The preregistered contract also planned replacement rules, sensitivity analyses, and a maximum of 300 model calls (one per condition per item) if independent generations were used.

Verifier and scoring rule. The verification oracle is deterministic_netops_validator_v5. For intent→configuration tasks the validator applies parse→state-update→property checks and returns a binary pass/fail along with detailed traces (observed_touched_field_count, parsed_command_count, required_command_count, violation codes). The preregistered scoring rule is full_intent_constraint_validation: outputs must achieve the intended final state and respect workflow constraints. Validator traces are archived for every episode and provide the empirical basis for the binary outcomes used in the paired analysis.

Planned sensitivity and power. The preregistration included sensitivity analyses: treating failed model calls as failures; evaluating fuzzy thresholds at 0.50, 0.65, 0.80; and subgrouping by templated vs realistic items. The power guidance documented that a multi-model or two-call per item design (300 independent generations) would detect smaller paired differences (≈4–6 percentage points under specified assumptions), while the single-model single-call shared-generation regime has reduced sensitivity. These planning notes are part of the preregistration archive and informed our interpretation of the executed null result.

IV. Results

Execution accounting and deterministic reconciliation. The archived execution manifest records planned_episode_count = 300 (150 items × 2 conditions) and completed_episode_count = 300 with failed_episode_count = 0. Because independent_condition_generation = false and no high-confidence exclusions occurred, the pipeline used model_calls_used = 150 (one shared generation per item) rather than two separate generations per condition; provider-call accounting and reconciliation artifacts confirm these counts and semantics (see execution/model_configuration.json, execution/task_manifest.jsonl, and analysis/deterministic_reconciliation.json). Stage-level audit fields show stage_counts: {"shared_initial_generation": 150} and condition_counts: {"shared": 150}, and provider_call_audit reports hosted_model_call_event_count = 150, completed_hosted_model_call_event_count = 150, cache_miss_count = 150, and cross_condition_cache_key_reuse_observed = false — all consistent with a single shared initial generation reused by both conditions for each item.

Contamination detectors: absence of preregistered high-confidence flags. The preregistered detectors produced zero high-confidence contamination flags for the executed sample; the analysis artifact analysis/contamination_summary.csv contains no flag rows. The detector workflow was implemented to persist per-item fuzzy overlap scores only when threshold crossings or replacements were required; because no items met the preregistered high-confidence threshold (exact OR 5-gram >= 0.80), the overlap pipeline did not materialize an aggregated per-item overlap table in the archive. This empty flag summary is an explicit artifact result (analysis/contamination_summary.csv) and not an absence inferred from silence: the detector logged no threshold events to persist. Consequently, no guarded exclusions or deterministic replacements were performed (replacements_performed = 0 in the execution manifest).

Validator outcomes and confirmatory statistics. The deterministic_netops_validator_v5 returned pass for all baseline episodes (baseline_success_count = 150/150; baseline_success_rate = 1.0) and for all guarded episodes (guarded_success_count = 150/150; guarded_success_rate = 1.0). The paired contingency table archived as analysis/paired_contingency_table.csv is n11 = 150, n10 = 0, n01 = 0, n00 = 0. The primary estimand (paired_success_rate_difference_guarded_minus_baseline) = 0.0 percentage points. The preregistered paired bootstrap percentile 95% CI (10,000 resamples, bootstrap_seed = 1729) is [0.0, 0.0]; the exact McNemar two-sided test yields p = 1.0. These numbers are computed by the archived paired_binary_analysis_v1 executor and are recorded in analysis/results.json and analysis/condition_summary.csv (condition,successes,failures,observed,success_rate: baseline,150,0,150,1.0; guarded,150,0,150,1.0).

Representative episode diagnostics and difficulty coverage. Representative archived episodes span the preregistered difficulty levels (medium, hard, very_hard) and illustrate validator behavior across task families. For example, pair-000001 (medium, pattern "shutdown_configure_restore") shows normalized_response identical to the reference and validator fields parsed_command_count = 4, required_command_count = 4, observed_touched_field_count = 3, and validation_after.valid = true; pair-000002 (hard, "make_before_break_uplink_switchover") shows parsed_command_count = 2 and validation_after.valid = true; pair-000004 (very_hard, "safe_access_service_transfer") shows required_command_count = 4 and validation_after.valid = true. These representative traces are included among the archived execution/scoring artifacts and demonstrate that the validator exercised per-task workflow constraints and recorded concrete parse→state metrics for inspection. Because all episodes were scored as successes (score = 1, scoring_reason_code = "VALID_CONFIGURATION") the global contingency counts and per-episode validator traces together indicate systematic validator agreement with the generated outputs for this synthetic sample.

Provider tracing, shared generation semantics, and implications. Provider trace fields recorded per-episode shared_initial_provider_trace_path entries pointing to execution/responses/*-shared-initial.txt and identical raw_response_sha256 values for baseline and guarded episodes in each pair. These traces, together with model_calls_used = 150 and the stage_counts described above, document the operational mechanism that produced identical outputs across conditions when no exclusions were triggered. In short: when detectors produced no high-confidence flags, the guarded pipeline performed no replacements and the shared initial generation was reused

for both conditions; this architecture produced perfect concordance (n11 = 150) and eliminated any generation-level discordance that an independent_generation design could have revealed.

Missing overlap summaries and interpretation. The preregistered fuzzy-overlap detector would have persisted per-item overlap scores and flags if thresholds were crossed. The archive contains no such per-item overlap summary because the detector implementation persisted overlap outputs only for items that triggered flags or replacements. The absence of persisted overlap scores is therefore an explicit artifact of the detector workflow combined with zero threshold crossings; it prevents reporting of min/median/percentile overlap statistics for this run, and we note this gap as an empirical limitation of the executed logging policy (see analysis/contamination_summary.csv and the analysis_log.jsonl entry describing detector events).

Synthesis of artifact evidence. The combination of (a) provenance-stable synthetic sampling (source_identifier values of the form "synthetic-netops:..." recorded in execution/task_manifest.jsonl), (b) shared_initial_candidate generation semantics (execution/model_configuration.json), (c) provider call accounting (analysis/deterministic_reconciliation.json), and (d) empty contamination_summary.csv together explain the degenerate confirmatory outcome: deterministic detectors produced no high-confidence flags and the reuse of a single shared generation per item produced identical baseline and guarded outputs, yielding complete concordance in validator outcomes. All numbers and traces cited above are archived in the execution bundle and referenced by the artifact paths given here (e.g., execution/raw_results.jsonl, analysis/paired_contingency_table.csv, analysis/condition_summary.csv, analysis/contamination_summary.csv, analysis/deterministic_reconciliation.json).

V. Discussion

Why the result is degenerate: artifact-grounded diagnosis. Two archived execution facts together explain the degenerate null outcome. First, the confirmatory sample was drawn entirely from provenance-stable synthetic tasks generated by deterministic_netops_task_generator_v5 (task entries carry the "synthetic-netops:" source_identifier). Synthetic provenance makes exact normalized token equality and high fuzzy 5-gram overlap with public corpus fingerprints unlikely by construction, and therefore the preregistered high-confidence detector rule (exact match OR 5-gram overlap >= 0.80) had little opportunity to trigger. Second, execution semantics set independent_condition_generation = false (shared_initial_candidate). When the detector produced no high-confidence exclusions, the pipeline reused a single model generation for both baseline and guarded episodes; this reuse is recorded in the model and provider accounting artifacts and reduces independent variance between conditions. Both facts are present in the archived execution/model_configuration.json and execution/task_manifest.jsonl artifacts and confirmed by the deterministic_reconciliation outputs. Together these facts deterministically produce identical baseline/guarded outputs per item and therefore make the paired test uninformative for generation-level contamination effects in this run.

Mechanics of the detector workflow and reporting consequences. The preregistered detector was implemented to persist detailed per-item overlap scores only when threshold crossings or replacement handling were invoked (to limit storage and to focus archival attention on resolved flags and replacements). Because no items crossed the high-confidence threshold, the detector produced no persisted per-item overlap table in analysis/contamination_summary.csv. This implementation choice—record only flagged events—yields a clear archival signal when contamination is detected but produces a blind spot for audits that later wish to inspect the full overlap distribution. The artifact evidence therefore shows an absence of high-confidence flags but does not let us characterize how close items were to the threshold in aggregate. Practically, audit pipelines should persist per-item detector scores at least as summary quantiles so a null flag set can still be interrogated post hoc; we recommend this change precisely because the current archive lacks the overlap distribution needed to rule out near-threshold cases.

Effect on estimand sensitivity and power. The preregistration explicitly contrasted two plausible execution regimes: (A) independent condition generation (two independent model calls per item, one for baseline and one for guarded) and (B) shared generation (single call reused across conditions when exclusions do not apply). Independent generation increases the number of independent model outputs and therefore raises the probability of observing cross-condition discordance arising from generation-level contamination, paraphrase variance, or randomness. The executed shared-generation regime collapses those independent draws into one and therefore reduces observable discordance unless the detector deterministically excludes an item. The execution manifest shows model_calls_used = 150 versus a planned maximum of 300; this arithmetic difference is the core practical reason the paired estimator had no discordant pairs. The preregistration's power guidance described this tradeoff qualitatively; the artifacts make the tradeoff concrete for this run and explain why the paired test could not detect generation-level leakage under the chosen semantics.

Validator coverage and semantic failure modes. The deterministic_netops_validator_v5 is a practical, deterministic oracle for the operational properties it encodes: parse→state updates, required command counts, dependency rules, and final-state checks. Validator traces archived per episode (validation_before/validation_after) demonstrate that the validator detected and recorded command counts, touched fields, and violation codes for every candidate. However, artifact evidence also makes clear the validator's scope is constrained to the preregistered intent constraints and modeled properties; it does not claim coverage of unmodeled semantic failures (for

example, subtle ordering effects outside the specified workflow_policies, vendor CLI idiosyncrasies not represented in the generator, or emergent cross-task interactions). Consequently, a passed validation trace is necessary evidence of meeting the modeled intent but not sufficient evidence of comprehensive production safety. This distinction is present in the archived validator_trace fields and motivates our recommendations for richer oracle coverage and human adjudication for borderline or high-risk items.

Concrete reproducibility and artifact inspection guidance. The execution bundle contains the artifacts that substantiate the diagnosis: execution/model_configuration.json documents the requested execution semantics (declared_model_version = "gpt-5-mini", independent_condition_generation setting), execution/task_manifest.jsonl records each sampled task with source_identifier and generation_seed, analysis/contamination_summary.csv is empty (no high-confidence flags), analysis/paired_contingency_table.csv records n11 = 150 and zero discordant pairs, and analysis/paired_difference_figure.svg contains the bootstrap distribution used to compute the CI. Inspecting these artifacts in the archive allows an auditor to verify the exact pipeline behavior without recomputing model outputs. We highlight that persistent per-item detector scores were intentionally omitted when flags were absent; auditors who require those scores should request pipelines that persist per-item summaries regardless of flag status.

Design prescriptions to increase audit sensitivity. Artifact-grounded lessons suggest specific, implementable changes: (1) sample provenance-rich, web-sourced items (with stable URLs/DOIs and timestamps) to exercise the fuzzy-overlap and provenance heuristics; (2) set independent_condition_generation = true and execute two independent model calls per item to create the independent comparisons that the paired test is designed to detect; (3) increase model diversity (multi-model replication) to surface model-specific memorization patterns; (4) persist per-item fuzzy-overlap scores and canonicalized fingerprints even when thresholds are not crossed so the overlap distribution is reconstructible; and (5) add paraphrase-robust detectors (semantic similarity beyond n-gram overlap) plus human adjudication for borderline results. Each prescription follows directly from the execution artifacts: the model_calls_used shortfall, the empty contamination_summary, and the recorded shared_initial_candidate semantics.

Operational implications and auditor tradeoffs. Auditors face tradeoffs between reproducibility and sensitivity. Deterministic synthetic sampling and shared generation maximize reproducibility and reduce compute, but those same choices can eliminate the observable signal needed to detect training-data leakage at generation time. Conversely, independent generations and multi-model designs increase sensitivity but require more calls, storage, and non-determinism management. The archived execution demonstrates that reproducibility-focused choices can produce a valid confirmatory test that nevertheless answers a narrower methodological question than originally targeted; explicit preregistration of those tradeoffs (as in our archive) is therefore essential for correct interpretation.

Broader methodological takeaways. Artifact-anchored preregistration remains a powerful practice: freezing detectors, thresholds, replacement rules, and analysis executors before observing outcomes prevents post hoc rule changes. However, the audit shows that preregistration must also commit to sampling and generation semantics that align with the estimand's sensitivity requirements. In short, auditors should preregister not only detectors and thresholds but also the independence structure of model draws and the persistence policy for detector outputs; the execution artifacts here show all three choices materially affected inferential ability.

VI. Limitations

- Single-model execution (gpt-5-mini) — results do not generalize across other LLM families, sizes, or training corpora.
- Sampled items were provenance-stable synthetic/deterministically generated tasks (source_identifier prefix "synthetic-netops:") for this run; absence of high-confidence flags here may not reflect contamination prevalence in real-world, web-crawled benchmarks.
- Generation semantics used shared_initial_candidate across baseline and guarded conditions (independent_condition_generation = false), producing identical model outputs when no exclusions occurred and thus limiting sensitivity to interventions that rely on independent re-generation.
- Contamination detection relied on preregistered conservative heuristics (exact match and fuzzy 5-gram overlap thresholds); subtler contamination patterns (paraphrased memorization, partial leakage) may escape these heuristics and require richer provenance or human adjudication.
- Passing the deterministic_netops_validator_v5 indicates satisfaction of the checked functional constraints but does not imply comprehensive production readiness or absence of semantic regressions outside the validator's scope.

VII. Conclusion

We executed a fully preregistered, artifact-anchored paired audit of conservative contamination detectors for LLM-for-NetOps evaluation. In the reproducible 150-pair synthetic sample we observed zero preregistered high-confidence contamination flags and identical validator pass rates (150/150 baseline, 150/150 guarded). The degenerate null outcome is artifact-supported and stems from provenance-stable synthetic sampling and shared generation semantics; it does not imply absence of contamination in web-sourced benchmarks or failure of the preregistered detectors in different sampling regimes. We provide an explicit preregistration→execution reconciliation, confirm deterministic reconciliation checks passed, and recommend concrete design changes—provenance-rich sampling, independent generation per condition, multi-model

scope, paraphrase-sensitive detectors, and matched power calculations—to increase audit sensitivity in future studies. All experiment artifacts, validator traces, analysis inputs/outputs, and execution manifests are archived and available as recorded in the Disclosure Statement below.

Disclosure Statement

This study was executed under the autonomous-run preregistration contamination-audit-netops-2026-v1 (preregistration SHA-256 recorded in the archived bundle). LLM used: gpt-5-mini (declared_model_version = "gpt-5-mini"). Orchestration adapter: hosted_netops_gvr_v1. Deterministic task generator: deterministic_netops_task_generator_v5 (version 5.0). Deterministic validator: deterministic_netops_validator_v5 (version 5.0). Representative executed artifacts (by path) include execution/model_configuration.json, execution/task_manifest.jsonl, execution/raw_results.jsonl, analysis/paired_contingency_table.csv, and analysis/contamination_summary.csv; the full execution bundle contains raw responses, validator traces, execution logs, and the transformation manifest. Exact immutable initial master-prompt reference: provenance/master_prompt.txt (SHA-256: 1872df1e1805d2d96940456ca016bd665d1d5196 add77f5acdf1582bb39b15ba). Topic constraints (Generative AI and LLMs for NetOps) were selected by the human Research Director prior to the autonomy lock. After the autonomy lock, the AI system autonomously performed literature search, gap selection, preregistration, experiment execution, analysis, manuscript drafting, peer-review simulation, and revision. All generation prompts, model outputs, validator traces, sampling seeds, replacement rules, execution logs, and analysis artifacts were produced and archived by the autonomous pipeline and are included in the execution bundle. No manual human editing of technical content, analyses, references, figures, tables, captions, or the Disclosure Statement occurred after the autonomy lock. Pre-lock development and rehearsal runs occurred (permitted) but pre-lock artifacts were not used as empirical evidence for this final study unless re-created under the locked pipeline. The execution followed preregistered missingness, retry, and exclusion policies; no deviations occurred.

References

[1] https://arxiv.org/abs/2604.22513
[2] https://doi.org/10.48550/arxiv.2606.06212
[3] https://doi.org/10.48550/arxiv.2605.22092
[4] http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.300.8659
[5] https://doi.org/10.1007/s11704-026-60308-3
[6] Buğra Alperen Uluırmak, Rıfat Kurban, "EvalSafetyGap: A Hybrid Survey and Conceptual Framework for LLM Evaluation-Safety Failures", 2026
[7] Nguyen Van Tu, Sukhyun Nam, James Won-Ki Hong, "Intent-Based Network Configuration Using Large Language Models", 2024, 10.1002/nem.2313
[8] Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, Ting Liu, "A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions", 2024, 10.1145/3703155