

Do Generate–Verify–Repair Guards Improve LLM NetOps Outputs? A Paired Emulator Study

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory experiment (n=200 natural-language NetOps intents; study_id cnd-001-5f3a) that compares ungated LLM generation to a bounded generate–verify–repair (GVR) pipeline applied to a single shared LLM candidate per intent. Generation used gpt-5-mini and the guarded arm applied a deterministic three-stage validator and a deterministic, rule-based repair operator with repair_budget ≤ 1 (no extra LLM calls). Artifact-backed primary outcomes: baseline successes = 199/200 (0.995), guarded = 200/200 (1.000); paired contingency n11=199, n10=1, n01=0, n00=0; paired difference (guarded - baseline) = 0.005. The preregistered exact two-sided McNemar (exact binomial on discordant pairs) yields $p = 1.0$; paired-bootstrap (10,000 resamples, seed 1729) 95% percentile CI = [0.0, 0.015]. The preregistered 0.15 absolute-effect target is excluded by the observed CI because the realized baseline was near-ceiling. We provide concise, auditable descriptions of the validator and repair operator, execution accounting (model_calls_used = 201; deterministic repair invocations = 1), exact-test justification appropriate for small discordant counts, and operationally grounded limitations. All prompts, provider traces, validator traces, repair traces (the single invoked repair trace), and analysis artifacts are archived in the study evidence bundle.

I. INTRODUCTION

Automating routine network-operations (NetOps) changes with large language models (LLMs) promises reduced manual effort but raises an operational-safety question: do LLM-produced configuration artifacts execute correctly when applied to control planes? Prior work demonstrates intent-to-configuration prototypes and documents structured-output failures (missing commands, misordering, protected-field modification) and proposes mitigations such as retrieval-augmentation, constrained decoding, verifier-assisted repair, and multi-stage tool-augmentation. Despite these proposals, there is limited execution-grounded evidence that quantifies the marginal benefit of applying a bounded deterministic repair to the same initial LLM candidate (a low-hosted-call operational estimand).

We preregistered and executed a paired experiment (study_id cnd-001-5f3a) that isolates the “repair-on-same-candidate” question by sharing a single initial LLM candidate across both arms (baseline ungated acceptance and guarded generate–verify–repair). The shared-candidate estimand is operationally relevant when hosted-model call budgets are constrained. Our contributions are: (1) a preregistered, artifact-backed paired evaluation under the shared-candidate estimand; (2) concise algorithmic descriptions of the deterministic three-stage validator and the deterministic, rule-based repair operator used in the guarded pipeline; and (3) complete execution

accounting and exact-test justification for small-discordant-pair inference.

II. RELATED WORK

Research on LLMs for NetOps has rapidly diversified along three complementary directions: intent-to-configuration translation and co-pilot prototypes, structured-output mitigation techniques, and formal/deterministic validators that can be integrated into tool-augmented pipelines. Prior prototype and survey work document practical promise—natural-language intents can be translated into device-level commands and assist operators in configuration generation—but also recurrent failure modes such as missing required commands, incorrect ordering, and unintended modification of protected fields (intent-translation prototypes summarized in recent literature) [doi:10.1002/nem.2313; <https://openalex.org/W7161165783>].

Mitigation efforts borrow from nearby domains. Retrieval-augmented generation (RAG), live schema grounding, and constrained decoding have shown material reductions in hallucinated or structurally invalid outputs in NL2SQL and document-grounded tasks; these approaches are natural candidates for NetOps but require careful adaptation to device CLI dialects, workflow policies, and operational constraints [doi:10.2139/ssrn.6909831]. Similarly, constrained or grammar-aware decoding methods and deterministic post-hoc canonicalizers reduce syntactic errors but do not by themselves guarantee correct control-plane behavior when applied to live or emulated networks.

Deterministic verification and formal-methods work (NetKAT variants, weighted NetKAT and ensemble verifiers) provide a complementary assurance axis by checking reachability, isolation, and quantitative properties of candidate configurations; these techniques can produce machine-readable violation codes that are actionable for automated repair or operator review [doi:10.1145/3808318]. However, verification tools face two practical constraints when applied in NetOps pipelines: (1) fidelity to dynamic control-plane phenomena (timing, convergence, vendor-specific idiosyncrasies) and (2) scalability when reasoning over larger multi-device changes. These limits motivate hybrid architectures that combine fast deterministic checks with auditable traces and bounded repair operations.

Security- and safety-oriented evaluations emphasize adversarial risks: human-readable jailbreak prompts and adversarial-context attacks remain an effective threat class for unconstrained generative pipelines and therefore constitute an op-

erational risk for NetOps automation unless mitigations and guards are deployed [doi:10.1145/3815174]. Architectural discussions of agentic NetOps emphasize that production reliability depends on the surrounding assurance machinery—auditable evidence traces, bounded tool use, rollback policies, and independent validators—rather than on raw model capability alone (see recent agentic NetOps position work) [https://openalex.org/W7161165783].

Across these literatures the recurring gap is empirical and execution-grounded: many studies show reduced failure modes in isolated subcomponents (generation, decoding, or verification) but there are relatively few reproducible, artifact-backed measurements of end-to-end generate-validate-repair pipelines under practical constraints such as limited hosted-model calls, deterministic repair budgets, and auditable validator traces. Our study targets this specific gap by measuring the marginal effect of a bounded deterministic repair applied to the same initial LLM candidate under an explicit shared-generation budget; this positions our work as an execution-grounded complement to prior prototype, mitigation, and verification literature [doi:10.1002/nem.2313; doi:10.2139/ssrn.6909831; doi:10.1145/3808318; https://openalex.org/W7161165783; doi:10.1145/3815174].

III. METHODOLOGY

Preregistration and high-level design: We preregistered a paired, shared-initial-generation confirmatory protocol (study_id cnd-001-5f3a). For each of $n = 200$ curated natural-language intents the hosted model produced one initial candidate configuration which was then scored directly as the baseline arm and also passed to the guarded arm. The guarded arm applied a deterministic validator; when the validator reported a repairable violation the pipeline invoked a deterministic, rule-based repair operator ($\text{repair_budget} \leq 1$ for the confirmatory run). The primary estimand is the paired success-rate difference (guarded - baseline); the preregistered primary test is the exact two-sided McNemar implementation (exact-binomial on discordant pairs). Planned maximum hosted-model calls = 400; realized $\text{model_calls_used} = 201$ ($\text{shared_initial_generation} = 200$; deterministic repair-stage calls = 1). The execution environment combined Dockerized Mininet+FRR with `deterministic_netops_validator_v5` as the canonical validator.

Task bank and operational scoring: The confirmatory bank contains 200 intent-expected-configuration pairs created by a deterministic task generator. Each task payload includes `initial_state`, `expected_final_state`, `required_sequence` (ordered per-task commands the authoritatively required), workflow policies (e.g., make-before-break, do-not-touch protected interfaces), and difficulty metadata. The scoring rule (`full_intent_constraint_validation`) requires that generated configurations (a) parse syntactically into canonical CLI-like statements, (b) include the task’s `required_sequence` (order and presence), (c) respect dependency and group-activity policies

(e.g., minimum-active constraints), (d) preserve protected-interface values, and (e) produce an emulator-observed `final_state` equal to `expected_final_state` when applied in the deterministic emulator. Provider/API generation failures are treated per the preregistered missingness plan as arm failures; none occurred in the confirmatory run.

Deterministic validator: The validator deterministically executes three ordered stages and emits a machine-readable trace that the scoring and repair logic consume. Stage 1 – syntactic parsing and canonical normalization: tokenizes CLI-like lines, canonicalizes variant synonyms, and produces `normalized_configuration`; syntax failures yield a `SYNTAX_PARSE_ERROR` and block further deterministic repair. Stage 2 – structural and workflow checks: the validator deterministically verifies presence and ordering of `required_sequence` entries and enforces dependency rules and group policies; it emits deterministic violation codes such as `MISSING_REQUIRED_COMMAND`, `DEPENDENCY_VIOLATION`, `PROTECTED_INTERFACE_MODIFICATION`, and `TRANSIENT_ORDER_VIOLATION`. Stage 3 – deterministic emulator application: the `normalized_configuration` is applied to the emulator to produce `observed_final_state`, and the validator reports `validation_after.valid` based on exact equality with `expected_final_state`. The validator trace fields archived per episode (and used by analysis) include `parsed_command_count`, `observed_touched_field_count`, `normalized_configuration`, `violation_count`, `violation_codes`, `validator_id` and `validator_version`, and `validation_before/after` summaries. These artifacts support auditable, post-hoc inspection of every failure mode.

Repair operator: The preregistered repair operator is a deterministic, rule-based transformer invoked only when the validator reports `violation_count > 0` and the violation family is in a preregistered repairable set. The operator enforces the post-lock constraints: (i) no additional LLM calls (repairs are implemented by deterministic code); (ii) bounded budget (≤ 1 repair attempt per episode in confirmatory runs); and (iii) unambiguous mapping requirement (repairs are applied only when the mapping from violation code to transformation is unique given the task payload). Repair decisions are computed from the task manifest (not inferred by an LLM): the operator consults `required_sequence`, `initial_state`, `expected_final_state`, and `workflow_policies` to deterministically generate corrections. The operator implements three canonical deterministic transformations: - Insertion: when a `MISSING_REQUIRED_COMMAND` violation exists and the missing step appears in the task’s `required_sequence`, the operator inserts the missing command(s) at canonical positions defined by the `required_sequence`. Positioning rules are deterministic: insert before the first command that references the same interface/field group or at the canonical step index derived from the `required_sequence` ordering in the manifest. - Reordering: for `DEPENDENCY_VIOLATION` or `TRANSIENT_ORDER_VIOLATION` the operator performs a stable reordering of contiguous, annotated com-

mand blocks to satisfy dependency constraints while preserving intra-block lexical order. The reordering algorithm deterministically identifies minimal block moves that resolve the dependency according to the manifest’s dependency rules. - Preservation enforcement (redaction/restore): for `PROTECTED_INTERFACE_MODIFICATION` or when `preserve_unspecified_state` would be violated, the operator deterministically replaces the offending command(s) with the original value obtained from `initial_state` (or, if required by `workflow_policies`, with the canonical preserved value from `expected_final_state`).

Ambiguity and nonrepairable cases: If a violation admits multiple equally plausible deterministic repairs (ambiguous insertion indices, multiple candidate values), the operator abstains (no change) and the episode remains a guarded failure; ambiguous families and policy-contradictory violations are flagged in the `repair_trace` and counted in analysis as nonrepairable. All repair decisions and the resulting normalized `configuration` are recorded in a repair trace archived with the run artifacts. In the confirmatory run the deterministic repair operator was invoked exactly once and that single repair converted its episode to a validator-passing configuration.

Statistical analysis: The primary confirmatory test is McNemar’s paired binary test implemented in its exact-binomial (exact McNemar) form appropriate for small discordant counts. Let n_{10} be the number of pairs where guarded = 1 and baseline = 0, and n_{01} where baseline = 1 and guarded = 0; $n = n_{10} + n_{01}$ is the total number of discordant pairs. The exact two-sided p-value is computed from the binomial distribution on n trials assuming $p = 0.5$ (symmetry under the null): the two-sided p equals twice the smaller of the lower and upper tail probabilities for observing $k = n_{10}$ successes under $\text{Binomial}(n, 0.5)$, truncated at 1.0. For our observed counts ($n_{10} = 1$, $n_{01} = 0$, $n = 1$) this exact-binomial two-sided computation yields $p = 1.0$. Complementary uncertainty quantification uses a paired-bootstrap percentile 95% CI for the guarded - baseline proportion difference with 10,000 resamples (`bootstrap_seed = 1729`) as preregistered; these parameters and the bootstrap result are archived in analysis artifacts. The preregistered power analysis assumed a baseline ≈ 0.5 and targeted a 15 percentage-point absolute effect at $n = 200$; the realized near-ceiling baseline substantially reduces the expected discordant counts and consequent sensitivity for that target effect size. All analysis scripts, seeds, resample counts, and the paired contingency CSV are archived in the evidence bundle for reproducibility.

Operational accounting and failure rules: Provider/API transient errors were treated per the preregistered missingness plan (single automatic retry for transient timeouts with 60s backoff); no such retries were consumed in the confirmatory run. Episodes failing syntactic parse are unrepared by design (repair operator requires a valid normalized `configuration`); such parse failures were zero in the confirmatory bank. Episodes where repair abstains are reported as guarded failures and counted in the paired contingency table. The deterministic pipeline’s audit traces (per-episode validator traces and any

repair traces) preserve the exact sequence of decisions for post-hoc inspection and are referenced by analysis outputs.

IV. RESULTS

Execution and primary outcomes: The confirmatory run completed all planned episodes (400 episodes = 200 paired intents) with no provider/API failures. Condition-level artifact-backed counts (analysis/condition_summary.csv) show `baseline_successes = 199/200 (0.995)` and `guarded_successes = 200/200 (1.000)`. The paired contingency (analysis/paired_contingency_table.csv) is $n_{11} = 199$, $n_{10} = 1$, $n_{01} = 0$, $n_{00} = 0$; paired difference (guarded - baseline) = 0.005.

Exact-test and CI details: The preregistered exact two-sided McNemar (exact binomial on discordant pairs) was computed on the observed discordant count $n = n_{10} + n_{01} = 1$; the resulting two-sided p-value reported by the analysis executor is $p = 1.0$ (analysis/results.json). Complementary uncertainty quantification used a paired-bootstrap percentile CI with 10,000 resamples (`bootstrap_seed = 1729`) yielding a 95% CI = [0.0, 0.015]. Because the preregistered primary effect-size target (absolute 0.15) lies well outside the observed CI, the confirmatory data do not support that magnitude of improvement under the shared-candidate estimand.

Repair and failure-accounting diagnostics: The validator reported `validation_before.valid == false` in exactly one confirmatory episode; deterministic repair invocations = 1 and the single deterministic repair converted that episode into a validator-passing final configuration, producing the sole guarded-only success ($n_{10} = 1$). Across the confirmatory bank syntactic/parse failures = 0 and no provider-call timeouts or API errors were recorded. Provider-call accounting (deterministic_reconciliation.json and execution/model_configuration.json) records `hosted_model_call_event_count = 201`, `completed_hosted_model_call_event_count = 201`, `cache_miss_count = 201`, `stage_counts = {shared_initial_generation: 200, repair: 1}`.

Representative exemplars and difficulty context: To illustrate task difficulty and typical intents we publish six representative exemplars with the validator traces and difficulty annotations (the execution/task_manifest.jsonl archive contains the full 200-entry manifest). The six exemplars included in the evidence bundle show difficulty labels: 1 medium (shutdown_configure_restore pattern), 2 hard (make-before-break uplink switchover; paired atomic MTU change), and 3 very_hard (safe-access/service-transfer and rolling migration patterns). These exemplars expose the kinds of dependency and group-activity constraints present in the confirmatory set; the full per-episode task manifest (execution/task_manifest.jsonl) contains complete per-task difficulty metadata for re-analysis.

Artifact-grounded diagnostics: All per-episode scoring JSONL artifacts include `validator_trace` and `scoring_reason` fields (e.g., execution/scoring/task-00000X-*.json). The `paired_contingency_table.csv` and `condition_summary.csv`

(analysis/) provide the exact counts used by the primary analysis. The `analysis_log.jsonl` records bootstrap parameters (resamples = 10,000, seed = 1729) and the exact McNemar computation event. The `deterministic_reconciliation.json` documents end-to-end episode accounting and verifies that `completed_episode_count` = 400, `complete_pair_count` = 200, and that no episodes were excluded for contamination.

Interpretation of the small-discordant outcome: With only one discordant pair ($n = 1$), the exact McNemar test conveys that the observed marginal improvement (one converted episode) is not unlikely under the null in two-sided terms given the small n ; the paired-bootstrap CI conveys the scale of uncertainty around the paired difference (upper bound 0.015). Practically, these diagnostics show that deterministic repair-on-the-same-candidate delivered an auditable, low-cost correction in exactly one case in this curated confirmatory bank; in higher-error regimes or under independent-resampling estimands the repair yield and statistical sensitivity could differ substantially.

Sensitivity notes and archived artifacts for reanalysis: The preregistered missingness plan treated provider/API failures as failures for the arm; because none occurred the primary analysis is not affected by missingness imputation. The preregistered power plan targeted a 0.15 absolute improvement assuming baseline ≈ 0.5 ; the realized near-ceiling baseline compresses discordant counts and reduces power to detect large absolute effects under this estimand. All relevant artifacts for re-analysis are archived (`execution/*`, `analysis/*`, and the full task manifest), enabling investigators to recompute alternative estimands (e.g., independent-resampling) or to inspect the single repaired episode and validator traces for mechanistic insight.

V. DISCUSSION

Operational interpretation: Under the shared-initial-generation estimand and the curated confirmatory distribution the ungated gpt-5-mini generator produced near-perfect candidates (199/200). Deterministic, auditable validators produce structured violation codes that enable targeted deterministic repairs; in low-error regimes the marginal yield of a conservative rule-based repair on the same candidate is small but auditable and low-cost. For operators constrained by model-call budgets, the guarded shared-candidate design preserves low hosted-call cost while providing an assurance gate; for contexts prioritizing maximal success, independent resampling or LLM-invoked repair loops merit exploration.

Design-space trade-offs and practitioner guidance: The confirmatory run’s execution accounting (`model_calls_used` = 201; `repair_invocations` = 1; `completed_episodes` = 400) enables simple cost-versus-yield comparisons. If operators accept higher hosted-call expense, independent-resampling (two or more independent LLM generations per intent) or LLM-guided iterative repair could raise success rates at the cost of greater provider usage and more complex governance. Deterministic validators should be considered mandatory assurance machin-

ery where auditable, repeatable checks and low-latency failure explanations are required.

Threats to validity: Internal validity is strong for the paired estimand and deterministic scoring rules; construct validity depends on emulator fidelity—our deterministic emulator models control-plane state deterministically but cannot fully capture runtime effects such as BGP convergence timing or specific device firmware idiosyncrasies. External validity is limited by the single LLM (gpt-5-mini) and a synthetic curated bank focused on small changes; extrapolation to multivendor CLI diversity or adversarial/out-of-distribution intents is limited. Conclusion validity is constrained by the near-ceiling baseline that reduced discordant-pair counts; the exact McNemar is the correct small-sample test for this observed pattern but offers limited sensitivity when n is small. These limits are documented in the preregistration and archived manifests.

Reproducibility and auditability: All artifacts needed to reproduce analysis and inspection are archived in the evidence bundle: per-episode prompts and responses, provider-call traces, validator traces, deterministic repair trace (the single invoked repair), task manifest entries, analysis scripts, paired contingency CSV, bootstrap parameters (seed = 1729; resamples = 10,000), and execution manifest (preregistration_sha256: e5c5fdbd717c...). The exact-test implementation is the exact-binomial McNemar formulation applied to discordant-pair counts (n_{10}, n_{01}) and is described above; the analysis log records the p-value and bootstrap CI used in the manuscript.

VI. LIMITATIONS

- Synthetic, curated confirmatory task bank focused on small single- and multi-device changes; not comprehensive of production NetOps workloads (multivendor CLI dialects, hardware idiosyncrasies, dynamic control-plane timing).
- Single LLM (gpt-5-mini) evaluated; findings do not imply multi-model generality.
- Deterministic emulator + validator model control-plane behavior but cannot fully capture real-world runtime dynamics such as BGP convergence timing or vendor-specific runtime effects.
- Preregistered repair budget limited to a single deterministic repair iteration; different repair strategies or additional iterations might yield different outcomes.
- Shared-initial-generation design reduces model-call cost and isolates repair-on-same-candidate effectiveness but reduces external generalizability compared with independent-resampling estimands.
- Near-ceiling baseline success limited statistical power to analyze failure-mode distributions in the confirmatory set; exploratory failure-taxonomy conclusions are therefore constrained.
- Adversarial or out-of-distribution intents (e.g., jailbreaks, malformed ambiguous intents) were not included in the confirmatory set and require separate evaluation.

- Full per-task difficulty label counts for all 200 confirmatory intents are recorded in `execution/task_manifest.jsonl` in the archived evidence bundle; the manuscript references that manifest rather than enumerating the full 200-line distribution inline to preserve concise presentation while preserving auditable reproducibility.

VII. CONCLUSION

We executed an autonomy-locked, preregistered paired experiment to measure whether a bounded generate–verify–repair guard improves emulator-executable success for LLM-generated NetOps configurations under a shared-candidate generation budget. Ungated gpt-5-mini generation succeeded in 199/200 confirmatory cases; the deterministic guarded pipeline succeeded in 200/200. The observed paired difference = 0.005 (exact McNemar two-sided $p = 1.0$; paired-bootstrap 95% CI [0.0, 0.015]) does not support the preregistered 0.15 absolute-effect target because the baseline was near-ceiling. For this estimand and task distribution, deterministic repair-on-same-candidate yielded negligible marginal improvement, but deterministic validators remain valuable, auditable assurance gates for operational NetOps pipelines. Full artifacts and analysis code are archived with the study for independent audit and re-analysis.

DISCLOSURE STATEMENT

This study was executed autonomously after an autonomy lock defined by the initial master prompt archived at `provenance/master_prompt.txt` (SHA-256: 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ba).

Human involvement prior to the lock was limited to selecting the topic family (Generative AI and LLMs for NetOps) and approving the master prompt. No manual human editing of technical content, figures, captions, references, results, or manuscript text occurred after the lock. The autonomous pipeline performed literature search and verification, research-gap selection, preregistration (study_id cnd-001-5f3a), code generation, experiment execution, artifact capture, analysis, internal critique, peer-review checks, and manuscript drafting.

Models and tooling: generation requested gpt-5-mini (declared_model_version in `execution/model_configuration.json`). Validation used deterministic_netops_validator_v5 implementing syntactic parsing, structural/workflow checks, and deterministic emulator application. Repairs in the confirmatory run were applied by deterministic code only (no additional LLM calls). The execution environment used Dockerized Mininet+FRF images orchestrated by adapter family hosted_netops_gvr_v1. Preregistered and enforced settings (not altered post-lock) include: primary estimand = paired_success_rate_difference_guarded_minus_baseline; confirmatory sample = $n=200$ paired intents; maximum model-call budget = 400; repair budget ≤ 1 deterministic repair iteration; primary analysis = exact McNemar two-sided (exact-binomial on discordant pairs); bootstrap CIs = 10,000 resamples, seed = 1729. Execution accounting (artifact-backed): the run completed 400 episodes with

model_calls_used = 201 (shared_initial_generation = 200; repair-stage calls = 1). All prompts, provider-call traces, responses, validator traces, repair traces (the single invoked repair trace), analysis artifacts, and the execution manifest are archived in the study evidence bundle.

REFERENCES

- [1] Nguyen Tu, Sukhyun Nam, James Won-Ki Hong, "Intent-Based Network Configuration Using Large Language Models", 2024, 10.1002/nem.2313
- [2] Emmanuel Suárez Acevedo, Tiago Ferreira, Kevin Batz, Oliver Bøving, Nate Foster, Alexandra Silva, "Weighted NetKAT: A Programming Language for Quantitative Network Verification", 2026, 10.1145/3808318
- [3] Sanjay Mishra, Divya Chukkapalli, Ganesh R. Naik, "Schema-Aware Localisation (SAL): Live Schema Grounding and Hallucination Validation for Oracle NL2SQL", 2026, 10.2139/ssrn.6909831
- [4] Muhammad Bilal, Jon Crowcroft, Ruizhi Wang, Xiaolong Xu, Shahram Dustdar, "Large Language Models for Agentic NetOps and AIOps: Architectures, Evaluation, and Safety", 2026, 10.48550/arxiv.2605.12729
- [5] Nilanjana Das, Edward Raff, Aman Chadha, Manas Gaur, "Human-Readable Adversarial Prompts: An Investigation into LLM Vulnerabilities Using Situational Context", 2026, 10.1145/3815174