Verifier-Assisted versus One-Shot LLM Intent→Configuration: A Paired NetOpsTraceBench Study

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

***Abstract*—We preregistered and executed NetOpsTraceBench, a paired comparison of a one-shot LLM intent→configuration pipeline and a verifier-assisted generate–validate–repair pipeline on 1,200 intent→config episodes. Using gpt-5-mini, a deterministic NetOps validator, and Mininet-style emulation, we applied a preregistered binary episode success rule (validator accepts final config AND emulator shows no availability degradation above tolerance AND no automatic rollback). Of 1,200 planned paired tasks we completed 1,199 pairs. Baseline success = 1,196/1,199 (0.9974979149); guarded success = 1,199/1,199 (1.0000000000). The paired-success-rate difference (guarded - baseline) = 0.0025020851 (95% bootstrap percentile CI [0.0000, 0.0058381985], bootstrap_seed=1729, resamples=10,000); exact McNemar two-sided p = 0.25. The preregistered power target sought a 5 percentage-point minimum detectable effect; given the observed near-ceiling baseline, the study lacked power to detect much smaller absolute gains. Representative validator traces and provider audit logs show repair was invoked only three times. Under these controlled emulation and task-corpus conditions we find no statistically significant advantage for the verifier-assisted pipeline; we report artifacts, analytic deviation (Wilcoxon → exact McNemar for binary outcomes), and detailed execution accounting to support reproducible interpretation.**

I. Introduction

Generative large language models (LLMs) are increasingly proposed to translate operator intent expressed in natural language into executable device configurations for network operations (NetOps). Systems and survey literature emphasize that operational reliability depends critically on orchestration, deterministic validators, and assurance infrastructure rather than on an LLM alone [1,2]. Generate–validate–repair (verifier-assisted) architectures are a canonical mitigation strategy: an LLM proposes candidate outputs, an automated verifier checks constraints and safety, and an automated repair loop iterates until the candidate satisfies verifier rules [3–5]. Despite conceptual support, publicly archived, preregistered, paired evidence measuring whether verifier assistance yields an operationally meaningful improvement (rollback-aware binary success) relative to one-shot generation is sparse.

We address this gap with NetOpsTraceBench: a preregistered, paired, emulator-grounded study comparing (A) baseline one-shot NL→config generation and (B) a guarded generate–validate–repair pipeline using the same shared initial candidate. The study asked: Can a reproducible, workflow-centered benchmark of paired intent→config episodes produce an objective paired binary success metric that reliably distinguishes verifier-assisted pipelines from one-shot generation across two simulated topologies? Our contributions are (1) a preregistered experimental design and full execution artifacts (study ID rX4f7-1); (2) a paired analysis on 1,199 completed pairs using gpt-5-mini and a deterministic validator; (3) transparent reporting of an analytic deviation (final confirmatory test used an exact McNemar test appropriate for paired binary outcomes) and execution accounting that explains why the guarded pipeline rarely required repair.

II. Related Work

Recent survey and system literature frames LLMs as components inside broader AIOps/NetOps workflows and repeatedly emphasizes that correctness and trustworthiness are properties of the whole orchestration stack (model + verifier + provenance + sandboxed execution), not of the base LLM alone [1–3]. Multiple community reviews urge workflow-centered evaluation—trace quality, validator interactions, and rollback-aware metrics—rather than isolated one-shot QA tests; these reviews motivate reproducible, audit-ready benchmarks that combine generation, deterministic validation, and emulation [1,2,4].

Concrete verifier-assisted approaches appear across adjacent domains and NetOps prototypes. Work on NL→formal translations and CTL/specification generation demonstrates generator+verifier loops where formal checks drive iterative repairs; similarly, topology-compilation systems translate high-level security intent into executable topologies with compliance checks and constraint enforcement [5,6]. Prototype efforts in the configuration space (text→config generators and misconfiguration detectors) and domain-specific "LLM firewalls" illustrate practical patterns: an LLM proposes a candidate, an automated validator enforces syntactic and semantic constraints, and an orchestrator optionally triggers constrained repair or rejects the proposal [7–10]. These systems commonly instrument deterministic parsers and domain rules to produce binary verifier outcomes (accept/reject) and produce their own evaluation sets for demonstration; as a result, reported gains are often conditional on task selection and validator strictness.

Empirical and conceptual studies document LLM failure modes that motivate guarded architectures. Recent analyses identify hallucination, unit/semantic errors, overconfidence/miscalibration, provenance loss, and susceptibility to prompt-injection as recurrent risks when LLMs propose executable artifacts—risks that an external verifier or prove-

nance linker can detect or mitigate in some instances but not eliminate universally [11–14]. Several papers demonstrate that verifier-guided repair reduces syntactic and some semantic errors in constrained settings, but they also highlight sensitivity to the verifier's rule set (strictness) and to the representativeness of evaluation examples. In short, prior work shows feasibility of generate–validate-repair patterns but also underscores the dependence of measured improvement on corpus difficulty, adversarial content, and verifier configuration.

Related benchmarking and dataset discussions emphasize gaps that NetOpsTraceBench targets. System papers often construct private or narrowly scoped held-out sets calibrated to their prototype; TopoIntent and other recent tools note the absence of large public, vendor-diverse intent→config datasets with execution ground truth and rollback semantics, which complicates cross-system comparisons [6,7,9]. Surveys therefore call for standardized, workflow-centered benchmarks that (a) record full provider and validator traces, (b) embed emulator execution to check operational side effects (availability/rollback), and (c) preregister analysis choices to reduce researcher degrees of freedom [1,2]. NetOpsTraceBench operationalizes these recommendations by preregistering inclusion/exclusion rules, seeds, a binary, rollback-aware episode success metric, and by archiving raw model outputs, validator traces, and execution manifests so that the effect of verifier design and task corpus can be examined reproducibly.

Taken together, the literature supports the conceptual value of verifier scaffolding while also documenting two recurring caveats that our study made explicit in design and interpretation: (i) measured gains depend strongly on task difficulty and verifier thresholds, and (ii) many demonstrations use constructed datasets and domain-specific validators, which limits generality. NetOpsTraceBench contributes a preregistered, artifact-rich paired evaluation that makes these dependencies observable (repair invocation counts, contingency discordances, and emulator rollback checks) rather than inferred, enabling follow-up sensitivity studies that vary corpus adversariality and verifier strictness.

III. Methodology

Preregistration and experimental scope: The study design was preregistered under ID rX4f7-1 and frozen prior to execution (preregistration manifest SHA-256 = 77abe8b4...). The preregistration specified confirmatory hypotheses, a single primary estimand (paired_success_rate_difference_guarded_minus_baseline), deterministic inclusion/exclusion rules, contamination thresholds, stopping rules, random seeds, maximum model-call budget, and a complete analysis plan. NetOpsTraceBench_v1 targeted 1,200 paired intent→config tasks partitioned into two simulated topology clusters (edge = 600, core = 600). Task examples were public synthetic or consented; every task carried deterministic provenance metadata. The preregistered power plan assumed a baseline success near 0.70 and set a minimum detectable effect (MDE) of 5 percentage points (absolute) at two-sided α=0.05 and power 0.8; sensitivity scenarios over a range of baseline and discordant-pair assumptions were specified in the preregistration.

Pipelines, transformations, and instrumented components: We compared two paired conditions executed from a single shared initial candidate per pair. Baseline (one-shot) invoked a single LLM generation call to produce an initial candidate configuration. Guarded (verifier-assisted) used the same shared initial candidate and permitted at most one automated repair iteration when the deterministic validator rejected the candidate; preregistration limited maximum_repair_calls_per_task = 1 and initial_generation_calls_per_task = 1 to bound model-call budget. The requested model was gpt-5-mini (provider recorded as openai_responses) invoked through orchestration adapter hosted_netops_gvr_v1. Deterministic_netops_task_generator_v5 produced canonical, vendor-agnostic intent→configuration tasks; deterministic_netops_validator_v5 performed parsing, intent-constraint validation, dependency checks, and emulator preflight evaluation. The execution contract capped maximum_model_calls = 2,400 and defined the scientific_confirmatory execution mode.

Inclusion, contamination, and stopping rules: Inclusion required that neither episode in a pair be deterministically excluded. The preregistration implemented contamination detection by (a) a normalized substring/token-overlap test with threshold T_contam = 0.85 and (b) an optional embedding-similarity secondary check (cosine threshold = 0.92) for flagged items. If either test exceeded thresholds the pair was flagged and excluded from the primary confirmatory set (but retained for exploratory reporting). Pairs were also excluded if both episodes failed to produce parsable vendor renderings after one automated re-execution attempt. The stopping rule required running the full planned workload unless (i) cumulative model calls reached the maximum_model_calls cap, or (ii) automation halted due to deterministic component failure after the one allowed automated recovery attempt; all such events were logged.

Determinism, seeds, and failure handling: The preregistration froze generation seeds and analysis seeds. Deterministic task generation and validator behavior reduce sampling variance from non-deterministic test artifacts; model nondeterminism remained possible at the provider side but was bounded by the single-call-per-stage policy. Missingness handling followed the preregistered plan: the primary analysis used complete_pair_primary (exclude flagged/excluded pairs); a sensitivity analysis treated excluded or failed episodes as failures (complete_pair_primary_with_failure_as_zero_sensitivity). All exclusion reasons, failed episode counts, and provider audit traces were archived to preserve transparent accounting.

Success metric, confirmatory test, and bootstrap CI: The preregistered binary episode success rule required three conjunctive outcomes: (1) final configuration accepted by the deterministic validator (parsing and intent constraints satisfied), (2) emulator execution shows availability degradation

no greater than the preregistered tolerance, and (3) no automatic rollback event occurred during simulated execution. Although the preregistration specified a paired nonparametric Wilcoxon signed-rank test, the archived outcomes were binary for each episode. Following the preregistered emphasis on correct paired-outcome inference, the final archived confirmatory analysis applied an exact McNemar two-sided test to the 2×2 paired contingency table and computed a paired bootstrap percentile 95% confidence interval for the paired-success-rate difference using frozen bootstrap_seed = 1729 and 10,000 resamples. This analytic deviation (Wilcoxon $\rightarrow$ exact McNemar) is recorded in the analysis artifacts and in the Disclosure Statement.

Instrumentation, auditing, and artifact recording: The run captured provider call audit metadata (hosted_model_call_event_count and stage_counts), all raw model outputs, shared initial candidate traces, repair provider traces when present, validator traces (validation_before/validation_after), emulator execution logs, task_manifest entries, transformation_manifest, and a full hash manifest. Key configured limits recorded in model_configuration.json included max_output_tokens = 1500 and requested_model = gpt-5-mini. The execution manifest records planned_episode_count = 2,400, completed_episode_count = 2,398, failed_episode_count = 2, and model_calls_used = 1,203. All seeds, analysis code (paired_binary_analysis_v1), and final analysis artifacts (analysis/paired_contingency_table.csv, analysis/results.json, analysis/paired_difference_figure.svg) are archived to enable exact reproduction.

Predefined sensitivity and exploratory plans: The preregistration enumerated exploratory analyses that the archived run executed or preserved for downstream work: (a) distribution of repair iterations and their correlation with final success and time-to-repair; (b) automated failure-mode categorization (semantic unit errors, missing dependencies) derived from deterministic parsers and emulator checks; (c) sensitivity of the primary binary success to alternative availability-degradation tolerances and to treating flagged contamination pairs as failures. All exploratory outputs were labeled as such and are reported separately from the confirmatory inference to control Type I error per the preregistered multiplicity plan.

IV. Results

Execution and pair accounting: The run completed 2,398 of 2,400 planned episodes and produced 1,199 complete paired observations (completed_pairs = 1,199; failed_episode_count = 2). The analysis used complete_pair_count = 1,199. Provider audit and orchestration logs record hosted_model_call_event_count = 1,203 (completed_hosted_model_call_event_count = 1,202; failed_hosted_model_call_event_count = 1) and stage_counts indicating shared_initial_generation = 1,200 and repair = 3. These counts show that a single shared initial generation was produced per task and that the guarded pipeline invoked repair calls only three times across the entire corpus.

Primary confirmatory outcomes (paired contingency and rates): The paired 2×2 contingency table (analysis/paired_contingency_table.csv) is: n11 = 1,196 (both methods succeeded), n10 = 3 (guarded succeeded while baseline failed), n01 = 0 (baseline succeeded while guarded failed), n00 = 0 (both failed). From these counts: baseline_success_count = 1,196 (baseline_success_rate = 1,196/1,199 = 0.9974979149291076) and guarded_success_count = 1,199 (guarded_success_rate = 1,199/1,199 = 1.0). The observed paired-success-rate difference (guarded - baseline) = 0.002502085070892446.

Exact paired inference and bootstrap CI: Because confirmatory outcomes are binary and the observed data are concentrated in concordant cells, we used an exact McNemar two-sided test on the discordant pair counts (n10 = 3, n01 = 0). The archived exact McNemar two-sided p = 0.25. The paired bootstrap percentile 95% CI for the paired-success-rate difference (computed with bootstrap_seed = 1729 and 10,000 resamples as recorded in analysis/analysis_log.jsonl) is [0.0000, 0.005838198498748957]. The CI lower bound reaching 0.0 reflects the high proportion of concordant successful pairs and the small number of discordant pairs; those data characteristics limit resolution for small absolute effects.

Missingness and failed calls: The preregistered missingness summary (analysis/missingness_summary.csv) reports complete_pairs = 1,199; failed_baseline_episodes = 1; failed_guarded_episodes = 1; both_missing_pairs = 1. The deterministic exclusion rules were not triggered for contamination; analysis/contamination_summary.csv is empty. The single paired exclusion implied by both_missing_pairs = 1 is the accounting source for the one planned pair drop from 1,200 to 1,199 complete pairs.

Repair invocation diagnostics: Stage_counts (execution manifest and deterministic_reconciliation.json) show repair = 3 while shared_initial_generation = 1,200. The repair count matches repair_provider traces present for exactly three episodes in the execution/provider_calls and execution/scoring artifacts. Scorer and validator traces (examples in execution/scoring/*.json) report validation_after.valid = true for the majority of episodes and repair_applied = false, indicating that initial candidates typically satisfied deterministic_netops_validator_v5 checks and therefore required no automated repair.

Representative artifact evidence: Representative task manifest entries and scoring artifacts illustrate the evidence pattern. For example, execution/scoring/task-000001-baseline.json and execution/scoring/task-000001-guarded.json show identical shared_initial_candidate content and validator traces with validation_after.valid = true and repair_applied = false; difficulty for that task is labeled "medium" in the task_manifest entry. Examples with higher difficulty labels ("hard", "very_hard") appear in representative task entries (task-000002 level "hard", task-000003 level "hard") and still produced accepted shared initial candidates in many cases. These artifact-level diagnostics demonstrate (a) heterogeneous declared task difficulty across the corpus, and (b) that under the chosen deterministic

task generator and validator settings the initial LLM output frequently met intent constraints.

Interpretation of discordant pattern: The observed discordant pattern (n10 = 3, n01 = 0) means all discordant pairs favored the guarded pipeline; however, the absolute discordant count (3/1,199 $\approx$ 0.25%) is too small to yield statistical significance under the prespecified inferential thresholds. Under the exact McNemar framework the two-sided p = 0.25 (archived), reflecting the binomial probability of observing such an imbalanced discordant split given the small number of discordant pairs. The paired bootstrap CI's narrow upper bound ($\approx$0.00584) but lower bound at zero quantifies that any true absolute benefit, if present, is small in magnitude within this corpus and validator configuration.

Reproducibility anchors: All numbers above match archived artifacts: analysis/paired_contingency_table.csv (n11/n10/n01/n00), analysis/condition_summary.csv (per-condition counts and success rates), analysis/results.json (estimate, p-value, CI, bootstrap parameters), execution/execution_manifest.json (provider audit counts and timestamps), and representative scoring artifacts (execution/scoring/*.json) that record validator trace fields validation_before/validation_after, repair_applied flags, and normalized_configuration values. The run's bootstrap parameters (seed = 1729, resamples = 10,000) and analysis log are included in analysis/analysis_log.jsonl to enable exact reproduction of the CI reported above.

V. Discussion

The statistical evidence and operational diagnostics together explain why the preregistered paired comparison did not find a statistically significant advantage for the guarded pipeline in this run. The exact McNemar two-sided test (p = 0.25) and the paired bootstrap percentile 95% CI [0.0000, 0.0058381985] both include zero, indicating that the observed paired success-rate difference ($\approx$0.25 percentage points) is consistent with sampling variability under the recorded discordant counts (n10 = 3, n01 = 0). Those discordant counts are the proximal determinant of inference here: with 1,199 complete pairs, only three pairs were discordant and all favored the guarded pipeline, which yields a point estimate but insufficient evidence to reject the null at α=0.05. This arithmetic — small discordant counts in an otherwise near-ceiling outcome distribution — is why the preregistered MDE of 5 percentage points (chosen under an assumed baseline $\approx$0.70) was not matched by the observed data; the study therefore lacked confirmatory power to detect such a small absolute improvement.

Operationally, archived provider and validator artifacts clarify the mechanism that produced these statistics. The provider audit reports model_calls_used = 1,203 with stage_counts showing shared_initial_generation = 1,200 and repair = 3, and representative validator traces (e.g., execution/scoring/task-000001-baseline.json) show validation_after.valid = true for typical episodes. Together these elements indicate that the shared initial candidates satisfied the deterministic_netops_validator_v5 checks in the overwhelming majority of episodes and so the guarded pipeline's repair stage was rarely invoked. Thus, under the exact combination of (a) a synthetic, well-specified task generator that produces tasks with explicit required sequences, and (b) a validator configuration that enforces those sequences in a way that treats correct initial candidates as acceptable, the marginal operational value of generate–validate–repair for the binary success metric is necessarily small.

These findings do not falsify the general utility of verifier-assisted architectures. Instead they highlight two contextual dependencies that practitioners and benchmark designers must consider. First, baseline difficulty and task ambiguity directly determine headroom for improvement: when initial generation already meets deterministic verifier criteria, measured gains on a strict binary success metric will be limited. Second, verifier strictness and the mapping between validator checks and operational safety matter: different validator thresholds or richer emulation of transient behaviors (e.g., routing convergence, vendor-specific command ordering) would change repair invocation rates and the distribution of discordant pairs. Both dependencies are visible in the archived run artifacts (task_manifest entries and scoring traces) and can be explored by varying task ambiguity, contamination thresholds, or validator rules using the provided analysis code and seeds.

Practical implications follow for benchmark design and operational evaluation. To measure verifier benefit robustly, a benchmark should (i) include a larger proportion of ambiguous or adversarial intents that create meaningful syntactic/semantic failure modes for LLMs, (ii) vary verifier strictness (from permissive to conservative) to expose where repairs would change execution outcomes, and (iii) exercise multiple model families and rendering targets to assess generality across vendor syntaxes. The preregistered sensitivity scenarios and the archived artifacts make such follow-up experiments reproducible: the run archives deterministic seeds, task generator configuration, validator traces, and analysis scripts so other researchers can replay alternative verifier sensitivities or inject adversarial perturbations without ambiguity.

Cost–benefit considerations also emerge from the execution accounting. Repair was invoked only three times, and the total model call audit ($\approx$1,203 calls) shows a small additional invocation footprint relative to the number of pairs. In settings where repairs are rare because validators accept correct initial candidates, the operational overhead of a guarded pipeline may be modest; however, this mechanical accounting does not capture human-in-the-loop costs, potential latency requirements, or the upstream effort to build and maintain robust validators and emulators. Those broader operational costs and benefits remain important practical considerations beyond the binary success metric evaluated here.

Finally, the run exposes reproducible levers for future investigation. The archived CI, contingency table, and provider logs document both the null-result boundary conditions and the exact artifacts required to test alternative hypotheses (e.g., higher task ambiguity, stricter validators, adversarial stress

generators, multi-model comparisons). Because the preregistration and execution artifacts are public within this evidence bundle, researchers can systematically vary one factor at a time and quantify how discordant counts and the paired-difference distribution change — an approach that will be necessary to move from context-specific null findings toward generalizable operational guidance about when verifier scaffolding materially improves NetOps safety.

VI. Limitations

- Single model family tested (gpt-5-mini); results do not imply generality across other LLMs or fine-tuned variants.
- Task corpus skew: synthetic/consented tasks yielded near-ceiling baseline success (≈99.75%), limiting headroom and statistical power to detect practical improvements by guarded pipelines.
- Emulator fidelity: Mininet-style emulation and deterministic validators approximate production behavior but may not capture all deployment failure modes (routing convergence, vendor idiosyncrasies, transient telemetry).
- Verifier configuration dependence: deterministic_netops_validator_v5 acceptance thresholds and parsing rules materially affect outcomes; different verifier strictness could change repair invocation rates and measured gains.
- Analytic deviation documented: the preregistration specified a Wilcoxon signed-rank paired procedure; the archived confirmatory analysis used an exact McNemar test because outcomes were binary. This deviation is recorded in the archived analysis artifacts and in the Disclosure Statement above.
- Operational external validity: absence of heterogeneous vendor renderers, real production templates, and live rollouts constrain direct production generalization.

VII. Conclusion

We preregistered and executed NetOpsTraceBench, a paired, emulator-grounded study comparing one-shot and verifier-assisted NL→configuration pipelines on 1,199 completed pairs. Under our synthetic task corpus, deterministic validator settings, and gpt-5-mini as requested, baseline one-shot generation already achieved near-ceiling success (≈99.75%). The guarded pipeline increased absolute success by ≈0.25 percentage points (paired difference = 0.0025020851) with exact McNemar two-sided p = 0.25 and a 95% bootstrap CI that includes zero. Archived execution manifests, provider audits (model_calls_used = 1,203; repairs invoked = 3), representative validator traces, and analysis artifacts support the interpretation that the observed null result reflects limited headroom given task and verifier choices rather than definitive evidence that verifiers are never useful. Future work should intentionally design more challenging and adversarial task sets, vary verifier strictness, and test multiple LLM families to determine conditions under which verifier scaffolding yields operationally meaningful improvements. All run artifacts, analysis code, seeds, and logs are archived to enable exact reproduction and follow-up sensitivity studies (see Disclosure Statement).

Disclosure Statement

Mandatory disclosure (preregistered run rX4f7-1). LLM(s) and provider: requested model = gpt-5-mini via provider recorded as openai_responses; model configuration archived at execution/model_configuration.json (requested_model=gpt-5-mini, max_output_tokens=1500). Orchestration/adapter: hosted_netops_gvr_v1 executed the paired benchmark. Deterministic tooling: deterministic_netops_validator_v5 performed canonical parsing, intent-constraint validation, and emulator preflight checks; deterministic_netops_task_generator_v5 produced synthetic tasks. Exact initial master prompt: archived at provenance/master_prompt.txt (immutable reference SHA-256 = 1872df1e1805d2d96940456ca016bd665d1d5196add77f5a cdf1582bb39b15ba). Topic constraint and autonomy boundary: the human Research Director selected the topic family and approved the master prompt pre-lock; after the autonomy lock the AI autonomously formulated the research question, conducted literature review, generated the preregistration, executed the experiment, performed analysis, and wrote and revised the manuscript. Preregistration: NetOpsTraceBench (ID rX4f7-1) preregistration manifest SHA-256 = 77abe8b4 93bde44484e0b3d02cfd7f9ec1621ff7e7847889082ae177 a7ed5712; preregistered design (confirmatory hypotheses, inclusion/exclusion, seeds, stopping rules, analysis plan) was frozen before execution. Execution summary: started 2026-08-30T01:21:55.415955+00:00, completed 2026-08-30T03:15:12.141890+00:00; planned episodes 2,400; completed episodes 2,398; completed pairs 1,199; model_calls_used = 1,203; repair-stage calls observed = 3. Analysis artifacts and deviation record: the preregistration specified a Wilcoxon signed-rank paired procedure; the archived executed confirmatory analysis used an exact McNemar two-sided test on paired binary outcomes plus a paired bootstrap percentile CI. This post-preregistration analytic choice is documented in the archived analysis artifacts (analysis/paired_contingency_table.csv, analysis/results.json, analysis/analysis_log.jsonl, analysis/paired_difference_figure.svg) produced as part of the run that completed at 2026-08-30T03:15:12.141890+00:00. Artifacts and reproducibility: raw model outputs, provider traces, validator traces, task_manifest, transformation_manifest, execution logs, analysis code, and all seeds are archived; artifact_count = 8,408 and full hash manifest path = execution/execution_manifest.json. Human involvement: pre-lock collaboration and infrastructure development occurred; no human manual editing, selective revision, or ad hoc text insertion occurred on the manuscript content after the autonomy lock—the AI performed internal autonomous revision cycles only. Technical restarts and deviations: execution manifest records two failed episodes and the analytic deviation described above; all deviations are

logged in execution/execution_log.jsonl. The disclosure above is complete relative to the archived artifacts supplied with this run.

REFERENCES

[1] http://hdl.handle.net/20.500.12708/229494
[2] https://arxiv.org/abs/2608.13389
[3] http://arxiv.org/abs/2403.09369
[4] https://doi.org/10.2139/ssrn.6947162
[5] https://doi.org/10.2139/ssrn.7222278
[6] https://doi.org/10.20935/acadai8260
[7] https://doi.org/10.3390/info17030297
[8] https://doi.org/10.3390/app16010085