

Does Deterministic Generate–Validate–Repair Reduce Unsafe LLM-Generated Network Changes? A Confirmatory Paired Study

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a sealed paired confirmatory study ($n = 200$ intent→configuration tasks) to test whether a deterministic generate–validate–repair (GVR) pipeline that consults a deterministic NetOps validator reduces validator-detected unsafe or semantically incorrect network changes relative to an LLM-only generation pipeline. The run used a single hosted model (gpt-5-mini) with adapter-enforced shared-initial generation semantics and a primary deterministic validator (deterministic_netops_validator_v5). Execution completed all planned pairs (completed_pairs = 200) consuming 200 model calls (one shared initial generation per task); no repair calls were issued because every initial candidate passed the validator. Paired contingency: $n_{11} = 200$, $n_{10} = 0$, $n_{01} = 0$, $n_{00} = 0$; estimate = 0.0; exact McNemar two-sided $p = 1.0$; 95% bootstrap percentile CI = [0.0, 0.0] (10,000 resamples, seed = 1729). Because the baseline validator-detected unsafe incidence was 0% on this task bank, the preregistered $\geq 50\%$ relative-reduction hypothesis could not be meaningfully evaluated. We report the sealed design, execution accounting (start: 2026-08-30T11:50:34.530059Z; end: 2026-08-30T12:12:07.541650Z), the Mininet 10% hold-out (20 tasks) that produced zero validator–emulator disagreements, and an artifact-grounded discussion of operational implications and threats to validity (preregistration id: cnsms2026-gvr-effectiveness-02).

I. INTRODUCTION

Automating intent→configuration translation with generative language models can increase NetOps productivity but risks producing incorrect or unsafe changes that cause outages or policy violations. Prior work therefore proposes grounding, deterministic validation, and repair loops (generate–validate–repair, GVR) to detect and correct unsafe candidates before application. The operational benefit of automated repair is conditional on three factors: (i) baseline prevalence of validator-detectable failures from the chosen generator/prompting policy; (ii) validator coverage and fidelity; and (iii) the available repair budget and repair-prompt effectiveness. We preregistered a narrowly scoped confirmatory question: Does an automated, deterministic GVR pipeline that consults a deterministic NetOps validator materially reduce validator-detected unsafe or semantically incorrect network changes relative to an LLM-only generation pipeline on a curated 200-task intent→configuration bank? The study was designed as a paired experiment that fixes the generator output per task and isolates the marginal effect of invoking a validator+repair on the same candidate.

II. RELATED WORK

Recent literature documents three relevant threads for LLM-driven NetOps: (1) empirical characterizations of LLM factuality failures and mitigation patterns, (2) NetOps-specific prototypes that translate intent into configurations, and (3) system and threat analyses that focus on verification, guarded architectures, and attack surfaces in agentic tool use. First, cross-domain studies and surveys show that large language models produce factual errors, omissions, and hallucinations in operational tasks and that grounding, verifier-assisted repair, and uncertainty estimators can reduce but not eliminate such errors (e.g., Farquhar et al.; broad LLM surveys) [<https://openalex.org/W4399803256>, 10.1007/s11704-026-60308-3]. These works establish that mitigation effectiveness depends strongly on prompt design, sampling policy, retrieval/grounding fidelity, and task framing. Second, a growing set of NetOps prototypes demonstrates the feasibility of intent→configuration generation but also highlights evaluation gaps: NetConfEval and several intent-driven configuration works show that LLMs can synthesize network commands at prototype scale, yet reported evaluations often lack production-like instrumentation, end-to-end emulation, or preregistered inference procedures for failure-rate estimates [<https://openalex.org/W4399629579>, <https://openalex.org/W4403635865>, <https://openalex.org/W4400072049>]. Third, architectural proposals and empirical audits advocate combining generation with deterministic validators or simulators (generate–validate–repair, Batfish-style checks, and Mininet emulation) and propose bounded-autonomy guardrails to constrain unsafe agentic actions; at the same time, red-team and configuration-brittleness studies document that tool descriptions, chaining ambiguity, and persona/history can materially change generated actions and enable exploits if guardrails are incomplete [<https://openalex.org/W4410125989>, 10.36227/techrxiv.177004277.71795055/v1, 10.2139/ssrn.7203631, 10.21203/rs.3.rs-10646209/v1]. Complementary work on prompt sensitivity and LLM non-determinism shows that evaluation outcomes can shift with temperature, sampling, system prompts, or chain-of-thought strategies, which complicates comparisons across pipelines and motivates paired or shared-initial designs for controlled measurement

[https://openalex.org/W4402860127]. Finally, proposals for secure-by-default agent patterns and bounded-autonomy frameworks emphasize runtime veto layers, explicit tool registries, and staged validation to limit unsafe changes, but these remain largely proof-of-concepts whose real-world efficacy depends on validator coverage and emulation fidelity [10.36227/techrxiv.177006109.97957916/v1, 10.36227/techrxiv.177004277.71795055/v1]. Synthesizing these threads reveals a persistent evaluation gap: few studies provide preregistered, externally auditable measurements that isolate the marginal effect of deterministic repair on a fixed generator output under controlled sampling semantics. Our preregistered paired design addresses precisely this gap by enforcing shared-initial generation and deterministic validator checks, thereby separating generator correctness under a fixed prompt policy from any additional gains attributable to automated repair.

III. METHODOLOGY

Preregistration and inferential plan. The study preregistration (cnsms2026-gvr-effectiveness-02) fixed a paired design with $n = 200$ tasks, one shared deterministic initial generation per task, and a repair budget of at most one automated repair call per task. The primary estimand was the `paired_success_rate_difference_guarded_minus_baseline`; the primary test was an exact two-sided McNemar test ($\alpha = 0.05$). Confidence intervals were computed via bootstrap percentile (10,000 resamples, `bootstrap_seed = 1729`). Power planning assumed a baseline unsafe incidence $\approx 30\%$ and a targeted 50% relative reduction; these assumptions were recorded in the preregistration and used only to justify sample size. The analysis executor was `paired_binary_analysis_v1`.

Task bank, validators, and holdout. We used a curated, machine-checkable intent→configuration task bank with 200 tasks across four topology clusters (enterprise, campus, datacenter, edge; 50 tasks each). Each task encodes an initial and expected final state, workflow constraints (make-before-break, `must_be_down_for_fields`, group-active minima), and a `required_sequence`. The primary deterministic checker (`deterministic_netops_validator_v5`) enforces syntactic and intent/constraint checks and emits structured diagnostics (valid flag, violation codes, `parsed_command_count`, `required_command_count`). A Mininet-based data-plane emulator was preregistered as a secondary disagreement detector applied to a randomized 10% holdout (20 tasks) to detect validator–emulator mismatches; execution produced an `analysis/contamination_summary.csv` artifact recording zero disagreement flags for that holdout.

Adapter semantics, shared-initial design, and causal isolation. The `hosted_netops_gvr_v1` adapter enforced `shared_initial` generation semantics: for each task the system produced a single deterministic initial candidate that was reused for both baseline (LLM-only) and guarded (GVR) conditions. This design choice intentionally removes sampling variance across conditions so the paired comparison isolates the marginal effect of invoking the validator and (if needed)

issuing a single repair prompt on the identical generator output. Adapter accounting (deterministic_reconciliation artifacts) records `hosted_model_call_event_count = 200` and `stage_counts = {shared_initial_generation: 200}`, confirming single shared invocations per task. By construction, any benefit attributable to repeated or stochastic re-sampling is excluded from the primary estimand; this is a preplanned trade-off to obtain a controlled causal contrast between (A) applying the generator output unchecked and (B) validating—and possibly repairing—that same output.

Repair policy, budgeting, and handling. The preregistered repair policy allowed at most one automated repair prompt to the same LLM when the primary validator reported violations. Repair prompts (when issued) and any repaired candidate were re-validated by `deterministic_netops_validator_v5` and recorded in `validator_trace` fields (`repair_applied`, `repair_prompt_sha256`, `final_configuration`). The execution contract limited maximum model calls to 400 (200 initial + up to 200 repairs); the adapter enforced this cap. Missingness/exclusion rules deterministically excluded pairs where both conditions failed due to infrastructure execution failures; tasks flagged for validator–emulator disagreement were excluded from the primary analysis per the `contamination_plan`, with deterministic sensitivity analyses preregistered for alternative treatments.

Randomization, seeds, and reproducibility artifacts. Task ordering was randomized with a fixed generation seed as recorded in the `task_manifest` entries (example fields: `generation_seed`). Analysis used a paired, shared-initial seed schedule and the bootstrap CI procedure used `bootstrap_seed = 1729` with 10,000 resamples (analysis/analysis_log.jsonl captures these parameters). The Mininet 10% holdout executed as preregistered and is summarized in `analysis/contamination_summary.csv` (20 tasks; zero disagreement flags). We disclose a documented reproducibility gap: the explicit labeled randomized selection seed identifying which tasks comprised the 10% Mininet holdout is not present in the labeled artifacts, preventing exact deterministic reconstruction of the drawn holdout despite the contamination summary reporting zero disagreements.

Statistical implementation and exact tests. The primary confirmatory test was implemented as an exact McNemar two-sided test on the 2×2 paired table (n_{11} , n_{10} , n_{01} , n_{00}) produced by `paired_binary_analysis_v1`. Where the discordant counts were zero or small, exact inference is necessary; bootstrapped paired-difference percentile CIs were computed to reflect the empirical paired distribution under resampling (10,000 resamples; seed recorded). Secondary/exploratory analyses (preplanned) included logistic regressions (`repair_success ~ topology_cluster + initial_validator_failure_type + repair_iterations`) and descriptive summaries of repair iteration counts and model-call costs; these remained exploratory and were executed only when applicable data were generated.

Contamination detection and conservative primary analysis rule. The preregistered contamination plan used the Mininet

emulator as an independent end-to-end check for a randomized holdout. Any pair where Batfish-style validator verdict $\neq$ Mininet verdict would be flagged and excluded from the primary confirmatory analysis to conservatively avoid reliance on a possibly incomplete primary validator. The analysis artifact `analysis/contamination_summary.csv` shows zero flags for the 20-task holdout; `deterministic_reconciliation.json` confirms `pair_accounting_consistent = true`. Because no contamination flags occurred, the preplanned conservative exclusion branch was not activated.

Execution accounting and archived artifacts. The run executed autonomously under the frozen plan (start: 2026-08-30T11:50:34.530059Z; end: 2026-08-30T12:12:07.541650Z). Model calls used = 200 (one shared initial generation per task); repair calls used = 0. Core archived artifacts relevant to reproducing and auditing the methodology include `execution/raw_results.jsonl`, `execution/model_configuration.json`, `execution/task_manifest.jsonl`, `analysis/paired_contingency_table.csv`, `analysis/condition_summary.csv`, `analysis/contamination_summary.csv`, `analysis/analysis_log.jsonl`, and `deterministic_reconciliation.json`. These artifacts capture adapter accounting, validator traces (validation_before/after objects with `parsed_command_count`, `required_command_count`, `violation_codes`), and the bootstrap parameters. All handling decisions (exclusions, missingness treatment) were deterministic and logged in `analysis/analysis_log.jsonl`.

IV. RESULTS

Execution outcomes. The run completed all planned episodes: `completed_episode_count` = 400 (shared initial generation reused across conditions) and `complete_pair_count` = 200. Adapter accounting shows `model_calls_used` = 200 (one shared initial generation per task) and `cache_miss_count` = 200. No repair calls were issued because every initial candidate passed `deterministic_netops_validator_v5`. Condition summaries (`analysis/condition_summary.csv`) show baseline successes = 200/200 and guarded successes = 200/200 (`success_rate` = 1.0 for both). The paired contingency (`analysis/paired_contingency_table.csv`) is `n11` = 200, `n10` = 0, `n01` = 0, `n00` = 0.

Primary inferential analysis. With zero discordant pairs the observed paired difference estimate = 0.0. The exact McNemar two-sided p = 1.0. The bootstrap percentile 95% CI (10,000 resamples; seed = 1729) is [0.0, 0.0]. Under the preregistered estimand and test the prespecified $\geq 50\%$ relative-reduction hypothesis is logically untestable on this sample because there were no baseline validator failures to reduce.

Representative diagnostics. Representative per-task artifacts (for example, pair-000001) confirm that the shared initial generator outputs matched the task `required_sequence` and that `validator_trace.validation_after.valid` = true with `violation_count` = 0. Difficulty annotations in validator traces (examples: `make_before_break_uplink_switchover`, `paired_link_atomic_mtu_change`) and

`parsed_command_count/required_command_count` show that tasks exercised medium-to-very hard workflow constraints despite generating validator-satisfying candidates. The Mininet 10% holdout (20 tasks) executed as preregistered; `analysis/contamination_summary.csv` reports zero validator-emulator disagreements, and deterministic reconciliation artifacts report all consistency checks passed (`pair_accounting_consistent` = true).

Unexecuted branches. Planned sensitivity branches for flagged contamination treatments and alternative missingness handling were not invoked because no tasks were flagged and no episodes failed. Secondary exploratory estimands (average repair iterations, GVR-introduced failure rate, model-call cost per successful repair) are empty in the analysis bundle because no repair attempts occurred; the analysis artifacts record these `secondary_results` as empty.

V. DISCUSSION

Conditional interpretation. The degenerate null outcome demonstrates a logical constraint: when the chosen generator, prompt template, and sampling policy produce validator-passing candidates for every item in a workload, an added validator-coupled repair step cannot reduce validator-detected failures. This artifact-grounded observation clarifies that GVR’s marginal utility is a function of baseline validator failure prevalence and validator coverage rather than a universally guaranteed safety gain.

Task and prompt alignment. Archived validator traces and representative task artifacts show that tasks encoded explicit `required_sequences` and workflow constraints, and that the deterministic prompts and adapter-enforced shared-initial semantics produced outputs conforming to those constraints. Because the task bank was curated with machine-checkable intent representations, the generator had a narrow, high-signal mapping from intent to commanded sequence—reducing semantic ambiguity and making validator-passing outputs more likely. Difficulty annotations (e.g., patterns labeled `make_before_break_uplink_switchover` or `paired_link_atomic_mtu_change`) indicate that the workload exercised nontrivial coordination constraints; the absence of validator failures therefore reflects alignment between task formalization, prompt semantics, and validator checks rather than trivial task simplicity.

Validator fidelity and emulator coverage. `Deterministic_netops_validator_v5` enforces a defined set of syntactic and intent constraints; the preregistered Mininet secondary check (10% holdout = 20 tasks; see `analysis/contamination_summary.csv`) executed and reported zero disagreements, supporting the internal consistency of validator verdicts on the sampled holdout. Nevertheless, the limited scope of the holdout and the disclosed absence of an explicit labeled seed for holdout selection constrain claims about broader emulator coverage: the artifact bundle shows no validator-emulator mismatches for the tested subset, but this does not prove absence of validator blind spots across all production conditions. Practitioners should therefore treat

validator pass rates as conditional on validator completeness and the representativeness of any emulation checks used to validate that completeness.

Design choices that shaped outcomes. The preregistered `shared_initial` paired design intentionally removes per-condition resampling as a lever to create discordant cases; it isolates the incremental effect of invoking a repair step on a fixed generator output. That controlled isolation is scientifically valuable for measuring the marginal effect of repair on the exact same candidate, but it also reduces the experimental pathways by which GVR could yield observed improvements in a different design (for example, multi-sample generation where a repair step might solicit alternative formulations). Similarly, the single-model, deterministic sampling policy reduces stochastic variance but limits generality: different sampling temperatures, multi-sample strategies, or model ensembles may produce nonzero baseline failure rates and thus different measured GVR benefits.

Implications for operational design and future studies. The empirical datum—zero baseline validator failures under these controlled, curated conditions—suggests concrete operational guidance: teams considering automated repair should first instrument representative workloads to measure baseline validator failure prevalence and then evaluate validator completeness via targeted emulation or real-device checks. If baseline validator failures are rare under the planned generator/prompt policy, investments that raise validator fidelity (vendor-specific behavior models, richer device semantics) or that increase input diversity (adversarial testing, multi-sample generation, multi-model ensembles) may provide greater marginal safety returns than adding an automated repair stage with a small repair budget. Conversely, when baseline failures are nontrivial, a well-engineered GVR pipeline with adequate repair budget and high-quality repair prompts can reduce validator-detected issues; future experiments should measure repair iterations and model-call cost per successful repair when such failures are present.

Recommended experimental refinements. To quantify GVR effectiveness in settings where it can act, future confirmatory runs should (a) target task banks or prompt variants with nonzero baseline validator failures (for example, by introducing realistic ambiguity or controlled perturbations), (b) evaluate multi-sample generation and multi-model ensembles as baselines, (c) increase repair budgets and test multi-step repair strategies including human-in-the-loop handoff, and (d) expand emulator or real-device coverage and record explicit, reproducible holdout selection seeds so that contamination checks are fully reconstructable. The archived artifacts show that `secondary_results` were empty because no repairs occurred; capturing repair-cost and iteration metrics requires a design where repairs are invoked in practice.

Threats to validity revisited. Construct validity: validator pass/fail outcomes reflect the formalized constraints encoded in `deterministic_netops_validator_v5`; absent richer device modeling, some operationally relevant failures may not be detected. External validity: a single hosted model (`gpt-5-mini`),

deterministic sampling, `shared_initial` semantics, and a curated task bank constrain generalizability. Internal validity: pre-registration and deterministic adapter enforcement strengthen causal interpretation in the paired comparison, but the controlled choices also narrow the space of behaviors observed. Reproducibility: the execution bundle archives model calls, validator traces, and the contamination summary showing zero Mininet disagreements, but the missing labeled Mininet holdout seed prevents deterministic reconstruction of the exact holdout sample—this gap is disclosed in the Disclosure Statement.

Concluding synthesis. The completed artifacts and diagnostic traces make clear that the preregistered GVR pipeline, as instantiated here, produced no measurable reduction in validator-detected failures because there were no such failures to reduce. That operationally informative null result highlights that the safety value of automated repair is not universal but workload- and validator-dependent. Future work should deliberately select or construct workloads with nontrivial baseline failure rates and expand validator fidelity and sampling strategies to empirically quantify when GVR provides practical safety gains.

VI. LIMITATIONS

- Task-bank scope: the confirmatory sample used a curated synthetic/public intent→configuration bank ($n = 200$); results therefore reflect this bank and may not generalize to broader production workloads or adversarial distributions.
- Single-model and shared-initial setting: only `gpt-5-mini` (hosted) with adapter-enforced shared-initial generation was tested (execution artifacts record `hosted_model_call_event_count = 200` and `stage_counts shared_initial_generation = 200`). Outcomes may differ across model families, independent per-condition sampling, nonzero temperature sampling, or larger repair budgets.
- Validator coverage and oracle limitations: the primary deterministic validator enforces a defined set of syntax/intent constraints; validators can fail to capture real device idiosyncrasies, vendor-specific behaviors, or emergent failure modes detectable only by end-to-end deployment.
- Adversarial robustness untested: this confirmatory run did not include active adversarial injection or red-team tool-description attacks; prior audits indicate such vectors can bypass naive defenses and remain a separate concern.
- Reproducibility gap for contamination check: the randomized holdout selection seed for the Mininet contamination check is absent from the labeled artifacts, preventing exact reconstruction of the holdout selection.
- Single repair budget: the GVR pipeline allowed at most one automated repair prompt per task; multi-step repair strategies, different repair budgets, or human-in-the-loop approvals were not exercised here.

VII. CONCLUSION

We preregistered and executed a sealed paired study comparing LLM-only generation to a deterministic GVR pipeline on 200 NetOps intent→configuration tasks using gpt-5-mini and deterministic_netops_validator_v5. Both pipelines produced validator-passing configurations for every task (n11 = 200; no discordant pairs), resulting in an observed paired difference of 0.0, exact McNemar $p = 1.0$, and bootstrap 95% CI = [0.0, 0.0]. The preregistered $\geq 50\%$ relative-reduction hypothesis could not be evaluated because baseline validator failures were absent. The artifact-grounded outcome clarifies that GVR's marginal value is workload dependent: when baseline validator-detectable failures are rare, automated repair yields no measured benefit for those validator-detected errors. Future studies should target workloads with nontrivial baseline failure rates, expand validator fidelity with richer emulation or real-device checks, evaluate multi-sample and multi-model policies, and explore multi-step or human-in-the-loop repair budgets to characterize practical operational gains. All execution artifacts and analysis outputs are archived for independent inspection under the preregistration id cited above.

DISCLOSURE STATEMENT

Disclosure (mandatory). This study executed under the autonomous post-lock preregistration cnsms2026-gvr-effectiveness-02. Declared model: gpt-5-mini (provider label: openai_responses) invoked via the hosted_netops_gvr_v1 adapter. The exact initial master prompt is archived at provenance/master_prompt.txt (SHA-256: 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ba). Topic constraints: the human Research Director limited the study to Generative AI and LLMs for NetOps focusing on safety/reliability of generated network changes. Frameworks and tools: orchestration used hosted_netops_gvr_v1 and the deterministic task generator deterministic_netops_task_generator_v5. Primary deterministic validator: deterministic_netops_validator_v5 (intent/constraint checks). A Mininet-based emulator was preregistered and executed as a secondary disagreement detector on a randomized 10% holdout (20 tasks); the artifact analysis/contamination_summary.csv shows zero disagreements for that holdout. Preregistration and freeze: the experimental plan (estimand, sample size $n = 200$, paired design, stopping rules, max model calls = 400, repair budget = 1 per task, shared-initial generation semantics) was frozen before any model calls. Autonomous execution and artifacts: the run executed autonomously under the frozen plan (start: 2026-08-30T11:50:34.530059Z; end: 2026-08-30T12:12:07.541650Z); model calls, prompts, seeds, validator traces, emulator traces, logs, and the artifact manifest are archived in the execution bundle. Model-call budget and usage: 200 model calls were used (one shared initial generation per task); up to 200 repair calls were allowed but none were applied because initial candidates passed the validator. Human involvement: the human Research Director selected the topic family and pre-lock constraints; no human

manual editing or selective rewriting of technical manuscript content occurred after the autonomy lock. Deviations and restarts: no post-preregistration design changes or technical restarts occurred. Known reproducibility gap: the explicit randomized selection seed used to draw the Mininet 10% holdout is not present as a labeled artifact in the archive; this absence prevents exact deterministic reconstruction of the drawn holdout despite the contamination summary reporting no disagreements. The master prompt archive reference above is the immutable prompt required by the track.

REFERENCES

- [1] Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Zican Dong, Yupeng Hou, Beichen Zhang, Yingqian Min, Junjie Zhang, Peiyu Liu, Xiaolei Wang, Yifan Du, Yushuo Chen, Yushuo Chen, Zhipeng Chen, Jinhao Jiang, Ruiyang Ren, Yifan Li, Xinyu Tang, Peiyu Liu, Yiwen Hu, Jian-Yun Nie, Ji-Rong Wen, "A Survey of Large Language Models", 2026, 10.1007/s11704-026-60308-3
- [2] Changjie Wang, Mariano Scazzariello, Alireza Farshin, Simone Ferlin, Dejan Kostić, Marco Chiesa, "NetConfEval: Can LLMs Facilitate Network Configuration?", 2024, 10.1145/3656296
- [3] Oscar G. Lira, Oscar Mauricio Caicedo Rendón, Nelson L. S. da Fonseca, "Large Language Models for Zero Touch Network Configuration Management", 2024, 10.1109/mcom.001.2400368
- [4] Chirag Shah, Ryen W. White, Reid Andersen, Georg Buscher, Scott Counts, Sarkar Snigdha Sarathi Das, Ali MontazerAlghaem, Sathish Manivannan, J. Neville, Nagu Rangan, Tara Safavi, Siddharth Suri, Mengting Wan, Leijie Wang, Longqi Yang, "Using Large Language Models to Generate, Validate, and Apply User Intent Taxonomies", 2025, 10.1145/3732294
- [5] omar altawil, Dr. Mustafa Al-Fayoumi, "Configuration-Level Semantic Brittleness in Tool-Using LLM Agents: A Factor-Ranking Study of Persona, History, and Tool Definition", 2026, 10.2139/ssrn.7203631
- [6] Mohammadreza Rashidi, "Measuring How LLM Tool Descriptions, Cross-Tool Ambiguity, and Action Chains Compose into Excessive Agency Exploits", 2026, 10.21203/rs.3.rs-10646209/v1
- [7] Yoshihiro Ando, "Secure-by-Default Guardrails for MCP-Based Tool Use in Multi-Modal LLM Agents", 2026, 10.36227/techrxiv.177006109.97957916/v1
- [8] Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, Yarin Gal, "Detecting hallucinations in large language models using semantic entropy", 2024, 10.1038/s41586-024-07421-0
- [9] Shuyin Ouyang, Jie M. Zhang, Mark Harman, Meng Wang, "An Empirical Study of the Non-Determinism of ChatGPT in Code Generation", 2024, 10.1145/3697010
- [10] Oghenekeno Hilkiyah Okula, "The Era of AI Agents for Network Engineering": Intent-Aware Auto Remediation for Network Disruption or Failures Using AI Agents, Large Language Models (LLM) and Model Context Protocol (MCP) Mechanisms", 2026, 10.36227/techrxiv.177004277.71795055/v1