

Measuring Validator and Sandboxed-Execution Coverage for LLM-Generated NetOps Artifacts: A Paired Confirmatory Study

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory experiment (study_id f2c3d8b1-7a6e-4a9a-b2c3-5f8e1d2b6c7e) that measured the marginal detection benefit of adding sandboxed emulation to a deterministic validator for LLM-generated intent→configuration artifacts. Under shared_initial_candidate semantics using the declared model "gpt-5-mini" (provider: openai_responses) we evaluated Npairs = 720 confirmatory tasks. The deterministic validator (deterministic_netops_validator_v5) flagged 719/720 = 0.9986111111111111; the guarded pipeline (validator + sandboxed emulation) flagged 720/720 = 1.0. Paired difference (guarded - baseline) = 0.001388888888888884; exact McNemar two-sided p = 1.0; 95% bootstrap percentile CI (10000 resamples, seed = 1729) = [0.0, 0.004166666666666667]. Run accounting: model_calls_used = 721 (720 shared initial + 1 repair-stage call). These artifact-grounded results lie far below our preregistered minimum effect-of-interest (0.20); we report the preregistered design, exact quantitative outcomes, run-accounting diagnostics, artifact-grounded interpretation for the negligible guarded benefit in this regime, and reproducibility pointers into the archived execution bundle.

I. INTRODUCTION

Generative large language models (LLMs) are increasingly proposed to translate operator intent into device configuration and NetOps workflows. Operational deployment requires generated artifacts to satisfy syntactic correctness, workflow ordering, and higher-order safety/policy constraints. A common mitigation architecture is generate→validate→(repair)→execute, sometimes augmented by sandboxed emulation to reveal runtime-only failures that static validators miss. Prior surveys and system reports advocate such workflow-centred mitigations [1]–[3], but provide limited large-scale paired measures of the marginal value added by sandboxed execution. Motivated by this gap, we preregistered and executed a controlled paired study to measure the aggregate detection-rate difference contributed uniquely by an automated sandboxed-emulation stage when the same initial generated candidate is analyzed by both conditions.

Research question and contributions. Our preregistered research question was: for synthetic intent→configuration tasks under shared_initial_candidate semantics, what is the paired detection-rate difference (guarded - baseline) between (A) a deterministic validator-only pipeline and (B) the same deterministic validator followed by automated sandboxed emulation? Contributions: (1) a machine-readable preregistration and faithful autonomous execution; (2) an artifact-grounded

paired analysis on Npairs = 720 with complete accounting; (3) an artifact-grounded interpretation explaining a near-zero marginal effect; (4) concise reproducibility pointers into the archived bundle and prescriptive boundary conditions where guarded stages are likely to matter.

II. RELATED WORK

Workflow-centred mitigations and verification tools. Surveys and recent system reports document architectures that pair LLM generation with validators, verifiers, and emulators to improve correctness of generated NetOps artifacts and adjacent program-generation tasks [1], [2]. Tool-focused work demonstrates generate→validate→repair patterns (e.g., Batfish/NetKAT style verification, emulator-assisted validation) but reports heterogeneous gains depending on task, validator rule coverage, and emulator fidelity [2], [3].

Empirical findings motivating a paired design. Empirical studies in program/specification generation show external verifiers often catch model mistakes not found by the model alone, and some NetOps prototypes use emulation to expose runtime issues [3]. These prior findings motivated our preregistered paired design that isolates guarded-stage contributions: because we enforce shared_initial_candidate semantics, the only permissible sources of between-condition discordance are downstream guarded-stage emulation observations and permitted repair-stage invocations recorded in run accounting.

III. METHODOLOGY

Preregistration and confirmatory design. We preregistered a paired-binary confirmatory design (study_id f2c3d8b1-7a6e-4a9a-b2c3-5f8e1d2b6c7e). The primary estimand is the paired_success_rate_difference_guarded_minus_baseline aggregated across confirmatory samples. The preregistered sample plan targeted Npairs = 720 (planned episodes = 1440) with a minimum effect-of-interest of 0.20 (20 percentage points) for power planning. The preregistration fixed inclusion/exclusion rules, missingness handling (deterministic reseed up to two times for parser failures; deterministic replacement for irrecoverable failures), multiplicity handling for exploratory secondary tests, and contamination controls based on deterministic hashing and artifact-version checks.

Task generation, topology templates, and vendor abstractions. Confirmatory items were generated deterministically from the intent_configuration_repair_v1 task family using a

deterministic task generator. The bank covers three abstracted vendor-CLI families and three topology templates (leaf-spine, 3-node service chain, triangle). Each task manifest entry contains: `initial_state`, `expected_final_state`, an explicit intent, a `required_command_sequence`, `allowed_touched_fields`, and workflow policies (e.g., `make-before-break`, `must-be-down-for-fields`, `minimum_active_in_groups`). Inclusion required vendor-specific parser acceptance after up to two deterministic reseeds. In the confirmatory run there were zero irrecoverable parser exclusions and `complete_pair_count = 720`, consistent with the preregistered missingness policy.

Model, sampling semantics, and `shared_initial_candidate` enforcement. The declared/requested model was recorded as "gpt-5-mini" (provider: `openai_responses`) and the model configuration archived (`requested_max_output_tokens = 2000`; the client did not set an explicit temperature parameter—recorded as null in the model configuration). Important preregistered clarifications: (a) temperature recorded as null indicates the client supplied no explicit temperature value; it does not imply provider-side determinism or `temperature=0`, and it does not alter the `shared_initial_candidate` contract; (b) `shared_initial_candidate` semantics were enforced per the preregistration: for each pair we obtained exactly one sampled initial candidate (a single provider call) and that identical candidate was presented to both downstream conditions (baseline and guarded). This design isolates downstream guarded-stage effects by removing generation variability as a source of between-condition discordance.

Baseline validator specification. The baseline validator (`deterministic_netops_validator_v5`) implements deterministic, rule-based intent-constraint validation. Its core checks (applied to a normalized candidate) include: vendor-CLI syntactic parsing; per-command parsing and canonicalization; sequence-order checks against `required_command_sequence`; atomicity/workflow constraints (e.g., `make-before-break`, `all_down_before_any_field_change`); protected-interface invariants; group-level activity constraints (minimum/maximum active interfaces within specified groups); and final-state conformance tests (`expected_final_state`). The validator returns a binary detection flag (`detected/not-detected`) under the preregistered scoring rule `full_intent_constraint_validation` and emits a `validator_trace` object (`validation_before/after`, `parsed_command_count`, `required_command_count`, `violation_codes`).

Guarded pipeline and sandboxed-emulation harness. The guarded pipeline runs the baseline validator and, irrespective of the validator outcome, then executes the candidate in an automated sandboxed-emulation harness that instantiates the declared topology template and executes normalized CLI commands against vendor-abstracted virtual devices. The emulation harness records runtime observations relevant to operational safety: interface admin transitions and timing, transient downtime and concurrent-interface states, MTU/VLAN application effects, and observable connectivity or reachability consequences at the topology level. Emulation results are translated into a set of runtime-observation rules that can

trigger a guarded detection if they reveal policy or safety violations not encoded in the static validator rule set. The emulation harness is intentionally conservative in mapping observations to detections; its purpose in this study is to reveal runtime-only misconfigurations that the baseline validator might miss.

Repair policy and run accounting. The preregistered repair policy allowed at most one repair-stage model invocation per task; repair invocations are permitted only after the shared initial candidate is generated and after baseline validator and emulation observations are available for the guarded condition. All repair calls and their transcripts are recorded in the execution manifest and deterministic reconciliation artifacts. Provider-call auditing for the confirmatory execution records `model_calls_used = 721` with `stage_counts` showing `shared_initial_generation = 720` and `repair = 1`; thus only a single repair-stage call occurred across the entire confirmatory set. Because `shared_initial_candidate` semantics were enforced, any discordant pair (guarded-only detection) must be attributable to either an emulation-observed runtime violation or to the single documented repair invocation; both are traceable in the archived per-episode artifacts.

Inclusion, determinism, and contamination controls. The execution used deterministic artifact hashing, generator-seed recording, and explicit binary/configuration-version checks before confirmatory runs. Parser-failure handling used deterministic reseeding with stored PRNG seeds under the task manifest; irrecoverable failures would be replaced deterministically and recorded. Deterministic reconciliation verifies `complete_pair_count = 720`, `completed_episode_count = 1440`, and `failed_episode_count = 0`. Provider-call auditing reports `cache_miss_count = 721` and zero failed hosted model calls. The preregistered contamination rules also halted confirmatory execution if validator or emulator binary/version mismatches were detected; no such mismatches occurred.

Analysis procedures and pre-registered inferential methods. The primary analysis executor is a paired-binary procedure aligned with McNemar-style inference and cluster-robust variance estimation. For the primary aggregated estimand we compute the paired detection-rate difference (guarded - baseline) with a two-sided exact McNemar test and produce a 95% paired-bootstrap percentile confidence interval (bootstrap resamples = 10000, bootstrap_seed = 1729) as specified in the preregistration. Secondary and exploratory per-class/vendor paired analyses were pre-specified and treated with Bonferroni-adjusted confidence intervals in the preregistered multiplicity plan. The primary analysis used the preregistered complete-pair handling (exclude irrecoverable technical pairs) and a pre-registered sensitivity analysis treating irrecoverable failures as nondetections; in the confirmatory execution no irrecoverable pairs were present, so conclusions rest on complete paired data.

Artifact-locator guidance for independent verification. All per-episode raw scoring and trace artifacts were archived. The canonical input-results file is `execution/raw_results.jsonl` (recorded in the execution manifest). Independent verifiers can locate any discordant pair by searching that file for

an entry where `baseline_score = 0` and `guarded_score = 1`; `deterministic_reconciliation.json` in the analysis artifacts links pair accounting and `provider_call_audit` entries and identifies the repair-stage call(s) associated with discordant outcomes. Summary paired counts and contingency outputs are archived as `analysis/paired_contingency_table.csv` and `analysis/condition_summary.csv`. The analysis used these archived artifacts as the direct input to the paired-binary executor; the numeric outputs reproduced in Results are verbatim from `analysis/results.json`.

Design rationale and trade-offs. We chose `shared_initial_candidate` semantics to produce a conservative, direct estimate of guarded-stage marginal contributions: by reusing the same sampled initial candidate for both conditions, any observed discordance can only arise downstream of generation (emulation observation or repair). This isolates the guarded-stage value but intentionally omits any potential benefit that could accrue from independent re-generation in the guarded condition (e.g., prompting-based repair prior to validator execution). Our inclusion policy (parser-validated synthetic tasks) reduces noise and ensures task executability but also concentrates the task bank where a strong validator can achieve high coverage; this design choice is central to interpreting the near-zero aggregate effect observed in the confirmatory run.

IV. RESULTS

Primary confirmatory counts and test statistics (artifact-grounded). The preregistered paired analysis produced the following exact contingency counts (`analysis/paired_contingency_table.csv`, sha256 efe6e8e7016fa843a8226e911dbde829e943903d719101da8d956b3ff899133c): `n11 = 719` (detected by both), `n10 = 1` (guarded-only), `n01 = 0` (baseline-only), `n00 = 0` (neither). Marginal detection rates are: `baseline = 719/720 = 0.9986111111111111` and `guarded = 720/720 = 1.0`. The paired estimate (`guarded - baseline`) = `0.0013888888888888884`. The exact McNemar two-sided test reported `p = 1.0`. Paired-bootstrap percentile 95% CI (10000 resamples, `bootstrap_seed = 1729`) = `[0.0, 0.004166666666666667]`. All numeric values above are reproduced verbatim from `analysis/results.json` and analysis artifacts.

Bootstrap and test details. The bootstrap CI was computed with 10000 resamples (`seed = 1729`) using paired resampling of episode-pairs; the analysis executor used a paired-bootstrap-percentile approach (artifact: `analysis/analysis_log.jsonl`). The McNemar exact two-sided p-value complements the bootstrap CI by testing symmetry in the discordant-pair counts; here the observed discordant counts (`n10 = 1`, `n01 = 0`) produce no evidence of a systematic guarded advantage at conventional significance thresholds (`p = 1.0` under the exact test reported by the executor).

Run accounting and discordant-pair provenance. Provider-call accounting and deterministic reconciliation confirm `model_calls_used = 721` (stages: `shared_initial_generation = 720`; `repair = 1`) and `complete_pair_count = 720`

(`analysis/deterministic_reconciliation.json`, sha256 49731a2e52e049ccf5b53c35ddac2351ad75e7e95504dc4a42bcb096d0dec3aa). Because `shared_initial_candidate` semantics were enforced, the single guarded-only detection (the unique `n10` entry) must arise from downstream guarded-stage activity (sandboxed emulation) or from the single recorded repair-stage invocation.

Explicit archival pointer to the discordant episode (artifact-grounded). The unique discordant episode can be located in the archived raw-results file: `execution/raw_results.jsonl` (input-results SHA-256 = `df367ecc40a6f0eaf402129260e11d8450884a4ea5fdb94a602e0d2513aef0e8`). Within that file, filter for the entry where `baseline_score = 0` and `guarded_score = 1`; that `raw_results.jsonl` record contains the episode's `pair_id` and the linked `validator_trace` and any `repair_provider_trace` fields that document the repair-stage invocation. The deterministic reconciliation artifact (`analysis/deterministic_reconciliation.json`, sha256 49731a2e52e0...) and the `provider_call_audit` block therein corroborate that the repair-stage call recorded in provider accounting maps to that unique discordant pair. We provide these exact artifact identifiers so readers can unambiguously locate the discordant pair and its repair/emulation traces in the archived bundle (see Reproducibility pointers below).

Discordant attribution and repair linkage. The archived `raw_results` entry for the discordant pair includes `validator_trace` (`validation_before/after`), any emulation observations, and fields capturing whether a repair prompt was issued and its trace path. The `deterministic_reconciliation.json` `provider_call_audit` documents `stage_counts` and confirms exactly one repair-stage call; matching timestamps and trace-path fields in `raw_results.jsonl` link that repair-stage call to the guarded-only detection. Thus the run artifacts substantiate that the single guarded-only detection corresponds to downstream guarded-stage activity (repair/emulation) rather than to generation differences, consistent with `shared_initial_candidate` semantics.

Representative per-episode diagnostics. Representative scoring traces archived under `execution/scoring/task-000001-*.json ... task-000006-*.json` include detailed `validator_trace` objects with `normalized_configuration`, `validation_before/after` states, `parsed_command_count`, `required_command_count`, `violation_count`, and `validator_version = 5.0`. Those representative artifacts illustrate that the vast majority of `shared_initial` candidates complied with validator rules (`violation_count = 0`) and that sandboxed emulation generally corroborated validator findings in this curated task bank. The `analysis/condition_summary.csv` artifact (sha256 394bc690671b7...) summarizes successes/failures by condition and reproduces the marginal success rates reported above.

Synthesis of numeric outcomes. In short, the guarded pipeline detected one more successful item than the baseline (720 vs. 719) across 720 paired episodes. The paired estimate (`~0.00139`) and its 95% bootstrap CI (`[0.0, 0.0041667]`) are both orders of magnitude smaller than our preregistered minimum effect-of-interest (0.20). Provider-call accounting

(model_calls_used = 721, repair = 1) further shows that the observed discordance is attributable to recorded downstream activity rather than to between-condition generation variance. These concrete, artifact-grounded outcomes are the basis for the interpretations and operational boundary conditions we discuss elsewhere in the manuscript.

V. DISCUSSION

Why was the guarded-stage aggregate benefit negligible? The archived evidence supports three artifact-grounded factors and suggests concrete boundary conditions where guarded stages remain important.

1) Validator ceiling effect (artifact-grounded). The deterministic validator covered the explicit syntactic and semantic constraints exercised by the synthetic tasks: most shared_initial candidates parsed successfully and produced validator traces with violation_count = 0. Analysis artifacts record n11 = 719 and n10 = 1, so the baseline validator already detected or accepted nearly all conformant configurations (baseline_success_rate $\approx$ 0.9986). With validator coverage at this level there is little statistical headroom for the guarded stage to contribute aggregate detections at scale. Representative validator_trace fields (parsed_command_count and required_command_count in execution/scoring/*.json) confirm that tasks were largely within the validator’s rule set.

2) Task curation and inclusion policy concentrate cases inside validator coverage. The confirmatory bank used parser-validated synthetic items that explicitly encoded workflow patterns (shutdown→configure→restore, make-before-break up-link switchover, paired MTU changes). The inclusion rule required vendor-parser acceptance before pairing, and the missingness artifact shows zero irrecoverable parser failures for the confirmatory set (analysis/missingness_summary.csv). This curation intentionally constrained the evaluation to cases where the validator’s deterministic rules apply; it therefore reduces the prevalence of subtle runtime-only failure modes (race conditions, unexpected protocol interactions, or stateful side-effects) that emulation would be uniquely positioned to reveal.

3) Shared_initial_candidate semantics and observed repair activity. Per the preregistration we enforced shared_initial_candidate: the same single sampled candidate was analyzed by both conditions for each pair. Provider-call accounting and deterministic_reconciliation.json show model_calls_used = 721 (720 shared initial + 1 repair-stage call) and stage_counts repair = 1. Under these semantics, between-condition discordance cannot arise from different model outputs; it can only arise from (a) additional runtime observations surfaced by emulation that are not encoded as validator rules, or (b) permitted downstream repair actions. The single n10 discordant count is explicitly mapped in the archived raw_results.jsonl record to the recorded repair-stage invocation, so the observed discordance is attributable to downstream guarded-stage activity rather than generation variability.

Practical interpretation and policy implications. The artifact-grounded conclusion is conditional: when a high-coverage static validator exists and inputs are constrained and parser-validated, automated sandboxed emulation may add negligible marginal detection at scale. Operationally, this suggests a cost/benefit trade-off: expensive emulation and canary infrastructure will have the most value when (i) validator coverage is known to be incomplete for target workflows, (ii) inputs include unconstrained or noisy natural-language intents, (iii) stateful multi-step changes introduce emergent runtime behaviors, or (iv) adversarial or fuzzed inputs are expected. In those regimes emulation can expose violations of invariants not expressible in the validator’s static rule set (for example, transient traffic blackholing under failure-induced routing changes or vendor-specific CLI side-effects). The archived per-episode traces (execution/scoring/*.json) provide concrete examples and can be used to build such adversarial testbeds.

Diagnostics useful to operators and evaluators. The run artifacts enable targeted diagnostics rather than blanket conclusions. For example: - Coverage auditing: compare validator_rule_inventory to parsed_command_count statistics across tasks to locate missing rules; analysis/condition_summary.csv and execution/scoring/* contain the necessary fields. - Discordant attribution: locate the unique discordant raw_results.jsonl entry (baseline_score = 0 and guarded_score = 1) and follow its validator_trace and repair_provider_trace to see whether the guarded detection arose from an emulation observation or an applied repair; deterministic_reconciliation.json confirms the single repair-stage call linkage. - Emulation fidelity checks: examine the emulation observation summaries in validator_trace.validation_after and compare them to expected_final_state in task_manifest entries to identify classes of runtime behaviors the emulator notices but the validator cannot express. These diagnostics are artifact-grounded and reproducible from the archived bundle referenced in the execution manifest.

Threats to validity and scope (artifact-grounded). Internal: shared_initial_candidate reduces variation in generation and isolates guarded-stage effects but constrains external generality — in practice operators may allow multiple candidate generations or stochastic sampling to improve correctness, which could change the relative benefit of guarded stages. Construct: our tasks are synthetic and parser-validated; they target representative NetOps patterns but do not exhaust the space of runtime faults (e.g., vendor implementation bugs, timing-sensitive state races). External: a single declared/requested model family (gpt-5-mini) and one validator implementation (deterministic_netops_validator_v5) limit generalization across models, validator families, and larger production topologies. Conclusion: power planning targeted a 0.20 minimum detectable paired difference; the observed paired difference ($\sim$ 0.0014) lies far below that sensitivity and therefore should be interpreted as a narrowly scoped near-zero finding for this experimental regime rather than a general negative result for all NetOps workflows. All of these limitations are documented in the preregistration and in analysis artifacts (missingness and rec-

conciliation files).

Operational recommendations and next steps grounded in the archived evidence. To meaningfully evaluate guarded-stage benefits where they matter operationally, future controlled studies (using the same artifact-oriented framework) should: (1) include adversarially generated tasks and fuzzed intents designed to evade validator rules; (2) broaden vendor-CLI families and emulate vendor-specific OS behaviors with higher-fidelity toolchains; (3) allow multiple sampled initial candidates per task to measure interplay between generation variability and guarded-stage detection; (4) instrument canary/rollback-aware scoring that penalizes dangerous but syntactically valid changes; and (5) measure cost and latency trade-offs for emulation at scale. The archived execution artifacts and per-episode traces supplied by this run provide a reproducible starting point and concrete templates for these extensions (see `execution/raw_results.jsonl` and `analysis/deterministic_reconciliation.json`).

Summary. The evidence archived with this run shows that, in a curated parser-validated synthetic task bank and under `shared_initial_candidate` semantics, a very strong deterministic validator already achieves near-complete coverage and the guarded sandboxed-emulation stage produced only a single additional detection tied to a documented repair-stage call. This narrow, artifact-grounded finding highlights circumstances in which validator-centric approaches may suffice, and it precisely identifies the empirical and methodological axes (task diversity, adversarial inputs, emulator fidelity, generation variability) where guarded stages are likely to provide substantive operational value.

VI. LIMITATIONS

- Synthetic task bank and deterministic task generator produce limited ecological diversity and create a validator-ceiling effect; results may not generalize to noisier, adversarial, or production distributions.
- Single declared/requested model family (gpt-5-mini) and single validator implementation (deterministic_netops_validator_v5) limit generalizability across models and validator families.
- Preregistered `shared_initial_candidate` semantics (one initial generation reused across paired conditions) isolate guarded-stage contribution but remove between-condition generation variability; this design choice constrains discordance sources to downstream guarded-stage observations or repair calls.
- Sandboxed-emulation fidelity relative to vendor/OS production behaviors and large-scale stateful interactions is limited by the emulation harness used; higher-fidelity emulators could change outcomes in other regimes.
- Study power planning targeted a minimum detectable paired difference of 0.20; the observed aggregate effect (~0.0014) lies far below that design sensitivity and should be interpreted as a narrowly scoped near-zero finding for this experimental regime rather than a general negative result for all NetOps workflows.

- Provider-side nondeterminism and hidden caching remain residual uncertainties despite deterministic reconciliation procedures; these are documented in `analysis/deterministic_reconciliation.json` and `provider_call_audit` artifacts.

VII. CONCLUSION

We preregistered and executed a paired confirmatory study comparing a strong deterministic validator with the same validator augmented by sandboxed emulation on a curated synthetic intent→configuration task bank. Over 720 paired tasks the guarded pipeline produced a negligible aggregate detection gain (0.001388888888888884) with exact McNemar $p = 1.0$ and 95% bootstrap CI [0.0, 0.004166666666666667]. Run accounting shows one repair-stage invocation corresponding to the single guarded-only detection; because `shared_initial_candidate` semantics were enforced, archived per-episode traces permit independent verification of that mapping via `execution/raw_results.jsonl` and `analysis/deterministic_reconciliation.json`. We interpret the result as arising from validator ceiling effects and curated task design; follow-up work should focus on adversarial and production-like task banks, multi-validator ensembles, and higher-fidelity emulation to identify regimes where guarded stages add substantive operational value.

Reproducibility pointers (compact). The canonical execution bundle archived by the run contains: `execution/execution_manifest.json` (execution summary and preregistration SHA-256), `execution/raw_results.jsonl` (input-results SHA-256 `df367ecc40a6f0eaf402129260e11d8450884a4ea5fdb94a602e0d2513aef0e8`), `analysis/deterministic_reconciliation.json` (sha256 `49731a2e52e049ccf5b53c35ddac2351ad75e7e95504dc4a42bcb096d0dec3aa`), `execution/model_configuration.json` (sha256 `a40e96e212ab87da251ac0d6dbb75bc543afa13438b5a6b9ebd073597ffdd820`), `analysis/paired_contingency_table.csv` (sha256 `efe6e8e7016fa843a8226e911dbde829e943903d719101da8d956b3ff899133c`), and `analysis/condition_summary.csv` (sha256 `394bc690671b731194f9b42b4dcbe28fded6e2ec8cf898b79831225343c22c28`).

The discordant episode is the unique `raw_results.jsonl` entry with `baseline_score = 0` and `guarded_score = 1`; that record contains the `pair_id` and validator/emulation traces that map it to the single repair-stage invocation recorded in `deterministic_reconciliation.json`. These artifacts are archived as the canonical reproducibility bundle referenced by the execution manifest and the preregistration record. (See Disclosure for exact manifest references.)

DISCLOSURE STATEMENT

Disclosure: Autonomous post-lock research run executed under the frozen master prompt supplied to the system. Declared/requested model: "gpt-5-mini" (provider recorded as `openai_responses`). Deterministic validator: `deterministic_netops_validator_v5`. Orchestration adapter family: `hosted_netops_gvr_v1`. Preregistration id:

f2c3d8b1-7a6e-4a9a-b2c3-5f8e1d2b6c7e (preregistration SHA-256: 0169919fb8c5aedb2c5a78e031f1ab5c3aeda405cde1ae80fb2f864768595c1b; recorded in the execution manifest). Initial master-prompt immutable reference: provenance/master_prompt.txt, SHA-256 = 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582b b39b15ba. Execution summary (artifact-grounded): completed_episode_count = 1440, complete_pair_count = 720, model_calls_used = 721 (720 shared initial + 1 repair-stage call); stage_counts and provider-call audit are archived in analysis/deterministic_reconciliation.json (sha256 49731a2e52e0...). Human role and autonomy boundary: a human Research Director selected the broad NetOps topic family and constrained the pre-lock master prompt; after the autonomy lock the agentic pipeline autonomously performed literature selection, preregistration, code and prompt generation, experiment execution, analysis, peer review, and manuscript drafting. No manual human editing of technical manuscript content occurred after the autonomy lock. The execution manifest and archived artifacts named above serve as the canonical reproducibility bundle.

REFERENCES

- [1] Muhammad Bilal, Jon Crowcroft, Ruizhi Wang, Xiaolong Xu, Schahram Dustdar, "Large Language Models for Agentic NetOps and AIOps: Architectures, Evaluation, and Safety", 2026, 10.48550/arxiv.2605.12729
- [2] Matt Brown, Ari Fogel, Daniel Halperin, Victor Heorhiadi, Ratul Mahajan, Todd Millstein, "Lessons from the evolution of the Batfish configuration analysis tool", 2023, 10.1145/3603269.3604866
- [3] <http://arxiv.org/abs/2407.08249>