Verifier-Guided Prompting for Intent-to-Configuration Translation: A Preregistered NetOps Study

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

***Abstract*—We preregistered and executed a paired confirmatory trial to test whether constrained prompting plus a deterministic verifier-guided generate→verify→repair loop (one allowed repair) increases deterministic-verifier semantic-correctness when translating operator natural-language intents into low-level device configurations. Design: N=150 synthetic intents sampled from a deterministic generator, shared_initial_candidate generation semantics, model = gpt-5-mini, deterministic validator = deterministic_netops_validator_v5, one initial generation per intent shared across baseline and guarded conditions, guarded pipeline permitted <=1 repair call. Results (artifact-grounded): baseline = 149/150 (99.333%); guarded = 150/150 (100.0%); paired absolute difference = 0.006667 (≈0.67 percentage points); exact McNemar two-sided p = 1.0; 95% bootstrap percentile CI = [0.00, 0.02] (10,000 resamples, seed = 1729). Provider accounting: completed episodes = 300; model_calls_used = 151; repair-stage invocations = 1. Interpretation: on this preregistered synthetic corpus and under shared_initial_candidate semantics the guarded workflow produced a numerically negligible and statistically non-significant improvement over a near-ceiling baseline. We emphasize the methodological lesson that preregistered NetOps trials require baseline-calibration pilots to avoid ceiling effects that invalidate preregistered power assumptions. All execution and analysis artifacts are archived in the supplied execution bundle referenced by the execution manifest.**

I. Introduction

Natural-language intent interfaces aim to reduce operator burden by letting humans express desired network changes as free text that is automatically translated into executable device configurations. Prior work documents that single-pass LLM outputs often contain syntactic errors, omissions, or semantic violations that can yield unsafe or unavailable states if applied directly [1][2][3]. A commonly proposed defense is a generate→verify→repair architecture in which an LLM produces a candidate, a deterministic verifier checks intent-level invariants or formalized constraints, and an automated repair (or human gate) corrects violations before execution [4][5]. Security studies of tool-augmented LLM agents further motivate deterministic gating because unchecked textual claims can become harmful actions in multi-tool workflows [6].

Research question. After an autonomous literature review and preregistration we posed a confined confirmatory question: does constrained prompting (structured templates with explicit semantic slots) combined with deterministic verifier-guided iterative feedback materially increase verifier-level semantic correctness of NL→configuration translation versus an unconstrained baseline under a paired design? The trial (study_cand_2026_b2_v1) was preregistered and frozen before any model outputs were observed.

Contributions. (1) A fully preregistered, artifactized paired confirmatory trial measuring deterministic-verifier pass/fail on a programmatically generated synthetic corpus (N=150 paired intents). (2) Artifact-backed outcomes showing that, under shared_initial_candidate semantics and a one-repair budget, the guarded pipeline raised verifier pass rate from 149/150 to 150/150 (+0.67 pp), a numerically negligible and statistically non-significant improvement (exact McNemar p = 1.0; 95% bootstrap CI = [0.00, 0.02]). (3) A methodological recommendation that preregistered NetOps trials include baseline-calibration pilots and explicitly specify generation semantics (shared_initial_candidate vs independent sampling) because semantics materially change the estimand and interpretation.

II. Related Work

Related literature falls into three strands. First, intent-to-configuration work demonstrates that converting free-text intents into verifier-sound artifacts is difficult and reports modest semantic-correctness on benchmarks, motivating post-generation checks [1][2]. Second, generate→verify→repair architectures—ranging from verifier-in-the-loop fine-tuning to policy-guided verifier feedback—have shown correctness improvements for infrastructure-as-code and configuration repair, but reported gains vary with baseline model competence, repair budget, and sampling semantics [4][7]. Third, security and systems analyses have demonstrated claim-to-action vulnerabilities where multi-tool LLM pipelines can convert false textual claims into concrete harmful actions, arguing for auditable gating and verifier-aware planning [6][8].

How this study situates. Our trial operationalizes a narrowly specified guarded architecture and emphasizes preregistration, paired inference, and deterministic-verifier endpoints. Key contrasts with prior reports that observed larger guarded-vs-baseline gains are: (a) shared_initial_candidate execution semantics (one initial sample shared across conditions), (b) a strict one-repair budget per intent, (c) a single tested model (gpt-5-mini), and (d) a synthetic verifier-capable task bank restricted to small topologies. These design choices reduce the maximal observable marginal benefit of the guarded repair under our estimand and explain why previously reported larger gains can coexist with our bounded result [4][7].

III. Methodology

Overview and preregistration. We preregistered study_cand_2026_b2_v1 and froze an analysis plan prior to any model calls. The design is paired-binary: each sampled intent yields one shared initial candidate used to evaluate both the baseline and guarded conditions (shared_initial_candidate semantics). The primary estimand was the paired_success_rate_difference_guarded_minus_baseline and the confirmatory test was an exact two-sided McNemar on the paired contingency table. The preregistration specified N=150 confirmatory intents, inclusion/exclusion rules, a one-repair budget for guarded episodes, a deterministic retry policy for provider timeouts (one retry), and an analysis plan that included a 10,000-resample paired bootstrap (bootstrap_seed = 1729) to produce percentile CIs. These preregistered elements are archived; the registered preregistration SHA-256 is recorded in the execution manifest.

Task generation, intent families, and prechecks. Confirmatory tasks were produced deterministically by deterministic_netops_task_generator_v5. The generator emits intent payloads with: (i) an explicit natural-language intent string, (ii) an initial_state and expected_final_state, (iii) an ordered required_sequence of field-level operations (e.g., shutdown→mtu→restore), (iv) policy constraints (allowed_touched_fields, protected_interfaces, final_group_constraints), and (v) a declared difficulty label. The confirmatory corpus was balanced across four intent families: safe sequence reachability/isolation, ACL-like access changes, VLAN/MTU sequence manipulations, and uplink switchover patterns. Topologies were restricted to verifier-capable small networks (4–32 nodes) to ensure deterministic validator coverage. Prior to model calls we ran deterministic unit-tests on the mapping/transformation code and a rule-based ambiguity detector; items flagged as ambiguous or failing parser/unit-tests were excluded per the preregistered exclusion rules. In this execution none of the confirmatory items were excluded.

Execution adapter, sampling semantics, and explicit baseline interpretation. The hosted_netops_gvr_v1 adapter enforced the shared_initial_candidate semantics required by the preregistration: for each intent the adapter performed exactly one initial model generation (the shared candidate) and used that identical candidate when scoring both the baseline and guarded conditions. The guarded pipeline could issue at most one additional repair call only if the deterministic validator failed the shared initial candidate. Because of this architecture, the baseline success rate measures the validity of the single initial sampled candidate; the guarded vs baseline contrast therefore quantifies the marginal effect of a single allowed repair applied to the identical sampled candidate rather than the effect of independently re-sampling first-pass candidates per condition. We include this precise sampling-to-estimand mapping in Methods to eliminate interpretive ambiguity about whether different first-pass generations were drawn for each condition.

Model configuration and nondeterminism. Declared model: gpt-5-mini (provider record: openai_responses). The archived execution/model_configuration.json records declared_model_version = "gpt-5-mini" and explicitly records temperature = null (unset). Importantly, the adapter and preregistration did not enforce a numeric temperature value or a fixed random seed; therefore model outputs may be nondeterministic across independent API sessions. We state this here in Methods and repeat it in the Disclosure Statement to comply with preregistration/runtime-parameter reporting requirements.

Repair budget, session isolation, and provider auditing. The guarded repair policy allowed at most one repair invocation per intent and the adapter maintained strict session isolation: each model call ran in a fresh API session (no cross-call state reuse). Provider auditing tracked hosted_model_call_event_count and cache statistics; the deterministic reconciliation artifact reports completed_hosted_model_call_event_count = 151, cache_miss_count = 151, and stage_counts = {shared_initial_generation: 150, repair: 1}. The adapter implementation enforces the preregistered retry rule (one retry after transient provider failure). All provider-call metadata and per-call runtime parameters were archived for audit.

Deterministic validator, scoring rule, and diagnostic trace structure. We used deterministic_netops_validator_v5 (validator_version = 5.0) as a deterministic oracle for the primary binary endpoint. The validator implements a full_intent_constraint_validation rule that: (a) parses the LLM-produced candidate into normalized command sequences, (b) computes parsed_command_count and compares against required_command_count derived from the task's required_sequence, (c) enforces workflow and group constraints (e.g., make-before-break, must_be_down_for_fields), and (d) returns a boolean valid flag plus structured diagnostics (violation_codes, observed_touched_field_count, normalized_configuration, final_state snapshot). Scoring for the confirmatory endpoint used the validator's boolean pass/fail; the richer diagnostics were used for failure-mode analysis and are archived as validator traces for each episode.

Inclusion/exclusion, contamination controls, and pre-execution unit-tests. The preregistered inclusion rules required that an intent pass the deterministic ambiguity detector and that mapping code successfully produce verifier inputs. We executed deterministic unit-tests for mapping code and verifier-input transformations before batch evaluation; these tests passed. Contamination checks (session-leak detection, duplicate-response entropy checks) were run; contamination_summary.csv recorded zero flags. Deterministic reconciliation confirms episode accounting consistency and reports no provider-call failures.

Statistical analysis posture (confirmatory vs exploratory). The methodology adheres to the preregistered multiplicity and missingness rules: the primary confirmatory contrast is guarded vs baseline at α=0.05 (exact McNemar two-sided) using complete_pair_primary (exclude incomplete pairs). Secondary contrasts were preregistered (constrained vs baseline,

guarded vs constrained) with Bonferroni correction; in practice shared_initial_candidate semantics and adapter limits constrained the realized secondary contrasts. Analysis artifacts include paired_contingency_table.csv, condition_summary.csv, and a paired bootstrap used to compute the 95% CI (bootstrap_resamples = 10000, bootstrap_seed = 1729). All analysis code, seeds, and artifacts are archived for reproducibility.

Reproducibility artifacts and archival pointers. The full execution manifest, raw model-call logs, per-call runtime metadata, task-generator seeds, verifier inputs/verdicts, and analysis outputs are archived in the execution bundle referenced by the execution manifest. Deterministic reconciliation and provider accounting artifacts record model_calls_used = 151 and stage-level counts; these artifacts support replication of the paired accounting and the precise sampling semantics used in this trial.

IV. Results

Execution and provider audit. The run completed as planned: completed_episode_count = 300 and complete_pair_count = 150. Provider-call accounting (audited at the provider-call level) shows model_calls_used = 151; completed_hosted_model_call_event_count = 151; failed_hosted_model_call_event_count = 0; cache_miss_count = 151. Stage-level counts recorded by the execution adapter were: shared_initial_generation = 150 and repair = 1. Contamination and mapping checks reported zero flags and no pre-execution unit-test failures.

Primary numerical outcomes and paired contingency. The paired 2×2 contingency (guarded $\times$ baseline) reported by the deterministic analysis executor is: n_11 = 149 (both-pass), n_10 = 1 (guarded-only-pass), n_01 = 0 (baseline-only-pass), n_00 = 0 (both-fail). Per-condition totals: baseline_success_count = 149/150 = 0.9933333333333333 and guarded_success_count = 150/150 = 1.0. The paired absolute difference (guarded - baseline) = 0.00666666666666671 ($\approx$0.67 percentage points). These counts are the basis for the confirmatory tests and are reproduced verbatim from the archived paired contingency artifact.

Statistical tests and bootstrap inference. The preregistered exact two-sided McNemar test on the observed discordant pair counts (n_10 = 1, n_01 = 0) returned p = 1.0 under the exact test implementation used by the paired_binary_analysis_v1 executor. Paired bootstrap percentile 95% confidence intervals were computed with bootstrap_resamples = 10000 and bootstrap_seed = 1729; the 95% CI for the paired absolute difference is [0.00, 0.02]. These analysis parameters and results are recorded in the archived analysis artifacts and analysis log and are reported here without modification.

Repair diagnostics and revision iterations. Consistent with the near-ceiling initial success rate, only a single guarded repair invocation occurred (1 of 150 guarded episodes). That single repair converted one initially failing candidate into a passing configuration under the deterministic validator; all other guarded episodes accepted the shared initial candidate without invoking repair. Summary statistics for revision iterations in guarded episodes: median = 0, IQR = 0. Validator traces for each episode include both validation_before and validation_after snapshots, normalized_configuration strings, parsed_command_count, required_command_count, observed_touched_field_count, and any violation_codes; these traces substantiate the single repair event and the distributional summary above.

Representative per-task diagnostics (artifact-grounded). Representative validator diagnostics for tasks 000001–000006 (selected from the archived representative set) illustrate variation in declared difficulty and concrete parsing outcomes while confirming consistent validator behavior: - task-000001 (declared difficulty: medium): parsed_command_count = 4, required_command_count = 4, observed_touched_field_count = 3, valid = true. - task-000002 (declared difficulty: hard): parsed_command_count = 2, required_command_count = 2, observed_touched_field_count = 2, valid = true. - task-000003 (declared difficulty: hard): parsed_command_count = 6, required_command_count = 6, observed_touched_field_count = 4, valid = true. - task-000004 (declared difficulty: very_hard): parsed_command_count = 4, required_command_count = 4, observed_touched_field_count = 3, valid = true. - task-000005 (declared difficulty: very_hard): parsed_command_count = 3, required_command_count = 3, observed_touched_field_count = 3, valid = true. - task-000006 (declared difficulty: very_hard): parsed_command_count = 6, required_command_count = 6, observed_touched_field_count = 4, valid = true. These per-task metrics and 'valid' flags are drawn directly from archived validator traces for the representative tasks and show that the verifier checked field-level counts, required sequences, and produced normalized configurations used for scoring.

Reproducibility parameters and analysis artifacts. Key analysis seeds and parameters used in inference are published in the artifacts: bootstrap_seed = 1729 and bootstrap_resamples = 10000. The complete paired contingency table, condition summaries, contamination summary (empty), and the full set of validator traces and raw model responses are archived. Provider accounting (model_calls_used = 151; repair invocations = 1) and execution manifest entries are recorded as part of the archived execution bundle to allow direct recomputation of the exact reported analyses.

Missingness, exclusions, and sensitivity checks. The preregistered missingness and exclusion rules were not triggered: failed_hosted_model_call_event_count = 0 and unscorable episodes = 0; no ambiguity-detection exclusions, parser failures, or verifier-unsupported-feature exclusions occurred. Consequently the primary analysis used complete_pair_primary and did not require sensitivity treatments; the alternate preregistered sensitivity (treating missing guarded outputs as failures) was not invoked because there were no missing or failed calls. The analysis artifacts include missingness_summary and contamination_summary CSVs confirming these counts.

Operational interpretation of the numeric outcome. The

observed baseline success rate (≈99.33%) leaves essentially no headroom for the preregistered detectible absolute improvement (target = +15 pp), producing a classical ceiling effect relative to the preregistered power assumptions. Under the shared_initial_candidate semantics used here the baseline measures the validity of the single sampled initial candidate and the guarded contrast isolates the marginal value of at most one repair applied to that identical candidate; hence, a near-perfect initial sample reduces the potential marginal utility of a single-repair guarded pipeline. The single observed repair shows the guarded pipeline can correct rare initial failures, but the aggregate effect on the preregistered estimand is numerically negligible and statistically unsupported (McNemar p = 1.0, bootstrap CI includes zero). These numeric conclusions are strictly artifact-grounded and follow directly from the archived contingency counts and analysis parameters.

Contextual diagnostics for practitioners. Beyond the binary pass/fail metric, the archived validator traces demonstrate additional operational artifacts that a guarded system would produce in deployment: normalized configuration text suitable for human audit, parsed_command_count and required_command_count to quantify task complexity, and violation_codes that explain specific invariant violations. Even when a marginal increase in pass rate is small on a given corpus, these deterministic diagnostics are operationally useful for triage, human review, and auditable gating; the execution bundle contains representative validator traces illustrating these affordances.

V. Discussion

We expand the discussion with concrete, artifact-grounded diagnostics, methodological implications, and operational interpretations that directly follow from the archived execution and analysis artifacts.

Ceiling effect, discordance structure, and inferential consequence. The contingency structure (n_11=149, n_10=1, n_01=0) makes explicit why the exact McNemar test returns p=1.0: almost all paired episodes agree and only a single discordant pair favors the guarded condition. Under McNemar's paired-proportion logic the test's power derives from the number of discordant pairs; with only one discordant pair there is effectively no evidence for a systematic directional effect. The paired bootstrap CI ([0.00,0.02], bootstrap_seed=1729, resamples=10000) quantifies the same intuition by showing that plausible values for the absolute improvement under resampling include zero and are bounded above by ≈2 percentage points. Both diagnostics are reported verbatim from analysis/paired_contingency_table.csv and analysis/condition_summary.csv and therefore trace directly to archived artifacts rather than post hoc interpretation.

Mechanistic explanation via execution semantics. The shared_initial_candidate semantics materially reduce the pathway by which the guarded pipeline can improve outcomes: because the identical initial candidate is used to evaluate both baseline and guarded conditions, any potential advantage of the guarded pipeline depends entirely on the single allowed repair call converting that specific initial candidate into a verifier-pass. This is reflected in provider accounting: stage_counts show shared_initial_generation=150 and repair=1, and model_calls_used=151. Practically, when baseline initial candidates are already highly likely to pass the verifier (here ≈99.33%), the repair budget is rarely exercised and the guarded condition has almost no opportunity to improve aggregate pass rates. This mechanism explains the observed near-ceiling result and illustrates how sampling semantics and repair budgets interact with baseline competence to determine observable effect sizes.

Design consequences for preregistration and power calibration. The preregistration targeted detection of a 15 percentage-point absolute improvement with N=150. That calculation assumed substantially lower baseline competence and therefore a larger expected discordant mass. Our artifact-grounded outcome demonstrates the risk of committing to a confirmatory N without an empirical baseline check: baseline calibration reduces the chance of an underpowered or effectively vacuous confirmatory test caused by ceiling effects. Concretely, a practical pilot procedure (consistent with our recommendations and implementable from the archived deterministic task generator) is: sample 20–40 tasks from the planned generator with the intended prompt styles; measure baseline verifier-pass rate on those pilot tasks; if baseline >~85–90% either increase task difficulty or enlarge N and/or adopt independent-sampling semantics or a larger repair budget. These procedural steps require no new experimental artifacts and can be implemented with the provided generator seeds and unit-tested mapping code archived in the execution bundle.

Operational affordances beyond the binary endpoint. Even when aggregate verifier-pass improvements are negligible, the guarded pipeline provides operationally valuable artifacts that the binary pass/fail metric does not capture. The deterministic validator returns structured diagnostics (normalized_configuration strings, parsed_command_count, required_command_count, violation_codes) for each episode; these are visible in validator traces for representative tasks (e.g., task-000001..000006). In practice these strings enable: (1) auditable diffs for operator review, (2) concise failure taxonomies to prioritize human triage, and (3) deterministic gating rules implementable as fail-stop or auto-remediation policies. The run-level audit (completed_hosted_model_call_event_count=151; failed_hosted_model_call_event_count=0) proves these artifacts were produced consistently across the sample and are therefore operationally available in the archived bundle for downstream tooling.

Trade-offs among sampling semantics, repair budgets, and estimands. Our evidence clarifies three non-identical estimands that practitioners must choose between and that should be explicit in preregistrations: (A) shared_initial_candidate marginal-repair effect (our estimand), (B) independent-sampling expected-performance difference (draw independent initial samples per condition), and (C) repair-robustness with expanded repair budgets (allow multiple iterative repairs).

Estimand A is conservative with respect to repair effectiveness because it fixes the initial draw; estimand B typically yields larger observable differences when model output variance is high; estimand C quantifies diminishing returns from additional repair attempts. The archived artifacts make it possible to re-run variants B or C (subject to provider budget and adapter constraints) starting from the same deterministic task generator seeds and the stored model_configuration metadata; this preserves reproducibility while acknowledging that different, equally defensible operational choices lead to different scientific answers.

Reproducibility notes tied to concrete artifact fields. Key numeric analysis parameters are archived and should be used by any reproducer: bootstrap_resamples=10000 and bootstrap_seed=1729 (analysis/analysis_log.jsonl), model_calls_used=151 and stage_counts (analysis/deterministic_reconciliation.json), and declared_model_version="gpt-5-mini" with model_configuration.json recording temperature=null (unset). Because temperature was unset and no fixed random seed was enforced, repeated independent provider sessions can produce nondeterministic outputs; reproductions seeking exact token-level matches must therefore control sampling parameters or accept equivalent statistical replications rather than byte-identical replays. All of these facts are documented in the execution manifest and model_configuration.json and are cited here to avoid ambiguity about what archive-grounded reproduction entails.

Threats to validity revisited with artifact grounding. Internal validity: the adapter's session-isolation and provider-call audit (no cross-condition cache hits observed) reduce concerns about prompt-history leakage; these checks are recorded in deterministic_reconciliation.json. Construct validity: deterministic validator pass/fail reflects the predeclared invariants from the generator's task_payload; mapping-code unit-tests passed pre-execution but any gaps in verifier coverage remain a construct threat—this is why we limited claims to the verifier-capable invariants and reported representative validator violation fields when present. External validity: the single-model constraint (gpt-5-mini) and small-topology synthetic corpus restrict generalization; these are enumerated in the preregistration and in the limitations section and are supported by the model_configuration and task-manifest artifacts. Conclusion validity: the mismatch between preregistered power assumptions and the observed baseline is a concrete illustration of how inferential conclusions depend on realized effect sizes and discordant-pair counts, both archived in the paired_contingency_table.csv.

Practical recommendations for practitioners and experimenters. (1) Always run a short baseline-calibration pilot using the final prompt templates and generator to estimate baseline competence and select N, difficulty, and repair budgets accordingly. (2) Make generation semantics explicit in preregistration: shared_initial_candidate vs independent sampling lead to different operational decisions and must match intended deployment semantics. (3) When verifier coverage is partial, report both binary pass/fail and the structured validator diagnostics to allow operators to triage and to enable metric-based comparisons across verifier designs. (4) Consider multi-arm designs that vary repair budgets and sampling semantics to estimate diminishing returns and to align statistical power with expected discordance rates. The provided execution bundle contains deterministic generator seeds, unit-tested mapping code, and validator traces sufficient to implement each recommendation without inventing new artifacts.

Summary. The expanded discussion ties our artifact-grounded numerical outcomes to concrete methodological mechanisms (sampling semantics and repair budget), diagnostic artifacts (validator traces and audit counts), reproducibility parameters (bootstrap seed, model metadata), and explicit experimental recommendations. These points are strictly supported by the archived execution and analysis artifacts and do not alter empirical claims reported elsewhere in the manuscript.

VI. Limitations

- Single-model constraint: experiments used only gpt-5-mini; results do not generalize across models or sizes.
- Synthetic corpus and scale: tasks were programmatically generated and limited to verifier-capable toy-to-small topologies (4–32 nodes); external validity to production WANs and heterogeneous vendor CLIs is untested.
- Shared_initial_candidate semantics and execution contract limited repair budget to at most one guarded repair call per intent; different generation semantics or multiple-repair budgets may yield different outcomes.
- Ceiling effect and preregistration mismatch: baseline correctness was far higher than the pilot assumptions used for power calculations (planned detectable effect = 15 percentage points), reducing the trial's ability to demonstrate benefit and illustrating sensitivity to baseline calibration.
- Verifier coverage and specification translation: the study measures deterministic validator pass/fail on the preregistered invariants but does not claim safety guarantees beyond the deterministic validator's modeled properties.
- Security and adversarial robustness: the experiment did not evaluate adversarial prompt-injection or adversary-driven intents; separate targeted studies are required to assess claim-to-action vulnerabilities in agentic NetOps workflows.

VII. Conclusion

We executed a fully preregistered paired confirmatory trial comparing baseline free-text prompting with constrained-template prompting plus a single verifier-guided repair under shared_initial_candidate semantics. On a deterministic synthetic corpus with gpt-5-mini and deterministic_netops_validator_v5, the guarded pipeline increased verifier pass rate from 149/150 to 150/150 ($\approx$0.67 pp), a numerically tiny and statistically non-significant difference (exact McNemar p = 1.0; 95% bootstrap CI = [0.00, 0.02]). The

principal scientific lesson is methodological: preregistered NetOps trials should include baseline calibration to avoid ceiling effects and preserve statistical power. Technically, our artifacts bound the marginal benefit of a one-repair guarded pipeline under shared_initial_candidate semantics; evaluating independent-sampling semantics, larger repair budgets, model diversity, and production-scale topologies remains important future work.

Disclosure Statement

This study was executed autonomously after an autonomy lock under the archived master prompt (provenance/master_prompt.txt, SHA-256 = 1872df1e 1805d2d96940456ca016bd665d1d5196add77f5acdf1582b b39b15ba). Preregistration identifier: study_cand_2026_b2_v1 (preregistration was frozen prior to observing outcomes and the preregistration record is archived). Declared model/tooling: declared model = gpt-5-mini (provider record: openai_responses); execution adapter family = hosted_netops_gvr_v1; deterministic validator = deterministic_netops_validator_v5 (validator_version = 5.0). Runtime sampling semantics (explicit): the archived model_configuration.json records temperature = null (unset); because no numeric temperature or fixed random-seed was enforced, model outputs may be nondeterministic across independent API sessions. Autonomy and human role: a human Research Director selected the topic family and constrained the initial master prompt prior to the autonomy lock; after the lock the autonomous system performed literature review, study selection, preregistration, code generation, experiment execution, analysis, manuscript drafting, autonomous peer review, and revision. No human manually edited technical content, analyses, figures, captions, or references after the autonomy lock. Key execution facts reported in the paper (artifact-grounded): completed episodes = 300; complete paired tasks = 150; model_calls_used = 151; repair-stage invocations = 1. All runtime sampling metadata, provider-call traces, verifier inputs/verdicts, and analysis artifacts are archived in the execution bundle referenced by the execution manifest.

References

[1] https://doi.org/10.1002/nem.2313
[2] https://doi.org/10.1145/3786583.3786898
[3] https://arxiv.org/abs/2604.22513
[4] https://doi.org/10.1145/3605764.3623985
[5] https://doi.org/10.1145/3656454
[6] https://doi.org/10.1002/widm.70111