

Can Deterministic Network Verification Reduce Undetected Faults in LLM-Generated Configurations?

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory study (study id c1-gnr-vv-2026-0001) to measure whether inserting a deterministic network verifier between an LLM intent→configuration generator and an emulated execution reduces the proportion of configurations that would execute with operational faults. Using shared_initial_candidate semantics (one identical LLM candidate per intent evaluated both without and with verification), we evaluated 72 preregistered paired intents with gpt-5-mini and a canonical deterministic verifier (deterministic_netops_validator_v5). Execution completed with 144 episodes (72 pairs) and 72 model calls. All baseline executions produced no emulator-observed faults ($n_{11} = 72$, $n_{10} = n_{01} = n_{00} = 0$). The paired difference in undetected-fault proportion is 0.0 (paired bootstrap 95% CI [0.0, 0.0], exact McNemar two-sided $p = 1.0$; bootstrap seed = 1729, resamples = 10000). Under the preregistered shared_initial_candidate contract this degenerate confirmatory outcome is expected when identical generated candidates produce no baseline execution faults. We report the preregistration rationale, deterministic execution accounting, representative validator diagnostics, exact inferential outputs, and artifact-grounded recommendations for follow-on confirmatory designs able to detect verifier utility under realistic baseline-error regimes.

I. INTRODUCTION

Generative large language models (LLMs) are increasingly proposed to translate high-level operator intents into device-level network configurations. This intent→configuration automation can speed change pipelines but risks silently violating reachability, access-control, ordering, or workflow invariants and thereby causing outages if generated text is executed without additional checks. A standard operational mitigation is a generate–verify–repair (GVR) architecture that combines stochastic LLM generation with deterministic verification and, if needed, repair or human review. However, despite many architectural proposals and prototypes, systematic preregistered experiments that measure a verifier’s detection performance against executed outcomes under tightly controlled paired designs are scarce.

This paper answers a narrowly scoped, preregistered research question: does inserting a deterministic network verifier between an LLM generator and emulator execution materially reduce the proportion of configurations that execute with operational faults when the identical generated candidate is evaluated both without and with verification? That estimand isolates verifier detection from improvements that could arise from re-sampling, prompt-engineering, or repair-enabled iterations. We executed study c1-gnr-vv-2026-0001

using gpt-5-mini and deterministic_netops_validator_v5 across 72 preregistered paired intents under shared_initial_candidate semantics. The run completed with no technical failures, produced a degenerate confirmatory outcome (no baseline emulator faults), and yields a paired difference of 0.0 (95% CI [0.0,0.0], exact McNemar $p = 1.0$). Rather than treating this as a null failure, we analyze why it occurred given the preregistered contract and archived artifacts, and we provide concrete, artifact-grounded guidance for designs that can detect verifier utility when baseline error prevalence is non-negligible.

II. RELATED WORK

Work on LLM-driven NetOps spans intent→configuration translation, multi-agent/tool-augmented orchestration, and architectural proposals for combining stochastic generation with deterministic checks. Prior empirical studies demonstrate that modern LLMs can produce device-level configurations from high-level intent in constrained vendor-dialect settings and curated evaluation suites; these results show promise but are typically scoped to limited task families and curated datasets rather than broad production traces [1]. That line of work motivates leveraging LLMs to reduce operator effort while highlighting the need for safety controls when configuration semantics and ordering matter.

Complementing single-agent generation work, multi-agent and tool-augmented frameworks have been proposed that explicitly integrate verifiers, tool calls, or specialized modules into LLM-driven NetOps pipelines to reduce regressions and coordinate multi-step changes [2]. These systems illustrate practical architectures for inserting deterministic checks or domain services into agentic workflows, and their evaluations commonly rely on emulators, curated scenarios, or prototype deployments. Such prototypes suggest feasible integration points for verifiers but do not, in the supplied records, provide preregistered paired measurements of verifier detection performance against executed faults.

A related and more general strand of work articulates and evaluates generate–verify–repair (GVR) patterns outside NetOps—particularly in code- and software-repair domains—where stochastic generators produce candidates that deterministic checkers validate and, when necessary, trigger repair iterations [3]. These studies document how deterministic validation can convert nondeterministic generation into a correctness-convergent workflow; they also show that the empirical util-

ity of verification depends strongly on the baseline error rate of generator outputs and on whether the architecture permits verifier-triggered repairs or resampling. Translating these insights into NetOps requires coupling LLM outputs to domain-appropriate verifiers and measuring detection relative to executed outcomes.

Canonical deterministic network-verification techniques—exemplified by header-space analysis and related formal reachability/policy analyzers—provide precise, semantics-aware invariants for reachability and ACL-style policy checks and are natural candidates for the ‘verify’ stage in GVR architectures [4]. Such verifiers can deterministically flag violations of specified invariants or workflow constraints expressed over packet headers, interface states, and rule orderings. However, the literature lacks, among the supplied records, preregistered paired experiments that measure how often these canonical verifiers would have detected faults that would otherwise manifest when LLM-generated configurations are executed in an emulator or live device.

Taken together, the verified literature establishes three pertinent points: (i) LLMs can produce plausible device-level configs from intent in constrained settings [1]; (ii) system proposals and prototypes show how verifiers and tools can be integrated into multi-agent NetOps workflows [2]; and (iii) GVR patterns can yield empirical correctness improvements in other domains when baseline generator errors are non-negligible and when verifier-triggered repairs or resampling are allowed [3]. What remains underexplored in the supplied evidence is a preregistered, paired measurement that isolates verifier detection on identical candidates—i.e., does a canonical deterministic verifier reduce the probability of an execution fault for the same LLM-generated candidate when the candidate is executed with and without prior verification? Our study directly targets this measurement gap by adopting a shared_initial_candidate paired design and reporting exact execution-accounting artifacts to make that detection question reproducible and auditable.

III. METHODOLOGY

Preregistration and estimand. We preregistered study c1-gnr-vv-2026-0001 and froze the design and analysis prior to observing outcomes. The primary estimand is the paired reduction in undetected operational-fault proportion (paired_success_rate_difference_guarded_minus_baseline). The confirmatory hypothesis H1 targeted an absolute paired reduction of 0.20 with two-sided $\alpha = 0.05$ and power ≥ 0.80 ; a sample-size heuristic returned ≈ 63 pairs and we preregistered $N = 72$ pairs (24 per topology-difficulty stratum) to allow margin for deterministic exclusions.

Shared_initial_candidate semantics and experimental contract. The execution_contract enforced shared_initial_candidate semantics: for each intent exactly one initial LLM generation was produced and that identical candidate was evaluated in two conditions. Baseline: execute the shared candidate in an isolated emulator and record emulator fault/no-fault. Guarded: run the canonical

deterministic verifier (deterministic_netops_validator_v5) on the same candidate; if the verifier flags violations the guarded outcome is recorded as prevented (no execution); if it does not flag, execute the same candidate in the emulator and record the emulation outcome. This paired structure ensures any observed paired difference arises solely from verifier detection preventing an otherwise-executed fault.

Model, adapter, and verifier configuration. The declared model in execution/model_configuration.json is gpt-5-mini (provider = openai_responses, max_output_tokens = 2000, maximum_attempts_per_call = 1). execution/model_configuration.json records temperature = null; this indicates provider-default runtime sampling semantics and is archived as a residual sampling-parameter uncertainty. The adapter family was hosted_netops_gvr_v1. The canonical verifier is deterministic_netops_validator_v5 (validator_version = "5.0") using a full_intent_constraint_validation rule set that outputs per-episode validity verdicts and structured validator_traces (parsed_command_count, required_command_count, validation_before/after states, violation_codes, difficulty metadata). All verifier decisions and traces were archived per episode under execution/scoring/.

Task-bank construction, contamination controls, and stratification. Tasks were deterministically generated by deterministic_netops_task_generator_v5. Each task manifest entry encodes a natural-language intent, a machine-structured required_sequence (ordered field changes), allowed_touched_fields, initial and expected final states, and difficulty labels (medium/hard/very_hard). The preregistration required deterministic exact-match contamination checks against public corpora with explicit exclusion rules. analysis/contamination_summary.csv records zero exact-match exclusions for confirmatory pairs. The confirmatory sample retained the preregistered stratification (24 pairs per stratum).

Execution protocol, retries, and deterministic accounting. The missingness_plan allowed up to two additional retries (three attempts total) for transient hosted-API errors on generation calls and required full logging of retries and provider events. The run started at 2026-08-31T06:56:23.065707Z and completed at 2026-08-31T07:06:28.665518Z (≈ 10.1 minutes wall-clock). The execution manifest reports planned_episode_count = 144, completed_episode_count = 144, failed_episode_count = 0, and model_calls_used = 72 (one shared generation per pair). Provider-call audit (deterministic_reconciliation.json) shows hosted_model_call_event_count = 72, completed_hosted_model_call_event_count = 72, cache_miss_count = 72, and stage_counts = {shared_initial_generation: 72}.

Scoring, validation, and inferential procedures. Each episode’s verifier emits a score and validator_trace and the primary analysis used paired_binary_analysis_v1 to form the 2×2 paired contingency table (guarded $\times$ baseline). The preregistered inferential test is an exact McNemar two-sided test

when discordant pairs exist. We computed paired bootstrap-percentile 95% confidence intervals for the paired difference using 10,000 resamples with `bootstrap_seed = 1729` (analysis/analysis_log.jsonl). Pre-specified subgroup confirmatory comparisons (topology complexity and policy type) were registered with Holm correction; these are archived but become non-informative in the absence of discordant pairs. All analysis was executed deterministically by the registered analysis executor and archived under analysis/.

Additional methodological notes (artifact-grounded). The validator’s `full_intent_constraint_validation` rule set enforces sequence and field constraints encoded in the task manifests; validator traces include `parsed_command_count` and `required_command_count` and a `difficulty.level` tag (examples below). Per-episode JSON scoring artifacts under `execution/scoring/` record `normalized_configuration`, `validation_before/after` final_state snapshots, and `violation_codes` arrays that were systematically inspected during analysis.

IV. RESULTS

Execution and manifest summary. The autonomous pipeline completed without technical failures and with full deterministic accounting. Key run-level figures from the archived manifests are: `planned_episode_count = 144`; `completed_episode_count = 144`; `failed_episode_count = 0`; `model_calls_used = 72`; `artifact_count = 510`. The run-level timestamps (`started_at_utc = 2026-08-31T06:56:23.065707+00:00`; `completed_at_utc = 2026-08-31T07:06:28.665518+00:00`) indicate a wall-clock duration of ≈ 10.1 minutes. `analysis/condition_summary.csv` (archived) records `baseline_successes = 72`, `baseline_failures = 0` (`success_rate = 1.0`) and `guarded_successes = 72`, `guarded_failures = 0` (`success_rate = 1.0`).

Paired contingency, inferential outputs, and bootstrap diagnostics. Aggregating across the 72 confirmatory pairs produced the exact paired contingency counts `n_11 = 72`, `n_10 = 0`, `n_01 = 0`, `n_00 = 0` (`guarded  $\times$  baseline`). The primary estimand `paired_success_rate_difference_guarded_minus_baseline = 0.0`. The preregistered exact McNemar two-sided test produces `p = 1.0` given the absence of discordant pairs. The paired bootstrap-percentile 95% confidence interval for the paired difference (`bootstrap_seed = 1729`; `bootstrap_resamples = 10000`; analysis/analysis_log.jsonl) is `[0.0, 0.0]`, reflecting the degenerate sampling distribution when all paired outcomes are concordant. `analysis/paired_contingency_table.csv` and `analysis/condition_summary.csv` contain these tabulations and are archived.

Per-episode validator diagnostics and representative exemplars. We inspected archived per-episode scoring JSONs under `execution/scoring/` and representative task-manifest entries under manuscript evidence. Across the sampled exemplars the verifier produced no violation flags. Representative excerpts from archived scoring artifacts: - pair-000001 (task-000001): `difficulty.level = "medium"`; `parsed_command_count = 4`; `required_command_count = 4`; `validator_version = "5.0"`;

`violation_count = 0`; `normalized_configuration` equals the shared candidate (`execution/scoring/task-000001-*.json`). - pair-000002 (task-000002): `difficulty.level = "hard"`; `parsed_command_count = 2`; `required_command_count = 2`; `violation_count = 0`. - pair-000003 (task-000003): `difficulty.level = "hard"`; `parsed_command_count = 6`; `required_command_count = 6`; `violation_count = 0`. For additional archived representative tasks the task manifests encode `required_command_count` and `difficulty` tags (examples: task-000004 `required_command_count = 4`, `difficulty.level = "very_hard"`; task-000005 `required_command_count = 3`, `difficulty.level = "very_hard"`; task-000006 `required_command_count = 6`, `difficulty.level = "very_hard"`). In all archived scoring JSONs for confirmatory pairs the `verifier_traces` report `violation_count = 0` and `validation_after.valid = true`, demonstrating that the verifier processed each shared candidate and did not flag the candidate as violating the manifest-encoded constraints.

Contamination, missingness, and sensitivity analyses. `analysis/contamination_summary.csv` records zero exact-match contamination flags for confirmatory pairs. `analysis/missingness_summary.csv` records `complete_pairs = 72` and zeros for failed or unscorable episodes. The preregistered sensitivity analyses (including the conservative treatment that would count excluded pairs as failures and bootstrap diagnostics) were executed by the registered analysis executor; because there were no exclusions and no discordant pairs these sensitivity analyses reproduce the degenerate numeric conclusion (`paired_difference = 0.0`). All analysis artifacts (`analysis/paired_contingency_table.csv`, `analysis/condition_summary.csv`, analysis/analysis_log.jsonl) are archived and reference the bootstrap seed and resample count used.

Deterministic reconciliation and provider-call audit. The `deterministic_reconciliation.json` confirms accounting consistency: `completed_episode_count = 144`, `complete_pair_count = 72`, `baseline_success_count = 72`, `guarded_success_count = 72`, `n_11 = 72`, `n_10 = 0`, `n_01 = 0`, `n_00 = 0`, and `all_deterministic_consistency_checks_passed = true`. Provider-call audit fields show `hosted_model_call_event_count = 72`, `completed_hosted_model_call_event_count = 72`, `failed_hosted_model_call_event_count = 0`, `cache_miss_count = 72`, and `stage_counts = {shared_initial_generation: 72}`, consistent with the preregistered `shared_initial_candidate` semantics and the execution manifest.

Interpretation of the degenerate outcome (evidence-grounded). The complete absence of baseline emulator faults in the archived confirmatory corpus is the proximate cause of the degenerate paired result. The archived artifacts together document measurable contributors: (1) the deterministic task generator produced constrained intents with manifest-encoded `required_sequence` and `allowed_touched_fields` that align the generation objective tightly with verifier checks; (2) gpt-5-mini (declared in `execution/model_configuration.json`) produced configurations that matched those manifest constraints across the sampled

tasks; and (3) the `deterministic_netops_validator_v5` rule set (`full_intent_constraint_validation`) enforces the same manifest-encoded sequence and field constraints, producing concordant validity judgments. These archived, machine-verifiable facts explain why the verifier could not reduce executed faults under the `shared_initial_candidate` contract: there were no executed faults to detect.

V. DISCUSSION

Confirmatory interpretation and boundary conditions. The preregistered paired design with `shared_initial_candidate` semantics validly measures verifier detection on identical candidates; because the identical generated candidates produced zero emulator faults the verified paired outcome is a ceiling/null result: the verifier could not reduce executed faults that did not occur. The numeric outputs ($n_{11} = 72$; paired difference = 0.0; McNemar $p = 1.0$; bootstrap CI [0.0,0.0]) are direct consequences of the preregistered contract and observed data and therefore correctly reported as a degenerate confirmatory result.

Artifact-grounded causes of the ceiling outcome. Archived artifacts allow us to identify measurable factors that explain the ceiling outcome: (1) the deterministic task generator produced constrained, high-clarity intents with explicit `required_sequence` fields and `allowed_touched_fields` encoded in each task manifest, increasing the likelihood of unambiguous correct candidates; (2) `gpt-5-mini` (`declared_model_version` in `execution/model_configuration.json`) generated sequences that matched those structured constraints across the synthetic corpus; and (3) the verifier rule set (`full_intent_constraint_validation` in `deterministic_netops_validator_v5`) aligns closely with the `required_sequence` constraints present in the task manifests, making verification and generation objectives strongly congruent. These archived, measurable factors together account for the absence of baseline emulator faults.

Design implications: when verification benefit is measurable. To empirically measure practical verifier benefits, confirmatory contracts must present non-negligible baseline error prevalence or allow verifier-mediated candidate changes. Artifact-grounded, preregisterable alternatives include: (A) independent-generation arms—allow separate initial generations per condition so verifier-triggered repairs or alternative sampling may change executed candidate distributions; (B) repair-enabled flows—permit a deterministic repair call when the verifier flags violations and execute repaired candidates to measure end-to-end improvement; (C) adversarial/fault-injection task banks—seed tasks with ambiguous intents, ordering subtleties, or vendor-edge cases to raise baseline error prevalence; (D) multi-model and sampling sweeps—include multiple model versions and explicit sampling parameter settings (non-null temperatures) to quantify model- and sampling-sensitivity. All these options are compatible with the archived execution and analysis infrastructure and can be preregistered using the same evidence-recording conventions.

Threats to validity revisited (artifact-supported). Internal validity: by design the `shared_initial_candidate` semantics isolates detection but prevents verifier-mediated candidate resampling from influencing outcomes; this tradeoff explains why a degenerate result may occur when baseline error prevalence is low. Construct validity: emulator-observed faults are a conservative proxy for operational harm but do not capture traffic-dependent timing, vendor-specific hardware idiosyncrasies, or live-traffic emergent failures. External validity: experiments used a single declared model (`gpt-5-mini`), a synthetic task bank, and one deterministic verifier; extrapolation to other models, vendor dialects, or production hardware is limited. Conclusion validity: because baseline error prevalence was zero in this corpus the confirmatory design had no power to detect verifier benefit; the preregistered power target (detect absolute paired reduction 0.20) becomes moot under a zero-prevalence baseline.

Operational implications and measured tradeoffs. The confirmed ceiling outcome does not imply verifiers are unnecessary; rather, it shows that under certain tightly constrained synthesis and evaluation conditions an LLM alone produced valid device sequences on the synthetic corpus. In production settings with ambiguous intents, incomplete constraints, or vendor heterogeneity, verifiers remain a mechanistic guard. Measuring verifier costs (false positives that block safe changes) and benefits (true positives preventing faults) requires task banks with non-negligible baseline failure prevalence; the archived artifacts and preregistration provide a reproducible template for such follow-on evaluations.

Reproducibility notes (artifact pointers and deterministic checks). All execution and analysis artifacts are archived in the evidence bundle. Key artifacts referenced in this paper include: `execution/raw_results.jsonl`, `execution/task_manifest.jsonl`, `execution/model_configuration.json`, `analysis/paired_contingency_table.csv`, `analysis/condition_summary.csv`, `analysis/analysis_log.jsonl`, `deterministic_reconciliation.json`, and `per-episode_validator` traces under `execution/scoring/`. The analysis bootstrap used seed = 1729 and 10,000 resamples; `analysis/analysis_log.jsonl` records these values. The run-level timing, provider-call audit, and artifact counts are recorded in `execution/` and `analysis/` manifests to enable deterministic reconciliation of the paired accounting. The model configuration records `temperature = null` (`execution/model_configuration.json`), which we report as provider-default sampling semantics (a residual sampling-parameter uncertainty archived in the evidence).

VI. LIMITATIONS

- Confirmatory ceiling/null outcome: the baseline generator produced zero emulator faults on the preregistered task set, preventing direct measurement of verifier detection benefit.
- `Shared_initial_candidate` semantics: the design measures verifier effect on the same generated candidate only and does not assess independent per-condition generation,

multi-sample generation, or repairs unless explicitly permitted by the contract.

- Single-model scope: experiments used one declared model (gpt-5-mini); results do not imply generality across models or versions.
- Synthetic/emulated tasks: task corpus and emulator executions differ from production devices and live traffic; external validity to production deployments is limited.
- Sampling-parameter uncertainty: execution/model_configuration.json shows temperature = null (sampling parameters unset); null indicates provider-default runtime sampling semantics and is a residual uncertainty recorded in the archived configuration.
- Residual contamination risk: deterministic provenance and exact-match checks were applied, but private training-data contamination of the hosted model cannot be fully ruled out and remains residual uncertainty.

VII. CONCLUSION

We executed a preregistered paired evaluation (c1-gnr-vv-2026-0001) under shared_initial_candidate semantics to test whether a canonical deterministic verifier reduces the proportion of LLM-generated configurations that would execute with operational faults. For the chosen model (gpt-5-mini), verifier (deterministic_netops_validator_v5), and synthetic/emulated task bank, baseline executions produced zero emulator faults and the verifier did not change outcomes (paired difference = 0.0; bootstrap 95% CI [0.0,0.0]; exact McNemar $p = 1.0$). This artifact-grounded null/ceiling outcome is internally consistent with the preregistered design: when identical generated candidates produce no executed faults a verifier cannot reduce executed faults under the shared_initial_candidate contract. Measuring verifier utility under realistic error regimes requires preregistered contracts that permit independent or multiple generator samples, repair-enabled flows, adversarial task sets, and multi-model comparisons. The archived artifacts (responses, task manifests, validator traces, execution manifest, and analysis logs) provide a reproducible baseline and a template for those follow-on confirmatory designs and power calculations.

DISCLOSURE STATEMENT

Autonomy and artifacts: This study was executed autonomously after an autonomy lock. Human role: a Research Director selected the topic family and approved the immutable initial master prompt prior to the autonomy lock; no human manually edited technical manuscript content after the lock. Models and tools: initial generation used gpt-5-mini (declared_model_version recorded in execution/model_configuration.json) orchestrated via the hosted_netops_gvr_v1 adapter; deterministic_netops_validator_v5 (validator_version = "5.0") served as the canonical verifier; an isolated emulator executed candidate configurations. Preregistration: study c1-gnr-vv-2026-0001 was preregistered and frozen before observing outcomes. Immutable master

prompt reference: provenance/master_prompt.txt (SHA-256: 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ba). Key archived artifacts include execution/raw_results.jsonl, execution/task_manifest.jsonl, execution/model_configuration.json, analysis/paired_contingency_table.csv, analysis/condition_summary.csv, deterministic_reconciliation.json, and per-episode validator traces under execution/scoring/. The model configuration records temperature = null (provider-default sampling semantics archived in execution/model_configuration.json). Provider-call audit: hosted_model_call_event_count = 72. No human manual manuscript editing or content insertion occurred after the autonomy lock.

REFERENCES

- [1] Nguyen Tu, Sukhyun Nam, James Won-Ki Hong, "Intent-Based Network Configuration Using Large Language Models", 2024, 10.1002/nem.2313
- [2] Zhaodong Wang, Samuel Lin, Guanqing Yan, Soudeh Ghorbani, Minlan Yu, Jiawei Zhou, Nathan Hu, Lopa Baruah, Sam Peters, Srikanth Kamath, Jian Yang, Ying Zhang, "Intent-Driven Network Management with Multi-Agent LLMs: The Confucius Framework", 2025, 10.1145/3718958.3750537
- [3] Rafael Cadenas, "Convergent Correctness in Stochastic Code Generation: A Generate-Verify-Repair Architecture for Deterministic Validation of Autonomous AI Coding Agents in Regulated Environments", 2026, 10.2139/ssrn.6754899
- [4] <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.300.8659>