

Measuring Validator Utility in Generate–Verify–Repair NetOps Pipelines: a Confirmatory Paired Study with Shared Initial Candidates

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory study (AC-2026-03) to evaluate whether a multidimensional validator metric suite predicts end-to-end agentic NetOps success better than a single verifier pass rate. Using a synthetic intent→configuration generator and a hosted generate–validate–repair (GVR) adapter under shared_initial_candidate semantics, we ran 200 paired tasks (one shared initial generation per task reused for baseline and guarded). The deterministic validator (deterministic_netops_validator_v5, v5.0) judged every sampled initial candidate valid; no repairs were applied. Baseline = 200/200, guarded = 200/200, paired difference = 0.00 (bootstrap 95% CI [0.00, 0.00], exact McNemar two-sided $p = 1.0$; bootstrap_seed = 1729, resamples = 10000). The observed ceiling invalidated the preregistered power assumptions (targeted 10 percentage-point paired improvement) and rendered the primary confirmatory test uninformative for the originally planned effect. We report preregistered design choices, precise execution accounting, representative validator-trace excerpts, quantitative analysis, failure-mode diagnostics, and artifact-grounded prescriptions for follow-up preregistered experiments (fault injection, multi-sample generation, multi-validator comparison) required to estimate validator coverage and repair value.

I. INTRODUCTION

Automating NetOps workflows with generative models requires reliable mechanisms to avoid harmful, silent, or wrong-but-plausible actions. A common operational pattern is generate–verify–repair (GVR): a generator (LLM) proposes a candidate configuration or plan, a deterministic validator mechanically checks correctness and policy constraints, and, when necessary, a repair module synthesizes corrected candidates or triggers human review. Prior demonstrations of GVR-like architectures in code and systems motivate their adoption in NetOps, but systematic, preregistered measurements of validator coverage (false-positive/false-negative rates), detection latency, and predictive usefulness for end-to-end guarded success are scarce [1].

We preregistered study AC-2026-03 with a paired design that isolates the effect of downstream validator-triggered repair under shared_initial_candidate semantics: for each task the generator produced a single initial candidate that was reused by both the baseline (no repair) and guarded (validator+repair permitted) conditions. Under these semantics any paired difference can only arise down-

stream from validator-triggered repair, not from independent generator sampling. The primary estimand was the paired_success_rate_difference_guarded_minus_baseline and the preregistered confirmatory hypothesis (H1) expected that a multidimensional validator-metric suite would explain more variance in end-to-end success than a single verifier pass rate. Our main contribution is a fully preregistered, artifact-archived experiment and an exact, transparent report of a degenerate confirmatory outcome (ceiling) with concrete, artifact-grounded prescriptions for follow-up designs that can measure validator utility.

II. RELATED WORK

This section synthesizes prior verified work that motivates our measurement focus and clarifies limits relevant to validator-metric evaluation in NetOps.

Generate–verify–repair (GVR) architectures and deterministic validators. Several recent works articulate and prototype GVR-style architectures that interpose mechanistic checks between stochastic generation and acceptance. Cadenas [1] formalizes the generate–verify–repair pattern in the context of code-generation for regulated environments and argues for deterministic acceptance criteria; independently, recent arXiv/system reports describe self-healing orchestrators and verifier-guided recovery in tool-augmented LLM systems [4]. These records consistently demonstrate that verifier layers can detect mechanically checkable defects (compilation failures, structural constraint violations) and can provide structured repair feedback to generators. However, the cited demonstrations focus on defect classes that are mechanically testable and do not establish comprehensive empirical coverage for complex NetOps semantics; as a result, these works motivate architecture adoption while leaving open quantification of validator false-negative coverage on semantic or intent-alignment failures.

Benchmarks and operational failure modes. Controlled benchmark efforts in AIOps/NetOps document substantive, operationally meaningful failure modes that complicate end-to-end automation. OpsEval and Fault-Bench-style suites highlight sensitivity to noisy or unreliable user tickets, wrong premises, and wrong-but-plausible outputs that may trigger

incorrect actions [2],[3]. These benchmark records show that end-to-end agentic success depends on (a) the generator’s factual and procedural correctness, (b) the validator’s capacity to detect the relevant defect classes, and (c) the repair mechanism’s ability to produce acceptable candidates. Crucially, benchmark findings imply that verifier effectiveness is conditional on validator coverage relative to the failure modes present in the task bank; controlled improvements reported in prior work emerge when validation coverage aligns with the dominant mechanical defect classes in the evaluated tasks, not necessarily when semantic intent mismatches dominate.

Measured benefits and demonstrated limits of verifier-augmented strategies. System reports and experimental prototypes indicate that tool-augmented orchestration (verifier-guided repair, auditable decision traces, self-healing orchestrators) reduces silent failures and improves guarded success in those controlled contexts where validators detect actionable faults [4]. The archived literature and system reports support an important qualification: verifier-driven gains depend strongly on validator completeness for the defect classes that actually occur. When validators are narrowly focused on syntactic or structural checks, they will miss correct-by-structure but incorrect-by-intent outputs; when validators are broad they may increase false positives and operational triage costs. Our experiment and artifact record treat these prior demonstrations as architectural motivation and emphasize that a standardized validator-metric suite (coverage, FP/FN rates, detection latency, operational-impact proxy) is required to quantify these trade-offs in NetOps GVR pipelines [1],[4].

Design and measurement gaps in prior work. The verified literature converges on two practical gaps motivating this study. First, while GVR prototypes show the pattern’s usefulness for mechanically testable errors, there is no widely adopted, empirically estimable validator-metric suite tailored to NetOps GVR flows that reports FP/FN rates and predictive power for operational impact (see CG2 in our synthesis). Second, existing benchmarks and prototypes often use single-sample or narrow sampling regimes; without controlled generator variability or explicit fault-injection, empirical estimates of validator FN-rate and repair utility remain unstable. These gaps explain why prior reports can demonstrate verifier benefits in specific controlled settings but cannot by themselves establish general-purpose validator effectiveness across heterogeneous NetOps failure modes [2],[3],[5].

Practical implications drawn from prior records. The literature implies specific experimental controls needed to measure validator utility: deterministic task banks with explicit fault classes (to estimate FN/FPR under known perturbations), multi-sample generator draws or explicit injection of generator errors (to reintroduce variance), and multi-validator or multi-model comparisons (to bound external validity). Prior survey and benchmark records recommend quantifying verification cost as part of evaluation—reporting verification coverage and detection latency alongside success rates—rather than treating a single pass/fail rate as the sole measure of safety or reliability [5]. Our preregistered metric suite and the follow-up

prescriptions (fault-injection, multi-sample generation, multi-validator comparison) are designed specifically to address these literature-identified needs without presuming that validator layers alone universally deliver operational improvements.

III. METHODOLOGY

Preregistered design and confirmatory intent. We preregistered study AC-2026-03 as a paired confirmatory experiment to isolate downstream validator/repair effects from generator variability. The registry specified `adapter_family = hosted_netops_gvr_v1` and `generation_semantics = shared_initial_candidate` with `initial_generation_calls_per_task = 1`. The primary estimand was the `paired_success_rate_difference_guarded_minus_baseline`; the primary test was an exact two-sided McNemar with paired bootstrap percentile confidence intervals (`bootstrap_resamples = 10000`, `bootstrap_seed = 1729`). The design targeted detection of an absolute 0.10 paired improvement ($\alpha = 0.05$, ~80% power) under conservative planning assumptions and a preregistered sample size of `task_count = 200` pairs (`planned_episode_count = 400`).

Task generation and constraints. Tasks were drawn deterministically from `deterministic_netops_task_generator_v5` (`task_family = intent_configuration_repair_v1`). Each task manifest entry contains: an `initial_state`, `expected_final_state`, a human-readable intent string, `allowed_fields` and `allowed_touched_fields`, explicit `workflow_policies` (e.g., `minimum_active_in_groups`, `must_be_down_for_fields`), and a `required_command_sequence`. Task selection followed the preregistered deterministic index rule `challenging_workflow_stress_v1` to produce a stress-profiled suite (levels labeled ‘very_hard’ and ‘extreme’ in difficulty fields). Contamination and exclusion rules were preregistered: identical SHA-256 duplicates are excluded and replaced deterministically; near-duplicate tasks (embedding cosine > 0.95) are deterministically excluded and replaced. The execution/analysis artifacts (`analysis/contamination_summary.csv`) record no contamination exclusions.

Model, sampling, and provider accounting. The declared model in `execution/model_configuration.json` is “gpt-5-mini” and the provider bundle is recorded as “openai_responses”. Preregistered `model_scope = gpt-5-mini` and `maximum_model_calls = 400`. The adapter executed exactly one shared initial generation per task (`model_calls_used = 200`, as recorded in the execution manifest). Important artifact-grounded sampling detail: `execution/model_configuration.json` records `temperature = null` (`temperature = null/unset`). Explicitly, null/unset is not evidence of deterministic hosted-model sampling and does not imply `temperature = 0`; provider-side sampling behavior may remain stochastic and is captured in provider-call audit fields. Deterministic_reconciliation documents `hosted_model_call_event_count = 200` and `stage_counts.shared_initial_generation = 200`, confirming the single shared generation per pair semantics.

Validator, scoring rule, and repair policy. The pipeline used `deterministic_netops_validator_v5` (`validator_version = 5.0`) as the mechanistic validator. The validator implements rule-based intent-constraint checks and emits per-episode fields used by scoring: `normalized_configuration`, `parsed_command_count`, `required_command_count`, `observed_touched_field_count`, `violation_count` and `violation_codes`, `valid` (boolean), `repair_applied` (boolean), and `repair_provider_trace` fields when a repair is invoked. The preregistered binary scoring rule `full_intent_constraint_validation` maps `validator.valid` (true/false) to the confirmatory success label. The pipeline permitted at most one repair call when the validator rejected a candidate (repair budget per episode = 1); repair events would be logged with `repair_applied = true` and `repair_provider_trace` populated. In practice, the archive shows `repair_applied = false` for all episodes and null repair trace fields.

Execution-level controls, exclusions, and missingness. Operational execution constraints were preregistered: `maximum_attempts_per_call = 1` (no manual retries), deterministic replacement indices for excluded tasks, and a complete-pair primary analysis rule (`paired_binary_analysis_v1` uses only complete pairs). The manifested execution ran to completion with `completed_episode_count = 400`, `failed_episode_count = 0`, and `complete_pair_count = 200`. Missingness artifacts (`analysis/missingness_summary.csv`) show no baseline-only, guarded-only, or both-missing pairs. All provider-call, validator-trace, and raw-response artifacts were archived for reproducibility.

Analysis plan and inferential checks. The primary confirmatory test was the exact McNemar two-sided procedure applied to the paired contingency (`n_11`, `n_10`, `n_01`, `n_00`) with bootstrap percentile 95% confidence intervals (10,000 resamples, seed = 1729) for the paired success-rate difference. Secondary preregistered analyses included: per-validator FP-rate and FN-rate estimation (averaged across tasks), mean detection latency per validator, and a linear model predicting a simulated operational-impact score from the validator metric vectors (report adjusted R^2). Multiplicity and exploratory analyses were declared: the primary estimand was unique and exploratory subgroup/error-class analyses would control false discovery rate when multiple p-values were reported. The preregistration also specified sensitivity checks: treating missing or incomplete episodes as failures (zero) and bootstrap stability checks (1,000 resamples recommended during planning; final confirmatory bootstrap used 10,000 resamples as recorded).

Why these choices. The paired shared_initial_candidate design was chosen to provide a clean scientific isolation: any improvement observed in the guarded condition must arise from deterministic validator detection and repair (or from repair-induced acceptance differences), not from independent generator sampling differences. The preregistered contamination rules, single-attempt policy, and deterministic task indexing were selected to reduce confounds and keep the confirmatory estimand interpretable. The trade-off is explicit and documented: `shared_initial_candidate` deliberately suppresses generator variability and therefore increases the risk

of ceiling/floor outcomes when the single sampled candidate systematically agrees with the validator, a risk that materialized in the present run.

Archive and artifact linkage for reproducibility of methods. All method-level configuration objects and runtime traces referenced above are present in the archived execution bundle: `execution/model_configuration.jsonl` (`declared_model_version = "gpt-5-mini"`, `temperature = null`, `maximum_attempts_per_call = 1`), `execution/task_manifest.jsonl` (deterministic_netops_task_generator_v5 payloads and required sequences), `execution/raw_results.jsonl` (raw model responses), `execution/scoring/*.json` (per-episode validator traces), and analysis artifacts (`analysis/paired_contingency_table.csv`, `analysis/condition_summary.csv`, `analysis/missingness_summary.csv`). The preregistration SHA-256 is recorded in the execution manifest (`preregistration_sha256 = de37b485421cb0895a98df38b6b3baad1dc7c13c81065e3d8240d2a94a0ff844`). These artifacts support independent verification of method-level choices, confirm that the single-sample-per-task and shared_initial_candidate semantics were enforced, and document the validator interface used to produce the binary success labels.

IV. RESULTS

Execution accounting and primary outcomes. The adapter completed the run exactly as preregistered: `planned_episode_count = 400`, `completed_episode_count = 400`, `failed_episode_count = 0`, and `model_calls_used = 200` (one shared initial generation per pair). Provider-call audit recorded `hosted_model_call_event_count = 200`, `completed_hosted_model_call_event_count = 200`, `cache_hit_count = 0`, `cache_miss_count = 200`, and `stage_counts.shared_initial_generation = 200` (`deterministic_reconciliation.json`). Contamination and missingness artifacts confirm no exclusions or missing episodes (`analysis/contamination_summary.csv` empty; `analysis/missingness_summary.csv` reports `complete_pairs = 200` and zeros for all missingness categories).

Primary confirmatory result and exact accounting. The deterministic validator judged every sampled shared initial candidate as valid; no repairs were applied in any episode. Analysis artifacts report `baseline_success_count = 200/200` (`success_rate = 1.0`) and `guarded_success_count = 200/200` (`success_rate = 1.0`). The paired contingency table (`analysis/paired_contingency_table.csv`) is `n_11 = 200`, `n_10 = 0`, `n_01 = 0`, `n_00 = 0`. Consequently the `paired_success_rate_difference_guarded_minus_baseline = 0.00`. The prespecified exact McNemar two-sided test yields `p-value = 1.0`. The paired bootstrap percentile 95% confidence interval (`bootstrap_resamples = 10000`, `bootstrap_seed = 1729`; `analysis/analysis_log.jsonl`) is `[0.00, 0.00]`. These numeric artifacts (`analysis/condition_summary.csv`, `analysis/paired_contingency_table.csv`, `analysis/analysis_log.jsonl`) are archived and match the reported values.

Representative per-episode validator diagnostics (artifact-grounded). To illustrate why no repairs were triggered, Table

and scoring artifacts contain the validator trace fields for every episode; a compact example (execution/scoring/task-000004-guarded.json) shows the key mechanistic outputs that were recorded for each episode:

```
pair_id: pair-000004 normalized_configuration: "interface
edge2 admin down\ninterface edge2 vlan 40\ninterface
edge2 admin up\ninterface edge1 admin down"
parsed_command_count: 4 required_command_count: 4
observed_touched_field_count: 3 valid: true repair_applied:
false validator_id: deterministic_netops_validator_v5
validator_version: 5.0
```

Equivalent validator_trace.validation_after.valid = true and repair_applied = false hold for every archived episode; the repair_provider_trace fields are null for all episodes (execution/scoring/*.json entries list repair_provider_trace_path = null). Representative tasks in the artifact_examples (task-000004, task-000005, task-000006) likewise show parsed_command_count equal to required_command_count and valid = true, demonstrating that, for these sampled seeds, generator outputs satisfied the validator’s mechanistic constraints exactly.

Secondary and exploratory diagnostics (uninformative under ceiling). Because repair_applied = false for all episodes, secondary metrics that depend on rejection or repair events are degenerate in this execution: per-validator false-positive or false-negative rate estimates cannot be computed from observed rejections; detection-latency and repair-invocation distributions are empty by design. The archived raw_results.jsonl and per-episode validator traces nevertheless provide the full data required to compute these metrics under alternative scoring rules or after applying the follow-up prescriptions (fault injection or multi-sample generation) described elsewhere in this paper.

Artifact-grounded diagnostic interpretation. Under shared_initial_candidate semantics any paired difference could only arise from validator-triggered repair. The degenerate all-valid outcome observed here is explicable from artifact-grounded evidence in three linked facts: (1) the deterministic task generator and the recorded generation seeds produced initial candidates that satisfied the validator’s mechanistic intent-constraint checks (as shown by per-episode validation_after summaries); (2) the single-sample-per-task sampling regime removed generator variability as a source of disagreement by design (execution/model_configuration.json and deterministic_reconciliation.json record one shared generation per task); (3) mechanistic validators operate on structural and policy checks and can be blind to semantic intent misalignment (correct-by-structure yet incorrect-by-intent) absent a separate semantic-intent model or operator feedback. These jointly explain why no repairs occurred and why the confirmatory test could not detect the preregistered 10-percentage-point improvement target.

Reproducibility pointers and archived artifacts. All numeric claims above are reproducible from the archived artifacts: execution/raw_results.jsonl (input_results_sha256 = 23f4cd010f3a7419eb21ef9b8caf4dbe7730552eaa3f3eaa22d2068cf3b5d4e3), execution/model_configuration.json

(sha256 = a40e96e2...), analysis/paired_contingency_table.csv (sha256 = 7bc1bd2a...), analysis/condition_summary.csv (sha256 = 1cb072b4...), analysis/analysis_log.jsonl (sha256 = dd2dbf18...), and representative per-episode scoring files (execution/scoring/task-*.json). These files show the per-episode validator.valid flags, repair_applied markers, shared_initial_candidate content and SHA-256s, and provider-call audit entries used to produce the summary statistics above. Because the observed ceiling outcome is entirely recorded in these artifacts, reanalysis under modified scoring rules or preregistered follow-up arms (fault injection, multi-sample generation, multi-validator comparison) can be performed without re-executing the original shared-initial generation stage.

V. DISCUSSION

Design implications of shared_initial_candidate semantics. The preregistered shared_initial_candidate choice isolates downstream validator effects by ensuring both conditions evaluate the same initial candidate for each pair. This isolation is scientifically valuable because any observed paired difference must derive from validator-triggered repair. It also concentrates the risk of a single-sample ceiling: if the generator sample aligns with the validator for all tasks, the design yields zero observed variance in the primary estimand, as occurred here. We emphasize this property because it determines what the paired estimator can and cannot measure: it cannot detect differences attributable to generator sampling or to alternative prompt constraints when the sampling semantics were shared.

Operational interpretation and validator scope. The deterministic validator (deterministic_netops_validator_v5) reliably enforces mechanistic intent-constraint checks (parsed_command_count, required_command_count, violation_codes, valid flag). However, artifact-grounded evidence here shows that correctness-by-structure does not imply semantic intent alignment under operator expectations. That limitation is intrinsic to validators that lack an external intent model or operator feedback; detecting semantic mismatches requires either richer intent models, human-in-the-loop signals, or tailored semantic validators. The present execution does not imply validators are unnecessary—rather, it demonstrates that in a synthetic bank and single-sample regime the validator had nothing mechanical to correct.

Prescriptions for follow-up preregistered experiments (artifact-grounded). The archived infrastructure and artifacts directly support concrete next-step preregistered arms that would produce estimable validator FP/FN and repair utility while preserving preregistration rigor: 1) Fault-injection arm (preregistered): deterministically inject defined fault classes into a controlled fraction f of reference answers (syntactic errors, configuration-logic faults, semantic-intent swaps). This increases baseline invalidity and permits measurement of repair invocation rates and post-repair success. 2) Multi-sample generation: set initial_generation_calls_per_task = $k \geq 3$ independent draws per task with independent seeds recorded. Allow the guarded pipeline to accept any validator-approved

draw or to perform repair. This reintroduces generator variability and creates opportunities for validator rejections and repairs. 3) Multi-validator and multi-model comparison: add alternative validator implementations and at least one additional model family. This will quantify sensitivity to validator design and model diversity and improve external validity.

Each prescription is directly supported by the archived adapter semantics and task-generation determinism; implementing them preregistered would produce non-degenerate data suitable for estimating per-validator FP/FN, repair utility, detection latency, and predictive R^2 against an operational-impact proxy as originally planned.

Threats to validity. Internal validity: `shared_initial_candidate` intentionally isolates downstream effects but increased single-sample ceiling risk. Construct validity: `validator.valid` maps to mechanistic correctness but may fail to capture semantic intent misalignment. External validity: synthetic task bank, single-model (gpt-5-mini), and a single deterministic validator limit generalization to heterogeneous production networks and additional model families. Conclusion validity: the degenerate outcome yields zero variance in the primary estimand so inferential conclusions about validator benefit are not possible from this execution alone. These limitations are recorded in the execution and analysis artifacts and motivate the artifact-grounded follow-up arms described above.

VI. LIMITATIONS

- `Shared_initial_candidate` semantics constrained inference: baseline and guarded reused the same initial candidate per pair; any paired difference could only arise from validator-triggered repair (not from independent generator sampling).
- Synthetic task bank (difficulty 'very_hard'/'extreme') is not equivalent to heterogeneous production networks (multi-vendor devices, real ticket noise); external validity is limited.
- Single-model experiment: only gpt-5-mini was executed; no multi-model generality claims are made.
- No repairs were applied because all initial candidates passed the deterministic validator; therefore the study could not estimate repair effectiveness or validator false-negative behavior in this execution.
- Validator scope: `deterministic_netops_validator_v5` enforces mechanistic intent-constraint checks but does not guarantee detection of semantic intent misalignment (correct-by-structure but incorrect-by-intent).
- Recorded `execution/model_configuration.json` temperature = null/unset does not imply deterministic sampling; provider-side sampling behavior may vary and is documented in execution manifests.

VII. CONCLUSION

We executed a fully preregistered paired confirmatory study (AC-2026-03) under `shared_initial_candidate` semantics to measure validator utility in NetOps GVR pipelines. For the

synthetic task bank, the sampled gpt-5-mini generations, and `deterministic_netops_validator_v5` used here, every sampled initial candidate passed the validator's mechanistic checks so no repairs were applied and guarded success equaled baseline (200/200). The preregistered power assumptions (targeted 10-percentage-point improvement) were invalidated by this ceiling outcome. The result is artifact-grounded and reproducible from the archived bundle. We publish the exact execution and analysis artifacts to support replication and provide concrete, preregistered experiment prescriptions (fault injection, multi-sample generation, multi-validator comparison) that are required to produce estimable validator FP/FN behavior and repair value in operationally meaningful conditions.

DISCLOSURE STATEMENT

Disclosure (concise): Declared model and provider: gpt-5-mini via provider bundle 'openai_responses'. Orchestration/adapter: `hosted_netops_gvr_v1`. Deterministic validator: `deterministic_netops_validator_v5` (`validator_version` = 5.0). Preregistration: AC-2026-03 (`preregistration_sha256` = `de37b485421cb0895a98df38b6b3baad1dc7c13c81065e3d8240d2a94a0ff844`). Master prompt immutable reference: `provenance/master_prompt.txt`, SHA-256 = `1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ba`. Autonomy boundary: a human Research Director selected the broad topic family and approved the master prompt before the autonomy lock; after the lock the AI system autonomously performed literature review, preregistration, experiment design, execution, analysis, internal peer review and manuscript drafting without manual editing of technical content. Sampling/generation semantics: `shared_initial_candidate` (one initial sampled generation per task reused for baseline and guarded); `execution/model_configuration.json` recorded temperature = null/unset (null/unset is not evidence of deterministic sampling and does not imply temperature = 0). Call budget and usage (preregistered): `maximum_model_calls` = 400; `actual_model_calls_used` = 200 (one shared initial generation per pair). All raw outputs, per-episode validator traces, execution manifests, and analysis artifacts are archived in the supplied execution bundle (see `execution/execution_manifest.json`). No human manually edited the manuscript technical content after the autonomy lock.

REFERENCES

- [1] [1] R. Cadenas, "Convergent Correctness in Stochastic Code Generation: A Generate-Verify-Repair Architecture for Deterministic Validation of Autonomous AI Coding Agents in Regulated Environments," SSRN Electronic Journal, 2026. doi:10.2139/ssrn.6754899.
- [2] [2] K.-H. Tseng, N. Bogahawatta, Y. Ginige, K. Patel, K. Dacic, S. Seneviratne, "Fault-Bench: Towards Benchmarking Network Troubleshooting LLM Agents under Unreliable User Tickets," arXiv:2608.27021, 2026.
- [3] [3] Y. Liu, C. Pei, L. Xu, B. Chen, M. Sun, Z. Zhang, Y. Sun, S. Zhang, K. Wang, H. Zhang, J. Li, G. Xie, X. Wen, X. Nie, M. Ma, P. Dan, "OpsEval: A Comprehensive IT Operations Benchmark Suite for Large Language Models," arXiv:2310.07637, 2023.
- [4] [4] R. S. Babu and A. Agrawal, "Self-Healing Agentic Orchestrators for Reliable Tool-Augmented Large Language Model Systems," arXiv:2606.01416, 2026.

- [5] [5] H. Zhou, C. Hu, Y. Yuan, Y. Cui, Y. Jin, C. Chen, H. Wu, D. Yuan, L. Jiang, D. Wu, X. Liu, J. Zhang, X. Wang, J. Liu, "Large Language Model (LLM) for Telecommunications: A Comprehensive Survey on Principles, Key Techniques, and Opportunities," *IEEE Commun. Surveys & Tutorials*, 2024. doi:10.1109/COMST.2024.3465447.