Verifier-Grounded Generate-Verify-Repair for Intent→Configuration NetOps Tasks: A Paired Mininet Emulation Study

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

***Abstract*—We preregistered and executed a provenance-traced paired confirmatory study (N = 200 intents; 400 paired episodes) to evaluate whether a deterministic verifier-grounded generate–verify–repair (GVR) loop reduces operationally detectable policy-violation errors when a hosted LLM produces device configuration sequences from operator intent in a Mininet emulation. Execution used shared_initial_candidate semantics (one sampled initial generation per intent reused by both baseline and guarded). Baseline success (no detected policy violation) = 196/200 (98%); guarded success = 200/200 (100%). Paired risk difference (guarded - baseline) = 0.02; preregistered exact McNemar two-sided p = 0.125 (primary confirmatory test). An exploratory paired bootstrap (seed = 1729, 10,000 resamples) yields a 95% percentile CI = [0.005, 0.04]. Four verifier-triggered repairs occurred (total hosted model calls = 204), giving mean extra model calls per intent = 0.02 and mean extra calls per avoided violation = 1.0. All reported numeric values, provider-call accounting, validator traces, and analysis artifacts are drawn verbatim from the archived execution and analysis bundle. We interpret the preregistered confirmatory result conservatively and discuss construct, internal, and external validity limits and operational implications.**

I. Introduction

Automating intent→configuration translation is an active NetOps research direction because natural-language or intent descriptions are a natural operator interface but can be transformed incorrectly by large language models (LLMs). Mechanically checkable faults introduced by LLMs—syntactic errors, topology/reachability violations, and ACL/policy mistakes—can lead to availability loss or exposure if applied to devices without validation [1], [4]. Generate–Verify–Repair (GVR) patterns pair an LLM generator with a deterministic verifier: the verifier checks mechanizable constraints and, if violations occur, issues localized diagnostics used to request corrective edits from the model [2], [3]. Prototype work across code and domain-specific program generation suggests verifier-grounded repair reduces mechanically verifiable errors; however, questions remain about scaling such pipelines to heterogeneous multi-vendor template families, the concrete hosted-LLM cost per avoided violation, and how to interpret paired comparisons when sampling semantics reuse the initial model candidate.

Research question and contribution. Our preregistered research question is: Does a verifier-grounded generate–verify-repair loop significantly reduce operationally detectable policy-violation errors produced by a hosted LLM when generating network configurations for operator intents across heterogeneous multi-vendor template families and topology patterns in a Mininet-based emulation? Contributions: (1) a provenance-traced confirmatory paired experiment (C1-GVR-multivendor-2026-v1) with shared_initial_candidate semantics executed on N = 200 intents (400 episodes); (2) audited provider and verifier accounting showing repair costs and concrete hosted-model calls; (3) artifact-grounded statistical inference (exact McNemar test) and exploratory bootstrap interval (seed = 1729, 10,000 resamples); (4) a focused analysis of failure modes, threats to validity, and operationally grounded recommendations consistent with archived artifacts.

II. Related Work

Recent literature on LLMs for NetOps and AIOps converges on two observations that are directly relevant to verifier-grounded pipelines. First, surveys and domain reviews document rapid uptake of LLMs for network tasks while highlighting recurring gaps in evaluation practice: missing provenance, limited verifier integration, and scarce multi-vendor or emulation-scale evidence [5], [6]. These reviews synthesize heterogeneous reports (research prototypes, bench demonstrations, and applied surveys) and conclude that while LLMs reduce human effort for drafting configurations, the community lacks standardized, provenance-traced benchmarks that pair generators with deterministically auditable verifiers. This limits cross-study comparability and weakens claims about operational safety or generalization across vendors and topologies [5], [6].

Second, targeted configuration studies and benchmarks show feasibility but also expose common failure modes. Prior NetOps evaluations (for example, work investigating LLM requirements for correct router configuration synthesis and the NetConfEval benchmark) document frequent classes of mechanically checkable failures—syntax/grammar mistakes, incorrect sequencing that violates workflow invariants, and topology/reachability errors when configurations contradict the network state [1], [4]. These studies typically evaluate generator accuracy on curated tasks and emphasize the need for automated checking; however, many prototype benchmarks either omit a closed-loop repair stage or do not provide provenance traces linking initial outputs, verifier diagnostics, and any applied repairs. In practice, reported improvements in such studies are therefore difficult to attribute to deterministic validator interventions versus differences in prompting or sampling.

A cross-domain body of work motivates verifier-grounded mitigation architectures (generate–verify–repair). Recent methodological proposals and prototype demonstrations in code and domain-specific program generation argue that a verifier capable of expressing mechanizable constraints plus localized diagnostic feedback enables targeted repairs that reduce mechanically detectable errors [2], [3]. These works provide conceptual and single-case empirical support: they show that when a verifier can (i) deterministically check the artifact against explicit constraints and (ii) return actionable, localized diagnostics, a bounded repair policy can converge failing samples to valid artifacts. Importantly, these prior results also identify the core dependency of GVR effectiveness on verifier coverage: if a verifier cannot encode an error class, that class remains a residual risk even with repeated repair attempts [2], [3].

How our study relates. Building on the empirical observations and methodological proposals above, our experiment addresses a concrete gap: scalable, provenance-traced paired evaluations that measure the marginal effect of deterministic verification and bounded repair across heterogeneous template families and topology patterns. Unlike prior NetOps benchmarks that either evaluate single-pass generation or lack explicit provenance between generator outputs and verifier actions [1], [4], or like GVR proposals that focus on single-domain demonstrations [2], [3], our paired design (with shared_initial_candidate semantics) isolates the verifier/repair contribution by reusing the identical initial LLM sample for both conditions and then measuring the final guarded outcome after permissive, verifier-triggered repair. This combination—multi-vendor synthetic templates exercised in Mininet, deterministic verifier traces, provable provider-call accounting, and preregistered paired inference—directly responds to the literature's call for provenance and auditable assessments of verifier-assisted generation [5], [6]. At the same time, we adopt the caution emphasized in the surveyed and prototype literature: verifier expressivity bounds the classes of errors that can be detected and repaired, and mechanical verification cannot by itself guarantee absence of higher-order semantic misinterpretations outside the verifier's rule set [2], [3], [5].

III. Methodology

Preregistration and estimand. The experiment was preregistered as C1-GVR-multivendor-2026-v1 before any model calls (preregistration SHA-256: 9e0fc992665835975b01289f7841926ef6abf358 4252c5d0459ba0ecf166bb50). Primary estimand: paired_success_rate_difference_guarded_minus_baseline (probability that an intent yields zero operationally detectable violations under guarded minus baseline), tested with an exact McNemar two-sided test at $\alpha = 0.05$. Predeclared exploratory estimators include a paired bootstrap percentile CI (bootstrap_seed = 1729; resamples = 10,000) and descriptive cost metrics (mean extra hosted model calls per intent and mean extra calls per avoided violation).

Sampling, strata, and shared_initial_candidate semantics. Task selection was deterministic and stratified across three synthetic vendor-template families and four topology patterns (12 strata) to exercise template heterogeneity. N = 200 unique operator intents were fixed prior to model calls. Execution used shared_initial_candidate semantics: exactly one initial sampled generation per intent was produced and reused by both baseline and guarded episodes. Under these semantics baseline outcome is the deterministic validator verdict on that shared initial candidate; the guarded outcome equals the final guarded pipeline result after deterministic validation and any permitted repair. The preregistration explicitly declared that independent constrained-prompt arms were not executed; therefore paired differences can arise only from deterministic validator or repair steps, not from different first-pass samples or prompt templates.

Model configuration, sampling notation, and budget. Declared model (as archived) is gpt-5-mini (provider recorded as openai_responses). The archived model_configuration.json records temperature = null/unset; we explicitly state that null/unset is not evidence of deterministic provider sampling and does not imply temperature = 0. Planned budget enforced in the adapter: planned_episode_count = 400, initial_generation_calls_per_task = 1, maximum_repair_calls_per_task = 1, maximum_model_calls = 400.

Deterministic verifier and repair policy. deterministic_netops_validator_v5 is a rule-based verifier combining (i) syntactic parsing and command pattern checks, (ii) required-command sequencing and workflow policy checks (per-task payload), (iii) reachability/topology checks against the emulated Mininet state, and (iv) encoded ACL/policy checks derived from the task payload. The scorer rule full_intent_constraint_validation yields a binary success = "no detected violations" if the final configuration satisfies all mechanizable constraints. Guarded pipeline behavior: if the validator flags violations, a single verifier-triggered repair call may be issued; repair prompts contain the verifier's localized diagnostics requesting corrective edits. Repair calls were permitted only when the validator reported violations, and up to one repair call per intent was allowed per the preregistration.

Scoring, missingness, and analysis plan. Primary analysis: exact McNemar two-sided test on paired binary success indicators. Missingness and exclusions were predeclared: duplicate tasks (detected via deterministic hashing) or unrecoverable infra failures would be excluded and reported; ambiguous validator outputs were conservatively treated as violations in the primary analysis. Exploratory analyses included the paired bootstrap percentile CI (seed = 1729, 10,000 resamples) and descriptive cost accounting. Analysis code and seeds are archived.

IV. Results

Execution completeness and timing. The preregistered run executed to completion without excluded pairs or unrecoverable infrastructure failures. The execution manifest records start and completion times (started_at_utc = 2026-09-01T12:52:23.426943+00:00; completed_at_utc = 2026-09-01T13:19:15.308203+00:00) and shows planned_episode_count = 400 with completed_episode_count = 400 and failed_episode_count = 0. Deterministic reconciliation confirms complete_pair_count = 200 and that all deterministic consistency checks passed.

Audited provider accounting and stage decomposition. Provider auditing in the execution manifest reports hosted_model_call_event_count = 204 and decomposes these calls into stage_counts = {shared_initial_generation: 200, repair: 4}. The manifest further records cache_hit_count = 0 and cache_miss_count = 204. These audited counts are the authoritative provider-call accounting for the run and were used directly in cost summaries and downstream arithmetic.

Per-condition tallies and raw contingency. The archived condition summary (analysis/condition_summary.csv) reports: baseline successes = 196, baseline failures = 4 (observed baseline success rate = 196/200 = 0.98); guarded successes = 200, guarded failures = 0 (guarded success rate = 200/200 = 1.0). The paired 2×2 contingency (analysis/paired_contingency_table.csv) is reproduced exactly from the archived analysis artifact: n11 = 196 (both succeed), n10 = 4 (guarded succeed, baseline fail), n01 = 0 (guarded fail, baseline succeed), n00 = 0 (both fail). These contingency counts are the direct inputs to the preregistered confirmatory test and to exploratory interval estimation.

Confirmatory inference (preregistered). Using the preregistered exact McNemar two-sided test on the paired binary success indicators yields p = 0.125; under the preregistered $\alpha = 0.05$ threshold this result does not reject the null of no difference. The point estimate for the paired risk difference (guarded - baseline) computed from the contingency counts is 0.02 (2 percentage points). We report the exact test p-value and the risk difference as primary, preregistered outputs and treat any auxiliary intervals as exploratory per the preregistration multiplicity plan.

Exploratory bootstrap interval and caution. An exploratory paired bootstrap percentile interval was computed using the archived analysis procedure (bootstrap_seed = 1729; bootstrap_resamples = 10,000) and produced a 95% percentile CI for the paired risk difference = [0.005, 0.04]. We report this interval as supplementary magnitude information only and note the known instability of bootstrap percentile intervals when discordant pair counts (n10 + n01) are small; with only four discordant pairs the bootstrap estimate provides limited additional inferential resolution and should not supplant the exact McNemar inference for confirmatory claims. The analysis log (analysis/analysis_log.jsonl) contains the bootstrap seed and resample count used to produce this interval.

Repair events, cost accounting, and arithmetic derivations. Four verifier-triggered repair calls were invoked across the N = 200 intents, consistent with stage_counts.repair = 4 in the provider audit. Total hosted model calls therefore equal 200 shared initial calls + 4 repair calls = 204 audited calls. From these audited counts we compute mean hosted model calls per intent = 204 / 200 = 1.02 and mean extra calls per intent attributable to repairs = 0.02. The archived contingency shows that each repair converted a failing baseline outcome into a passing guarded outcome (the four n10 discordant pairs), so mean extra calls per avoided violation = 1.0 (4 repair calls / 4 avoided violations). All these arithmetic results are direct derivations from the archived provider and contingency artifacts and are reported verbatim.

Representative validator trace behavior and substantiation. The archived validator traces and scoring artifacts show the operational pattern that produced the contingency entries: for each pair, the shared_initial_candidate was passed to the deterministic verifier; when the verifier found mechanizable violations it produced localized diagnostics and (for four intents) the guarded pipeline issued a single repair call using those diagnostics. For those discordant pairs the pre-repair validator output recorded violation_count > 0 and the repair decision was accompanied by a subsequent validator validation_after that recorded violation_count = 0 and a final scorer decision of success = 1. Representative examples of valid shared_initial candidates and their validator traces are provided in the archive (execution/scoring and execution/responses artifacts) and illustrate the exact shared_initial_candidate reuse semantics: the same initial generated candidate was scored as baseline and was the starting point for any guarded repair sequence. The reconciliation log confirms cross_condition_cache_key_reuse_observed = false; therefore the guarded conversions arose from validator/repair activity rather than from independent new initial samples.

Power context and interpretation. The preregistered power calculation targeted detecting a 10 percentage point absolute reduction with approximate required $N \approx 157$; N = 200 was chosen to provide margin and permit stratified description. The observed baseline failure prevalence (4/200 = 2%) is substantially lower than the planning scenario used for power targeting, which reduces power to detect small absolute improvements and yields a small discordant count (4) that underlies the non-significant McNemar result. We therefore present the exact McNemar p = 0.125 as the formal confirmatory outcome, report the exploratory bootstrap interval (seeded and archived), and interpret the cost accounting and repair effectiveness as artifact-grounded operational diagnostics rather than as definitive evidence of broad efficacy beyond the constrained emulation and verifier scope.

V. Discussion

Causal interpretation under shared_initial_candidate semantics. Because execution used shared_initial_candidate semantics (one sampled initial generation per intent reused by baseline and guarded), paired improvements arise solely from deterministic validator/repair actions. The preregistration explicitly declared that independent constrained-prompt arms were not executed; thus the observed n10 = 4 discordant pairs reflect validator-triggered repairs that converted failing shared initial candidates into passing final configurations. Any statements attributing improvements to prompt-engineering or differing first-pass samples would be incorrect for this run.

Provider accounting and manifest labels. The archived execution manifest and deterministic reconciliation provide audited provider_call counts that clarify how model events map to pipeline stages. The manifest records hosted_model_call_event_count = 204 and decomposes that total into stage_counts = {"shared_initial_generation": 200, "repair": 4}. Condition or stage labels observed in the manifest (for example entries labeled "shared":200 and "guarded":4) reflect this decomposition: the 200 "shared_initial_generation" events correspond to the single initial sampled generation produced and reused for both baseline and guarded episodes, while the 4 repair events are additional guarded-stage model calls that occurred only when the deterministic verifier reported violations. This audited decomposition underpins the cost accounting reported in Results (mean extra model calls per intent = 0.02; mean extra calls per avoided violation = 1.0) and supports the causal claim that paired gains derive from verifier-triggered repairs rather than from independent re-sampling.

Statistical interpretation and bootstrap caution. The paired bootstrap percentile CI we report (95% CI = [0.005, 0.04]; bootstrap_seed = 1729; resamples = 10,000) is explicitly exploratory. Small discordant pair totals (n10 + n01 = 4 here) make bootstrap percentile intervals sensitive to resampling variability and to the discrete structure of paired binary data; consequently such intervals can underrepresent sampling uncertainty or be unstable in edge cases. For this reason we rely on the preregistered exact McNemar two-sided test ($p = 0.125$) as the formal confirmatory inference and present the bootstrap interval only to provide a complementary magnitude estimate archived with its seed and resampling count. Researchers planning follow-ups should (a) anticipate discordant-count regimes when designing paired tests, (b) consider planning for larger N or higher baseline failure prevalence to increase discordant counts, and (c) report both exact conditional inferences and multiple interval estimators to better triangulate effect size in sparse discordance settings.

Verifier coverage and residual semantic risks. deterministic_netops_validator_v5 checks syntactic correctness, required sequencing, reachability against the emulated state, and encoded ACL/policy constraints derived from task payloads. These mechanizable checks were sufficient to detect and localize the failures corrected by the four repairs. However, the verifier can only detect violations that are expressible in its rule set and encoded policies; higher-order semantic misinterpretations of intent that do not manifest as a mechanizable constraint violation remain possible false negatives. The observed complete elimination of mechanically detectable baseline violations by the bounded repair budget applies only to the verifier's detectable fault classes; extending verifier expressivity (for example, formal policy ontologies or richer semantic specifications) is necessary to reduce residual semantic risk.

Operational implications and bounded generality. Within the constrained Mininet emulation, synthetic template families, and validator rule set used here, a small repair budget removed all mechanically detectable baseline violations at very low hosted-model cost. This artifact-grounded observation suggests that, in settings where a deterministic verifier can express relevant constraints and produce localized, actionable diagnostics, a bounded GVR policy can eliminate template-driven mechanically checkable faults in similar emulated contexts. We emphasize bounded generality: these results do not establish production readiness, multi-model generality, or coverage for semantic errors outside the verifier. Re-deploying the approach in operational environments requires (i) richer verifier coverage tuned to production policies, (ii) paired designs that also execute independent multi-sample initial generations to separate sampling/prompting effects from repair effects, and (iii) replication across multiple hosted models and more realistic hardware testbeds.

Recommendations for future evaluations. Based on the archived evidence and the statistical constraints observed, we recommend that future confirmed-design studies: (1) preregister whether shared_initial_candidate semantics or independent-sample semantics will be used and plan estimands accordingly; (2) select sample sizes or baseline scenarios that anticipate sufficient discordant pairs for confirmatory power, or else use study designs that enlarge expected discordance (e.g., harder task strata); (3) broaden verifier expressivity to capture more semantic/policy classes and measure verifier false-negative rates against annotated ground truth; and (4) report audited provider_call decompositions and seeds for bootstrap/replication so others can unambiguously reconcile cost and resampling diagnostics. These recommendations are directly motivated by patterns in the archived execution manifest, deterministic reconciliation, and analysis artifacts.

VI. Limitations

- Scope limited to Mininet-style emulation with synthetic, provenance-traced intent→configuration tasks; results do not generalize directly to production hardware or live operations.
- Single declared model (gpt-5-mini) evaluated; multi-model generality is not established.
- Deterministic validator coverage limited to mechanically checkable errors (syntax, reachability, ACL/policy encodings); higher-level semantic or specification misinterpretations outside the validator may remain undetected.
- Shared_initial_candidate semantics imply observed paired differences derive only from verifier-triggered repair; this design does not measure effects of alternative prompting strategies or independent multi-sample candidate selection.

- Observed confirmatory effect (2 percentage points) did not reach statistical significance at α = 0.05 for N = 200 (exact McNemar p = 0.125); low baseline failure prevalence (2%) reduced power to detect small absolute improvements under some preregistered scenarios.

VII. Conclusion

We preregistered and executed a provenance-traced paired confirmatory study (C1-GVR-multivendor-2026-v1) using shared_initial_candidate semantics to measure whether a deterministic verifier plus bounded repair reduces mechanically detectable policy violations for intent→configuration tasks in a Mininet emulation. The preregistered exact McNemar two-sided test did not reject at α = 0.05 (p = 0.125). Guarded GVR invoked four repairs and corrected the four observed baseline violations at low model-call overhead (model calls = 204; mean extra calls per intent = 0.02), yielding a paired risk difference estimate of 0.02 and an exploratory paired bootstrap CI = [0.005, 0.04] (bootstrap_seed = 1729, resamples = 10,000). These artifact-grounded results indicate that when a deterministic verifier can express and check relevant constraints and provide localized feedback, a small repair budget can eliminate mechanically detectable policy violations in similar template-based emulated settings. The results do not establish production readiness or multi-model generality. All reported numbers, provider accounting, validator traces, and analysis artifacts are drawn verbatim from the archived execution and analysis bundles referenced in the preregistration and execution manifests.

References [1] R. Mondal, A. Tang, R. Beckett, T. Millstein, and G. Varghese, "What do LLMs need to Synthesize Correct Router Configurations?", Proceedings of the ACM, 2023, doi: 10.1145/3626111.3628194. [2] R. Cadenas, "Convergent Correctness in Stochastic Code Generation: A Generate-Verify-Repair Architecture for Deterministic Validation of Autonomous AI Coding Agents in Regulated Environments", SSRN, 2026, doi: 10.2139/ssrn.6754899. [3] Z. Ding, "Executable but Wrong: Verifier-Grounded Contracts and Counterexamples for LLM-Generated Music Programs", 2026, doi: 10.31224/7995. [4] C. Wang, M. Scazzariello, A. Farshin, S. Ferlin, D. Kostić, and M. Chiesa, "NetConfEval: Can LLMs Facilitate Network Configuration?", Proceedings of the ACM, 2024, doi: 10.1145/3656296. [5] S. Y. Long et al., "A Survey on Intelligent Network Operations and Performance Optimization Based on Large Language Models", IEEE Communications Surveys & Tutorials, 2025, doi: 10.1109/comst.2025.3526606. [6] L. Zhang et al., "A Survey of AIOps in the Era of Large Language Models", Proceedings of the ACM, 2025, doi: 10.1145/3746635.

Disclosure Statement

Declared model (archived): gpt-5-mini (provider recorded as openai_responses). Adapter family: hosted_netops_gvr_v1. Deterministic validator: deterministic_netops_validator_v5. Task generator: deterministic_netops_task_generator_v5. Preregistration: C1-GVR-multivendor-2026-v1 (SHA-256 = 9e0fc992665835975b01289f7841926ef6abf3584252c5d0 459ba0ecf166bb50), recorded prior to any model calls. Initial master prompt archived at provenance/master_prompt.txt (SHA-256 = 1872df1e1805d2d96940456ca016bd665d1d5196 add77f5acdf1582bb39b15ba). Human role: a human Research Director selected the topic family and constrained the initial master prompt prior to the autonomy lock; after the lock the AI autonomously performed literature verification, preregistration, experiment execution, analysis, peer review, revision, and manuscript drafting; no human manually edited technical manuscript content after the autonomy lock. Sampling semantics: execution used shared_initial_candidate (exactly one sampled initial generation per intent was produced and reused by both baseline and guarded); archived model configuration records temperature = null/unset (null/unset is not evidence of deterministic provider sampling and does not imply temperature = 0). Provider accounting (audited): hosted_model_call_event_count = 204 decomposed as 200 shared initial generations + 4 repair calls. Reported numbers, provider-call accounting, validator traces, and analysis artifacts are drawn verbatim from the archived execution and analysis bundles referenced in the preregistration and execution manifests.

References

[1] Rajdeep Mondal, Alan Tang, Ryan Beckett, Todd Millstein, George Varghese, "What do LLMs need to Synthesize Correct Router Configurations?", 2023, 10.1145/3626111.3628194

[2] Rafael Cadenas, "Convergent Correctness in Stochastic Code Generation: A Generate-Verify-Repair Architecture for Deterministic Validation of Autonomous AI Coding Agents in Regulated Environments", 2026, 10.2139/ssrn.6754899

[3] Zhengyang Ding, "Executable but Wrong: Verifier-Grounded Contracts and Counterexamples for LLM-Generated Music Programs", 2026, 10.31224/7995

[4] Changjie Wang, Mariano Scazzariello, Alireza Farshin, Simone Ferlin, Dejan Kostić, Marco Chiesa, "NetConfEval: Can LLMs Facilitate Network Configuration?", 2024, 10.1145/3656296

[5] S.Y. Long, Jingjing Tan, Bomin Mao, Fengxiao Tang, Yangfan Li, Ming Zhao, Nei Kato, "A Survey on Intelligent Network Operations and Performance Optimization Based on Large Language Models", 2025, 10.1109/comst.2025.3526606

[6] Lingzhe Zhang, Tong Jia, Mengxi Jia, Yifan Wu, Aiwei Liu, Yong Yang, Zhonghai Wu, Xuming Hu, Philip S. Yu, Ying Li, "A Survey of AIOps in the Era of Large Language Models", 2025, 10.1145/3746635