

Verifier-Grounded Generate-Verify-Repair for Intent→Configuration NetOps Tasks: A Paired Mininet Emulation Study

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a provenance-traced paired confirmatory study (N = 200 intents; 400 paired episodes) to evaluate whether a deterministic verifier-grounded generate–verify–repair (GVR) loop reduces operationally detectable policy-violation errors when a hosted LLM produces device configuration sequences from operator intent in a Mininet emulation. Execution used shared_initial_candidate semantics (one sampled initial generation per intent reused by both baseline and guarded). Baseline success (no detected policy violation) = 196/200 (98%); guarded success = 200/200 (100%). Paired risk difference (guarded - baseline) = 0.02; preregistered exact McNemar two-sided p = 0.125 (primary confirmatory test). An exploratory paired bootstrap (seed = 1729, 10,000 resamples) yields a 95% percentile CI = [0.005, 0.04]. Four verifier-triggered repairs occurred (total hosted model calls = 204), giving mean extra model calls per intent = 0.02 and mean extra calls per avoided violation = 1.0. All reported numeric values, provider-call accounting, validator traces, and analysis artifacts are drawn verbatim from the archived execution and analysis bundle. We interpret the preregistered confirmatory result conservatively and discuss construct, internal, and external validity limits and operational implications.

I. INTRODUCTION

Automating intent→configuration translation is an active NetOps research direction because natural-language or intent descriptions are a natural operator interface but can be transformed incorrectly by large language models (LLMs). Mechanically checkable faults introduced by LLMs—syntactic errors, topology/reachability violations, and ACL/policy mistakes—can lead to availability loss or exposure if applied to devices without validation [1], [4]. Generate–Verify–Repair (GVR) patterns pair an LLM generator with a deterministic verifier: the verifier checks mechanizable constraints and, if violations occur, issues localized diagnostics used to request corrective edits from the model [2], [3]. Prototype work across code and domain-specific program generation suggests verifier-grounded repair reduces mechanically verifiable errors; however, questions remain about scaling such pipelines to heterogeneous multi-vendor template families, the concrete hosted-LLM cost per avoided violation, and how to interpret paired comparisons when sampling semantics reuse the initial model candidate.

Research question and contribution. Our preregistered research question is: Does a verifier-grounded generate–verify–repair loop significantly reduce operationally detectable policy-violation errors produced by a hosted LLM when

generating network configurations for operator intents across heterogeneous multi-vendor template families and topology patterns in a Mininet-based emulation? Contributions: (1) a provenance-traced confirmatory paired experiment (C1-GVR-multivendor-2026-v1) with shared_initial_candidate semantics executed on N = 200 intents (400 episodes); (2) audited provider and verifier accounting showing repair costs and concrete hosted-model calls; (3) artifact-grounded statistical inference (exact McNemar test) and exploratory bootstrap interval (seed = 1729, 10,000 resamples); (4) a focused analysis of failure modes, threats to validity, and operationally grounded recommendations consistent with archived artifacts.

II. RELATED WORK

Surveys and NetOps-specific reviews document rapid growth in LLM-for-NetOps work and recurring evaluation gaps—benchmark provenance, verifier integration, and limited multi-vendor emulation evidence [5], [6]. Prior prototypes show LLMs can draft device configurations but often introduce syntax, topology, or policy errors; verifier-guided repair appears as a recurrent mitigation pattern in code and domain program generation [2], [3]. NetConfEval and related benchmarks demonstrate feasibility of evaluating LLMs on configuration tasks but commonly lack provenance-traced verifier pipelines or multi-vendor emulation at scale [4]. Our study builds on these findings by providing a provenance-traced paired design across synthetic vendor templates in Mininet and by executing the study under explicit shared_initial_candidate semantics to isolate verifier/repair effects.

III. METHODOLOGY

Preregistration and estimand. The experiment was preregistered as C1-GVR-multivendor-2026-v1 before any model calls (preregistration SHA-256: 9e0fc992665835975b01289f7841926ef6abf3584252c5d0459ba0ecf166bb50). Primary estimand: paired_success_rate_difference_guarded_minus_baseline (probability that an intent yields zero operationally detectable violations under guarded minus baseline), tested with an exact McNemar two-sided test at $\alpha = 0.05$. Predeclared exploratory estimators include a paired bootstrap percentile CI (bootstrap_seed = 1729; resamples = 10,000) and descriptive cost metrics (mean extra hosted model calls per intent and mean extra calls per avoided violation).

Sampling, strata, and shared_initial_candidate semantics. Task selection was deterministic and stratified across three synthetic vendor-template families and four topology patterns (12 strata) to exercise template heterogeneity. $N = 200$ unique operator intents were fixed prior to model calls. Execution used shared_initial_candidate semantics: exactly one initial sampled generation per intent was produced and reused by both baseline and guarded episodes. Under these semantics baseline outcome is the deterministic validator verdict on that shared initial candidate; the guarded outcome equals the final guarded pipeline result after deterministic validation and any permitted repair. The preregistration explicitly declared that independent constrained-prompt arms were not executed; therefore paired differences can arise only from deterministic validator or repair steps, not from different first-pass samples or prompt templates.

Model configuration, sampling notation, and budget. Declared model (as archived) is gpt-5-mini (provider recorded as openai_responses). The archived model_configuration.json records temperature = null/unset; we explicitly state that null/unset is not evidence of deterministic provider sampling and does not imply temperature = 0. Planned budget enforced in the adapter: planned_episode_count = 400, initial_generation_calls_per_task = 1, maximum_repair_calls_per_task = 1, maximum_model_calls = 400.

Deterministic verifier and repair policy. deterministic_netops_validator_v5 is a rule-based verifier combining (i) syntactic parsing and command pattern checks, (ii) required-command sequencing and workflow policy checks (per-task payload), (iii) reachability/topology checks against the emulated Mininet state, and (iv) encoded ACL/policy checks derived from the task payload. The scorer rule full_intent_constraint_validation yields a binary success = "no detected violations" if the final configuration satisfies all mechanizable constraints. Guarded pipeline behavior: if the validator flags violations, a single verifier-triggered repair call may be issued; repair prompts contain the verifier's localized diagnostics requesting corrective edits. Repair calls were permitted only when the validator reported violations, and up to one repair call per intent was allowed per the preregistration.

Scoring, missingness, and analysis plan. Primary analysis: exact McNemar two-sided test on paired binary success indicators. Missingness and exclusions were predeclared: duplicate tasks (detected via deterministic hashing) or unrecoverable infra failures would be excluded and reported; ambiguous validator outputs were conservatively treated as violations in the primary analysis. Exploratory analyses included the paired bootstrap percentile CI (seed = 1729, 10,000 resamples) and descriptive cost accounting. Analysis code and seeds are archived.

IV. RESULTS

Execution completion and deterministic accounting. The preregistered run completed without infrastructure exclusions: planned_episode_count = 400, completed_episode_count =

400, complete_pair_count = 200, failed_episode_count = 0 (execution manifest). Audited provider accounting (deterministic_reconciliation) reports hosted_model_call_event_count = 204; this decomposes deterministically into stage_counts = {shared_initial_generation: 200, repair: 4} and condition_counts recorded as {shared: 200, guarded: 4} in the manifest. Cache statistics show cache_hit_count = 0 and cache_miss_count = 204, confirming that every required hosted call executed and that the four repair calls were distinct provider calls. All deterministic consistency checks passed in the reconciliation artifacts.

Paired binary outcomes and primary confirmatory test. The preregistered paired contingency (analysis/paired_contingency_table.csv) is $n_{11} = 196$, $n_{10} = 4$, $n_{01} = 0$, $n_{00} = 0$. These cells yield baseline_success_rate = $196/200 = 0.98$ and guarded_success_rate = $200/200 = 1.00$; the paired risk difference (guarded - baseline) = 0.02. The preregistered exact McNemar two-sided test on these paired indicators returned $p = 0.125$ (analysis results). Under the preregistered inferential contract this p -value means we do not reject the null at $\alpha = 0.05$.

Repair events, candidate reuse, and per-intent cost accounting. The stage decomposition (200 shared initial generation events + 4 repair events) matches the model-call audit and demonstrates that the guarded pipeline invoked repairs for exactly four intents. From these audited counts we compute total_model_calls = 204, mean_model_calls_per_intent = $204/200 = 1.02$, mean_extra_calls_per_intent_due_to_repairs = $4/200 = 0.02$. Because each repair converted one failing baseline into a passing guarded outcome in the archived contingency, the observed mean_extra_calls_per_avoided_violation = $4 \text{ repairs} / 4 \text{ avoided violations} = 1.0$. These arithmetic derivations follow directly from the archived execution manifest and deterministic reconciliation artifacts.

Shared_initial_candidate evidence and representative artifact substantiation. Representative task artifacts included in the evidence bundle show explicit reuse of the same sampled initial generation by both baseline and guarded episodes. For example, the representative tasks (task-000004, task-000005, task-000006, etc.) exhibit identical shared_initial_candidate contents and identical shared_initial response file SHA-256 values across their baseline and guarded response artifacts. The archived validator_trace fields for the discordant pairs record validation_before entries showing violations and validation_after entries showing repair_applied = true with final validation_after.valid = true; the bundle therefore substantiates that the four discordant cases were resolved by verifier-triggered repairs operating on the identical initial candidate.

Diagnostic pattern observed in discordant pairs. In every archived discordant pair the validator_trace shows a mechanizable violation detected in validation_before (violation_count > 0 and violation_codes returned by deterministic_netops_validator_v5) and a subsequent repair application that produced a final configuration satisfying full_intent_constraint_validation. The manifest and per-task

scoring artifacts indicate that repairs were only issued for intents where the verifier produced explicit, localized diagnostics; no repair calls occurred for intents without detected violations. The archived traces therefore support the causal pathway: identical initial candidate $\rightarrow$ verifier detects mechanizable violation $\rightarrow$ single permitted repair call $\rightarrow$ final guarded configuration passes deterministic validation.

Exploratory bootstrap interval and interpretation caution. The archived exploratory paired bootstrap (bootstrap_seed = 1729; resamples = 10,000) produced a 95% percentile bootstrap CI for the paired risk difference = [0.005, 0.04]. We report this interval as exploratory only. Practically, the bootstrap percentile CI in paired binary settings relies on resampling paired outcomes and can be unstable when the total number of discordant pairs ($n_{10} + n_{01}$) is small; here $n_{10} + n_{01} = 4$ is small, which reduces the bootstrap’s reliability as a standalone uncertainty summary. Accordingly, we treat the exact McNemar test as the prespecified confirmatory inference ($p = 0.125$) and use the bootstrap interval only to convey plausible effect-size magnitudes suggested by the archived data.

Power context and observed baseline prevalence. The pre-registered power calculation targeted a 10 percentage-point absolute reduction with $N \approx 157$ (McNemar approximation) and noted that with $N = 200$ the detectable absolute risk difference under the same assumptions would be ≈ 8 percentage points. The observed baseline failure prevalence (2%) was substantially lower than scenarios that motivated the target effect size, which reduced the study’s power to detect small absolute improvements and produced the small discordant count (4) observed. This lower than anticipated baseline error rate is an empirical fact recorded in the archived condition summary and underlies the conservative interpretation of the null result from the pre-registered test.

Audit reconciliation and reproducibility pointers. All numerical claims above (paired cell counts, model_call totals, decomposition into shared initial and repair stages, bootstrap parameters, and seeds) are directly traceable to archived artifacts: execution/raw_results.jsonl, analysis/paired_contingency_table.csv, deterministic_reconciliation.json, execution/model_configuration.json, and analysis/analysis_log.jsonl. The execution manifest records started_at_utc and completed_at_utc timestamps for the run and the preregistration SHA-256; the analysis artifacts record the exact bootstrap_seed = 1729 and bootstrap_resamples = 10000 used to produce the exploratory interval. Reproducibility of the primary contingency counts does not require re-sampling because shared_initial_candidate semantics ensured exactly one initial sampled generation per intent was reused by both conditions; the archived response files and validator traces substantiate this reuse.

Summary of result-section evidence. In short: (1) the formal pre-registered confirmatory test did not reject the null hypothesis at $\alpha = 0.05$ (exact McNemar $p = 0.125$); (2) deterministic verifier-triggered repairs converted four failing shared initial candidates into passing final configurations, as shown by

per-task validator_trace entries; (3) the total hosted model-call overhead for these repairs was low (4 additional calls across 200 intents, mean extra calls per intent = 0.02, mean extra calls per avoided violation = 1.0); (4) the exploratory bootstrap CI [0.005, 0.04] suggests small plausible effect magnitudes but must be interpreted cautiously given the sparse discordant pair count. These findings are reported verbatim from the archived execution and analysis bundle and are discussed further in the Discussion section.

V. DISCUSSION

Causal interpretation under shared_initial_candidate semantics. Because execution used shared_initial_candidate semantics (one sampled initial generation per intent reused by baseline and guarded), paired improvements arise solely from deterministic validator/repair actions. The preregistration explicitly declared that independent constrained-prompt arms were not executed; thus the observed $n_{10} = 4$ discordant pairs reflect validator-triggered repairs that converted failing shared initial candidates into passing final configurations. Any statements attributing improvements to prompt-engineering or differing first-pass samples would be incorrect for this run.

Provider accounting and manifest labels. The archived execution manifest and deterministic reconciliation provide audited provider_call counts that clarify how model events map to pipeline stages. The manifest records hosted_model_call_event_count = 204 and decomposes that total into stage_counts = {"shared_initial_generation": 200, "repair": 4}. Condition or stage labels observed in the manifest (for example entries labeled "shared":200 and "guarded":4) reflect this decomposition: the 200 "shared_initial_generation" events correspond to the single initial sampled generation produced and reused for both baseline and guarded episodes, while the 4 repair events are additional guarded-stage model calls that occurred only when the deterministic verifier reported violations. This audited decomposition underpins the cost accounting reported in Results (mean extra model calls per intent = 0.02; mean extra calls per avoided violation = 1.0) and supports the causal claim that paired gains derive from verifier-triggered repairs rather than from independent re-sampling.

Statistical interpretation and bootstrap caution. The paired bootstrap percentile CI we report (95% CI = [0.005, 0.04]; bootstrap_seed = 1729; resamples = 10,000) is explicitly exploratory. Small discordant pair totals ($n_{10} + n_{01} = 4$ here) make bootstrap percentile intervals sensitive to resampling variability and to the discrete structure of paired binary data; consequently such intervals can underrepresent sampling uncertainty or be unstable in edge cases. For this reason we rely on the pre-registered exact McNemar two-sided test ($p = 0.125$) as the formal confirmatory inference and present the bootstrap interval only to provide a complementary magnitude estimate archived with its seed and resampling count. Researchers planning follow-ups should (a) anticipate discordant-count regimes when designing paired tests, (b) consider planning for larger N or higher baseline failure prevalence to increase discordant

counts, and (c) report both exact conditional inferences and multiple interval estimators to better triangulate effect size in sparse discordance settings.

Verifier coverage and residual semantic risks. `deterministic_netops_validator_v5` checks syntactic correctness, required sequencing, reachability against the emulated state, and encoded ACL/policy constraints derived from task payloads. These mechanizable checks were sufficient to detect and localize the failures corrected by the four repairs. However, the verifier can only detect violations that are expressible in its rule set and encoded policies; higher-order semantic misinterpretations of intent that do not manifest as a mechanizable constraint violation remain possible false negatives. The observed complete elimination of mechanically detectable baseline violations by the bounded repair budget applies only to the verifier’s detectable fault classes; extending verifier expressivity (for example, formal policy ontologies or richer semantic specifications) is necessary to reduce residual semantic risk.

Operational implications and bounded generality. Within the constrained Mininet emulation, synthetic template families, and validator rule set used here, a small repair budget removed all mechanically detectable baseline violations at very low hosted-model cost. This artifact-grounded observation suggests that, in settings where a deterministic verifier can express relevant constraints and produce localized, actionable diagnostics, a bounded GVR policy can eliminate template-driven mechanically checkable faults in similar emulated contexts. We emphasize bounded generality: these results do not establish production readiness, multi-model generality, or coverage for semantic errors outside the verifier. Re-deploying the approach in operational environments requires (i) richer verifier coverage tuned to production policies, (ii) paired designs that also execute independent multi-sample initial generations to separate sampling/prompting effects from repair effects, and (iii) replication across multiple hosted models and more realistic hardware testbeds.

Recommendations for future evaluations. Based on the archived evidence and the statistical constraints observed, we recommend that future confirmed-design studies: (1) preregister whether `shared_initial_candidate` semantics or independent-sample semantics will be used and plan estimands accordingly; (2) select sample sizes or baseline scenarios that anticipate sufficient discordant pairs for confirmatory power, or else use study designs that enlarge expected discordance (e.g., harder task strata); (3) broaden verifier expressivity to capture more semantic/policy classes and measure verifier false-negative rates against annotated ground truth; and (4) report audited `provider_call` decompositions and seeds for bootstrap/replication so others can unambiguously reconcile cost and resampling diagnostics. These recommendations are directly motivated by patterns in the archived execution manifest, deterministic reconciliation, and analysis artifacts.

VI. LIMITATIONS

- Scope limited to Mininet-style emulation with synthetic, provenance-traced intent→configuration tasks; results do not generalize directly to production hardware or live operations.
- Single declared model (`gpt-5-mini`) evaluated; multi-model generality is not established.
- Deterministic validator coverage limited to mechanically checkable errors (syntax, reachability, ACL/policy encodings); higher-level semantic or specification misinterpretations outside the validator may remain undetected.
- `Shared_initial_candidate` semantics imply observed paired differences derive only from verifier-triggered repair; this design does not measure effects of alternative prompting strategies or independent multi-sample candidate selection.
- Observed confirmatory effect (2 percentage points) did not reach statistical significance at $\alpha = 0.05$ for $N = 200$ (exact McNemar $p = 0.125$); low baseline failure prevalence (2%) reduced power to detect small absolute improvements under some preregistered scenarios.

VII. CONCLUSION

We preregistered and executed a provenance-traced paired confirmatory study (`C1-GVR-multivendor-2026-v1`) using `shared_initial_candidate` semantics to measure whether a deterministic verifier plus bounded repair reduces mechanically detectable policy violations for intent→configuration tasks in a Mininet emulation. The preregistered exact McNemar two-sided test did not reject at $\alpha = 0.05$ ($p = 0.125$). Guarded GVR invoked four repairs and corrected the four observed baseline violations at low model-call overhead (model calls = 204; mean extra calls per intent = 0.02), yielding a paired risk difference estimate of 0.02 and an exploratory paired bootstrap CI = [0.005, 0.04] (bootstrap_seed = 1729, resamples = 10,000). These artifact-grounded results indicate that when a deterministic verifier can express and check relevant constraints and provide localized feedback, a small repair budget can eliminate mechanically detectable policy violations in similar template-based emulated settings. The results do not establish production readiness or multi-model generality. All reported numbers, provider accounting, validator traces, and analysis artifacts are drawn verbatim from the archived execution and analysis bundles referenced in the preregistration and execution manifests.

References [1] R. Mondal, A. Tang, R. Beckett, T. Millstein, and G. Varghese, "What do LLMs need to Synthesize Correct Router Configurations?", *Proceedings of the ACM*, 2023, doi: 10.1145/3626111.3628194. [2] R. Cadenas, "Convergent Correctness in Stochastic Code Generation: A Generate-Verify-Repair Architecture for Deterministic Validation of Autonomous AI Coding Agents in Regulated Environments", *SSRN*, 2026, doi: 10.2139/ssrn.6754899. [3] Z. Ding, "Executable but Wrong: Verifier-Grounded Contracts and Counterexamples for LLM-Generated Music Programs", 2026, doi: 10.31224/7995. [4] C. Wang, M. Scazzariello, A.

Farshin, S. Ferlin, D. Kostić, and M. Chiesa, "NetConfEval: Can LLMs Facilitate Network Configuration?", Proceedings of the ACM, 2024, doi: 10.1145/3656296. [5] S. Y. Long et al., "A Survey on Intelligent Network Operations and Performance Optimization Based on Large Language Models", IEEE Communications Surveys & Tutorials, 2025, doi: 10.1109/comst.2025.3526606. [6] L. Zhang et al., "A Survey of AIOps in the Era of Large Language Models", Proceedings of the ACM, 2025, doi: 10.1145/3746635.

DISCLOSURE STATEMENT

Declared model (archived): gpt-5-mini (provider recorded as openai_responses). Adapter family: hosted_netops_gvr_v1. Deterministic validator: deterministic_netops_validator_v5. Task generator: deterministic_netops_task_generator_v5. Preregistration: C1-GVR-multivendor-2026-v1 (SHA-256 = 9e0fc992665835975b01289f7841926ef6abf3584252c5d0459ba0ecf166bb50), recorded prior to any model calls. Initial master prompt archived at provenance/master_prompt.txt (SHA-256 = 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ba). Human role: a human Research Director selected the topic family and constrained the initial master prompt prior to the autonomy lock; after the lock the AI autonomously performed literature verification, preregistration, experiment execution, analysis, peer review, revision, and manuscript drafting; no human manually edited technical manuscript content after the autonomy lock. Sampling semantics: execution used shared_initial_candidate (exactly one sampled initial generation per intent was produced and reused by both baseline and guarded); archived model configuration records temperature = null/unset (null/unset is not evidence of deterministic provider sampling and does not imply temperature = 0). Provider accounting (audited): hosted_model_call_event_count = 204 decomposed as 200 shared initial generations + 4 repair calls. Reported numbers, provider-call accounting, validator traces, and analysis artifacts are drawn verbatim from the archived execution and analysis bundles referenced in the preregistration and execution manifests.

REFERENCES

- [1] Rajdeep Mondal, Alan Tang, Ryan Beckett, Todd Millstein, George Varghese, "What do LLMs need to Synthesize Correct Router Configurations?", 2023, 10.1145/3626111.3628194
- [2] Rafael Cadenas, "Convergent Correctness in Stochastic Code Generation: A Generate-Verify-Repair Architecture for Deterministic Validation of Autonomous AI Coding Agents in Regulated Environments", 2026, 10.2139/ssrn.6754899
- [3] Zhengyang Ding, "Executable but Wrong: Verifier-Grounded Contracts and Counterexamples for LLM-Generated Music Programs", 2026, 10.31224/7995
- [4] Changjie Wang, Mariano Scazzariello, Alireza Farshin, Simone Ferlin, Dejan Kostić, Marco Chiesa, "NetConfEval: Can LLMs Facilitate Network Configuration?", 2024, 10.1145/3656296
- [5] S.Y. Long, Jingjing Tan, Bomin Mao, Fengxiao Tang, Yangfan Li, Ming Zhao, Nei Kato, "A Survey on Intelligent Network Operations and Performance Optimization Based on Large Language Models", 2025, 10.1109/comst.2025.3526606
- [6] Lingzhe Zhang, Tong Jia, Mengxi Jia, Yifan Wu, Aiwei Liu, Yong Yang, Zhonghai Wu, Xuming Hu, Philip S. Yu, Ying Li, "A Survey of AIOps in the Era of Large Language Models", 2025, 10.1145/3746635