

Verifier-Guided Evaluation of LLM-Generated Network Configurations: a Paired Study with Shared Initial Candidates

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory study (CAND-2026-001) that measured whether a verifier-triggered guarded pipeline (generate → validate → repair, ≤1 repair) reduces deterministic-invariant violations relative to single-shot initial generation for intent→configuration NetOps tasks. Using a deterministic synthetic task generator and a deterministic network verifier, we ran 200 preregistered tasks under shared_initial_candidate semantics with the hosted model gpt-5-mini (execution/model_configuration.json records temperature = null/unset). Baseline configurations passed the verifier in 199/200 tasks (99.5%); the guarded pipeline passed 200/200 (100%). Paired absolute risk difference (guarded - baseline) = 0.005; contingency counts n11=199, n10=1, n01=0, n00=0; exact two-sided McNemar $p = 1.0$; paired-bootstrap 95% percentile CI = [0.00, 0.015] (10,000 resamples, seed 1729). The single guarded improvement arose from one verifier-triggered repair that corrected a multi-step ordering violation. We report preregistration, full execution accounting (201 model calls: 200 shared initial + 1 repair), paired analysis, representative diagnostic artifacts, and bounded limitations grounded in the archived execution bundle.

I. INTRODUCTION

Generative large language models (LLMs) are increasingly proposed to translate high-level operator intent into device configurations for network operations (NetOps). Operational correctness for such workflows requires generated configurations to satisfy deterministic invariants (reachability, policy compliance, workflow ordering, and group-level safety). Prior prototype systems and benchmarks illustrate feasibility but leave open questions about how much deterministic verification and bounded repair alter acceptance rates when the identical initial sample is reused across comparison conditions.

This paper reports a preregistered paired experiment (CAND-2026-001) designed to isolate the marginal effect of a verifier-triggered automated repair on an identical generated candidate. Under shared_initial_candidate semantics the same single initial generation is deterministically reused by both baseline and guarded episodes; any paired difference therefore arises solely from the permitted validator-triggered repair. The primary estimand is the paired difference in per-task binary acceptance by a deterministic verifier (paired_success_rate_difference_guarded_minus_baseline). Because shared_initial_candidate semantics were enforced, this estimand measures a repair-only effect and does not capture potential benefits from independent guarded first-pass generations, constrained prompts, or multi-sample selection.

Contributions. (1) A preregistered, fully archived paired experiment executing 200 intent→configuration tasks with a deterministic verifier and a hosted LLM (gpt-5-mini) under shared_initial_candidate semantics. (2) Complete execution accounting and deterministic reconciliation showing 200 shared initial generations and 1 repair call (201 model calls total). (3) Artifact-grounded confirmatory analysis reporting baseline pass 199/200 and guarded pass 200/200 with contingency counts (n11=199, n10=1, n01=0, n00=0), paired risk difference 0.005, exact two-sided McNemar $p = 1.0$, and bootstrap 95% CI = [0.00, 0.015] (10,000 resamples, seed 1729). (4) A localized failure diagnosis for the single repaired pair using archived validator traces. (5) Detailed reproducibility guidance and bounded limitations tied to archived artifacts and reconciliation records.

II. RELATED WORK

We conducted a fresh literature investigation and verified records covering intent→configuration systems, generate→validate→repair patterns, and agentic-AI safety discussions relevant to NetOps. Prior intent-to-configuration prototypes and task-bench evaluations demonstrate that LLMs can synthesize device configurations from operator intents and that verification is widely advocated as a guardrail [1]. Work such as NetConfEval provides task-oriented benchmark evidence that LLMs can facilitate configuration synthesis, but reported evaluations commonly focus on syntactic or task-level correctness metrics rather than verifier-backed invariant enforcement across deterministic topologies [1].

Deterministic verification and formal checking of configuration invariants have a separate, longer literature in networking. Beckett et al. present a general approach to network configuration verification that emphasizes encoding device semantics and network-wide properties into analyzable specifications; such verifier-focused methods clarify the kinds of invariants a downstream oracle must implement to make generator→validator pipelines meaningful in practice [2]. Our study situates itself at the intersection of these bodies of work: rather than benchmarking raw generative capability, we measure how an explicit verifier gate and a bounded repair step change acceptance rates when the identical initial candidate is reused.

Cross-domain empirical work on generate→validate→repair and test-in-the-loop LLM repair

in software engineering provides a methodological precedent. Test-driven or test-in-loop program-repair demonstrations show that pairing generation with a deterministic test oracle (failing test $\rightarrow$ patch $\rightarrow$ regression test) yields reproducible corrections for many classes of faults, and those studies emphasize artifact archiving to permit independent reproduction [4]. Analogously, tool-augmented agent studies and analyses of LLM agents that call external analyzers or solvers argue that coupling models with deterministic tools materially changes task outcomes and creates audit points amenable to reproducible evaluation [6]. However, these demonstrations are primarily in code or program-synthesis domains; adapting their empirical lessons to NetOps requires careful mapping of network-specific invariants (ordering, group constraints, dependency rules) and verifier semantics.

Survey and reporting-guideline work underline evaluation and reporting challenges for LLM studies. TRIPOD-LLM and broad LLM surveys call for rigorous preregistration, honest reporting of sampling and randomness settings, and preservation of execution artifacts to support reproducibility and avoid selective reporting [3],[5]. These calls motivated our preregistration, deterministic task selection, and archival of provider-call accounting and verifier traces. In particular, TRIPOD-style guidance stresses that reporting model sampling parameters is necessary even when fields are unset; execution bundles must record temperature/null semantics and the reuse policy for initial candidates so readers can correctly interpret paired designs [3].

Comparative synthesis. Across the verified records there is a consistent pattern: (a) promising LLM-based intent $\rightarrow$ config prototypes and benchmarks exist, (b) generate $\rightarrow$ validate $\rightarrow$ repair architectures are advocated and empirically effective in adjacent domains, and (c) end-to-end verifier-integrated NetOps experiments with full artifact grounding and preregistered paired estimands are scarce. Our paired shared_initial_candidate design directly addresses this methodological gap by (i) enforcing identical initial candidates across baseline and guarded episodes, (ii) using a deterministic verifier whose per-episode validator_trace fields provide fine-grained evidence about the nature of failures (parsed_command_count, violation_codes, required_command_count), and (iii) archiving provider-call accounting and stage_counts so that the precise degree of model interaction (shared_initial_generation=200, repair=1) is auditable.

Limitations in the prior literature that motivate our approach. Many prior works evaluate generation quality via held-out references or human judgement rather than deterministic invariant enforcement; others integrate emulators or dataplane tests but do not archive the exact paired generation semantics needed to isolate repair-only effects. By contrast, the verifier-centric studies and test-in-the-loop program-repair literature demonstrate that a deterministic oracle can make repair effects measurable and reproducible, but these prior empirical results must be adapted to NetOps because network invariants include ordering and group-level constraints that differ from unit tests

used in program repair. Our contribution is therefore methodological and empirical: we apply a preregistered, verifier-grounded paired protocol to quantify the marginal repair effect under shared_initial_candidate semantics and provide the full artifact bundle needed to reproduce that narrowly scoped estimand.

III. METHODOLOGY

Preregistration and design freeze. The study design and analysis plan were frozen in preregistration CAND-2026-001 prior to any model calls. The preregistration fixed: (a) a paired binary estimand (paired_success_rate_difference_guarded_minus_baseline) comparing per-task binary verifier acceptance under baseline (no repair) versus guarded (verifier may trigger ≤ 1 repair) conditions; (b) deterministic task selection for $N=200$ tasks; (c) generation_semantics = shared_initial_candidate (one sampled initial generation per task deterministically reused by both episodes); (d) a maximum of one automated repair call per task; (e) primary failed-call treatment complete_pair_primary and a sensitivity treatment that counts persistent unparsables as failures; and (f) the confirmatory test (exact two-sided McNemar via exact binomial tail) plus a paired-bootstrap procedure for a percentile CI. The preregistration artifact (immutable) is referenced in the execution manifest.

Deterministic task sampling. Tasks were produced by a deterministic task generator (deterministic_netops_task_generator_v5) and selected by a fixed index rule documented in the preregistration. The selection rule is deterministic across the 200-task bank (the preregistration specifies the cycle-and-offset index mapping used to pick each task index); this ensures reproducible task composition and eliminates post-hoc selection. Each manifest encodes: initial_state, expected_final_state, required_sequence (ordered command list), allowed_touched_fields, workflow_policies (group-level constraints, minimum-active rules), and difficulty metadata (changed_interface_count, dependency_rule_count, required_command_count). Difficulty labels used in the study include very_hard and extreme; representative required_command_count values ranged from 3 up to 9 in the archived sample.

Shared_initial_candidate semantics and sampling notes. The adapter contract enforced shared_initial_candidate semantics: for each pair the exact same initial generation (a single sampled candidate) was provided to both the baseline and guarded episodes. Consequently, any observed paired difference can only arise from the action of the deterministic verifier and the single permitted repair; it cannot arise from independent first-pass guarded generations, different sampling calls, prompt-engineering differences, or selection among multiple generated candidates. The model-configuration record declares the requested model as gpt-5-mini and records temperature = null (unset). We explicitly note that temperature = null/unset in the configuration file is not evidence of deterministic sampling and does not imply temperature = 0; nondeterministic hosted-model sampling therefore remains a possible background prop-

erty even though the experiment used a single initial sample per pair by design.

Generation, sanitizer, and repair policy. For each task the pipeline executed a deterministic formatting sanitizer (`automatic_syntax_sanitizer_v1`) to correct minor syntactic normalization issues before verifier parsing; this sanitizer is deterministic and is applied prior to scoring per the preregistration. After the shared initial generation and sanitizer, the deterministic verifier parsed the candidate. If the verifier reported any invariant violations, the guarded logic was permitted to invoke at most one automated repair call to the hosted model with a repair prompt constructed from the validator trace (this repair policy and prompt format were specified in the preregistration). Repair calls were counted against the `maximum_model_calls` budget; the execution manifest reports exactly one repair call across all pairs in the run. The pipeline’s missingness plan treated persistent unparsable outputs (after sanitizer and an allowed repair attempt) as scorable failures in the sensitivity analysis; in the primary analysis pairs with complete episodes remained eligible per the preregistered `complete_pair_primary` rule.

Deterministic verifier as the oracle. The deterministic verifier (`deterministic_netops_validator_v5`) is the study’s operational oracle. The verifier implements: syntactic parsing into discrete commands; enforcement of `allowed_fields` and `allowed_touched_fields` per task manifest; comparison of parsed command sequences against the manifest’s `required_sequence` ordering; evaluation of group-level invariants (e.g., `final_group_constraints` and `minimum_active_in_groups`); and a final boolean `full_intent_constraint_validation = true` only if all encoded constraints are satisfied. The verifier emits structured tracer output for each episode that includes `parsed_command_count`, `required_command_count`, `observed_touched_field_count`, `normalized_configuration`, `validation_before` (violation codes and counts), `validation_after` (post-repair state and validity flag), `validator_id` and version, and a human-readable description of detected invariant breaches. These tracer fields are the authoritative evidence for per-episode scoring in the analysis.

Paired execution bookkeeping and provider accounting. The preregistered planned episodes were 400 (200 pairs $\times$ 2 conditions). The run completed all planned episodes without terminal provider failures. Provider accounting recorded hosted-model-call events and stage counts consistent with preregistration: the stage counts indicated 200 shared initial generation events and 1 repair event, for a total of 201 model calls. Cache and reuse checks verified that the shared initial candidate was reused per pair and that no cross-condition cache-reuse introduced additional independent generation calls. All runtime provider metadata (call counts, cache hits/misses, stage counts) were reconciled deterministically in the execution reconciliation artifact.

Statistical procedures — confirmatory and CI computation. The primary confirmatory test is the exact two-sided McNemar test computed from discordant pair counts n_{10} and n_{01} (where n_{10} = count of pairs with `guarded=1` & `baseline=0`, and n_{01}

= count with `guarded=0` & `baseline=1`). We implemented the exact two-sided McNemar p-value using the exact binomial-tail formulation: let $n = n_{10} + n_{01}$ and $k = n_{10}$; the two-sided p-value is $p = 2 * \min\{ \text{BinomCDF}(k; n, 0.5), 1 - \text{BinomCDF}(k - 1; n, 0.5) \}$ with the conventional handling when $k = 0$. This exact-binomial realization is documented in the analysis log and reproduces the $p = 1.0$ reported for the observed discordants ($n = 1, k = 1$). For the paired absolute risk difference we computed a paired-bootstrap 95% percentile confidence interval with 10,000 resamples and `bootstrap_seed = 1729`; bootstrap parameters and seed are recorded in the analysis log.

Missingness, exclusions, and contamination controls. The preregistration defined deterministic exclusion criteria prior to model calls (e.g., token-normalized identity to public vendor-source templates to avoid leakage) and required logging of any excluded tasks. The run recorded no preregistered exclusions. The sanitizer and missingness policy were applied deterministically; persistent unparsables did not occur in this execution. Contamination checks included token-level similarity scans between generated candidates and vendor-template sources and a contamination summary that would flag unusually high overlaps; no contamination flags were raised in this run.

Reproducibility and artifact grounding. All implementation details required to reproduce the methodology are archived in the execution bundle: the immutable preregistration, the deterministic task generator artifacts and manifests (task bank with `required_sequence` and `workflow_policies`), the model-configuration record (including declared model and the explicit `temperature = null/unset` entry), the verifier implementation and structured tracer outputs, provider-call accounting and deterministic reconciliation, the sanitizer and repair-prompt templates, and the analysis-executor configuration (exact McNemar code path and bootstrap seed/resample settings). The verifier’s tracer fields provide the operational link between generated text and the binary scoring used in the confirmatory analysis. Re-running the archived paired analysis executor on the raw results reproduces the contingency counts and CI parameters reported in the paper.

IV. RESULTS

Execution accounting and provider audit. The execution manifest records start and end timestamps (UTC) and provider-call accounting: the run started at 2026-09-01T16:03:05.491039+00:00 and completed at 2026-09-01T16:32:13.298536+00:00 (≈ 29 minutes wall-clock). `Planned_episode_count = 400` and `completed_episode_count = 400`; `failed_episode_count = 0`. Provider-call audit (`deterministic_reconciliation.json`) reports `hosted_model_call_event_count = 201` with `stage_counts = {shared_initial_generation: 200, repair: 1}`. These counts match the preregistered `shared_initial_candidate` semantics and the single observed repair.

Primary confirmatory outcomes (archived CSV artifacts). `analysis/condition_summary.csv` shows: Baseline successes = 199, failures = 1 (`success_rate = 0.995`); Guarded successes

= 200, failures = 0 (success_rate = 1.000). The paired contingency table (analysis/paired_contingency_table.csv) records $n_{11} = 199$, $n_{10} = 1$, $n_{01} = 0$, $n_{00} = 0$. The paired absolute risk difference (guarded - baseline) = 0.005. Exact two-sided McNemar $p = 1.0$ as computed above; paired-bootstrap percentile 95% CI (10,000 resamples, seed 1729) = [0.00, 0.015]. All numeric values are present in analysis/*.csv and analysis/results.json and reproduced by the archived paired_binary_analysis_v1 executor.

Representative repaired case and diagnostics. The guarded pipeline invoked exactly one automated repair. Validator traces for the repaired pair indicate an extreme composite task (required_command_count = 9, changed_interface_count = 3) where the shared initial candidate violated the task’s required_sequence ordering constraint. The repair applied a localized reordering that aligned the command sequence to required_sequence and yielded validation_after.valid = true. Per-episode validator_trace fields (validation_before.violation_codes, parsed_command_count, required_command_count, normalized_configuration) document both the pre-repair violation and the post-repair correction. These archived traces (execution/scoring/task-00000X-*.json) are the evidence for the localized repair effect and support the preregistered expectation that multi-device sequencing errors concentrate in higher-difficulty tasks.

Missingness, exclusions, and contamination. The run recorded no preregistered exclusions and no persistent unparsables; analysis/missingness_summary.csv shows complete_pairs = 200 and zeros for other missingness categories. analysis/contamination_summary.csv is empty (no flagged contamination). Deterministic_reconciliation.json marks all checks passed and pair_accounting_consistent = true. Provider-call audit shows cache_miss_count = 201 and cache_hit_count = 0.

Exploratory diagnostics by difficulty metadata. Representative task manifests illustrate required_command_count values from 3 to 9 and difficulty labels very_hard or extreme. The single baseline failure is associated with an extreme pattern and high required_command_count, consistent with the pre-registered exploratory question that multi-device sequencing and dependency complexity dominate residual failures. analysis/paired_difference_figure.svg (archived) visualizes condition rates and the bootstrap distribution used for the CI calculation.

V. DISCUSSION

Interpretation under shared_initial_candidate semantics. Because baseline and guarded conditions reuse the identical sampled candidate per pair, observed differences isolate the marginal effect of permitting an automated verifier-triggered repair. The experiment therefore quantifies a repair-only effect and does not evaluate other guarded variants that would permit independent guarded first-pass generations or constrained prompting. We reiterate this estimand interpretation explicitly so readers do not generalize to independent-generation contrasts without new experiments.

Statistical regime, power, and limitations of inference. Our preregistered power rationale targeted detection of moderate absolute paired increases (~10–15 percentage points) at $N = 200$ under mid-range baseline assumptions (~45–65%). The realized baseline success rate ($199/200 = 99.5\%$) places the analysis in a sparse-discordant regime where the exact McNemar p -value gives limited directional evidence ($p = 1.0$ for our observed discordants). The paired-bootstrap CI [0.00, 0.015] provides a bounded estimate of plausible effect sizes compatible with the data and shows the dataset is consistent with at most small absolute improvements under our repair-only design. This underscores that unexpectedly high baseline performance can substantially reduce realized inferential sensitivity for small marginal effects.

Operational implications and bounded claims. Artifact-grounded evidence supports a narrow operational conclusion: in this synthetic deterministic task bank using gpt-5-mini and deterministic_netops_validator_v5, a verifier-triggered bounded repair corrected a localized multi-step ordering fault in one extreme task. This is direct evidence that verifier-guided bounded repair can remedy certain ordering/dependency violations in generated configurations when those violations are expressible in the task manifest’s invariants. It is not evidence that bounded repair will correct omissions only observable at dataplane runtime, or that guarded repair will address adversarially poisoned contexts; those remain open evaluation axes requiring live emulation, multi-model studies, and adversarial contamination experiments.

Threats to validity and generalization. Internal validity: the paired shared_initial_candidate design isolates repair effects but does not address prompt sensitivity or model sampling variability. Construct validity: the verifier enforces encoded syntactic and ordering invariants but does not substitute for dataplane emulation or vendor-specific runtime semantics. External validity: the experiment used a single model (gpt-5-mini) and a synthetic, deterministic task bank; results do not imply cross-model or production-device generality. Conclusion validity: with only one discordant pair the statistical power to detect small repair-only effects is limited; the bootstrap CI quantifies that uncertainty.

Reproducibility, artifact grounding, and recommended reanalyses. All execution and analysis artifacts are archived in the execution bundle. Key artifacts for reanalysis include: execution/task_manifest.jsonl (task selection and manifests), execution/model_configuration.json (declared model and sampling fields), execution/responses/ (shared-initial and per-condition responses), execution/scoring/ (validator_trace per episode), analysis/*.csv (paired contingency and condition summaries), deterministic_reconciliation.json (provider accounting and stage_counts), and analysis/analysis_log.jsonl (bootstrap parameters and seed). The initial master prompt is archived (provenance/master_prompt.txt; SHA-256 provided in the Disclosure Statement). Re-executing paired_binary_analysis_v1 on the archived input_results_sha256 reproduces the confirmatory numbers reported above. For sensitivity, the preregistered alternative

treatment (persistent unparseables as failures) is supported by the artifacts but was not needed because such cases did not occur.

VI. LIMITATIONS

- Synthetic deterministic task bank and verifier: results reflect correctness under the chosen synthetic intent templates and deterministic verifier and do not guarantee similar performance on live heterogeneous production devices or telemetry.
- Single-model experiment (gpt-5-mini): findings do not demonstrate multi-model generality.
- Shared_initial_candidate semantics: baseline and guarded conditions began from the same initial generation; the study measures verifier/repair effects only and does not evaluate independent first-pass generation contrasts.
- Limited adversarial/contamination testing: the run did not include systematic adversarial retrieval or poisoned-context attacks; such threats remain an open evaluation axis.
- Oracle coverage: the deterministic verifier enforces encoded invariants but may omit runtime behaviors and device-specific semantics observable only through live execution; external validity to live deployments is therefore limited.

VII. CONCLUSION

We executed a preregistered paired experiment measuring the marginal effect of a verifier-triggered repair under shared_initial_candidate semantics for LLM-generated network configurations. In 200 synthetic deterministic tasks with gpt-5-mini and deterministic_netops_validator_v5, baseline acceptance was 199/200 and guarded acceptance was 200/200; a single automated repair corrected a multi-device ordering violation. Paired absolute risk difference = 0.005 with paired-bootstrap 95% CI [0.00, 0.015] (10,000 resamples, seed 1729) and exact two-sided McNemar $p = 1.0$. Artifact-grounded evidence therefore supports a constrained conclusion: verifier-guarded bounded repair can correct localized ordering faults in this synthetic verifier-backed setting. Broader operational claims require additional experiments: multi-model studies, independent-first-pass guarded semantics, adversarial contamination testing, and live-device or dataplane emulation to capture runtime semantics not modeled by the deterministic verifier.

DISCLOSURE STATEMENT

Disclosure: This run implemented preregistered study CAND-2026-001 and was executed autonomously after the autonomy lock. Declared model: gpt-5-mini (execution/model_configuration.json records temperature = null/unset; null/unset is not evidence of deterministic sampling or temperature = 0). Orchestration adapter: hosted_netops_gvr_v1. Exact initial master prompt archived at provenance/master_prompt.txt (SHA-256 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582b

b39b15ba). Preregistration fixed generation_semantics = shared_initial_candidate (one sampled initial generation per task was deterministically reused by baseline and guarded episodes) and allowed at most one automated repair call per task. Model-call accounting in the archived execution manifest reflects 200 shared initial generation calls and 1 repair call (201 model calls total). The deterministic verifier used was deterministic_netops_validator_v5. Humans selected the broad topic family and constrained the initial master prompt prior to the autonomy lock as permitted by the track; no human manually edited technical manuscript content after the autonomy lock. All provider traces, verifier logs, task manifests, model-configuration records, and analysis artifacts are archived in the execution bundle and indexed by the execution manifest.

REFERENCES

- [1] Changjie Wang, Mariano Scazzariello, Alireza Farshin, Simone Ferlin, Dejan Kostić, Marco Chiesa, "NetConfEval: Can LLMs Facilitate Network Configuration?", 2024, 10.1145/3656296
- [2] Ryan Beckett, Aarti Gupta, Ratul Mahajan, David Walker, "A General Approach to Network Configuration Verification", 2017, 10.1145/3098822.3098834
- [3] Jack Gallifant, Majid Afshar, Saleem Ameen, Yindalon Aphinyanaphongs, Shan Chen, Giovanni Cacciamani, Dina Demner-Fushman, Dmitriy Dligach, Roxana Daneshjou, Chrystinne Oliveira Fernandes, Lasse Hyldig Hansen, Adam Landman, Lisa Soleymani Lehmann, Liam G. McCoy, Timothy A. Miller, Amy C. Moreno, Nikolaj Munch, David Restrepo, Guergana Savova, Renato Umeton, Judy Wawira Gichoya, Professor Gary S. Collins, Karel G.M. Moons, Leo Anthony Celi, Danielle S. Bitterman, "The TRIPOD-LLM reporting guideline for studies using large language models", 2025, 10.1038/s41591-024-03425-5
- [4] Yunhe Li, "Test-in-the-loop LLM Repair: Verifiable Automated Program Repair on QuixBugs with a "Failing Test → Patch → Regression Test" Loop", 2024, 10.69987/jacs.2024.40206
- [5] Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Zican Dong, Yupeng Hou, Beichen Zhang, Yingqian Min, Junjie Zhang, Peiyu Liu, Xiaolei Wang, Yifan Du, Yushuo Chen, Yushuo Chen, Zhipeng Chen, Jinhao Jiang, Ruiyang Ren, Yifan Li, Xinyu Tang, Peiyu Liu, Yiwen Hu, Jian-Yun Nie, Ji-Rong Wen, "A Survey of Large Language Models", 2026, 10.1007/s11704-026-60308-3
- [6] Haoran Dong, "Tool-Augmented Large Language Model Agents for Analysis: A Focused Study", 2026, 10.2139/ssrn.6647800