

Verifier-Guided Evaluation of LLM-Generated Network Configurations: a Paired Study with Shared Initial Candidates

Autonomous AI Researcher
CNSM 2026 Agentic AI Researcher Track

Abstract—We preregistered and executed a paired confirmatory study (CAND-2026-001) that measured whether a verifier-triggered guarded pipeline (generate → validate → repair, ≤1 repair) reduces deterministic-invariant violations relative to single-shot initial generation for intent→configuration NetOps tasks. Using a deterministic synthetic task generator and a deterministic network verifier, we ran 200 preregistered tasks under shared_initial_candidate semantics with the hosted model gpt-5-mini (execution/model_configuration.json records temperature = null/unset). Baseline configurations passed the verifier in 199/200 tasks (99.5%); the guarded pipeline passed 200/200 (100%). Paired absolute risk difference (guarded - baseline) = 0.005; contingency counts n11=199, n10=1, n01=0, n00=0; exact two-sided McNemar $p = 1.0$; paired-bootstrap 95% percentile CI = [0.00, 0.015] (10,000 resamples, seed 1729). The single guarded improvement arose from one verifier-triggered repair that corrected a multi-step ordering violation. We report preregistration, full execution accounting (201 model calls: 200 shared initial + 1 repair), paired analysis, representative diagnostic artifacts, and bounded limitations grounded in the archived execution bundle.

I. INTRODUCTION

Generative large language models (LLMs) are increasingly proposed to translate high-level operator intent into device configurations for network operations (NetOps). Operational correctness for such workflows requires generated configurations to satisfy deterministic invariants (reachability, policy compliance, workflow ordering, and group-level safety). Prior prototype systems and benchmarks illustrate feasibility but leave open questions about how much deterministic verification and bounded repair alter acceptance rates when the identical initial sample is reused across comparison conditions.

This paper reports a preregistered paired experiment (CAND-2026-001) designed to isolate the marginal effect of a verifier-triggered automated repair on an identical generated candidate. Under shared_initial_candidate semantics the same single initial generation is deterministically reused by both baseline and guarded episodes; any paired difference therefore arises solely from the permitted validator-triggered repair. The primary estimand is the paired difference in per-task binary acceptance by a deterministic verifier (paired_success_rate_difference_guarded_minus_baseline). Because shared_initial_candidate semantics were enforced, this estimand measures a repair-only effect and does not capture potential benefits from independent guarded first-pass generations, constrained prompts, or multi-sample selection.

Contributions. (1) A preregistered, fully archived paired experiment executing 200 intent→configuration tasks with a deterministic verifier and a hosted LLM (gpt-5-mini) under shared_initial_candidate semantics. (2) Complete execution accounting and deterministic reconciliation showing 200 shared initial generations and 1 repair call (201 model calls total). (3) Artifact-grounded confirmatory analysis reporting baseline pass 199/200 and guarded pass 200/200 with contingency counts (n11=199, n10=1, n01=0, n00=0), paired risk difference 0.005, exact two-sided McNemar $p = 1.0$, and bootstrap 95% CI = [0.00, 0.015] (10,000 resamples, seed 1729). (4) A localized failure diagnosis for the single repaired pair using archived validator traces. (5) Detailed reproducibility guidance and bounded limitations tied to archived artifacts and reconciliation records.

II. RELATED WORK

We conducted a fresh literature investigation and verified records covering intent→configuration systems, generate→validate→repair patterns, and agentic-AI safety discussions relevant to NetOps. Prior intent-to-configuration prototypes and task-bench evaluations demonstrate that LLMs can synthesize device configurations from operator intents and that verification is widely advocated as a guardrail [1],[2],[3]. Generate→validate→repair has empirical support in program repair and synthesis-with-tests domains and motivates guarded NetOps pipelines, but those demonstrations are primarily in software-engineering contexts rather than artifact-backed NetOps verifier studies [4],[5],[6].

Survey and safety literature on agentic AI documents substantive failure modes (hallucination/confabulation) and argues for verifier gates and provenance; this conceptual grounding motivated our guarded design and strict oracle definition but does not supply the verifier-backed NetOps empirical outcomes reported here [7],[8]. Synthesis. The supplied, verified records indicate active prototyping and conceptual advocacy for verifier-augmented pipelines but leave an empirical gap: few reproducible paired experiments quantify the marginal repair effect when the identical initial candidate is reused. Our paired shared_initial_candidate design directly addresses that gap.

III. METHODOLOGY

Preregistration and design freeze. The study design and analysis plan were frozen in preregistration CAND-2026-001 prior to any model calls. The preregistration fixed: (a) a paired binary estimand (paired_success_rate_difference_guarded_minus_baseline) comparing per-task binary verifier acceptance under baseline (no repair) versus guarded (verifier may trigger ≤ 1 repair) conditions; (b) deterministic task selection for $N=200$ tasks; (c) generation_semantics = shared_initial_candidate (one sampled initial generation per task deterministically reused by both episodes); (d) a maximum of one automated repair call per task; (e) primary failed-call treatment complete_pair_primary and a sensitivity treatment that counts persistent unparseables as failures; and (f) the confirmatory test (exact two-sided McNemar via exact binomial tail) plus a paired-bootstrap procedure for a percentile CI. The preregistration artifact (immutable) is referenced in the execution manifest.

Deterministic task sampling. Tasks were produced by a deterministic task generator (deterministic_netops_task_generator_v5) and selected by a fixed index rule documented in the preregistration. The selection rule is deterministic across the 200-task bank (the preregistration specifies the cycle-and-offset index mapping used to pick each task index); this ensures reproducible task composition and eliminates post-hoc selection. Each manifest encodes: initial_state, expected_final_state, required_sequence (ordered command list), allowed_touched_fields, workflow_policies (group-level constraints, minimum-active rules), and difficulty metadata (changed_interface_count, dependency_rule_count, required_command_count). Difficulty labels used in the study include very_hard and extreme; representative required_command_count values ranged from 3 up to 9 in the archived sample.

Shared_initial_candidate semantics and sampling notes. The adapter contract enforced shared_initial_candidate semantics: for each pair the exact same initial generation (a single sampled candidate) was provided to both the baseline and guarded episodes. Consequently, any observed paired difference can only arise from the action of the deterministic verifier and the single permitted repair; it cannot arise from independent first-pass guarded generations, different sampling calls, prompt-engineering differences, or selection among multiple generated candidates. The model-configuration record declares the requested model as gpt-5-mini and records temperature = null (unset). We explicitly note that temperature = null/unset in the configuration file is not evidence of deterministic sampling and does not imply temperature = 0; nondeterministic hosted-model sampling therefore remains a possible background property even though the experiment used a single initial sample per pair by design.

Generation, sanitizer, and repair policy. For each task the pipeline executed a deterministic formatting sanitizer (automatic_syntax_sanitizer_v1) to correct minor syntactic normalization issues before verifier parsing; this sanitizer is deter-

ministic and is applied prior to scoring per the preregistration. After the shared initial generation and sanitizer, the deterministic verifier parsed the candidate. If the verifier reported any invariant violations, the guarded logic was permitted to invoke at most one automated repair call to the hosted model with a repair prompt constructed from the validator trace (this repair policy and prompt format were specified in the preregistration). Repair calls were counted against the maximum_model_calls budget; the execution manifest reports exactly one repair call across all pairs in the run. The pipeline’s missingness plan treated persistent unparseable outputs (after sanitizer and an allowed repair attempt) as scorable failures in the sensitivity analysis; in the primary analysis pairs with complete episodes remained eligible per the preregistered complete_pair_primary rule.

Deterministic verifier as the oracle. The deterministic verifier (deterministic_netops_validator_v5) is the study’s operational oracle. The verifier implements: syntactic parsing into discrete commands; enforcement of allowed_fields and allowed_touched_fields per task manifest; comparison of parsed command sequences against the manifest’s required_sequence ordering; evaluation of group-level invariants (e.g., final_group_constraints and minimum_active_in_groups); and a final boolean full_intent_constraint_validation = true only if all encoded constraints are satisfied. The verifier emits structured tracer output for each episode that includes parsed_command_count, required_command_count, observed_touched_field_count, normalized_configuration, validation_before (violation codes and counts), validation_after (post-repair state and validity flag), validator_id and version, and a human-readable description of detected invariant breaches. These tracer fields are the authoritative evidence for per-episode scoring in the analysis.

Paired execution bookkeeping and provider accounting. The preregistered planned episodes were 400 (200 pairs $\times$ 2 conditions). The run completed all planned episodes without terminal provider failures. Provider accounting recorded hosted-model-call events and stage counts consistent with preregistration: the stage counts indicated 200 shared initial generation events and 1 repair event, for a total of 201 model calls. Cache and reuse checks verified that the shared initial candidate was reused per pair and that no cross-condition cache-reuse introduced additional independent generation calls. All runtime provider metadata (call counts, cache hits/misses, stage counts) were reconciled deterministically in the execution reconciliation artifact.

Statistical procedures — confirmatory and CI computation. The primary confirmatory test is the exact two-sided McNemar test computed from discordant pair counts n_{10} and n_{01} (where n_{10} = count of pairs with guarded=1 & baseline=0, and n_{01} = count with guarded=0 & baseline=1). We implemented the exact two-sided McNemar p-value using the exact binomial-tail formulation: let $n = n_{10} + n_{01}$ and $k = n_{10}$; the two-sided p-value is $p = 2 * \min\{ \text{BinomCDF}(k; n, 0.5), 1 - \text{BinomCDF}(k - 1; n, 0.5) \}$ with the conventional handling when $k = 0$. This exact-binomial realization is documented in

the analysis log and reproduces the $p = 1.0$ reported for the observed discordants ($n = 1$, $k = 1$). For the paired absolute risk difference we computed a paired-bootstrap 95% percentile confidence interval with 10,000 resamples and `bootstrap_seed = 1729`; bootstrap parameters and seed are recorded in the analysis log.

Missingness, exclusions, and contamination controls. The preregistration defined deterministic exclusion criteria prior to model calls (e.g., token-normalized identity to public vendor-source templates to avoid leakage) and required logging of any excluded tasks. The run recorded no preregistered exclusions. The sanitizer and missingness policy were applied deterministically; persistent unparsables did not occur in this execution. Contamination checks included token-level similarity scans between generated candidates and vendor-template sources and a contamination summary that would flag unusually high overlaps; no contamination flags were raised in this run.

Reproducibility and artifact grounding. All implementation details required to reproduce the methodology are archived in the execution bundle: the immutable preregistration, the deterministic task generator artifacts and manifests (task bank with `required_sequence` and `workflow_policies`), the model-configuration record (including declared model and the explicit `temperature = null/unset` entry), the verifier implementation and structured tracer outputs, provider-call accounting and deterministic reconciliation, the sanitizer and repair-prompt templates, and the analysis-executor configuration (exact McNemar code path and bootstrap seed/resample settings). The verifier’s tracer fields provide the operational link between generated text and the binary scoring used in the confirmatory analysis. Re-running the archived paired analysis executor on the raw results reproduces the contingency counts and CI parameters reported in the paper.

IV. RESULTS

Execution accounting and provider audit. The run started at 2026-09-01T16:03:05.491039+00:00 and completed at 2026-09-01T16:32:13.298536+00:00 (≈ 29.2 minutes wall-clock). The execution manifest records `planned_episode_count = 400` and `completed_episode_count = 400` with `failed_episode_count = 0`. Provider-call accounting (`deterministic_reconciliation.json`) reports `hosted_model_call_event_count = 201`, with `stage_counts = {shared_initial_generation: 200, repair: 1}` and `condition_counts` reported as `{shared: 200, guarded: 1}` under the adapter’s stage/condition accounting semantics. Cache statistics show `cache_miss_count = 201` and `cache_hit_count = 0`, consistent with fresh shared-initial generations for all tasks plus a single repair call. Dividing the wall-clock duration by the 201 hosted-model calls yields an approximate average model-call latency of 8.7 seconds per call (simple elapsed-time / call count calculation from archived timestamps and `model_calls_used`).

Primary confirmatory outcomes (archived CSV artifacts). The archived condition summary (`analysis/condition_summary.csv`) records baseline successes

$= 199$, baseline failures $= 1$ (`success_rate = 0.995`) and guarded successes $= 200$, guarded failures $= 0$ (`success_rate = 1.000`). The paired contingency table (`analysis/paired_contingency_table.csv`) contains the exact cell counts used by the confirmatory test: $n_{11} = 199$ (both pass), $n_{10} = 1$ (guarded pass only), $n_{01} = 0$ (baseline pass only), $n_{00} = 0$ (both fail). From these counts the observed paired absolute risk difference (guarded - baseline) $= 0.005$.

Exact-test and interval quantification. Applying the exact two-sided McNemar (binomial-tail formulation) to the discordant counts ($n = n_{10} + n_{01} = 1$, $k = n_{10} = 1$) yields $p = 1.0$ as computed in Methods. To quantify plausible effect sizes we computed a paired-bootstrap percentile 95% confidence interval for the absolute risk difference using 10,000 resamples with `bootstrap_seed = 1729` (`analysis/analysis_log.jsonl`). The resulting 95% percentile CI is $[0.00, 0.015]$, giving a bounded estimate consistent with at most small absolute improvements under the repair-only estimand given the observed data.

Missingness, contamination, and deterministic reconciliation. The run recorded no preregistered exclusions and no persistent unparsables; `analysis/missingness_summary.csv` shows `complete_pairs = 200` and zeros for all other missingness categories. `analysis/contamination_summary.csv` contains no flagged contamination entries. The deterministic reconciliation artifact (`analysis/deterministic_reconciliation.json`) reports `pair_accounting_consistent = true` and `all_deterministic_consistency_checks_passed = true`, confirming internal accounting consistency between provider-call logs, `stage_counts`, and analysis inputs.

Localized failure diagnosis and repair evidence. The guarded pipeline invoked exactly one automated repair (`stage_counts` shows `repair = 1`). Validator traces for the repaired pair (archived under `execution/scoring/`) indicate an extreme composite task whose manifest metadata records `required_command_count = 9` and `changed_interface_count = 3`. The shared initial candidate for that pair violated the manifest’s `required_sequence` ordering constraint; `validation_before.violation_codes` and `validation_before.normalized_configuration` recorded the pre-repair mismatch. The repair interaction produced a localized reordering of commands that matched the manifest’s `required_sequence` and yielded `validation_after.valid = true`. The evidence supporting this sequence of events is entirely artifact-grounded: the per-episode `validator_trace` fields (`validation_before.violation_codes`, `parsed_command_count`, `required_command_count`, `normalized_configuration`, `repair_applied` flag, and `validation_after.*`) and the repair provider trace are archived and indexable. These traces show that the repair fixed an ordering/dependency violation expressible in the task manifest’s invariants rather than a syntactic parsing error.

Exploratory diagnostics by difficulty metadata. Representative task manifests available in the execution bundle document `required_command_count` values ranging from 3 up to 9 and difficulty levels labeled `very_hard` or `extreme`. The

single baseline failure (the repaired pair) is associated with the highest recorded complexity in the representative sample (`required_command_count` = 9, `changed_interface_count` = 3), consistent with the preregistered exploratory hypothesis that multi-device sequencing and dependency complexity concentrate residual failures. No other tasks in the archived representative subset exhibited violations requiring repair; many very_hard tasks (`required_command_count` in the mid-range) were validated successfully without repair.

Practical reproducibility pointers for confirming results. All artifacts needed to reproduce the primary numbers are archived and referenced by the execution manifest: `analysis/paired_contingency_table.csv` and `analysis/condition_summary.csv` contain the per-condition counts; `analysis/analysis_log.jsonl` contains the bootstrap parameters (`resamples`=10000, `seed`=1729); `deterministic_reconciliation.json` and `execution/model_configuration.json` record provider-call accounting and declared model fields (including `temperature` = null); `execution/responses/` and `execution/scoring/` contain the shared-initial responses and per-episode `validator_trace` objects demonstrating `validation_before` and `validation_after` states and the single repair interaction. Re-running the archived `paired_binary_analysis_v1` executor on the archived `input_results_sha256` reproduces the confirmatory outputs (`analysis/results.json`) without recomputation of model calls.

Summary of empirical evidence. The archived run shows that under `shared_initial_candidate` semantics (the same sampled candidate reused by both conditions) and with the chosen synthetic task bank and deterministic verifier, nearly all initial single-shot candidates already satisfied the encoded deterministic invariants (199/200). Allowing a single verifier-triggered bounded repair corrected one manifest-expressed ordering violation in an extreme composite task, producing the observed paired improvement ($n_{10} = 1$). The small number of discordants places the inferential regime in a sparse-discordant setting where the exact McNemar test yields $p = 1.0$ but the bootstrap CI bounds the plausible absolute improvement to at most ~1.5 percentage points under the paired repair-only estimand.

V. DISCUSSION

We expand the interpretive analysis and diagnostics here while preserving the paper’s existing supported content.

Estimand interpretation and scope (explicit). Because the experiment used `shared_initial_candidate` semantics, baseline and guarded episodes for each task began from the identical sampled candidate; consequently the primary estimand intentionally measures the marginal effect of permitting an automated verifier-triggered repair on that identical candidate (a repair-only effect). This design isolates repair mechanics from any potential benefits that could arise if the guarded condition were allowed to sample independently or draw multiple candidates.

Sparse-discordant regime and inferential consequences. The observed data (n_{11} =199, n_{10} =1, n_{01} =0) place the analysis

in a sparse-discordant regime, a setting where conventional asymptotic approximations are unreliable and exact discrete tests or resampling are appropriate. Our use of the exact binomial-tail McNemar calculation and a paired-bootstrap percentile CI (10,000 resamples, seed 1729 as archived in `analysis/analysis_log.jsonl`) is consistent with best practice for small discordant counts. Practically, the realized baseline success rate (99.5%) substantially reduced the experiment’s sensitivity to detect modest repair-only effects relative to our preregistered power target (designed for mid-range baseline rates ~45–65%). The bootstrap CI [0.00, 0.015] quantifies that the data are compatible with effects up to 1.5 percentage points in absolute terms under our paired repair-only estimand; the preregistered design could not reliably detect larger effects because none were present.

Mechanistic diagnostic of the repaired episode. The single repair call corrected a concrete sequencing violation expressible in the task manifest’s `required_sequence` field (`required_command_count` = 9 in the repaired manifest). Archived `validator_trace` fields show `validation_before.violation_codes` listing the ordering breach, `parsed_command_count` equal to the presented commands, and `validation_after.valid` = true after the automated reordering repair. These per-episode traces demonstrate: (1) the verifier’s capacity to detect ordering invariants encoded in the manifest; (2) the repair module’s ability to perform localized, semantics-preserving reorderings that the verifier accepts; and (3) that such faults concentrated in tasks with higher `required_command_count` and extreme difficulty labels in our deterministic task bank. All diagnostic fields cited above are available in `execution/scoring/task-00000X-*.json` in the archived bundle.

Construct validity of the deterministic verifier. `deterministic_netops_validator_v5` enforces a structured subset of operational invariants encoded in the task payloads (ordering, `allowed_touched_fields`, group-level up/down constraints, and explicit per-field prerequisites). This verifier therefore provides a high-fidelity oracle for manifest-encoded invariants but does not emulate dataplane behavior, vendor runtime idiosyncrasies, or external control-plane side-effects. Construct-valid conclusions should therefore be framed strictly in terms of manifest-specified, deterministically verifiable invariants; the experiment does not claim to validate dataplane reachability or vendor CLI execution semantics that lie outside the verifier’s encoding.

Design implications for future NetOps evaluations. The paired `shared_initial_candidate` design cleanly isolates repair capability; however, operational questions remain that require alternative designs: (a) independent guarded-first-pass sampling would quantify whether guarded prompting or constrained generation yields a higher baseline of valid samples before repair; (b) multi-sample selection (n -best) could reveal selection-improvement effects; and (c) multi-model comparisons would quantify model-dependent failure patterns. From an experimental-planning perspective our results highlight that preregistered power calculations must account for plausible

high baseline performance—if baseline rates turn out near 100% the discordant-pair approach will often be underpowered to detect small absolute repairs-only gains.

Reproducibility, artifact usage, and recommended reanalyses. All artifacts necessary to reproduce analysis are archived: `execution/task_manifest.jsonl` (deterministic task selection and `required_sequence` fields), `execution/model_configuration.json` (declared model, noting `temperature = null/unset`), `execution/responses/` (shared-initial and per-condition responses), `execution/scoring/` (per-episode validator traces), `analysis/paired_contingency_table.csv` and `analysis/condition_summary.csv`, `deterministic_reconciliation.json` (provider accounting and `stage_counts`), and `analysis/analysis_log.jsonl` (bootstrap parameters and seed). Re-running the `paired_binary_analysis_v1` executor on the archived `input_results_sha256` reproduces the confirmatory results. For targeted reanalysis, researchers can (i) recompute bootstrap intervals under alternative resampling schemes, (ii) treat the single repaired task as an outlier to study robustness, or (iii) reclassify manifest invariants to evaluate sensitivity of verifier coverage—all without altering original model-call artifacts.

Threats to validity revisited (artifact-grounded). Internal validity: `shared_initial_candidate` semantics intentionally eliminate differences due to separate sampling, but they also prevent measuring effects of guarded prompt variants; model sampling nondeterminism remains possible because `temperature` was unset (`execution/model_configuration.json` records `temperature = null`). Construct validity: the verifier enforces only encoded invariants; omissions in the manifest (for example, dataplane-dependent invariants) are not tested. External validity: a single hosted model (gpt-5-mini) and a synthetic deterministic task bank limit generalizability; results do not imply identical outcomes for other models, live devices, or telemetry-driven validation. Conclusion validity: with one discordant pair the statistical evidence for a nonzero repair-only effect is weak ($p = 1.0$), though the bootstrap CI constrains plausible magnitudes.

Operational implications and cautious hypotheses. Artifact-grounded evidence demonstrates that bounded verifier-triggered repairs can fix localized ordering violations that are detectable and expressible in task manifests. From an operational standpoint this suggests a practical pattern: require explicit, machine-checkable manifest invariants for change tasks and pair LLM synthesis with deterministic validation/repair gates before human or automated deployment. We emphasize these are actionable hypotheses supported by the archived diagnostics in this controlled synthetic setting; they are not operational claims that bounded repair will resolve dataplane-level or adversarially induced failures without further evaluation.

Summary of expanded diagnostics. The archived execution bundle provides the evidentiary chain: deterministic task manifests encoding the invariants, the shared initial generated candidates, per-episode validator traces showing before/after states, provider-call accounting confirming

`shared_initial_generation = 200` and `repair = 1`, and analysis artifacts (contingency table, bootstrap parameters) documenting the statistical regime. Together these artifacts ground the narrower, repair-only conclusion and illuminate concrete design decisions for future, broader evaluations (independent guarded sampling, multi-model studies, live emulation).

VI. LIMITATIONS

- Synthetic deterministic task bank and verifier: results reflect correctness under the chosen synthetic intent templates and deterministic verifier and do not guarantee similar performance on live heterogeneous production devices or telemetry.
- Single-model experiment (gpt-5-mini): findings do not demonstrate multi-model generality.
- `Shared_initial_candidate` semantics: baseline and guarded conditions began from the same initial generation; the study measures verifier/repair effects only and does not evaluate independent first-pass generation contrasts.
- Limited adversarial/contamination testing: the run did not include systematic adversarial retrieval or poisoned-context attacks; such threats remain an open evaluation axis.
- Oracle coverage: the deterministic verifier enforces encoded invariants but may omit runtime behaviors and device-specific semantics observable only through live execution; external validity to live deployments is therefore limited.

VII. CONCLUSION

We executed a preregistered paired experiment measuring the marginal effect of a verifier-triggered repair under `shared_initial_candidate` semantics for LLM-generated network configurations. In 200 synthetic deterministic tasks with gpt-5-mini and `deterministic_netops_validator_v5`, baseline acceptance was 199/200 and guarded acceptance was 200/200; a single automated repair corrected a multi-device ordering violation. Paired absolute risk difference = 0.005 with paired-bootstrap 95% CI [0.00, 0.015] (10,000 resamples, seed 1729) and exact two-sided McNemar $p = 1.0$. Artifact-grounded evidence therefore supports a constrained conclusion: verifier-guarded bounded repair can correct localized ordering faults in this synthetic verifier-backed setting. Broader operational claims require additional experiments: multi-model studies, independent-first-pass guarded semantics, adversarial contamination testing, and live-device or dataplane emulation to capture runtime semantics not modeled by the deterministic verifier.

DISCLOSURE STATEMENT

Disclosure: This run implemented preregistered study CAND-2026-001 and was executed autonomously after the autonomy lock. Declared model: gpt-5-mini (`execution/model_configuration.json` records `temperature = null/unset`; `null/unset` is not evidence of deterministic sampling or `temperature = 0`). Orchestration adapter:

hosted_netops_gvr_v1. Exact initial master prompt archived at `provenance/master_prompt.txt` (SHA-256 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acdf1582bb39b15ba). Preregistration fixed `generation_semantics = shared_initial_candidate` (one sampled initial generation per task was deterministically reused by baseline and guarded episodes) and allowed at most one automated repair call per task. Model-call accounting in the archived execution manifest reflects 200 shared initial generation calls and 1 repair call (201 model calls total). The deterministic verifier used was `deterministic_netops_validator_v5`. Humans selected the broad topic family and constrained the initial master prompt prior to the autonomy lock as permitted by the track; no human manually edited technical manuscript content after the autonomy lock. All provider traces, verifier logs, task manifests, model-configuration records, and analysis artifacts are archived in the execution bundle and indexed by the execution manifest.

REFERENCES

- [1] Changjie Wang, Mariano Scazzariello, Alireza Farshin, Simone Ferlin, Dejan Kostić, Marco Chiesa, "NetConfEval: Can LLMs Facilitate Network Configuration?", 2024, 10.1145/3656296
- [2] Ryan Beckett, Aarti Gupta, Ratul Mahajan, David Walker, "A General Approach to Network Configuration Verification", 2017, 10.1145/3098822.3098834
- [3] Jack Gallifant, Majid Afshar, Saleem Ameen, Yindalon Aphinyanaphongs, Shan Chen, Giovanni Cacciamani, Dina Demner-Fushman, Dmitriy Dligach, Roxana Daneshjou, Chrystinne Oliveira Fernandes, Lasse Hyldig Hansen, Adam Landman, Lisa Soleymani Lehmann, Liam G. McCoy, Timothy A. Miller, Amy C. Moreno, Nikolaj Munch, David Restrepo, Guergana Savova, Renato Umeton, Judy Wawira Gichoya, Professor Gary S. Collins, Karel G.M. Moons, Leo Anthony Celi, Danielle S. Bitterman, "The TRIPOD-LLM reporting guideline for studies using large language models", 2025, 10.1038/s41591-024-03425-5
- [4] Yunhe Li, "Test-in-the-loop LLM Repair: Verifiable Automated Program Repair on QuixBugs with a "Failing Test → Patch → Regression Test" Loop", 2024, 10.69987/jacs.2024.40206
- [5] Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Zican Dong, Yupeng Hou, Beichen Zhang, Yingqian Min, Junjie Zhang, Peiyu Liu, Xiaolei Wang, Yifan Du, Yushuo Chen, Yushuo Chen, Zhipeng Chen, Jinhao Jiang, Ruiyang Ren, Yifan Li, Xinyu Tang, Peiyu Liu, Yiwen Hu, Jian-Yun Nie, Ji-Rong Wen, "A Survey of Large Language Models", 2026, 10.1007/s11704-026-60308-3
- [6] Haoran Dong, "Tool-Augmented Large Language Model Agents for Analysis: A Focused Study", 2026, 10.2139/ssrn.6647800