

# Validator-Guided Repair on a Shared Initial LLM Candidate: A Preregistered Paired Evaluation for Network Configuration Safety

Autonomous AI Researcher  
CNSM 2026 Agentic AI Researcher Track

**Abstract**—We preregistered and executed a paired confirmatory experiment (PREREG-CAND-001-VER-GVR-2026-09-01) to measure whether permitting a deterministic network-configuration validator plus one automated repair changes the rate of validator-valid executable configurations relative to leaving the identical sampled LLM candidate unguarded (shared\_initial\_candidate semantics). Using the declared generator gpt-5-mini (hosted adapter openai\_responses) and deterministic\_netops\_validator\_v5 on  $N = 200$  paired intents (completed\_episode\_count = 400; model\_calls\_used = 200), every shared initial candidate passed validator checks and no repairs were attempted or applied (repair\_attempt\_count = 0; repair\_applied\_count = 0). Observed outcomes: baseline\_success = 200/200, guarded\_success = 200/200; paired contingency  $n_{11} = 200$ ,  $n_{10} = 0$ ,  $n_{01} = 0$ ,  $n_{00} = 0$ ; exact McNemar two-sided  $p = 1.0$ ; paired bootstrap percentile 95% CI (seed = 1729; 10,000 resamples) = [0.0, 0.0]. We present artifact-grounded methodology, execution accounting, representative validator diagnostics, compact contingency and repair summaries, and a focused discussion clarifying the causal limits of the executed estimand under shared\_initial\_candidate semantics and preregistered follow-on designs required to measure repair efficacy under independent sampling, higher difficulty, or adversarial inputs.

## I. INTRODUCTION

Generative large language models (LLMs) are increasingly proposed to translate operator intent into device configuration sequences in network operations (NetOps). Erroneous sequences, syntactic mistakes, or fabricated fields can cause service disruption or policy violations when applied to devices. A common architectural mitigation is a generate→validate→repair (GVR) pipeline: generate candidate configuration(s), deterministically validate them against intent/topology/workflow constraints, and trigger automated repair when the validator finds faults. Prior work documents verifier-guided improvements in infrastructure-as-code and related code-generation domains and recommends grounding strategies to reduce hallucination in operational tasks [1–3]. However, whether verifier-guided pipelines reduce execution-level operational risk for network configurations remains an open empirical question because prior demonstrations often measure textual correctness or IaC test suites rather than executed device-state safety in packet-level or emulator contexts [3,4].

This paper reports a fully preregistered paired confirmatory experiment that isolates the downstream causal contribution of a deterministic validator and any permitted repair acting on the

exact same sampled LLM candidate (shared\_initial\_candidate semantics). Under these semantics baseline and guarded episodes for each intent begin from the identical initial generation; any paired difference can therefore only arise from deterministic validation and a permitted repair of that same candidate. Our primary research question is: under shared\_initial\_candidate semantics, does permitting a deterministic validator plus at most one automated repair change the proportion of validator-valid executable configurations relative to leaving the same initial candidate unguarded? The contribution is an artifact-grounded report of the executed estimand with complete execution accounting, exact inferential statistics, representative validator diagnostics, and precise recommendations for preregistered follow-on designs that can measure repair efficacy when independent sampling, elevated difficulty, or adversarial inputs are employed.

## II. RELATED WORK

Hallucination and factual unreliability in LLM outputs are well documented and motivate grounding, retrieval, and verification strategies for operational tasks [1]. Retrieval-augmented and evidence-grounded pipelines can reduce hallucination in operations-style QA, though these works evaluate textual or documentation correctness rather than device-level executability [2]. Generate→verify→repair architectures have shown empirical improvements in IaC and program generation (for example, verifier-guided fine-tuning and policy-guided repair loops) and are a primary methodological inspiration for our work [3]. Recent NetOps evaluations examine intent translation and LLM capability but emphasize capability benchmarks rather than standardized execution-level safety measures that map generated commands to executable device states in a simulator/emulator [4]. Comprehensive AIOps surveys highlight the need for execution-aware benchmarks and validator fidelity uplift to bridge generation to operational safety [5].

From these verified sources we identified two methodological gaps that motivated the current paired design. First, verifier-guided demonstrations in adjacent domains do not directly measure whether corrections prevent unsafe executable states when enacted on networks; empirical transferability to NetOps therefore requires execution-level validation. Second, many prior studies mix sampling variability with downstream validator/repair effects; precisely defined estimands that separate these sources of variation are uncommon. We therefore

preregistered a shared\_initial\_candidate paired design that isolates the downstream validator/repair effect on identical generator outputs.

### III. METHODOLOGY

**Preregistration and primary estimand.** The study preregistration (PREREG-CAND-001-VER-GVR-2026-09-01) froze the primary estimand and analysis before any model calls. Primary estimand: paired\_success\_rate\_difference\_guarded\_minus\_baseline (per-intent paired difference in proportion of validator-valid executable configurations) under shared\_initial\_candidate semantics. Design: deterministic paired sample of task\_count = 200 intents producing planned\_episode\_count = 400 episodes. Generation semantics: exactly one sampled initial generation per pair (shared) and deterministic reuse by baseline and guarded episodes. Guarded episodes allowed up to one automated repair (maximum\_repair\_calls\_per\_task = 1). Primary test: exact two-sided McNemar test; paired bootstrap (seed = 1729; 10,000 resamples) for a percentile 95% CI.

**Task bank and contamination controls.** A deterministic task generator produced the manifest of 200 synthetic intent\_configuration\_repair\_v1 tasks. Each manifest entry records initial\_state, required\_sequence, difficulty annotations, and a generation\_seed for deterministic replay. Preexecution contamination checks (SHA-256 exact-match and n-gram similarity thresholds) were applied; no exact-match contamination exclusions were flagged and the confirmatory sample executed intact. Contamination and task manifest artifacts are archived in the execution evidence bundle.

**Model calls, generation semantics, and provider audit.** Execution used the declared generator labeled gpt-5-mini via the hosted adapter openai\_responses and consumed model\_calls\_used = 200 (one shared initial generation per pair). The archived model\_configuration.json records declared\_model\_version = "gpt-5-mini" and temperature = null/unset; we explicitly state that null/unset is not evidence of deterministic sampling and does not imply temperature = 0. Provider audit traces record hosted\_model\_call\_event\_count = 200, cache\_miss\_count = 200, and stage\_counts showing shared\_initial\_generation = 200 consistent with the preregistered shared\_initial\_candidate semantics.

**Deterministic validator and repair policy.** Scoring used deterministic\_netops\_validator\_v5 with scoring\_rule = "full\_intent\_constraint\_validation." The validator parses candidate commands, applies task workflow\_policies and required\_sequence checks, and produces structured validator traces with validation\_before and validation\_after subtables, parsed\_command\_count, required\_command\_count, normalized\_configuration, violation\_codes, and a repair\_applied flag. For guarded episodes one automated repair attempt was permitted; the repair, if invoked, would produce a new candidate which the validator would re-validate. Under shared\_initial\_candidate semantics only the guarded arm

could alter the identical initial candidate via the permitted repair; the baseline arm reused the same candidate unchanged.

**Mapping to outcomes and missingness.** Validator.valid = true mapped to a binary "executable/valid" outcome for each episode. Primary analysis adhered to the preregistered complete\_pair\_primary policy (exclude incomplete pairs). A sensitivity analysis (failure-as-zero) was preregistered to treat missing or failed episodes as unsafe. Multiplicity: single confirmatory McNemar test; exploratory subgroup analyses (e.g., by difficulty patterns) were labeled exploratory and corrective procedures (Benjamini–Hochberg) were specified if applied.

### IV. RESULTS

**Execution accounting and deterministic reconciliation.** The execution completed exactly as preregistered: planned\_episode\_count = 400, completed\_episode\_count = 400, failed\_episode\_count = 0, and model\_calls\_used = 200 (one shared initial generation per pair). Analysis missingness\_summary reports complete\_pairs = 200 with zero failed or unscorable episodes. Deterministic reconciliation records hosted\_model\_call\_event\_count = 200, completed\_hosted\_model\_call\_event\_count = 200, cache\_miss\_count = 200, and stage\_counts = {"shared\_initial\_generation": 200}; cross\_condition\_cache\_key\_reuse\_observed = false and all deterministic consistency checks passed.

**Primary binary outcomes and contingency.** The validator mapped each episode to a binary executable outcome (validator.valid). Condition summaries show baseline\_successes = 200/200 and guarded\_successes = 200/200 (success\_rate = 1.0 for both arms). The paired 2×2 contingency table is n11 = 200, n10 = 0, n01 = 0, n00 = 0; exact McNemar two-sided p = 1.0. The paired bootstrap (seed = 1729; 10,000 resamples) produced a percentile 95% CI for the paired difference of [0.0, 0.0]. These files are archived (analysis/condition\_summary.csv and analysis/paired\_contingency\_table.csv) and exactly reproduce the contingency and CI numbers reported here.

**Repair activity, validator diagnostics, and representative examples.** Across the confirmatory sample the validator flagged no violations on any shared initial candidate; as a result repair\_attempt\_count = 0 and repair\_applied\_count = 0. Representative per-episode scoring JSON artifacts illustrate the pattern. For example, task-000004 (pair\_id pair-000004) shared\_initial\_candidate and both condition responses were identical:

```
interface edge2 admin down interface edge2 vlan 40 interface edge2 admin up interface edge1 admin down
```

For task-000004 the archived validator trace shows parsed\_command\_count = 4, required\_command\_count = 4, observed\_touched\_field\_count = 3, and validation\_after.valid = true. Similarly, task-000005 shows parsed\_command\_count = 3 and required\_command\_count = 3 with validation\_after.valid = true; task-000006 shows parsed\_command\_count = 6 and required\_command\_count = 6 with validation\_after.valid = true. These per-episode diagnostics (execution/scoring/task-00000X-\*.json)

demonstrate that the validator parsed the expected command set, matched required\_sequence checks, and reported no violation codes for the sampled candidates. Because the validator never produced a violation, the guarded repair path was never invoked and could not alter any outcome.

Contamination and provider audit summary. Preexecution contamination checks found no exact-match contamination exclusions; analysis/contamination\_summary.csv is empty. Provider audit and execution manifests corroborate the sampling semantics: hosted\_model\_call\_event\_count = 200, cache\_miss\_count = 200, and stage\_counts indicating exactly one shared\_initial\_generation per pair. No infrastructure retries or failures occurred; analysis/analysis\_log.jsonl documents the paired analysis completion and bootstrap parameters (bootstrap\_resamples = 10000; bootstrap\_seed = 1729).

Compact repair summary (explicit). In concise preregistered terms: total\_pairs = 200; repair\_attempt\_count = 0; repair\_applied\_count = 0. The empty repair counts fully explain the degenerate paired contingency (all pairs valid before any repair opportunity). All artifacts supporting these counts are archived in the evidence bundle (notably the execution/scoring JSONs, execution/task\_manifest.jsonl, analysis/paired\_contingency\_table.csv, and analysis/condition\_summary.csv) and permit independent verification that no repairs were produced or applied in guarded episodes.

## V. DISCUSSION

Interpreting the executed estimand. The shared\_initial\_candidate semantics were intentionally chosen to isolate the causal effect of deterministic validation and any permitted repair acting on the exact same generator output. The executed estimand therefore answers a narrowly defined causal question: when baseline and guarded episodes begin from the identical sampled candidate, can deterministic validation and at most one automated repair change the validator-valid/executable outcome? The observed null effect (paired difference = 0.0) shows that, for this task bank and this single sampled candidate per pair, the validator found no violations and the permitted repair path was never exercised. Importantly, this does not evaluate whether guarded pipelines reduce first-pass generator error rates under independent sampling or different model temperatures; such claims would require a different preregistered estimand and execution.

Threats to validity (artifact-grounded). Internal validity: the paired shared\_initial\_candidate design eliminates sampling variability across arms and thus preserves attribution of any observed repair effect to validator/repair alone; however, this design precludes measurement of generator-level effects. Construct validity: deterministic\_netops\_validator\_v5 enforces syntactic and workflow-policy checks and maps results to a boolean valid flag; it is an abstraction that does not execute vendor CLI on physical hardware nor perform packet-level emulation—dynamic protocol/timing effects are therefore not represented. External validity: results are conditioned on one declared generator (gpt-5-mini via adapter openai\_responses),

a synthetic deterministic task bank, and one validator implementation; generalization beyond these artifacts is unsupported. Conclusion validity: zero discordant outcomes preclude inference about repair benefit magnitude under different sampling regimes because the guarded path was not invoked.

Operational implications and preregistered follow-on designs. The null result here should not be read as evidence that GVR pipelines are unnecessary in practice. Instead it demonstrates an estimand boundary: under shared\_initial\_candidate semantics and the executed task bank the guarded pipeline had no opportunity to change outcomes. To empirically evaluate repair efficacy in operationally relevant regimes, we preregistered and recommend the following estimand/design modifications (each reuses archived artifacts where applicable): (A) remove shared\_initial\_candidate semantics so baseline and guarded arms receive independent first-pass samples (this permits the repair path to act on generator faults but requires a revised paired/unpaired analysis plan to separate generation variability from repair effects); (B) select or synthesize tasks labeled "extreme" difficulty or inject adversarial perturbations to raise expected validator failure rates and thereby exercise repair logic; (C) preregister non-null temperature values and temperature sweeps to evaluate sample diversity versus validator-failure tradeoffs; (D) increase validator fidelity by integrating vendor CLI semantics or packet-level emulation so repairs are evaluated against dynamic runtime behavior; (E) expand to multi-model sampling to assess generality across generator families.

Reproducibility and artifact utility. All per-episode scoring JSON artifacts include validator traces with validation\_before/validation\_after subtables, normalized\_configuration, parsed\_command\_count, required\_command\_count, and repair\_applied flags permitting independent verification that repairs did not run. Execution and analysis artifacts (execution/raw\_results.jsonl, execution/task\_manifest.jsonl, analysis/paired\_contingency\_table.csv, analysis/condition\_summary.csv, analysis/analysis\_log.jsonl) are archived in the evidence bundle and enable replay of the exact shared\_initial\_candidate sampling. The model\_configuration.json records declared\_model\_version = "gpt-5-mini" and temperature = null; null/unset must not be interpreted as deterministic sampling. The master prompt immutable reference is archived (see Disclosure). These artifacts are sufficient to reproduce the executed estimand and deterministic reconciliation.

## VI. LIMITATIONS

- Shared\_initial\_candidate semantics: the design produced exactly one sampled initial generation per intent and reused it across baseline and guarded conditions; this measures validator/repair effects only and cannot detect differences caused by independent per-condition sampling or prompt engineering.
- Temperature unset (temperature = null) in the archived model configuration: null/unset is not evidence of deter-

ministic sampling and does not imply temperature = 0; sampling nondeterminism may still have occurred during the single sampled generation.

- Single model and single hosted provider: results reflect gpt-5-mini under the archived adapter; generalization to other models or model families is unsupported by these artifacts.
- Synthetic task bank and validator abstraction: tasks were synthetic 'intent\_configuration\_repair\_v1' items and the deterministic validator is an abstraction; realistic vendor CLI idiosyncrasies, emergent runtime interactions, and hardware effects were not executed on live devices.
- No observed validator failures: all initial candidates were validator-valid, so the experiment could not assess the repair step's effectiveness; the null effect therefore reflects the task bank and sampling semantics rather than evidence that GVR is unnecessary in general.

## VII. CONCLUSION

We executed a fully preregistered paired confirmatory experiment that measures the downstream effect of a deterministic validator plus an allowed single automated repair on the exact same sampled LLM candidate per intent (shared\_initial\_candidate semantics). Using the declared generator gpt-5-mini and deterministic\_netops\_validator\_v5 on 200 paired intents, every shared initial candidate passed validation and no repairs were attempted or applied; guarded and baseline outcomes were identical for the executed estimand (paired difference = 0.0; McNemar  $p = 1.0$ ; paired bootstrap CI = [0.0, 0.0]). This artifact-grounded null result demonstrates an estimand boundary: under shared\_initial\_candidate semantics and the executed task bank the guarded pipeline could not be exercised and therefore could not change outcomes. It does not imply that GVR pipelines are unnecessary generally. Preregistered follow-on experiments that (1) remove shared\_initial\_candidate semantics, (2) increase task difficulty or inject adversarial inputs, (3) set explicit sampling parameters, and (4) raise validator fidelity are required to empirically evaluate repair efficacy under realistic sampling variability and adversarial conditions. The archived artifacts and reconciliation files enable exact reproduction of the executed estimand and provide a secure base for precise preregistered extensions.

References [1] A. Alansari and H. Luqman, "Large Language Models Hallucination: A Comprehensive Survey," arXiv preprint arXiv:2510.06265, 2025. [2] B. Zhang, H. Rao, and D. Zhao, "Evidence-Grounded RAG for Cloud-Native DevOps: Hallucination-Resistant AIOps Question Answering over Private Operations Documents," J. Autom. Cloud-Native Syst., 2024. [3] P. Jana et al., "TerraFormer: Automated Infrastructure-as-Code with LLMs Fine-Tuned via Policy-Guided Verifier Feedback," Proc. ACM Symp. (2026). [4] Y. Miao et al., "An Empirical Study of NetOps Capability of Pre-Trained Large Language Models," arXiv:2309.05557, 2023. [5] L. Zhang et al., "A Survey of AIOps in the Era of Large Language Models," ACM Comput. Surv., 2025.

One-line reiteration of the primary estimand limit: the preregistered primary estimand intentionally used shared\_initial\_candidate semantics; evaluating per-condition independent sampling requires a different preregistration and analysis plan (explicitly recommended above).

## DISCLOSURE STATEMENT

Disclosure (concise): The human Research Director selected the broad topic family (Generative AI / LLMs for NetOps) before the autonomy lock. After the autonomy lock the autonomous system performed literature review, experiment selection, preregistration (PREREG-CAND-001-VER-GVR-2026-09-01), code generation, execution, deterministic validation, automated analysis, and manuscript drafting without manual human editing of technical content. Declared model used in execution: gpt-5-mini via hosted adapter 'openai\_responses' (execution used model\_calls\_used = 200; archived model\_configuration.json records temperature = null/unset; null/unset is not evidence of deterministic sampling and does not imply temperature = 0). Generation semantics executed: shared\_initial\_candidate — exactly one sampled initial generation per intent was produced and deterministically reused for both baseline and guarded conditions. Validator: deterministic\_netops\_validator\_v5 scored candidates using 'full\_intent\_constraint\_validation' and a single automated repair iteration was permitted in guarded episodes; in this run the validator flagged no violations and no repairs were applied (repair\_attempt\_count = 0; repair\_applied\_count = 0). Execution accounting: completed\_episode\_count = 400 (200 paired intents) completed without preregistration deviations. Master prompt immutable reference: provenance/master\_prompt.txt — SHA-256: 1872df1e1805d2d96940456ca016bd665d1d5196add77f5acd1582bb39b15ba. Pre-lock infrastructure development and rehearsal occurred but pre-lock artifacts were not used as empirical evidence in this final run. After the autonomy lock no human manually rewrote technical results, manuscript text, figures, or references.

## REFERENCES

- [1] Aisha Alansari, Hamzah Luqman, "Large Language Models Hallucination: A Comprehensive Survey", 2025, 10.48550/arxiv.2510.06265
- [2] Boning Zhang, Hengning Rao, Derek Zhao, "Evidence-Grounded RAG for Cloud-Native DevOps: Hallucination-Resistant AIOps Question Answering over Private Operations Documents", 2024, 10.69987/jacs.2024.40308
- [3] Prithwish Jana, Sam Davidson, Bhavana Bhasker, Andrey Kan, Anoop Deoras, Laurent Callot, "TerraFormer: Automated Infrastructure-as-Code with LLMs Fine-Tuned via Policy-Guided Verifier Feedback", 2026, 10.1145/3786583.3786898
- [4] Yukai Miao, Yu Bai, Li Chen, Dan Li, Haifeng Sun, Xizheng Wang, Ziqiu Luo, Ren, Yanyu, Dapeng Sun, Xiuting Xu, Qi Zhang, Chao Xiang, Xinchu Li, "An Empirical Study of NetOps Capability of Pre-Trained Large Language Models", 2023, 10.48550/arxiv.2309.05557
- [5] Lingzhe Zhang, Tong Jia, Mengxi Jia, Yifan Wu, Aiwei Liu, Yong Yang, Zhonghai Wu, Xuming Hu, Philip S. Yu, Ying Li, "A Survey of AIOps in the Era of Large Language Models", 2025, 10.1145/3746635